Synthesizing the Consensus: Generative AI, LLM Visibility, and the Integration of Multi-Provider Consensus Forecasts*

Leonard Yang Liu†

September 2026

[Newest Version]

Abstract

The consensus earnings forecast is not one number: forecast data providers (FDPs) construct economically distinct consensus numbers for the same firm-quarter. This paper asks whether generative AI (GenAI) changes how equity markets synthesize consensus forecasts across providers. I study a panel of S&P 1500 firm-quarters from 2021 to 2025 with consensus forecasts from five FDPs—FactSet, Bloomberg, Zacks, S&P Capital IQ, and I/B/E/S. After ChatGPT's release, the market reacts more strongly to forecasts from more-accurate FDPs: provider-specific earnings response coefficients align more closely with each FDP's accuracy rank. Post-ChatGPT, market reactions also align increasingly with an accuracy-weighted consensus across the five FDPs rather than with the I/B/E/S-only or equal-weighted heuristic alternatives. The shift concentrates where GenAI-enabled synthesis is most valuable: in firm-quarters with high cross-FDP disagreement and larger nonrecurring items, and—once GenAI tools can retrieve forecasts from the web—in firms with high large language model (LLM) visibility. A long–short strategy trading the gap between the accuracy-weighted and heuristic consensus measures earns economically meaningful returns. Google search data support a post-ChatGPT shift in investor search toward free forecast distributors, and tests using major ChatGPT outages as exogenous shocks point to GenAI as the driver of the shift. The evidence suggests that GenAI helps investors both gather and integrate forecast information—and changes how that information is impounded into prices.

Keywords: Generative AI, Analyst, Forecast Data Provider, Consensus Forecast, Earnings Surprise, Information Processing, Earnings Response Coefficient.

JEL classification: G11, G14, G41, M41.

*I am deeply grateful to my dissertation committee for their invaluable guidance and support: Musa Subasi (chair), Michael Kimbrough (co-chair), Rebecca Hann, Nick Seybert, Jingyi Qian, and Erkut Y. Ozbay. I gratefully acknowledge the Texas A&M University AI Dissertation Award in Accounting. All errors are my own.

†Department of Accounting and Information Assurance, Robert H. Smith School of Business, University of Maryland. Email: yliu337@umd.edu.

1 Introduction

"The consensus" earnings forecast is not one number. Sell-side analyst forecasts serve as the primary market benchmark against which realized earnings are evaluated (e.g., Bradshaw et al., 2017; Kothari, 2001). The earnings surprise relative to this consensus is a key determinant of announcement-window stock price reactions and post-announcement drift (e.g., Bartov et al., 2002; Bernard and Thomas, 1989; Livnat and Mendenhall, 2006), with measurable consequences for price discovery, trading volume, and firms' own disclosure choices. Yet multiple *forecast data providers* (FDPs) construct economically distinct consensus numbers for the same firm-quarter through proprietary aggregation rules. Larocque et al. (2025) document these differences across the five leading FDPs: I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks.[1] Although these FDPs are widely used and cited for investment decisions (Coleman et al., 2025; Xue et al., 2025), little is known about whether the market reacts differently across FDPs and, more importantly, how it selects and integrates across them.

The question matters because acquiring and integrating data across multiple providers is costly.[2] Since the release of ChatGPT, generative AI (GenAI) has diffused rapidly into investor research and analysis workflows, reshaping the economics of how financial information is gathered and combined. Industry surveys document widespread adoption (CFA Institute, 2024), and a growing body of research shows that these tools substantially lower the costs of acquiring, integrating, and synthesizing financial information (e.g., Bertomeu et al., 2025; Blankespoor et al., 2026; Chang et al., 2026; Cheng et al., 2025; Liu and Subasi, 2026; Xue et al., 2025). Prior technological advances such as EDGAR, XBRL, and social-media disclosure each lowered a specific component of investors' information-processing costs (e.g., Blankespoor, 2019; Blankespoor et al., 2014; Drake et al., 2015), but none automated what GenAI uniquely enables: cross-FDP synthesis of analyst consensus forecasts. This paper asks whether the post-ChatGPT market values an accuracy-weighted

[1] In a sample of 94,030 firm-quarters covered by all five FDPs over 2002–2020, Larocque et al. (2025) find that the FDPs report different earnings surprises (beyond rounding) for the same firm-quarter 58 percent of the time, and that at least one FDP signals a miss while another signals a beat in 9 percent of firm-quarters.

[2] Larocque et al. (2025) show that announcement-window returns align more with I/B/E/S and Capital IQ signals than with the other FDPs, consistent with acquisition cost being a binding factor.

multi-FDP consensus more than the conventional single-FDP consensus or the equal-weighted average across providers.[3]

The case for cross-FDP synthesis rests on the substantive disagreement among providers, who differ in their screening rules, treatment of unusual items, and analyst-coverage thresholds (Bochkay et al., 2022). Which FDP yields the most accurate consensus varies across firms and over time, with I/B/E/S ranking highest on average, consistent with its prevalence as the academic-research standard (Larocque et al., 2025). This variation implies meaningful gains from synthesizing across providers. Historically, such synthesis was impractical: investors had to recognize that the FDPs disagree, subscribe to each consensus separately, and reconcile differing methodologies.

Recognize, subscribe, reconcile—these three frictions correspond to the awareness, acquisition, and integration costs in the information-processing-cost framework of Blankespoor et al. (2020). GenAI attenuates all three, but the binding friction for accuracy-weighted multi-FDP aggregation has historically been the third: cross-source integration of heterogeneous consensus measures, which prior single-source technologies left intact.[4] The framework yields two empirical predictions: (i) the post-ChatGPT market should shift its implicit consensus toward an accuracy-weighted synthesis across multiple FDPs, and (ii) this shift should concentrate where GenAI synthesis is most accessible and most valuable.

Neither prediction is obvious ex ante. The strongest reason for doubt is theoretical: in the forecast-combination literature, combining heterogeneous, partially independent forecasts almost always dominates sub-selecting on accuracy (e.g., Ahlers and Lakonishok, 1983; Bates and Granger, 1969). Consistent with this logic, Michaely et al. (2024) show that constructing a consensus from high-quality analysts within a single provider does not improve investor outcomes; the frictions they identify carry over with even greater force to the cross-FDP setting. A separate concern is

[3]The accuracy-weighted construction is not a claim that the market literally applies this weighting rule. It serves as a tractable proxy for the broader hypothesis that the post-ChatGPT market draws on multiple FDPs and discriminates by source accuracy; if so, market reactions should align more closely with this benchmark than with the single-FDP or equal-weighted alternatives, even if the specific weights the market applies differ from mine.

[4]A single LLM query of the form "What is the consensus EPS forecast for [ticker] next quarter?" typically returns a synthesized response citing several free third-party distributors (e.g., Yahoo Finance, MarketBeat, TipRanks). Appendix B reports two worked examples (IBM and Microsoft) showing the cited distributors, their reported figures, and the disclosed upstream FDP where available.

that lower acquisition costs do not imply more selective integration: a richer information set need not improve forecast accuracy when the synthesis itself is unselective, and LLMs in particular can rely heavily on memorization rather than reasoning (e.g., Lopez-Lira et al., 2025). Whether GenAI nonetheless shifts the market toward a more selective synthesis is therefore an empirical question.

I test these predictions using a panel of S&P 1500 firm-quarters from 2021 to 2025 with consensus EPS forecasts for the same firm-quarter from all five FDPs.[5] After ChatGPT's release, the market reacts more strongly to forecasts from accurate FDPs and discounts those from less accurate ones. In a within-firm-quarter rank-based test that estimates provider-specific earnings response coefficients (ERCs) by accuracy rank (e.g., Collins and Kothari, 1989; Easton and Zmijewski, 1989), the ERC increment per one-rank accuracy gain roughly doubles after ChatGPT, from 0.31 to 0.62 (post-pre differential 0.32, $p < 0.01$). A complementary firm-level design partitions firms by whether I/B/E/S's accuracy rank rose or fell from pre- to post-ChatGPT. The result is asymmetric: the market response to I/B/E/S-based earnings surprises weakens significantly where I/B/E/S lost rank but not where it gained, consistent with pre-GenAI overweighting of the I/B/E/S consensus.

Given this evidence of selective use, I next test whether the market's aggregation rule itself has shifted toward an accuracy-weighted synthesis. I construct an accuracy-weighted standardized unexpected earnings (SUE) in which each FDP's SUE is weighted by its trailing forecast accuracy, and benchmark it against two heuristic alternatives: the single-FDP I/B/E/S consensus and the equal-weighted average across the five FDPs.[6] Post-ChatGPT, ERCs on the accuracy-weighted surprise rise sharply, while ERCs on the I/B/E/S and equal-weighted surprises remain flat or decline. The resulting accuracy premium widens significantly against both heuristic benchmarks. In economic terms, from pre- to post-ChatGPT, the three-day announcement CAR generated by a one-percentage-point standardized surprise rises 2.40 percentage points more under the

[5]The sample begins in 2021 for two reasons: (i) it yields a relatively balanced panel of pre- and post-ChatGPT firm-quarters around the November 30, 2022 cut, and (ii) Bloomberg's historical consensus data are accessible only for the past five years. Given the different methodologies used by FDPs, I compute each FDP's forecast accuracy from its own end-of-quarter consensus EPS and reported actual EPS.

[6]Specifically, the weight on FDP g is the softmax of the negative of its z-scored four-quarter rolling mean absolute forecast error, so providers with better recent track records receive larger weights. Figure 2 confirms that the accuracy-weighted SUE has the smallest mean absolute forecast error of the three benchmarks, with the difference statistically significant. All variable definitions are in Appendix A.

accuracy-weighted SUE than under the equal-weighted benchmark, and 2.67 percentage points more than under I/B/E/S.

The timing and the mechanism of this shift both point to GenAI. Re-estimated within four GPT-capability regimes, the premium against I/B/E/S climbs from negative in the pre-ChatGPT and text-only regimes to 1.52 under browsing and 4.91 under native search—a dose–response staircase that tracks the tools' retrieval capabilities. The gating condition is acquisition: a firm-level *LLM visibility* measure[7]—how many forecast sources a browsing-enabled model actually retrieves for the firm—is empirically distinct from classical prominence proxies, and the premium concentrates in high-visibility firms only once the model can browse. Together, these shifts imply that the consensus reflected in market prices has moved from heuristic alternatives toward an accuracy-weighted multi-FDP synthesis.

Two cross-sectional tests show that the accuracy premium concentrates in settings where GenAI-enabled synthesis is most valuable. Because the regime evidence localizes the premium to the browsing and native-search regimes, these tests use a post-scraping indicator—equal to one for announcements on or after the onset of GPT browsing (September 27, 2023), when web retrieval first made cross-provider synthesis feasible—in place of the post-ChatGPT indicator. First, the premium concentrates in firm-quarters with high cross-FDP disagreement, where accuracy-weighting offers the largest improvement over heuristic benchmarks. Second, it concentrates in firm-quarters with large nonrecurring items: because FDPs differ most in their treatment of unusual items (Bochkay et al., 2022), this is where multi-FDP synthesis most attenuates single-provider methodology noise.

I next examine whether the accuracy-weighted consensus is economically valuable to investors.[8] I construct a predicted earnings surprise—the price-scaled difference between the accuracy-weighted and each heuristic consensus—as an ex ante predictor of post-announcement returns. Regressing the four-day [-2, +1] announcement-window buy-and-hold abnormal return (BHAR) on the standardized predicted surprise yields a positive coefficient ($p < 0.01$ for both

[7]Operationally, LLM visibility is an EPS-extraction indicator times $\log(1 + \text{unique cited domains})$, averaged across ChatGPT and Perplexity responses to firm-specific EPS queries (from the DataForSEO LLM Responses API).

[8]This is not a foregone conclusion: Michaely et al. (2024) show that historically a consensus from high-quality analysts within a single provider did not generate significantly better investor outcomes.

heuristic benchmarks), indicating that the predicted surprise carries return-relevant information. Following the earnings-surprise trading-strategy literature (e.g., Battalio and Mendenhall, 2005; Michaely et al., 2024), a quintile long–short portfolio on the predicted surprise earns 0.63–0.73 percent per announcement against the equal-weighted benchmark and 1.35 percent against I/B/E/S (all $p < 0.01$).[9]

Three additional tests probe the GenAI channel. First, a forecast-based informativeness measure captures, firm by firm and in real time, how closely the consensus has anticipated the firm's year-ahead street numbers over the preceding twenty quarters. The post-ChatGPT rise in announcement-window reactions concentrates in firms with an informative consensus: the ERC–informativeness gradient roughly doubles, significant under both return adjustments; the increment is orthogonal to announcement composition; and it prices the common cross-provider consensus rather than any single provider. Second, GenAI might merely substitute for information investors previously acquired directly from free distributors. Google search-volume data, a standard proxy for investor information demand (e.g., Da et al., 2011; Drake et al., 2012), suggest otherwise: post-ChatGPT search share for free forecast-data distributors rises relative to paywalled FDPs, consistent with GenAI acting as a gateway to, rather than a substitute for, direct investor acquisition. Third, following Cheng et al. (2025) and Liu and Subasi (2026), I use major ChatGPT outages as an exogenous shock to GenAI availability. On announcement days following a major outage, the accuracy premium against I/B/E/S collapses while the equal-weighted contrast shows no significant change—an asymmetry consistent with GenAI as the integration technology underlying the premium: when multi-FDP synthesis is unavailable, investors revert specifically to the I/B/E/S consensus.[10]

This paper makes three contributions. First, I show that the market's reliance on FDPs is not static but technology-dependent, extending recent evidence on cross-FDP differences in analyst

[9] Specifically, the portfolio goes long firm-quarters with positive predicted surprise and short those with negative predicted surprise, holding each position over the four-day $[-2, +1]$ announcement window. The per-trade alpha is estimated by panel OLS with firm-clustered standard errors, and annualized by scaling by 252 trading days divided by the four-day holding window. Results are robust to alternative event windows, including $[-2, +2]$ and $[-1, +1]$.

[10] Outages are sourced from OpenAI's public status-page history; I require an incident of at least 240 minutes, yielding four incidents across five unique calendar days through 2024.

consensus information (Larocque et al., 2025), which documents substantial methodological and accuracy differences across providers. When information-processing costs fall, the market reweights its implicit consensus toward the more accurate provider's information, with measurable effects on earnings-announcement returns. For researchers, the implication is methodological: the FDP whose consensus best reflects market expectations shifts with information technology.

Second, I add a distinct mechanism to the literature on the capital-market consequences of generative AI (e.g., Bertomeu et al., 2025; Blankespoor et al., 2026; Chang et al., 2026; Cheng et al., 2025; Liu and Subasi, 2026; Lopez-Lira and Tang, 2023; Xue et al., 2025). Prior work documents that GenAI lowers investor processing costs and shifts prices around specific events such as outages or regulation. The mechanism here is synthesis: GenAI makes it cheap for the market to combine forecasts across multiple FDPs, and this combined consensus is now priced into earnings-announcement reactions.

Third, I sharpen the disclosure-processing-cost framework of Blankespoor et al. (2020) by isolating which cost component drives the market's accuracy-weighted use of consensus information. A long line of work shows how prior single-source technologies reduce specific cost components: EDGAR for filing acquisition (Drake et al., 2015; Loughran and McDonald, 2017), XBRL for structured-data processing (Blankespoor, 2019), social media for awareness latency (Blankespoor et al., 2014), and automated text processing for retail-investor processing (Blankespoor et al., 2018). The evidence here shows that the binding friction for accuracy-weighted multi-FDP aggregation was cross-source integration of heterogeneous consensus measures. The earlier technologies left that friction intact; GenAI relaxes it, with measurable consequences for earnings-announcement pricing.

The remainder of the paper is organized as follows. Section 2 reviews the related literature and presents the theoretical framework. Section 3 describes the empirical setting. Section 4 presents initial evidence that the market begins to rely selectively on more-accurate FDPs after ChatGPT. Section 5 tests whether the post-ChatGPT market prices the accuracy-weighted synthesis across FDPs, documents its tradeable returns, traces the premium across GPT-capability regimes, and identifies the cross-sectional drivers of the resulting accuracy premium. Section 6 shows that LLM visibility—a direct measure of GenAI forecast acquisition—is distinct from classical prominence proxies and gates the premium. Section 7 constructs a forecast-based informativeness measure and tests its post-ChatGPT pricing, examines investor information acquisition, and exploits major ChatGPT outages as an exogenous shock to GenAI availability. Section 8 concludes.

2 Related Literature and Theoretical Framework

2.1 Forecast Data Providers and the Construction of Analyst Consensus

Sell-side analyst earnings forecasts have long served as the benchmark proxy for market expectations in research and practice (Bartov et al., 2002; Beyer et al., 2010; Bradshaw et al., 2017; Kothari, 2001), in part because analysts integrate firm-specific guidance, industry knowledge, and macroeconomic conditions in ways that simple time-series extrapolation does not (Bradshaw et al., 2012; Schipper, 1991). Individual forecast accuracy varies systematically with an analyst's experience, employer resources, and portfolio complexity (Clement, 1999; Mikhail et al., 1997), and the market discriminates on these differences at the analyst level: revisions from more accurate or higher-status analysts trigger larger price reactions (Clement and Tse, 2003; Park and Stice, 2000; Stickel, 1992). At the consensus level, however, the analogous accuracy-weighted sub-selection has not been shown to dominate the pooled benchmark, even within a single provider (Michaely et al., 2024). This asymmetry — discrimination at the analyst level but not at the consensus level — organizes the discussion that follows.

The consensus number itself is not a primitive of nature; third-party forecast data providers (FDPs) construct it by pulling each underlying analyst forecast through proprietary screening, exclusion, and adjustment rules.[11] Bochkay et al. (2022) document the first-order capital-market consequences of these construction choices in the setting of Thomson Reuters' (now Refinitiv's) 2009 methodology change for I/B/E/S.[12] They show that an FDP's treatment of unexpected charges

[11]FDPs differ along several dimensions, including (i) the set of broker estimates each provider includes (some providers exclude forecasts deemed stale or non-comparable; others apply tighter recency cutoffs), (ii) the treatment of nonrecurring items (some FDPs report GAAP-based earnings, others a "street" earnings adjusted for one-time charges), and (iii) rounding and staleness conventions.

[12]In September 2009, Thomson Reuters discontinued its prior practice of relying on analysts' ex post consensus

and gains feeds directly into both the time-series predictive properties of street earnings and price discovery around announcements. The FDP layer, in other words, is an active participant in information production, not a passive conduit.

The same construction discretion generates pervasive disagreement across providers. Comparing the five FDPs in a large sample, Larocque et al. (2025) find that providers report different earnings surprises (beyond rounding) for the same firm-quarter 58 percent of the time, and that in 9 percent of firm-quarters at least one FDP signals a beat while at least one other signals a miss. These disagreements measurably affect price responsiveness, abnormal volatility, and liquidity around earnings announcements, and the accuracy ranking across FDPs is itself unstable across firms and over time.

Disagreement of this magnitude leaves substantial room, in principle, for investors to synthesize across providers rather than anchor on a single one.[13] If the market could anticipate each FDP's accuracy ex ante, accuracy-weighted aggregation across providers could outperform any single-provider benchmark. The historical evidence points the other way: Larocque et al. (2025) document that I/B/E/S ranks highest on their composite measures of FDP quality and that market reactions align most closely with the I/B/E/S surprise, consistent with I/B/E/S serving as the historical anchor.

The closest evidence on whether the market performs this accuracy-weighted aggregation comes from Michaely et al. (2024), who examine the within-provider analog. Constructing a consensus from high-quality analysts inside I/B/E/S, they find no meaningful improvement in investor outcomes over the full-analyst consensus, and they attribute the null to three frictions: error cancellation in the broader pool, imperfect persistence in analyst quality, and the cognitive cost of identifying high-quality forecasters.[14] The same logic that makes within-FDP sub-selection

to determine the I/B/E/S street-earnings actual when an earnings press release reported unexpected charges or gains, replacing it with a uniform rule that immediately excludes such items from GAAP earnings. Because the change was rolled out gradually, Bochkay et al. (2022) use the staggered variation in a difference-in-differences design.

[13]Classical forecast-combination theory (Bates and Granger, 1969; Clemen, 1989; Timmermann, 2006) establishes that the accuracy gain from combining multiple forecasts is increasing in the heterogeneity of their forecast errors and decreasing in their pairwise correlation; whether the gain is realizable in practice depends on the cost of identifying, acquiring, and integrating the underlying sources, a point developed formally below.

[14]This pattern echoes a broader theme in the analyst and disclosure literature, in which forecast dispersion is

costly should make cross-FDP sub-selection costlier still. The open question is whether a large enough reduction in the cost of identifying, acquiring, and integrating accurate sources can unlock the consensus-level discrimination that the historical environment failed to produce. The next subsection turns to the technologies that determine this cost.

2.2 Technological Advancements and Information Processing Costs in Capital Markets

Whether investors use available information depends on what it costs them to do so. Blankespoor et al. (2020) provide the canonical economic frame, decomposing the cost of using any piece of information into three components: *awareness* of its existence, *acquisition* of the information itself, and *integration* of that information into the investor's decision. Each component is sensitive to technology, and an exogenous reduction in any one of them, holding the underlying information fixed, can measurably change investor behavior, asset prices, and firms' own reporting choices.

An extensive empirical literature documents these effects component by component, with much of the identifying variation coming from single-source technological shocks introduced by regulators or by changes in market infrastructure. On the acquisition side, Drake et al. (2015) show that EDGAR search activity tracks firm-level information demand and that surges in retrieval precede price adjustments, while Loughran and McDonald (2017) use granular access logs to show that few investors look at most filings, but the filings that are accessed move prices substantively. On the integration side, the SEC's XBRL mandate forced firms to tag financial-statement items in machine-readable form, and Blankespoor (2019) documents that the resulting investor processing advantage shifted firms' own disclosure choices; on the textual margin, Tetlock (2007) and Loughran and McDonald (2011) establish that machine-readable sentiment and finance-specific dictionaries extract economically meaningful information from narrative disclosure. On the awareness side, Blankespoor et al. (2014) find that firm-initiated social-media dissemination broadens the investor base that responds to earnings news, and Blankespoor et al. (2018) demonstrate that robo-journalism (automated article generation) lowers the processing cost retail investors face when consuming

itself a priced characteristic (Diether et al., 2002) and disclosure clarity measurably affects investor behavior at the individual-disclosure level (Lawrence, 2013).

financial information.

A parallel literature on demand-side attention validates these supply-side mechanisms. Da et al. (2011) establish Google search as a measure of retail attention that predicts short-horizon returns, and Drake et al. (2012) document that earnings-window search proxies information demand more directly than trading-based measures. Attention is itself a binding constraint, with documented inattention along multiple margins: complicated firms (Cohen and Lou, 2012), competing simultaneous announcements (Driskill et al., 2020; Hirshleifer et al., 2009), Friday timing (DellaVigna and Pollet, 2009), strategic scheduling (deHaan et al., 2015), and high-fear states (Da et al., 2015).

Closest to the FDP-aggregation question, the disseminator-side literature shows that broader access to firm disclosures shifts price formation: opening conference calls to all investors alters retail-trading and volatility patterns (Bushee et al., 2003) and reduces post-announcement underreaction (Kimbrough, 2005), and business-press coverage narrows spreads and deepens markets around earnings releases (Bushee et al., 2010).

A common theme runs through this work: each technological advance lowered the cost of one component within a single information source. None enabled the low-cost cross-source integration of heterogeneous numerical outputs from multiple distinct producers — combining, say, I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks consensus estimates in real time. That friction survived every one of these technologies.

2.3 Generative AI in Capital Markets

GenAI is the first technology positioned to relax this cross-source friction. It extends a machine-learning literature in capital markets whose central finding is well established: nonlinear, high-dimensional methods (e.g., neural networks and tree-based models) extract incremental predictive signal from financial data that traditional linear regression-based asset-pricing models miss (Goldstein et al., 2021; Gu et al., 2020). What distinguishes the current generation of LLMs is the joint possession of three features prior technologies lacked: they process unstructured information from heterogeneous sources, they synthesize that information through natural-language queries without specialized retrieval infrastructure or programming expertise, and they do so at low marginal cost per query. Cross-source synthesis that once demanded time, specialized expertise, and dedicated computational infrastructure available only to institutional teams collapses to the marginal cost of a single natural-language query — placing the capability within reach of retail investors for the first time.

Early evidence on the capital-market consequences of this shift is consistent with a meaningful price-relevant effect. Lopez-Lira and Tang (2023) show that ChatGPT-derived sentiment signals predict cross-sectional returns; Cheng et al. (2025) use ChatGPT service outages as exogenous shocks and find that they degrade investor responses to information; Bertomeu et al. (2025) exploit Italy's 2023 ChatGPT ban for country-level identification; and Chang et al. (2026) show that retail traders' alignment with AI-derived earnings-call sentiment rose sharply after ChatGPT's release, narrowing the information-processing gap between less-sophisticated retail traders and more-sophisticated short sellers.

Most directly relevant, a separate strand examines AI's effects on the analyst research function. Kimbrough et al. (2026) find that analysts at brokerages with greater AI integration produce more accurate earnings forecasts and that their revisions are more informative to capital markets. Bradshaw et al. (2026) document that AI-generated content reached 13.5 percent of Seeking Alpha after ChatGPT's release, with AI-using authors covering more firms but producing less-informative pieces per article. Liu and Subasi (2026) find an analogous democratization on the crowdsourced side, with GenAI closing the accuracy gap between non-professional and professional forecasters. Together, these papers establish that AI is reshaping the production of analyst research. How GenAI reshapes the consumption of that research — in particular, how investors aggregate across the heterogeneous consensus numbers that competing FDPs construct — is the gap this paper fills.

2.4 Theoretical Framework and Predictions

Two opposing forces shape the framework. The first comes from the forecast-combination literature: combining heterogeneous, partially independent forecasts almost always dominates relying on any single source (Bates and Granger, 1969; Clemen, 1989; Granger and Ramanathan, 1984; Stock and Watson, 2004; Timmermann, 2006). The intuition is the wisdom-of-crowds principle popularized by Surowiecki (2004): averaging diversifies away idiosyncratic forecaster errors, so sub-selecting on individual accuracy can lose more from forgone diversification than it gains from individual precision. Applied to the analyst-consensus setting, this force rationalizes the historical equilibrium in which the market anchors on a single pooled benchmark (typically the I/B/E/S consensus) and treats cross-FDP differences as noise. It is also consistent with the within-FDP evidence in Ahlers and Lakonishok (1983) and Michaely et al. (2024), in which more sophisticated within-source sub-selection fails to dominate the broader pool consistently.

The opposing force is the documented cross-sectional and time-series variation in FDP accuracy (Larocque et al., 2025), together with the methodology-driven information content in Bochkay et al. (2022). When cross-provider dispersion is large, contains a non-trivial accuracy-relevant component, and is correlated with exploitable provider-specific characteristics, a sufficiently low-cost synthesis across providers can dominate the pooled benchmark. A parallel point holds in the model-based earnings-forecast literature, where cross-sectional models deliver more accurate forecasts than any single analyst for specific sub-samples (Hou et al., 2012). The cost of identifying, acquiring, and integrating accurate sources resolves the tension between the two forces: when this cost is high, the wisdom-of-crowds force dominates and the market anchors on a pooled single-FDP consensus; when it is low, accuracy-weighted synthesis becomes feasible and can dominate the pooled benchmark.

In the Blankespoor et al. (2020) framework, cross-FDP synthesis binds on all three cost components at once. Even when consensus data are freely redistributed through retail-facing platforms, an investor seeking the cross-FDP picture must know which platform carries which FDP's consensus (awareness), visit each individually to retrieve the number (acquisition), and reconcile inconsistent presentations of methodology-driven differences (integration). The time cost per firm-quarter scales with the number of firms followed, and performing the synthesis at scale further requires programming and data-engineering expertise. This combination has historically confined the capability to specialists, anchoring the market on a single source.

GenAI removes both barriers — the per-firm time cost and the expertise requirement — at once: a single LLM query identifies candidate FDP sources, retrieves their numerical content, and integrates them into a synthesized response that surfaces the implied consensus and the cited distributors, all without specialized infrastructure or programming expertise. The effective cost of cross-FDP synthesis collapses to that of issuing one query.[15] Figure 1 summarizes the framework, which yields three empirical predictions:

Prediction 1. After ChatGPT's release, the consensus reflected in earnings-announcement prices shifts away from heuristic aggregation (single-FDP anchoring or equal-weighting across FDPs) toward a more sophisticated synthesis that draws on multiple FDPs and places greater weight on the more accurate ones.

Prediction 2. The post-ChatGPT shift is stronger when and where GenAI-mediated synthesis is more accessible (LLMs effectively retrieve firm-level analyst information) and more valuable (cross-FDP disagreement and its methodology-driven component are large).

Before the final prediction, one scope condition deserves emphasis. The framework operates entirely on the demand side: falling awareness, acquisition, and integration costs change how investors use the FDPs' numbers, while the numbers themselves — each provider's construction rules and the resulting street earnings — are held fixed. This yields a falsifiable corollary that separates the framework from a supply-side alternative in the spirit of the 2009 I/B/E/S methodology change studied by Bochkay et al. (2022): if the post-ChatGPT patterns instead reflected a change in how the street numbers are constructed, the predictive quality of street earnings for future performance should shift at the ChatGPT boundary, and differentially so in announcements containing large unexpected items, where construction discretion binds most.

Prediction 3. The post-ChatGPT shift extends beyond trailing accuracy: announcement reactions load increasingly on providers whose consensus is more informative about the firm's future performance.

[15] Appendix B reports two worked examples (IBM and Microsoft) showing the actual LLM responses, the cited distributors, and the upstream FDPs where disclosed.

3 Institutional Background, Data, Sample, and Key Measures

This section describes the institutional setting that generates cross-FDP disagreement, the data and sample, the dependent variables, and the construction of the accuracy-weighted SUE.

3.1 Institutional Background

Five major data providers—I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks—construct and sell analyst-consensus EPS information to sell-side brokerages, buy-side investors, financial media outlets, and academic researchers. I/B/E/S is credited with creating the consensus-earnings-estimate industry in the early 1970s, and Zacks claims to have first distributed quarterly consensus forecasts in the early 1980s; Bloomberg, Capital IQ, and FactSet entered the segment from adjacent product lines (terminal data, investment-banking software, and printed "fact sets" delivered to Wall Street clients, respectively).[16]

Two sources of variation drive cross-FDP differences in consensus numbers. First, the underlying analyst set differs: each FDP includes a different subset of contributing sell-side analysts, applies its own recency cutoffs (e.g., Bloomberg drops estimates more than 30 days stale), and runs proprietary statistical checks to filter outliers. Second, the methodology for constructing "street" earnings varies: each FDP exercises discretion over how to treat unexpected charges, gains, and other adjustments to reported earnings (Bochkay et al., 2022). Zacks, for example, treats stock-based compensation as a recurring expense in street earnings, whereas I/B/E/S has historically excluded it. More broadly, 46.7 percent of firm-quarters in the 2002–2020 sample of Larocque et al. (2025) contain a large unexpected item that FDPs treat inconsistently. Together, these coverage and methodology differences generate the cross-FDP disagreement documented in Section 2; Larocque et al. (2025) provide a comprehensive FDP-by-FDP review.

No single FDP, however, strictly dominates the others at the firm-quarter level. Although I/B/E/S ranks highest on average across composite measures of FDP quality (Larocque et al., 2025) and has historically served as the de facto anchor for both research and trading, it is the most accurate

[16] Each FDP today operates within a broader financial-data platform: I/B/E/S under LSEG Data & Analytics (formerly Refinitiv); Capital IQ under S&P Global; FactSet and Bloomberg as standalone parents. Bloomberg's earnings estimates are bundled with its core terminal offering, while the other four sell estimates as standalone or modular products. See Larocque et al. (2025) for a detailed institutional review.

provider in only 16.9 percent of firm-quarters in my S&P 1500 sample and the least accurate in 12.1 percent; the remaining 71 percent of firm-quarters split across the other four FDPs. This variation across firms and over time is what makes the cross-FDP synthesis question economically meaningful: a static anchoring choice cannot exploit it; an adaptive multi-FDP rule can.

3.2 Data and Sample

Exploiting this firm-quarter variation requires parallel consensus histories from all five FDPs, which no single database supplies. I obtain quarterly consensus EPS forecasts and reported actuals from each provider through WRDS and direct terminal access. I/B/E/S and Zacks are available through WRDS (the I/B/E/S Summary History file and the Zacks Estimates database, respectively). The remaining three FDPs (FactSet, Bloomberg, and S&P Capital IQ) require proprietary access through their respective terminals or institutional portals, and I download their quarterly consensus histories firm by firm.[17]

The sample begins with all S&P 1500 firm-quarters with an earnings announcement over 2021–2025, totaling 29,395 firm-quarters across 1,502 firms.[18] I then require consistent coverage across all five FDPs, dropping 4,368 firm-quarters missing the consensus from at least one provider. After dropping firm-quarters with missing market-outcome variables or controls (1,179 observations), missing constructed return or predicted-surprise variables (10), and singleton firms (5), the final sample contains 23,833 firm-quarters across 1,334 firms; Table 1 reports the full selection procedure. Table 2 reports descriptive statistics for the main variables.[19]

Firm controls and market-outcome variables are drawn from WRDS-hosted databases: accounting fundamentals from Compustat, returns and trading data from CRSP, institutional ownership from the Thomson Reuters 13F Holdings database, and analyst-coverage characteristics from I/B/E/S. Variable definitions appear in Appendix A.

[17] Larocque et al. (2025) follow a similar collection procedure for their 2002–2020 study, downloading or hand-collecting data from each FDP's proprietary system. Bloomberg consensus data, in particular, must be retrieved via the Bloomberg Query Language (BQL) function on a ticker-by-ticker basis, and Bloomberg's historical consensus is accessible only for the most recent five years, which constrains the sample start date.

[18] The November 30, 2022 ChatGPT release falls roughly in the middle of this window, yielding a relatively balanced panel of pre- and post-ChatGPT firm-quarters. The five-year window is also the maximum permitted by Bloomberg's historical-retrieval constraint described above.

[19] All continuous variables in the analyses that follow are winsorized at the $(1,99)$ level.

3.3 Dependent Variables

Because the market's reliance on any given consensus is not directly observable, I infer it from the price response at the earnings announcement: a more-trusted consensus should generate a stronger reaction to its implied surprise. The main outcome variables are therefore short-window abnormal returns around earnings announcements: a three-day cumulative abnormal return $CAR_{[-1,+1]}$ for the immediate-reaction tests and a four-day buy-and-hold abnormal return $BHAR_{[-2,+1]}$ for the trading-strategy test.[20] For each window I use two risk adjustments: a Fama-French-Carhart four-factor (FFC) abnormal return and a market-adjusted return.[21]

3.4 Constructing the Accuracy-Weighted SUE

The paper's central construct is an accuracy-weighted aggregation of the five FDP consensus signals. Let $g \in \{\text{FactSet, Bloomberg, Zacks, CapIQ, IBES}\}$ index providers and let $SUE^{g}_{i,q}$ denote provider g's price-scaled standardized unexpected earnings for firm i in quarter q. The accuracy-weighted composite is

$$SUE^{Acc}_{i,q} = \sum_{g} w^{g}_{i,q} SUE^{g}_{i,q}, \tag{1}$$

with weights $w^{g}_{i,q}$ derived from each provider's recent forecast accuracy. The equal-weighted benchmark sets $w^{g}_{i,q} = 1/5$, yielding $SUE^{EW5}_{i,q}$. Entering both composites in the same regression tests whether the market loads more on the accuracy-tilted aggregate than on the naive equal-weighted aggregate.

Weight construction. I first compute each provider's price-scaled absolute forecast error, $err^{g}_{i,q} = |actual^{g}_{i,q}\text{-}consensus^{g}_{i,q}|/p_{i,q}$.[22] I then take a backward-looking four-quarter rolling mean

[20] The four-day BHAR window adds a pre-announcement day to capture any leakage and a post-announcement day to capture overnight drift in announcements made after the market close.

[21] FFC factor loadings are estimated on the [-260, -11] pre-announcement window; the market-adjusted return subtracts the contemporaneous CRSP value-weighted index from each daily return. All returns are expressed in percent. All inferences are robust to both adjustments: Tables 4, 5, and 11 report results under both adjustments side by side (Table 3 Panel B likewise), while Tables 3 (Panel A), 6, 7, 8, 9, 10 (Panel B), and 12 (Panel A) report only the market-adjusted return for brevity (FFC counterparts deliver identical inferences and are available on request).

[22] The reported actual EPS is itself FDP-specific: each provider computes its own "street" earnings actual under its own screening and adjustment rules. Pairing each FDP's consensus with its own actual ensures that the resulting forecast error reflects only ex ante predictive accuracy, not methodology mismatch between the consensus and the realized value. This is the approach taken by Larocque et al. (2025) and is consistent with how each FDP reports its own pre-announcement consensus relative to its own post-announcement actual.

to obtain a recent mean absolute forecast error, $MAFE^{g}_{i,q} = \frac{1}{4}\sum_{k=1}^{4} err^{g}_{i,q\text{-}k}$, and flip the sign so larger values denote greater accuracy: $acc^{g}_{i,q} = \text{-}MAFE^{g}_{i,q}$. The rolling window excludes quarter q itself, so the weights use only information available before the quarter-q announcement and cannot induce look-ahead bias. I then standardize across the five providers within firm-quarter, $acc\_z^{g}_{i,q} = (acc^{g}_{i,q}\text{-}\overline{acc}_{i,q})/\mathrm{sd}(acc_{i,q})$, and convert the z-scores into convex weights via the softmax (multinomial-logit) transform,

$$w^{g}_{i,q} = \frac{\exp(acc\_z^{g}_{i,q})}{\sum_h \exp(acc\_z^{h}_{i,q})}. \tag{2}$$

The resulting weights are strictly positive, sum to one within each firm-quarter, preserve the accuracy ordering, and reduce to the uniform 0.2 vector when all five major FDPs are equally accurate. Appendix C works through this construction for a single hypothetical firm-quarter, tracing how raw accuracy histories pass through the four layers to produce both the weights $w^{g}_{i,q}$ and the composite $SUE^{Acc}_{i,q}$.

Design choices. Three choices in the construction merit comment. First, the within-firm-quarter standardization (rather than cross-sectional) is essential because earnings noise differs across firms and across periods: what matters economically is relative accuracy for this firm this quarter, not absolute accuracy compared with other firms or other quarters. Second, the softmax transform, equivalent to the exponential weighting scheme in the forecast-combination literature (e.g., Bates and Granger, 1969; Timmermann, 2006), occupies a natural middle ground between equal weighting (which ignores accuracy) and winner-take-all selection (which would amplify the noise in the four-quarter MAFE estimate). It is also the maximum-entropy weight distribution consistent with the accuracy ordering—the least arbitrary mapping from relative accuracy to weights. Third, the backward-looking four-quarter window balances recency against statistical noise: shorter windows overweight transient accuracy shocks; longer windows blur structural changes in provider quality.

Empirical weight distribution. In practice, the weights tilt meaningfully—but not extremely—away from equal weighting. The average within-firm-quarter minimum weight is 0.028 and the corresponding maximum is 0.414, against a uniform benchmark of 0.200: a typical firm-quarter downweights its worst provider to roughly three percent of the composite and concentrates roughly forty-one percent on its best. SUE^{Acc} and SUE^{EW5} are highly correlated ($\rho \approx 0.82$) but not collinear.

4 GenAI and Selective Reliance Across FDPs: Initial Evidence

This section presents the first evidence on Prediction 1. If GenAI lets the market discriminate across FDPs by accuracy, the within-firm-quarter price reaction should rise with provider accuracy, and more steeply post-ChatGPT. Two tests implement this logic: a within-firm-quarter test of whether announcement returns track provider accuracy ranks, and a firm-level test of whether reliance on the historical anchor tracks its own accuracy.

Accuracy is the appropriate dimension along which to rank providers. A validation exercise in the Online Appendix (Table OA1) formalizes why: the accuracy-weighted consensus is strictly closer to realized earnings than either heuristic in both regimes (paired tests, $t > 10$), while providers' consensus slopes for future fundamentals are statistically indistinguishable, suggesting that the accuracy tilt sacrifices little fundamental informativeness. At the firm level, the accuracy and informativeness of the consensus — constructed identically at the current-quarter and year-ahead horizons — are themselves highly correlated (Pearson 0.72; Section 7), so operating on accuracy is econometric convenience, not substance.

Design. The first test asks whether the ERC rises with provider accuracy rank. Ranking the five FDPs each firm-quarter by ex ante trailing MAFE (rank 1 = least, 5 = most accurate) and pooling in long form, I estimate

$$CAR^{Mkt}_{i,q,[-1,+1]} \; = \; \alpha + \beta \, SUE^{g}_{i,q} + \delta \, Rank^{g}_{i,q} + \gamma \left(SUE^{g}_{i,q} \times Rank^{g}_{i,q} \right) + \varepsilon_{i,q,g}, \tag{3}$$

where γ is the ERC increment per one-rank accuracy gain.[23] Because rank is assigned within firm-quarter, identification comes from differential reactions to the same announcement, differencing out firm- and time-level confounders, including common provider-level methodology drift.

[23] The earnings response coefficient (ERC) is the slope of announcement returns on the earnings surprise — the abnormal return the market delivers per unit of surprise. A larger ERC means the market treats the surprise as more informative; γ measures how much that slope rises when the surprise comes from a one-rank-more-accurate provider.

The second, complementary test asks whether reliance on the historical anchor tracks its own accuracy at the firm level. Classifying firms by the change in I/B/E/S's mean rank between the pre- and post-ChatGPT windows ($IBES\,Up_i$/$IBES\,Down_i$ for moves of ≥ 2 positions), I estimate the triple-interaction

$$\begin{aligned} CAR_{i,q} = \alpha + \beta_1\, SUE^{IBES}_{i,q} + \beta_2\, SUE^{IBES}_{i,q} \times Post_q \times IBES\,Up_i \\ + \beta_3\, SUE^{IBES}_{i,q} \times Post_q \times IBES\,Down_i + \text{controls} + \text{FE} + \varepsilon_{i,q}, \end{aligned} \tag{4}$$

where the regression includes the lower-order terms, suppressed here. The triple interactions carry the test: selective reliance predicts $\beta_3 < 0$ and $\beta_2 \approx 0$.

Results. Table 3 reports both tests. Post-ChatGPT, the price reaction to a given surprise depends increasingly on which provider produced it. In Panel A, the per-rank ERC increment γ roughly doubles post-ChatGPT, from 0.306 to 0.621; the post-pre differential is 0.316 (both $p < 0.01$).[24] Panel B shows the adjustment is accuracy-conditioned: the response to I/B/E/S weakens where it *lost* rank — the post-ChatGPT triple interaction for I/B/E/S-down firms ranges from -2.3 to -2.8 across specifications ($p < 0.05$) — but not where it gained or held.[25] The reweighting is thus asymmetric, operating where the historical anchor's accuracy advantage has eroded.

5 GenAI and the Accuracy-Weighted Synthesis Across FDPs

Section 4 establishes that the post-ChatGPT market discriminates across FDPs by accuracy. This section asks a more demanding question: does the market price the accuracy-weighted synthesis of all five FDPs, as Prediction 1 asserts, beyond what the heuristic alternatives capture? The analysis proceeds in three steps. I first validate the construct itself: accuracy ranks persist, the composite is the most accurate of the three, and its ex ante difference from the heuristics earns an ex post trading return. I then horse-race the accuracy-weighted SUE against the two heuristic benchmarks and show that the accuracy premium is statistically and economically significant after ChatGPT. Finally, testing Prediction 2, I trace the premium across the four GPT-capability regimes and then show that

[24] Provider-by-rank cell-level versions of Panel A, estimated separately in absolute-value and signed form, yield the same inferences (untabulated).

[25] A generalized post-ChatGPT shift in risk premia or attention would move all providers' ERCs together; the within-firm-quarter, accuracy-conditioned patterns here do not.

it concentrates where the framework predicts: in firm-quarters with greater cross-FDP disagreement and greater methodological divergence across providers.

5.1 Validating the Accuracy-Weighted Construct

Premise and benefits. The design rests on one premise: weighting by *trailing* accuracy is informative only if accuracy *persists*. Given persistence, two benefits should follow: SUE^{Acc} should be more accurate than the heuristics, and it should carry trading value — its difference from a heuristic, known ex ante, earning an ex post return. For the latter I form the *predicted surprise*

$$Pred.Surp._{i,q}^{X} = \left(Consensus_{i,q}^{Acc}\text{-}Consensus_{i,q}^{X}\right) / \, p_{i,q}, \qquad X \in \{EW5, IBES\}, \tag{5}$$

and test whether it forecasts the four-day announcement BHAR and a quintile long–short return.

Results. All three conditions hold. Table 4, Panel A, shows that provider ranks are highly persistent (Spearman $\rho = 0.71$); the least- and most-accurate providers retain those ranks 76 and 65 percent of the time, roughly three times the uniform baseline. In Figure 2, SUE^{Acc} has the smallest mean absolute error of the three composites in both regimes. In Panel B, the standardized predicted surprise forecasts the four-day BHAR in all four columns ($p < 0.01$), and a long–short portfolio on its quintiles earns 0.63–0.73 percent per announcement against equal-weighting and 1.35 percent against I/B/E/S ($p < 0.01$).

5.2 The Accuracy Premium: Horse-Race Tests

Design. With the construct validated, the horse race asks whether the market prices it. I race the accuracy-weighted composite SUE^{Acc} against two natural benchmarks: the equal-weighted composite SUE^{EW5} treats each FDP as equally informative, and the I/B/E/S-only consensus SUE^{IBES} represents the historical anchor. The pooled specification is

$$\begin{aligned} CAR_{i,q} \;=\; & \alpha + \beta_1 SUE_{i,q}^{Acc} + \beta_2 SUE_{i,q}^{X} + \beta_3 SUE_{i,q}^{Acc} \times Post_q \\ & + \beta_4 SUE_{i,q}^{X} \times Post_q + \text{controls} + \text{FE} + \varepsilon_{i,q}, \end{aligned} \tag{6}$$

for $X \in \{EW5, IBES\}$. The pre-period accuracy premium is the contrast β_1-β_2; its post-pre change is β_3-β_4, the central test of Prediction 1. A positive estimate of β_3-β_4 means that after ChatGPT, the market loads more heavily on the accuracy-weighted consensus than on the heuristic benchmark.

This contrast is the estimand for every test in this section: the regime and cross-sectional analyses that follow re-estimate it within subperiods and subsamples.[26] I report both the three-day market-adjusted abnormal return and the Fama-French-Carhart four-factor abnormal return.

Results. Table 5 reports the estimates. In Panel A (equal-weighted benchmark), the pre-ChatGPT accuracy premium is statistically indistinguishable from zero (-0.972), consistent with the equal-weighted composite already capturing the pre-ChatGPT information set. After ChatGPT, the premium widens sharply: the accuracy-weighted slope rises from 1.376 to 3.420 while the equal-weighted slope falls from 2.324 to 1.767, yielding a post-pre change in the accuracy premium of 2.401 ($p < 0.01$) under the market-adjusted CAR and 2.584 ($p < 0.01$) under the Fama-French-Carhart four-factor abnormal return. Panel B (I/B/E/S benchmark) tells a complementary story: the pre-ChatGPT accuracy premium is itself negative (-0.623, statistically insignificant), reflecting the over-reliance on I/B/E/S documented in Section 4. Post-ChatGPT, the premium climbs to 2.049 ($p < 0.01$), implying a post-pre change of 2.672 ($p < 0.01$) under the market-adjusted CAR and 2.688 ($p < 0.01$) under the Fama-French-Carhart adjustment.

Economically, the magnitudes are large. After ChatGPT, a one-percentage-point standardized surprise generates a three-day announcement CAR roughly 2.4 percentage points larger under the accuracy-weighted consensus than under the equal-weighted benchmark, and 2.7 percentage points larger than under I/B/E/S — contrasts comparable to the entire pre-ChatGPT ERC on the I/B/E/S consensus itself, 2.2 in Panel B. The post-ChatGPT market loads on the accuracy-weighted synthesis with a weight that rivals the historical reliance on the I/B/E/S consensus. The shift tracks real informational content: Figure 2 shows that the accuracy-weighted consensus also delivers the smallest mean absolute forecast error of the three composites, and the difference is statistically significant.[27]

[26] Although SUE^{Acc} and the benchmark SUE^{X} are weighted averages of the same five per-FDP surprises and are therefore correlated by construction ($\rho \approx 0.82$ for $X = EW5$), they are not collinear: with $R_j^2 \approx 0.67$, the variance inflation factor $VIF_j = 1/(1\text{-}R_j^2) \approx 3.0$, so each individual slope's standard error is only $\sqrt{3.0} \approx 1.7$ times its single-regressor counterpart. The identifying quantity, however, is the contrast β_1-β_2, whose variance $\mathrm{Var}(\hat{\beta}_1) + \mathrm{Var}(\hat{\beta}_2)\text{-}2\,\mathrm{Cov}(\hat{\beta}_1, \hat{\beta}_2)$ is materially less inflated, because positively correlated regressors yield positively correlated slope estimates. The horse race is therefore the conservative specification: a reparameterization using the accuracy tilt SUE^{Acc}-SUE^{X} directly would recover the same contrast with cleaner individual slopes.

[27] As emphasized in the introduction, the accuracy-weighted construction is not a claim that the market literally

5.3 The Accuracy Premium Across GPT-Capability Regimes

Motivation. The pooled horse race dates the premium only coarsely, to the post-ChatGPT boundary. Prediction 2 implies a finer dose–response pattern: if the premium reflects GenAI-mediated synthesis, it should strengthen as the tools' retrieval capabilities expand, not merely jump once at the ChatGPT release. I therefore re-estimate the horse race of Equation (6) within four GPT-capability regimes: pre-ChatGPT (R0), text-only (R1, from November 30, 2022), browsing (R2, from September 27, 2023), and native search (R3, from October 31, 2024).

Results. Table 6 and Figure 3 sharpen the timing. Against I/B/E/S, the accuracy premium rises from -0.75 (pre-ChatGPT) and -1.23 (text-only) to 1.52 (browsing) and 4.91 (native search, $p < 0.01$) — the staircase that Prediction 2 predicts. Against the equal-weighted benchmark, the path is -0.95, 0.18, 2.43 ($p < 0.05$), and 2.55 ($p < 0.05$). The premium does not arrive with the chatbot; it arrives, step by step, with the ability to retrieve and integrate the underlying forecasts.

5.4 Cross-Sectional Tests of the Accuracy Premium

Design. Prediction 2 also specifies where the shift should appear: where GenAI-mediated synthesis is more valuable. Because the regime evidence localizes the accuracy premium to the browsing and native-search regimes, the cross-sectional tests use the post-scraping indicator $Post_q^{\mathrm{Scr}}$ — equal to one for announcements on or after the onset of GPT browsing (September 27, 2023), when web retrieval first made cross-provider synthesis feasible — in place of the post-ChatGPT $Post_q$. I re-estimate the horse-race specification in Equation (6) within the low and high subsamples of two partitions, each probing a distinct channel. (i) *Concurrent cross-FDP SUE dispersion* is the interquartile range of the five FDPs' SUE for the same firm-quarter; this is the setting in which provider weighting has the most discriminating work to do. (ii) *Special-items magnitude* is the absolute size of special items relative to the absolute size of income before extraordinary items; FDPs differ in their treatment of unexpected charges and gains, so cross-FDP disagreement on what

applies softmax weights. It is a tractable proxy for the broader hypothesis that the post-ChatGPT market weights providers by accuracy. Formally, the accuracy tilt $SUE_{i,q}^{Acc}$-$SUE_{i,q}^{EW5} = \sum_g (w_{i,q}^g\text{-}1/5)\, SUE_{i,q}^g$ is the inner product of each provider's deviation from the uniform weight and its standardized surprise; this is the object the horse-race contrast identifies.

to include in street earnings should be largest where such items are economically material.[28]

Results. Table 7 reports the partition-by-partition estimates; the browsing-era premium change is significant only in the predicted half of each partition. In Panel A (cross-FDP dispersion), the post-pre change in the accuracy premium is 2.211 ($p < 0.05$) against the equal-weighted benchmark and 2.427 ($p < 0.01$) against I/B/E/S in the high-disagreement subsample, versus insignificant estimates (2.037 and -0.691) in the low-disagreement subsample. In Panel B (special items), the concentration is the same: 2.518 ($p < 0.05$) and 3.638 ($p < 0.01$) under high special items, versus -0.552 and -0.219 under low. Two readings of the mechanism emerge. The concentration in high-dispersion firm-quarters supports the discrimination claim: the premium is largest where the five FDPs disagree most about whether the firm beat or missed. The concentration in high special-items firm-quarters supports the integration claim: the premium operates where the methodological choices distinguishing the FDPs most affect the reported surprise.

Together, Tables 4–7 document an accuracy premium that is statistically significant, economically meaningful, concentrated where the framework predicts, and tradeable in real time. The shift from heuristic to accuracy-weighted aggregation is the central empirical signature of Prediction 1. Section 6 turns to the natural follow-up: did this shift in market behavior arise from a change in how investors acquire and integrate forecast information?

6 LLM Visibility and Effective Synthesis Across FDPs

Motivation. This section develops the LLM-visibility pillar named in the paper's title: it introduces *LLM visibility* as a direct measure of GenAI's information acquisition, shows the measure is distinct from classical prominence proxies, and tests whether acquisition gates the accuracy premium. The preceding section establishes that the market prices the accuracy-weighted synthesis; the accessibility clause of Prediction 2 asks what makes that synthesis possible. Integration requires acquisition: the model cannot synthesize forecasts it cannot retrieve. LLM visibility measures

[28] A third moderator the framework predicts — whether GenAI actually retrieves the firm's forecasts — requires its own construct and is developed in Section 6. The plain post-ChatGPT-boundary variant of this table, which adds a median split on that LLM-visibility score, yields similar inferences, with the premium concentrating in the high half of each partition (untabulated).

acquisition directly — the number of source citations a browsing-enabled model returns when asked for a firm's forward consensus EPS, averaged across ChatGPT and Perplexity responses.[29] It captures what GenAI is *doing*, not what it *can* do.

A distinct construct. LLM visibility is empirically distinct from the classical prominence proxies. Table 9 shows that its high-visibility set overlaps those of firm size, analyst coverage, and institutional ownership only about 47–50 percent of the time — near the 50 percent coin-flip benchmark for independent median splits — whereas the classical proxies overlap one another up to 80 percent. In a pooled horse race entering each measure's premium-concentration triple interaction at once, LLM visibility survives against both benchmarks: 5.84 ($p < 0.01$) against I/B/E/S, where it is the only surviving measure, and 3.17 ($p < 0.10$) against the equal-weighted benchmark.

Acquisition as a precondition. Distinctness permits the sharper test: if acquisition gates synthesis, the premium should appear in high-visibility firms, and only once the model can browse. I re-estimate the horse race of Equation (6) within each regime, separately for above- and below-median visibility firms; the object of interest is the high-minus-low gap in the accuracy premium,

$$\Delta Premium_r \;=\; Premium_r^{\text{High vis}}\text{-}Premium_r^{\text{Low vis}}, \qquad r \in \{\text{R0}, \text{R1}, \text{R2}, \text{R3}\}. \tag{7}$$

Table 8 shows that $\Delta Premium_r$ is statistically insignificant at the five-percent level in the pre-ChatGPT and text-only regimes (R0, R1) but turns large and positive once the model can browse and search: 2.7 to 3.0 in the browsing regime (R2, $p < 0.05$) and 3.1 to 4.2 under native search (R3, $p < 0.01$ against I/B/E/S and $p < 0.10$ against equal-weighting). The gap opens only where forecasts are visible to the model *and* only once it can retrieve them — acquisition is a binding precondition, not an incidental correlate.

Ruling out prominence. The regime timing also disposes of the natural concern that LLM visibility merely proxies for firm prominence. A firm's visibility is essentially time-invariant over

[29] When the browsing-enabled model answers the query, it lists the web pages it consulted (e.g., Yahoo Finance, MarketBeat, TipRanks). Counting these citations measures how many independent forecast sources the model actually retrieves for a firm — the breadth of its information acquisition, not the content of any one estimate. Appendix B reports worked examples.

this window, so a prominence explanation would predict the premium gap in the pre-ChatGPT and text-only regimes as well; the gap instead appears only once the model can browse, which is difficult to reconcile with a static firm characteristic driving the gap.

7 Additional Tests

Sections 4 and 5 establish that the post-ChatGPT market prices an accuracy-weighted synthesis of consensus EPS forecasts. This section subjects that conclusion to three further tests, each addressing a distinct concern a skeptic could raise. The first asks whether the market's selectivity is specific to trailing accuracy or extends to a second quality dimension: a forecast-based informativeness measure built from the consensus's own predictive record. The last two identify the GenAI channel itself — first through investor information acquisition, then through major ChatGPT outages.

7.1 A Forecast-Based Informativeness Measure

Motivation. The accuracy dimension priced in Section 5 is one notion of quality; a complementary notion asks how *informative* the consensus is about the firm's future performance. If the post-ChatGPT market's selectivity reflects a general taste for consensus quality rather than a peculiarity of the trailing-accuracy construction, it should price this second dimension as well.

Design. For each firm, I score the consensus's predictive record in real time. The building block is the own-basis anticipation error — "own basis" because the consensus is paired with the same provider's street actual: the absolute distance between the quarter-q I/B/E/S consensus and the seasonally matched quarter-$(q+4)$ street actual, per share, with a stock-split guard. I average this error over up to the twenty most recent source quarters whose year-ahead actuals are already public at the announcement (at least twelve), scale it by the window mean $|actual|$ so it is comparable across firms, and winsorize. The measure is then simply $\textit{Predictive Value}^{IBES}_{i,q} = 1\text{-}error_{i,q}$, a continuous cross-firm variable (mean 0.56, standard deviation 0.30): higher values mean the firm's consensus has anticipated its year-ahead street numbers more closely. The measure covers 22,364 of the 23,833 firm-quarters (8,730 pre-ChatGPT); the remainder are recent listings with fewer than twelve scoreable history quarters. If post-ChatGPT investors read earnings through tools that weight the consensus by its demonstrated predictive value, the announcement-window response to the surprise should steepen in *Predictive Value*^{IBES} after ChatGPT.

Results. Table 10 reports the pricing. The ERC gradient on predictive value is strongly positive already before ChatGPT (2.17, $t = 4.0$) — firms with a predictable consensus carry higher ERCs, a familiar association — so the identified claim is the post-ChatGPT increment: $+2.05$ ($t = 2.68$) market-adjusted and $+2.10$ ($t = 2.82$) under the Fama-French-Carhart adjustment. The increment is comparable in size to the pre-period gradient, so the gradient roughly doubles after ChatGPT; dispersion, volatility, and the unexpected-item flag are among the controls, and the *Predictive Value*^{IBES} level terms are indistinguishable from zero. Columns (3)–(4) trace the increment across the three post-ChatGPT capability regimes: 0.92 ($t = 0.9$) under text-only, 1.43 ($t = 1.5$) under browsing, and 3.54 ($t = 3.1$) under native search — the same capability staircase as the accuracy premium of Table 6, and consistent with the retrieval-intensive nature of a forecast track record, which only browsing and native search made accessible.

Two alternative readings of the increment fail to absorb it. It is not announcement composition in disguise: jointly adding interactions of the large-unexpected-item flag of Bochkay et al. (2022) leaves the increment unchanged (2.04, $t = 2.67$), and the two measures are essentially uncorrelated (-0.02). Nor is it selection toward I/B/E/S specifically: in an untabulated decomposition of the surprise into the component common to all providers and the I/B/E/S-specific tilt, both the gradient (2.3, $t = 3.7$) and its post-ChatGPT increment (1.9, $t = 2.0$) price the common component. The post-ChatGPT market rewards the consensus signal of well-predicted firms, wherever it is sourced.

The provider level tells the same story, with one qualification. A within-firm-quarter four-provider rank version of the same score (each provider's own-basis error against its own year-ahead actual; Bloomberg excluded by its 2020Q2 native-history start) shows the same profile: no pre-period rank gradient, a post-period gradient of 0.320 ($t = 4.14$), and a shift of 0.242 ($t = 2.56$). But its significance concentrates in Capital IQ provider-events, and the informativeness and accuracy rankings are statistically indistinguishable once their low rank reliabilities are corrected, so accuracy weighting already embeds much of the provider-level informativeness signal.

Beyond the pricing test, the measure rationalizes the paper's operating choice. Constructed identically at the current-quarter horizon, the same procedure yields the consensus's trailing *accuracy*, and across firms the two are highly correlated (Pearson 0.72, Spearman 0.78; 0.82 correcting for split-half measurement error; stable pre- and post-ChatGPT and within announcement quarters). Informativeness and accuracy are thus two horizons of one underlying quality — how predictable the firm's street numbers are by its consensus. That the paper's analyses operate on the accuracy dimension is therefore a matter of econometric convenience rather than substance: the current-quarter error is realized within a quarter rather than after a four-quarter maturation, so accuracy is observable in real time at every announcement while carrying most of the informativeness content.

Two boundaries temper the reading. First, the own-basis pairing rewards internal consistency of street-earnings rules alongside genuine foresight about the firm, and the level gradient partly reflects earnings predictability — which is why the identified claim is confined to the post-ChatGPT increment. Second, the evidence concerns the consensus signal of well-predicted firms, not provider selection: the untabulated I/B/E/S tilt interactions are at best marginal, and the provider-level rank version's shift is, as noted, concentrated in Capital IQ provider-events (excluding Capital IQ, 0.099, $t = 0.75$; excluding any other provider, 0.26–0.31) rather than uniform across the four.

7.2 Additional Tests on Information Acquisition

Motivation. The remaining two tests turn from what the market prices to how the signal reaches it. The accuracy-weighted synthesis documented in Section 5 requires that investors actually access consensus information from multiple FDPs. As Appendix B documents, ChatGPT and Perplexity responses to direct EPS queries overwhelmingly cite free distributor platforms — Yahoo Finance, MarketWatch, Seeking Alpha, and similar aggregators that publish forecasts from several FDPs without paywalls. Whether GenAI substitutes for direct access to these distributors, or instead acts as a gateway that broadens it, is an open question.

The two interpretations make opposite predictions for direct Google search of free distributors. Under substitution, investors query ChatGPT in place of typing "Yahoo Finance AAPL earnings" into Google, so search for free distributors should fall post-ChatGPT and the cross-FDP access channel narrows rather than broadens. Under the gateway interpretation, GenAI surfaces these distributors to investors who would not otherwise visit them, so post-ChatGPT search for free distributors should rise even as ChatGPT itself becomes available. The direction of post-ChatGPT platform search therefore provides a sharp test between the two interpretations.

Design. I obtain daily Google search volume for 17 platform brands, classified into Free Distributors (twelve aggregator sites including Yahoo Finance, MarketWatch, and Seeking Alpha), Paywalled FDPs (Bloomberg, Capital IQ, FactSet, I/B/E/S), and the Free FDP (Zacks); the full classification is in Appendix A. The dependent variable is each brand's share of total platform-related search on the day.[30] I estimate

$$S_{p,t} = \alpha + \sum_{g} \delta_g \, Post_t \times Group_{g,p} + \lambda \, Post_t + \gamma X_t + \text{calendar-month FE} + \varepsilon_{p,t}, \quad (8)$$

where $Group_{g,p}$ indexes the three platform groups (Free Distributor in the pre-ChatGPT regime is the omitted baseline) and X_t contains controls for total platform-related search, earnings-announcement intensity, and trading-day status. Under the gateway interpretation, the post-ChatGPT shift should be positive for Free Distributors and negative for the two FDP groups.

Results. Table 11, Panel B, reports the platform-day estimates. Column (1) uses a single post-ChatGPT indicator; column (2) splits the post-ChatGPT period into three GenAI regimes. In column (1), the post-ChatGPT share of total platform-related search rises by 2.1 percentage points for Free Distributors and falls by 7.5 and 7.1 percentage points for the Free FDP and Paywalled FDPs respectively, all significant at the 1% level. Because shares sum to one across platforms each day, these movements describe substitution within the platform universe: search shifts from FDP brands to free aggregators. Column (2) confirms that the pattern is stable across the three regimes; the share gain for Free Distributors holds at roughly 2.0 to 2.2 percentage points throughout.

[30] Google Trends reports a relative search index normalized within a query window rather than an absolute volume (see, e.g., Da et al., 2011; Drake et al., 2012), so I use the DataForSEO Search Volume API, which produces absolute search counts from Google's keyword-level impression and click data. Scaling by total platform-related search captures general demand for financial-information platforms, the relevant baseline for the substitution-versus-gateway comparison; scaling instead by platform-plus-earnings combinations (e.g., "Yahoo Finance EPS forecast") yields volumes below DataForSEO's minimum detectability threshold for most platforms.

The pattern is consistent with the gateway interpretation. Post-ChatGPT, investor demand for FDP-branded search falls and demand for free aggregators rises, even though ChatGPT itself can answer the underlying queries directly. On this reading, GenAI operates as a gateway rather than a substitute: investors appear to learn about free distributors through ChatGPT and return to them through direct search. The result broadens the empirical foundation of the accuracy premium and is difficult to reconcile with the alternative reading in which GenAI merely intermediates existing investor demand.

7.3 Major ChatGPT Outages as an Exogenous Shock

Motivation. To identify the GenAI channel against alternative explanations for the post-ChatGPT accuracy premium, I exploit major ChatGPT outages as exogenous, short-duration shocks to GenAI availability, following Cheng et al. (2025) and Liu and Subasi (2026). On announcement dates immediately following a major outage, investors temporarily lose the integration capability that GenAI provides and should revert to the only heuristic that requires no multi-FDP synthesis: the single-FDP I/B/E/S consensus. The prediction is therefore an asymmetric outage effect across the two horse-race benchmarks. Against the I/B/E/S benchmark, the accuracy premium should weaken on outage days, because investors substitute toward I/B/E/S. Against the equal-weighted benchmark, it should be unchanged, because both the accuracy-weighted and equal-weighted aggregates require multi-FDP synthesis and become equally infeasible when GenAI is down.

Design. I collect ChatGPT outage incidents from OpenAI's public status-page history and define a major outage as any incident lasting at least 240 minutes; this yields four incidents across five unique calendar days through 2024. The treatment indicator $Outage_{i,q}$ equals one for firm-quarter announcement dates falling on the calendar day immediately following a major outage and zero otherwise. I re-estimate the horse-race specification of Equation (6) on the post-ChatGPT subsample, interacting both surprise measures with $Outage_{i,q}$, separately for the equal-weighted and I/B/E/S benchmarks. The small number of treated days limits power on the level of the outage effect; identification rests instead on the asymmetric sign predicted across the two benchmarks.

Results. Table 11, Panel A, reports the estimates. The two benchmarks deliver sharply different outage effects, exactly as the asymmetric prediction implies. Against the equal-weighted benchmark (column 1), the outage-day change in the accuracy premium is statistically insignificant (3.787, $t = 1.37$), indicating no detectable effect on the accuracy-weighted-versus-equal-weighted contrast. Against the I/B/E/S benchmark (column 2), the outage-day change is -3.173 ($p < 0.05$): the accuracy-weighted slope falls ($\beta_3 =$ -1.685, $p < 0.10$) and the I/B/E/S slope rises ($\beta_4 = 1.488$, $p < 0.05$) on outage days. Investors temporarily reweight toward I/B/E/S precisely when GenAI is unavailable.

The asymmetry across benchmarks is the identifying signature of the GenAI channel. Both the accuracy-weighted and equal-weighted aggregates require investors to draw consensus information from all five FDPs; when GenAI access is severed, neither weighting is operative and their relative loading is unchanged. The I/B/E/S consensus, by contrast, is the single-FDP fallback that requires no multi-source integration. Investors therefore revert specifically to I/B/E/S, the accuracy premium collapses against the I/B/E/S benchmark, and the equal-weighted contrast remains unaffected. This sign-and-asymmetry pattern is difficult to generate without a GenAI-mediated multi-FDP integration channel. Together with the cross-sectional concentration in Section 5 and the gateway evidence above, the outage test identifies GenAI as the channel through which the post-ChatGPT accuracy-weighted synthesis reaches market prices.

8 Conclusion

This paper shows that after the release of ChatGPT, market reactions to earnings announcements increasingly reflect an accuracy-weighted synthesis across the five major forecast data providers, departing from the equal-weighted or single-FDP heuristics that dominated prior decades. Four pieces of evidence support this conclusion. First, the post-ChatGPT market discriminates across FDPs by accuracy: the within-firm-quarter ERC steepens with accuracy rank between regimes, and the firm-level response to the I/B/E/S consensus surprise weakens specifically in firms where I/B/E/S has lost relative accuracy. Second, the accuracy-weighted multi-provider consensus carries a statistically and economically significant accuracy premium over both the equal-weighted alternative and the I/B/E/S-only consensus. The premium concentrates where GenAI synthesis has the most discriminating work to do (high cross-FDP disagreement) and is most integration-intensive (large special items); it translates into tradeable quintile long–short returns on the predicted surprise; it climbs across GPT-capability regimes to 4.91 under native search; and it concentrates in firms whose forecasts the model demonstrably retrieves (LLM visibility) — but only once retrieval is possible.

The remaining two pieces tie this shift to GenAI itself. Third, the market's selectivity extends to forecast-based informativeness: when each firm is scored in real time by how closely the consensus has anticipated its year-ahead street numbers over the preceding twenty quarters, the post-ChatGPT rise in announcement reactions concentrates in firms with an informative consensus. The increment is significant under both return adjustments, rises monotonically across GPT-capability regimes to its peak under native search, is orthogonal to announcement composition, and loads on the common cross-provider consensus rather than an I/B/E/S-specific tilt. Fourth, major ChatGPT outages produce an asymmetric pattern that ties the premium to GenAI-enabled integration: the accuracy premium against I/B/E/S collapses on outage days while the equal-weighted contrast shows no significant change — the signature a GenAI-mediated multi-FDP channel predicts. A corroborative search-data test shows that post-ChatGPT investor acquisition broadens, rather than substitutes for, direct access to free distributors of cross-FDP consensus.

These findings make three contributions. First, they reposition the choice of forecast data provider from a researcher's measurement decision to an endogenous market outcome that responds to the information environment, extending recent work on cross-FDP differences in consensus information (Larocque et al., 2025). The historical reliance on I/B/E/S reflects a technology-bound equilibrium that GenAI now relaxes, with measurable consequences for earnings-announcement returns. Second, they identify a specific channel through which GenAI reaches capital markets: low-cost multi-source synthesis of analyst forecasts, which the market now prices into earnings-announcement reactions. This isolates the analyst-information channel from the broader set of GenAI effects on investor behavior and asset prices documented in recent work (e.g., Bertomeu et al., 2025; Cheng et al., 2025; Liu and Subasi, 2026). Third, they sharpen the disclosure-processing-cost framework of Blankespoor et al. (2020) by identifying cross-source integration, distinct from awareness or acquisition within a single source (e.g., Blankespoor et al., 2014; Blankespoor, 2019; Drake et al., 2015), as the cost component that prior single-source technologies left intact and that GenAI uniquely relaxes.

Four caveats warrant mention. First, the accuracy-weighted construction is a tractable proxy; alternative weighting schemes consistent with the same hypothesis should yield similar inferences. Second, the outage test rests on few treated days, none spanning a full calendar day, so identification comes from the asymmetric sign across benchmarks rather than from level estimates, and the documented magnitude is likely a lower bound. Third, the LLM visibility measure is a point-in-time snapshot of inherently variable LLM outputs, which limits its reliability (e.g., Chen et al., 2023; Lopez-Lira et al., 2025). Fourth, the forecast-based informativeness measure of Section 7 rewards internal consistency of street-earnings rules alongside genuine foresight, and its level gradient partly reflects earnings predictability; the identified claim is therefore the post-ChatGPT increment, and the pricing operates on the common consensus rather than on provider selection. The provider-level rank version of the measure is noisy (split-half reliability below 0.35), with its shift concentrated in Capital IQ provider-events rather than uniform across the four providers.

Several questions remain open. First, whether the accuracy premium widens or contracts as GenAI tools evolve and as institutional investors integrate them into their workflows requires a longer post-ChatGPT panel to settle. Second, whether the cross-FDP synthesis advantage extends to other forecasted quantities, such as revenue, cash flow, or industry-specific metrics, is unexamined. Third, the framework here treats FDPs as fixed information producers, but the market's new discrimination across them gives FDPs an incentive to revisit their own methodologies; this supply-side response bears on both the disclosure-processing-cost framework and the structure of the FDP industry.

The broader lesson is that the market's earnings benchmark is technology-bound. For decades, a single provider's consensus stood in for the market's expectation because integrating across providers was costly; GenAI has relaxed that constraint, and the benchmark has moved. The market now weights providers by the accuracy they have earned, not the convenience of reaching them — what investors can cheaply integrate shapes what prices reflect.

References

Ahlers, D. and J. Lakonishok. 1983. A study of economists' consensus forecasts. *Management Science* 29 (10): 1113–1125. doi:10.1287/mnsc.29.10.1113.

Bartov, E., D. Givoly, and C. Hayn. 2002. The rewards to meeting or beating earnings expectations. *Journal of Accounting and Economics* 33 (2): 173–204. https://doi.org/10.1016/S0165-4101(02)00045-9.

Bates, J. M. and C. W. J. Granger. 1969. The combination of forecasts. *Operational Research Quarterly* 20 (4): 451–468. doi:10.1057/jors.1969.103.

Battalio, R. H. and R. R. Mendenhall. 2005. Earnings expectations, investor trade size, and anomalous returns around earnings announcements. *Journal of Financial Economics* 77 (2): 289–319. https://doi.org/10.1016/j.jfineco.2004.08.002.

Bernard, V. L. and J. K. Thomas. 1989. Post-earnings-announcement drift: Delayed price response or risk premium? *Journal of Accounting Research* 27: 1–36. https://doi.org/10.2307/2491062. Supplement.

Bertomeu, J., Y. Lin, Y. Liu, and Z. Ni. 2025. The impact of generative AI on information processing: Evidence from the ban of ChatGPT in Italy. *Journal of Accounting and Economics* 80 (1): 101782. https://doi.org/10.1016/j.jacceco.2025.101782.

Beyer, A., D. A. Cohen, T. Z. Lys, and B. R. Walther. 2010. The financial reporting environment: Review of the recent literature. *Journal of Accounting and Economics* 50 (2–3): 296–343. https://doi.org/10.1016/j.jacceco.2010.10.003.

Blankespoor, E. 2019. The impact of information processing costs on firm disclosure choice: Evidence from the XBRL mandate. *Journal of Accounting Research* 57 (1): 75–120. doi:10.1111/1475-679X.12237.

Blankespoor, E., G. S. Miller, and H. D. White. 2014. The role of dissemination in market liquidity: Evidence from firms' use of Twitter. *The Accounting Review* 89 (1): 79–112. https://doi.org/10.2308/accr-50576.

Blankespoor, E., E. deHaan, and C. Zhu. 2018. Capital market effects of media synthesis and dissemination: Evidence from robo-journalism. *Review of Accounting Studies* 23 (1): 1–36. https://doi.org/10.1007/s11142-017-9422-2.

Blankespoor, E., E. de Haan, and I. Marinovic. 2020. Disclosure processing costs, investors' information choice, and equity market outcomes: A review. *Journal of Accounting and Economics* 70 (2–3): 101344. https://doi.org/10.1016/j.jacceco.2020.101344.

Blankespoor, E., J. Croom, and S. M. Grant. 2026. Generative AI and investor processing of financial information. SSRN Working Paper No. 5053905; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5053905. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5053905.

Bochkay, K., S. Markov, M. Subasi, and E. H. Weisbrod. 2022. The roles of data providers and analysts in the production, dissemination, and pricing of street earnings. *Journal of Accounting Research* 60 (5): 1695–1740. https://doi.org/10.1111/1475-679X.12457.

Bradshaw, M. T., M. S. Drake, J. N. Myers, and L. A. Myers. 2012. A re-examination of analysts' superiority over time-series forecasts of annual earnings. *Review of Accounting Studies* 17 (4): 944–968. https://doi.org/10.1007/s11142-012-9185-8.

Bradshaw, M. T., Y. Ertimur, and P. C. O'Brien. 2017. Financial analysts and their contribution to well-functioning capital markets. *Foundations and Trends in Accounting* 11 (3): 119–191. https://doi.org/10.1561/1400000042.

Bradshaw, M. T., C. Ma, B. P. Yost, and Y. Zou. 2026. Generative AI use by capital market information intermediaries: Evidence from Seeking Alpha. *Journal of Accounting Research* . doi:10.1111/1475-679X.70053; forthcoming.

Bushee, B. J., D. A. Matsumoto, and G. S. Miller. 2003. Open versus closed conference calls: The determinants and effects of broadening access to disclosure. *Journal of Accounting and Economics* 34 (1–3): 149–180. https://doi.org/10.1016/S0165-4101(02)00073-3.

Bushee, B. J., J. E. Core, W. Guay, and S. J. W. Hamm. 2010. The role of the business press as an information intermediary. *Journal of Accounting Research* 48 (1): 1–19. https://doi.org/10.1111/j.1475-679X.2009.00357.x.

CFA Institute. 2024. AI industry survey: Investment industry employers' views on the responsible and ethical use of AI. CFA Institute Research and Policy Center Survey Report. https://www.cfainstitute.org/about/press-room/2024/ai-in-investment-sector-survey. https://www.cfainstitute.org/about/press-room/2024/ai-in-investment-sector-survey.

Chang, A. Y., X. Dong, X. Martin, and C. Zhou. 2026. AI democratization and trading inequality. *Journal of Accounting Research* . https://doi.org/10.1111/1475-679X.70063.

Chen, L., M. Zaharia, and J. Zou. 2023. How is ChatGPT's behavior changing over time? arXiv preprint arXiv:2307.09009; https://arxiv.org/abs/2307.09009.

Cheng, Q., P. Lin, and Y. Zhao. 2025. Does generative AI facilitate investor trading? early evidence from ChatGPT outages. *Journal of Accounting and Economics* 80 (2): 101807. https://doi.org/10.1016/j.jacceco.2025.101807.

Clemen, R. T. 1989. Combining forecasts: A review and annotated bibliography. *International Journal of Forecasting* 5 (4): 559–583. https://doi.org/10.1016/0169-2070(89)90012-5.

Clement, M. B. 1999. Analyst forecast accuracy: Do ability, resources, and portfolio complexity matter? *Journal of Accounting and Economics* 27 (3): 285–303. https://doi.org/10.1016/S0165-4101(99)00013-0.

Clement, M. B. and S. Y. Tse. 2003. Do investors respond to analysts' forecast revisions as if forecast accuracy is all that matters? *The Accounting Review* 78 (1): 227–249. https://doi.org/10.2308/accr.2003.78.1.227.

Cohen, L. and D. Lou. 2012. Complicated firms. *Journal of Financial Economics* 104 (2): 383–400. https://doi.org/10.1016/j.jfineco.2011.08.006.

Coleman, B., T. Dyer, and M. Lang. 2025. The influence of financial data subscriptions on analyst research. Working paper, University of Georgia, Brigham Young University, and University of North Carolina at Chapel Hill.

Collins, D. W. and S. P. Kothari. 1989. An analysis of intertemporal and cross-sectional determinants of earnings response coefficients. *Journal of Accounting and Economics* 11 (2–3): 143–181. https://doi.org/10.1016/0165-4101(89)90004-9.

Da, Z., J. Engelberg, and P. Gao. 2011. In search of attention. *Journal of Finance* 66 (5): 1461–1499. https://doi.org/10.1111/j.1540-6261.2011.01679.x.

Da, Z., J. Engelberg, and P. Gao. 2015. The sum of all FEARS: Investor sentiment and asset prices. *Review of Financial Studies* 28 (1): 1–32. https://doi.org/10.1093/rfs/hhu072.

deHaan, E., T. Shevlin, and J. Thornock. 2015. Market (in)attention and the strategic scheduling and timing of earnings announcements. *Journal of Accounting and Economics* 60 (1): 36–55. https://doi.org/10.1016/j.jacceco.2015.03.003.

DellaVigna, S. and J. M. Pollet. 2009. Investor inattention and Friday earnings announcements. *Journal of Finance* 64 (2): 709–749. https://doi.org/10.1111/j.1540-6261.2009.01447.x.

Diether, K. B., C. J. Malloy, and A. Scherbina. 2002. Differences of opinion and the cross section of stock returns. *Journal of Finance* 57 (5): 2113–2141. https://doi.org/10.1111/0022-1082.00490.

Drake, M. S., D. T. Roulstone, and J. R. Thornock. 2012. Investor information demand: Evidence from Google searches around earnings announcements. *Journal of Accounting Research* 50 (4): 1001–1040. https://doi.org/10.1111/j.1475-679X.2012.00443.x.

Drake, M. S., D. T. Roulstone, and J. R. Thornock. 2015. The determinants and consequences of information acquisition via EDGAR. *Contemporary Accounting Research* 32 (3): 1128–1161. https://doi.org/10.1111/1911-3846.12119.

Driskill, M., M. P. Kirk, and J. W. Tucker. 2020. Concurrent earnings announcements and analysts' information production. *The Accounting Review* 95 (1): 165–189. https://doi.org/10.2308/accr-52489.

Easton, P. D. and M. E. Zmijewski. 1989. Cross-sectional variation in the stock market response to accounting earnings announcements. *Journal of Accounting and Economics* 11 (2–3): 117–141. https://doi.org/10.1016/0165-4101(89)90003-7.

Goldstein, I., C. S. Spatt, and M. Ye. 2021. Big data in finance. *The Review of Financial Studies* 34 (7): 3213–3225. https://doi.org/10.1093/rfs/hhab038.

Granger, C. W. J. and R. Ramanathan. 1984. Improved methods of combining forecasts. *Journal of Forecasting* 3 (2): 197–204. https://doi.org/10.1002/for.3980030207.

Gu, S., B. Kelly, and D. Xiu. 2020. Empirical asset pricing via machine learning. *The Review of Financial Studies* 33 (5): 2223–2273. https://doi.org/10.1093/rfs/hhaa009.

Hirshleifer, D., S. S. Lim, and S. H. Teoh. 2009. Driven to distraction: Extraneous events and underreaction to earnings news. *Journal of Finance* 64 (5): 2289–2325. https://doi.org/10.1111/j.1540-6261.2009.01501.x.

Hou, K., M. A. van Dijk, and Y. Zhang. 2012. The implied cost of capital: A new approach. *Journal of Accounting and Economics* 53 (3): 504–526. https://doi.org/10.1016/j.jacceco.2011.12.001.

Kimbrough, M. D. 2005. The effect of conference calls on analyst and market underreaction to earnings announcements. *The Accounting Review* 80 (1): 189–219. https://doi.org/10.2308/accr.2005.80.1.189.

Kimbrough, M. D., M. Subasi, and L. Yang Liu. 2026. AI-powered analysts. Working paper, University of Maryland.

Kothari, S. P. 2001. Capital markets research in accounting. *Journal of Accounting and Economics* 31 (1–3): 105–231. https://doi.org/10.1016/S0165-4101(01)00030-1.

Larocque, S. A., J. Watkins, and E. H. Weisbrod. 2025. Consensus? an examination of differences in earnings information across forecast data providers. SSRN Working Paper No. 4640370; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4640370. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4640370.

Lawrence, A. 2013. Individual investors and financial disclosure. *Journal of Accounting and Economics* 56 (1): 130–147. https://doi.org/10.1016/j.jacceco.2013.05.001.

Liu, L. Y. and M. Subasi. 2026. The role of generative ai in democratizing investment research: Evidence from crowdsourced forecasts. Working paper, University of Maryland.

Livnat, J. and R. R. Mendenhall. 2006. Comparing the post-earnings announcement drift for surprises calculated from analyst and time series forecasts. *Journal of Accounting Research* 44 (1): 177–205. https://doi.org/10.1111/j.1475-679X.2006.00196.x.

Lopez-Lira, A. and Y. Tang. 2023. Can ChatGPT forecast stock price movements? return predictability and large language models. SSRN Working Paper No. 4412788; arXiv:2304.07619; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4412788. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4412788.

Lopez-Lira, A., Y. Tang, and M. Zhu. 2025. The memorization problem: Can we trust LLMs' economic forecasts? SSRN Working Paper No. 5217505; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5217505. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5217505.

Loughran, T. and B. McDonald. 2011. When is a liability not a liability? Textual analysis, dictionaries, and 10-Ks. *The Journal of Finance* 66 (1): 35–65. https://doi.org/10.1111/j.1540-6261.2010.01625.x.

Loughran, T. and B. McDonald. 2017. The use of EDGAR filings by investors. *Journal of Behavioral Finance* 18 (2): 231–248. https://doi.org/10.1080/15427560.2017.1308945.

Michaely, R., A. Rubin, D. Segal, and A. Vedrashko. 2024. Do differences in analyst quality matter for investors relying on consensus information? *Management Science* 70 (2): 751–772. doi:10.1287/mnsc.2023.4699.

Mikhail, M. B., B. R. Walther, and R. H. Willis. 1997. Do security analysts improve their performance with experience? *Journal of Accounting Research* 35: 131–157. https://doi.org/10.2307/2491456. Supplement.

Park, C. W. and E. K. Stice. 2000. Analyst forecasting ability and the stock price reaction to forecast revisions. *Review of Accounting Studies* 5 (3): 259–272. https://doi.org/10.1023/A:1009668711298.

Schipper, K. 1991. Analysts' forecasts. *Accounting Horizons* 5 (4): 105–121. December issue.

Stickel, S. E. 1992. Reputation and performance among security analysts. *Journal of Finance* 47 (5): 1811–1836. doi:10.1111/j.1540-6261.1992.tb04672.x.

Stock, J. H. and M. W. Watson. 2004. Combination forecasts of output growth in a seven-country data set. *Journal of Forecasting* 23 (6): 405–430. https://doi.org/10.1002/for.928.

Surowiecki, J. 2004. *The Wisdom of Crowds: Why the Many Are Smarter Than the Few and How Collective Wisdom Shapes Business, Economies, Societies and Nations*. New York, Doubleday. ISBN 0-385-50386-5.

Tetlock, P. C. 2007. Giving content to investor sentiment: The role of media in the stock market. *The Journal of Finance* 62 (3): 1139–1168. https://doi.org/10.1111/j.1540-6261.2007.01232.x.

Timmermann, A. 2006. Forecast combinations. In Elliott, G., C. W. J. Granger, and A. Timmermann, editors, *Handbook of Economic Forecasting*, volume 1, chapter 4, pages 135–196. Elsevier. https://doi.org/10.1016/S1574-0706(05)01004-9.

Xue, J., Q. Zhang, and W. Zhu. 2025. Generative AI for analysts. SSRN Working Paper No. 5821142; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5821142.

FIGURE 1. Theoretical Framework

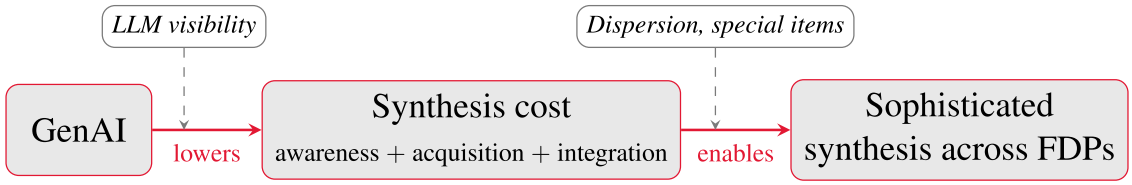

Notes. This figure summarizes the theoretical framework. Generative AI (GenAI) lowers the cost of synthesizing analyst consensus forecasts across the five forecast data providers (FDPs), the joint cost of *awareness*, *acquisition*, and *integration* in the Blankespoor et al. (2020) decomposition, which in turn enables the market to price a sophisticated, accuracy-weighted synthesis across providers in place of a single-FDP or equal-weighted heuristic. The strength of the channel is moderated by how *accessible* synthesis is for a given firm (LLM visibility, the information-acquisition channel) and by how *valuable* it is (cross-FDP disagreement and special-items magnitude). These mechanisms map to Predictions 1 and 2, developed in Section 2; Prediction 3 adds the demand-side corollary that the numbers themselves are unchanged.

FIGURE 2. Mean Magnitude of *SUE*: Accuracy-Weighted vs. Heuristic Models

(a) Quarterly mean *SUE* by consensus method

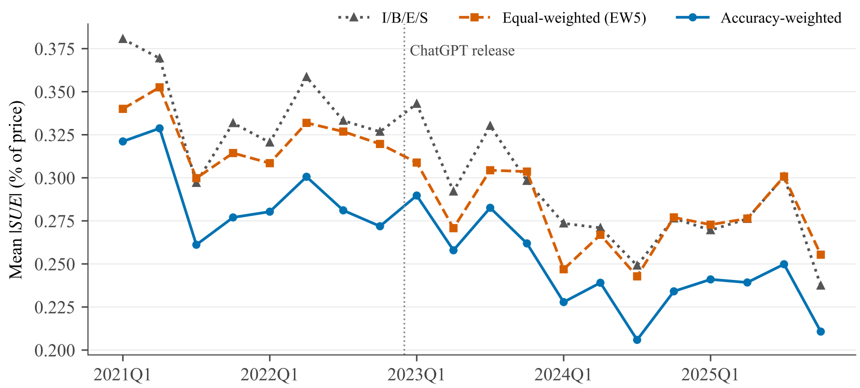

(b) Mean *SUE* by consensus method

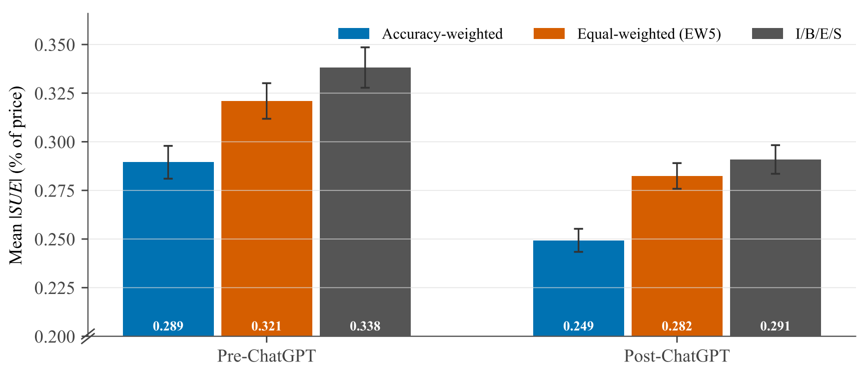

Notes. This figure compares the accuracy of the accuracy-weighted multi-vendor consensus SUE^{Acc} against the equal-weighted five-FDP consensus SUE^{EW5} and the I/B/E/S consensus SUE^{IBES} across the five-FDP sample (FactSet, Bloomberg, Zacks, Capital IQ, I/B/E/S) of S&P 1500 firm-quarters with announcement dates in 2021Q1–2025Q4. $|SUE|$ is the absolute standardized unexpected earnings scaled as a percentage of pre-announcement price, winsorized at the (1, 99) level. **Panel (a)** plots the quarterly mean $|SUE|$ for each of the three methods over 20 announcement quarters. The vertical reference line marks the public release of ChatGPT (November 30, 2022). **Panel (b)** shows the pooled mean $|SUE|$ within each Pre- vs. Post-ChatGPT regime, with 95% confidence intervals from the cell-level standard error of the mean; the *y*-axis is truncated below 0.20 (break mark on axis) to expand the visible range.

FIGURE 3. The Accuracy Premium Rises with GPT Retrieval Capability

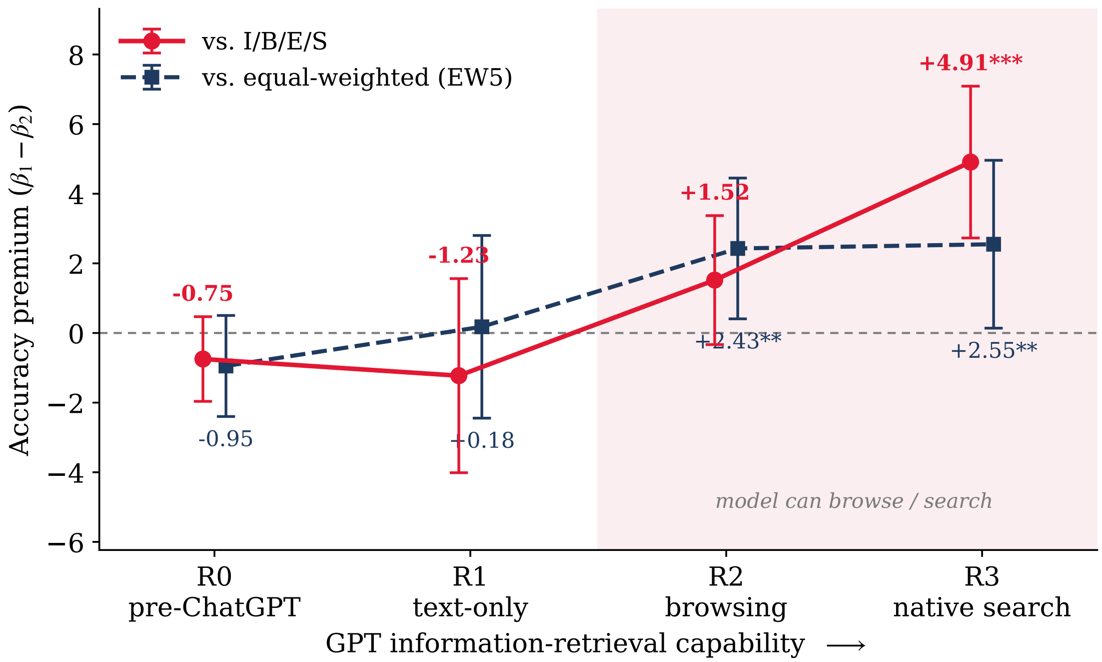

Notes. This figure plots the *accuracy premium*, the contrast β_1-β_2 between the announcement-return loading on the accuracy-weighted consensus surprise SUE^{Acc} and the loading on the benchmark surprise, estimated separately within each GPT-capability regime. The estimates are taken directly from Table 6: the red line uses the I/B/E/S benchmark (Panel B) and the navy line the equal-weighted five-FDP benchmark (Panel A). Regimes are ordered by ChatGPT's information-retrieval capability: **R0** (before 2022-11-30, pre-ChatGPT), **R1** (text-only), **R2** (browsing, from 2023-09-27), and **R3** (native search, from 2024-10-31); the shaded region marks the regimes in which the model can retrieve forecasts from the web. Whiskers are 95% confidence intervals; point labels report the premium and its significance ($^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$) from firm-clustered standard errors. The premium is statistically indistinguishable from zero before retrieval is available, turns positive once the model can browse, and grows monotonically thereafter, the dose-response Prediction 2 predicts.

TABLE 1. Sample Selection

Panel A. Main Sample (firm × quarter)	**Observations**	**# Firms**
S&P 1500 firm-quarters with an earnings announcement (announcement date 2021–2025)	29,395	1,502
Retain firm-quarters with non-missing consensus from all five FDPs (FactSet, Bloomberg, Zacks, Capital IQ, and I/B/E/S)	25,027	1,353
Drop firm-quarters missing market-outcome variables or required controls	(1,179)	(13)
Drop firm-quarters missing any constructed return or predicted-surprise variable	(10)	(1)
Drop firms with only one firm-quarter (no within-firm variation for firm fixed effects)	(5)	(5)
Final Sample	**23,833**	**1,334**

Panel B. Information Acquisition Tests (platform × day)	**Observations**	**# Platforms**
Identified platforms (5 FDPs + 15 distributors)	20	20
Platform-day observations, 2020–2025	43,820	20
Drop platform-days missing Ads scaling	(2,191)	(1)
Require consistent Trends coverage 2020–2025	(4,382)	(2)
Final Sample	**37,247**	**17**

Notes. This table reports the sample selection procedure for the two analytic samples in the paper. **Panel A** reports the firm-quarter sample used in the main across-FDP tests, built from CRSP, Compustat, I/B/E/S, and the five-FDP consensus panel. **Panel B** reports the platform-day sample used in the information-acquisition test.

TABLE 2. Descriptive Statistics

	N	Mean	SD	Min	p25	Median	p75	Max
Panel A. Main Sample (Firm $\times$ quarter).								
Dependent Variables: Abnormal Returns (%)								
$CAR^{Mkt}_{[-1,+1]}$	23,833	0.204	7.739	-22.532	-3.933	0.168	4.255	23.535
$CAR^{FF4}_{[-1,+1]}$	23,833	0.196	7.635	-22.779	-3.776	0.186	4.135	23.606
$BHAR^{Mkt}_{[-2,+1]}$	23,833	0.228	7.978	-22.803	-4.165	0.177	4.418	24.992
$BHAR^{FF4}_{[-2,+1]}$	23,833	0.193	7.847	-22.935	-3.966	0.111	4.255	24.819
Per-FDP Standardized Earnings Surprise (SUE^g, %)								
$SUE^{FactSet}$	23,833	0.132	0.534	-2.190	0.000	0.074	0.238	2.484
$SUE^{Bloomberg}$	23,833	-0.222	1.659	-10.049	-0.362	-0.027	0.162	5.780
SUE^{Zacks}	23,833	0.149	0.467	-1.606	0.000	0.075	0.240	2.291
SUE^{CapIQ}	23,833	0.137	0.530	-2.185	0.001	0.076	0.243	2.477
SUE^{IBES}	23,833	0.130	0.548	-2.271	0.000	0.075	0.241	2.502
Cross-FDP Composites and Predicted Surprise								
SUE^{Acc}	23,833	0.110	0.450	-1.656	-0.015	0.059	0.212	1.972
SUE^{EW5}	23,833	0.066	0.510	-1.936	-0.064	0.041	0.201	1.942
$Standardized\ Pred.Surp.^{EW5}$	23,833	0.606	0.116	0.000	0.590	0.619	0.631	1.000
$Standardized\ Pred.Surp.^{IBES}$	23,833	0.483	0.109	0.000	0.469	0.478	0.493	1.000
Firm and Event Controls								
Log Market Equity	23,833	8.992	1.448	6.390	7.923	8.769	9.942	12.883
Book-to-Market	23,833	0.440	0.361	-0.159	0.170	0.362	0.634	1.670
Leverage	23,833	0.314	0.218	0.001	0.133	0.301	0.444	1.008
Prior 12-Month Return	23,833	0.201	0.487	-0.589	-0.101	0.112	0.379	2.400
Log Price	23,833	4.166	0.980	1.880	3.491	4.178	4.831	6.580
Return Volatility	23,833	0.021	0.009	0.009	0.015	0.019	0.025	0.054
Analyst Dispersion	23,833	0.002	0.003	0.000	0.000	0.001	0.002	0.018
Busyness	23,833	4.944	0.973	1.946	4.431	5.124	5.684	6.299
$\lvert \Delta IB \rvert$	23,833	0.012	0.026	0.000	0.001	0.004	0.011	0.178
$\lvert \Delta Shares \rvert$	23,833	0.010	0.021	0.000	0.001	0.003	0.009	0.150
$\lvert SpecItems \rvert$	23,833	0.003	0.010	0.000	0.000	0.000	0.002	0.074
SBC High	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
Unexpected Item	23,833	0.686	0.464	0.000	0.000	1.000	1.000	1.000
Stock Split	23,833	0.003	0.051	0.000	0.000	0.000	0.000	1.000
Fiscal Q4	23,833	0.250	0.433	0.000	0.000	0.000	0.000	1.000
Institutional Ownership	23,833	0.850	0.238	0.011	0.813	0.940	1.000	1.000
# Analysts	23,833	10.914	6.536	2.000	6.000	10.000	15.000	31.000
Cross-Sectional Partitioning Indicators								
High Visibility	23,827	0.489	0.500	0.000	0.000	0.000	1.000	1.000

(continued on next page)

	N	Mean	SD	Min	p25	Median	p75	Max
High Cross-FDP Dispersion	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
High Special Items	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
IBES Accuracy-Rank Up	23,833	0.099	0.299	0.000	0.000	0.000	0.000	1.000
IBES Accuracy-Rank Down	23,833	0.059	0.236	0.000	0.000	0.000	0.000	1.000
ChatGPT Outage (prior day)	23,833	0.004	0.059	0.000	0.000	0.000	0.000	1.000
Panel B. Information Acquisition Sample (Platform × day).								
Dependent Variable								
Brand Share ($S_{\text{brand}}/S_{\text{all platforms}}$)	37,247	0.059	0.101	0.000	0.005	0.022	0.071	0.832
Macro and Calendar Controls								
Log Total Platform Searches	37,247	15.089	1.148	12.462	14.239	15.100	15.891	17.732
Log # Earnings Today	37,247	2.996	2.017	0.000	0.693	3.434	4.605	6.410
Trading Day	37,247	0.714	0.452	0.000	0.000	1.000	1.000	1.000
Vendor-Group Indicators								
Free Distributor	37,247	0.706	0.456	0.000	0.000	1.000	1.000	1.000
Free FDP	37,247	0.059	0.235	0.000	0.000	0.000	0.000	1.000
Paywalled FDP	37,247	0.235	0.424	0.000	0.000	0.000	0.000	1.000
Regime Indicators								
Post-ChatGPT	37,247	0.514	0.500	0.000	0.000	1.000	1.000	1.000
Regime 1	37,247	0.137	0.344	0.000	0.000	0.000	0.000	1.000
Regime 2	37,247	0.183	0.386	0.000	0.000	0.000	0.000	1.000
Regime 3	37,247	0.194	0.396	0.000	0.000	0.000	0.000	1.000

Notes. This table reports descriptive statistics for the two analytic samples. **Panel A** reports the across-FDP firm-quarter sample used in the main tests; **Panel B** reports the platform-day sample used in the information-acquisition test. All continuous variables are winsorized at the (1, 99) level. Variable definitions are in Appendix A.

TABLE 3. Initial Evidence on Selective Reliance

Panel A. Earnings-response coefficients by forecast-provider accuracy rank

	Pre-GPT		Post-GPT	
Provider	β	γ	β	γ
IBES	2.591***	+0.022	2.611***	+0.295**
	(6.60)	(0.17)	(5.49)	(2.09)
FactSet	2.316***	+0.050	2.059***	+0.380**
	(4.78)	(0.36)	(3.80)	(2.42)
Bloomberg	0.985***	-0.058	0.789***	+0.162**
	(5.82)	(-0.80)	(5.01)	(2.10)
Zacks	3.603***	-0.139	4.757***	-0.020
	(8.35)	(-1.04)	(8.57)	(-0.13)
CapIQ	1.107**	+0.543***	2.483***	+0.349**
	(2.40)	(3.94)	(4.94)	(2.40)
Stacked	**1.206*** **	**+0.306*** **	**0.973*** **	**+0.621*** **
	(9.21)	**(6.93)**	**(8.09)**	**(11.65)**
Observations	47,130		72,035	
Stacked Δ (Post-Pre)			**-0.232**	**+0.316*** **
			(-1.38)	**(4.91)**

Panel B. Heterogeneous post-ChatGPT effects by firm-level IBES accuracy-rank change

Dep. Var.:	Market-adjusted $CAR_{[-1,+1]}$		FF4 $CAR_{[-1,+1]}$	
	(1)	(2)	(3)	(4)
SUE^{IBES}	2.859***	2.845***	2.775***	2.776***
	(13.91)	(14.03)	(13.67)	(13.84)
$SUE^{IBES} \times Post$	1.029***	0.978***	1.136***	1.065***
	(3.56)	(3.32)	(4.11)	(3.76)
$SUE^{IBES} \times IBES\,Up$	0.693	0.628	0.803	0.737
	(1.24)	(1.16)	(1.46)	(1.39)
$SUE^{IBES} \times IBES\,Down$	1.397*	1.582**	1.913***	2.080***
	(1.81)	(2.00)	(2.61)	(2.74)
$SUE^{IBES} \times Post \times IBES\,Up$	**-1.248**	**-1.135**	**-1.291**	**-1.185**
	(-1.44)	**(-1.30)**	**(-1.40)**	**(-1.26)**
$SUE^{IBES} \times Post \times IBES\,Down$	**-2.278** **	**-2.344** **	**-2.758*** **	**-2.791** **
	(-2.27)	**(-2.15)**	**(-2.70)**	**(-2.51)**
$Post \times IBES\,Up$	0.353	0.232	0.437	0.297
	(1.13)	(0.68)	(1.37)	(0.85)
$Post \times IBES\,Down$	0.544	0.713	0.678*	0.859*
	(1.34)	(1.61)	(1.68)	(1.95)
Controls	No	Yes	No	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	23,833	23,833	23,833	23,833
Adjusted R^2	0.052	0.070	0.054	0.072

Notes. Both panels test whether the market relies more on more accurate forecasts. **Panel A**: within each firm-quarter the five FDPs are ranked by their ex-ante trailing four-quarter MAFE (Rank 1 = least accurate, Rank 5 = most accurate); pooling across firm-quarters, $CAR^{Mkt}_{[-1,+1]} = \alpha + \beta\, SUE^g + \delta\, Rank^g +$

$\gamma(SUE^g \times Rank^g)$, where γ is the ERC increment per one-rank accuracy improvement. Each provider row is its own regression; the *Stacked* row pools all five FDPs; the *Stacked* Δ row reports the $SUE \times Post$ and $SUE \times Rank \times Post$ coefficients. **Panel B**: firms are split by the change in IBES's rounded mean rank among the five FDPs between the pre- and post-ChatGPT windows ($IBES\,Up = 1$ if the rank improved by ≥ 2 positions; $IBES\,Down = 1$ if it worsened by ≥ 2); the bolded triple interactions are the key coefficients, and Panel B includes firm and year fixed effects. Pre-GPT (Post-GPT) covers announcements before (on or after) 2022-11-30. Standard errors are clustered by firm; continuous variables are winsorized at $(1, 99)$. t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$. Variable definitions are in Appendix A.

TABLE 4. Accuracy Persistence and the Predicted-Surprise Trading Strategy

Panel A. Persistence of FDP accuracy ranks (trailing-MAFE rank, quarter t to $t+1$)

	$\Pr(rank_{t+1} = j \mid rank_t = i)$, in %				
$rank_t$	$j = 1$ *(least)*	$j = 2$	$j = 3$	$j = 4$	$j = 5$ *(most)*
1 *(least)*	**76.1**	12.1	4.9	3.4	3.6
2	12.1	**54.1**	19.7	8.9	5.2
3	4.7	19.5	**46.7**	21.1	8.0
4	3.4	9.1	20.8	**48.3**	18.5
5 *(most)*	3.8	5.2	7.9	18.4	**64.7**

Spearman $\rho = 0.71$; overall same-rank rate = 58%; 109,915 quarter-to-quarter pairs.

Panel B. Trading strategy from predicted surprise

Dep. Var.:	Market-adj. $BHAR_{[-2,+1]}$		FF4 $BHAR_{[-2,+1]}$	
Benchmark:	Equal-Weighted	IBES	Equal-Weighted	IBES
	(1)	(2)	(3)	(4)
Standardized Pred. Surp.	**2.493*****	**3.433*****	**2.440*****	**3.357*****
	(4.10)	**(5.35)**	**(4.08)**	**(5.25)**
Long-Short Return (Q5-Q1)	**+0.73%*****	**+1.35%*****	**+0.63%*****	**+1.35%*****
Controls / Firm FE / Year FE	Yes	Yes	Yes	Yes
Observations (firm-q)	23,833	23,833	23,833	23,833
Adjusted R^2	0.016	0.017	0.018	0.019

Notes. This table motivates the accuracy-weighted consensus. **Panel A** reports the quarter-to-quarter persistence of forecast-provider accuracy ranks. Within each firm-quarter the five FDPs are ranked by their trailing four-quarter mean absolute forecast error (Rank 1 = least accurate, Rank 5 = most accurate, the same ex-ante measure used to build SUE^{Acc}); cell (i, j) is the share of provider-quarters at Rank i that move to Rank j the following quarter (rows sum to 100). **Panel B** horse-races two predicted-surprise signals as predictors of post-announcement abnormal returns: for benchmark $X \in \{EW5, IBES\}$ the predicted surprise is $(Consensus^{Acc}\text{-}Consensus^{X})/\text{price}$, min-max standardized to $[0,1]$; columns (1)–(2) use the market-adjusted $BHAR_{[-2,+1]}$ and (3)–(4) the FF4 four-factor $BHAR_{[-2,+1]}$. The Long-Short row reports the abnormal return from going long the top quintile and short the bottom quintile of predicted surprise. All Panel B columns include firm and year fixed effects, the full control set, and firm-clustered standard errors; continuous variables winsorized at $(1,99)$. t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$. Variable definitions are in Appendix A.

TABLE 5. Accuracy-Weighting: Does the Market Place Higher Weight on More Accurate FDPs?

Panel A. Horse race: SUE^{Acc} vs SUE^{EW5}.

Dep. Var.:	Mkt CAR$_{[-1,+1]}$			FF4 CAR$_{[-1,+1]}$		
	(1) Pre	(2) Post	(3) Pooled	(4) Pre	(5) Post	(6) Pooled
SUE^{Acc} (β_1)	1.376***	3.420***	1.307***	1.401***	3.519***	1.290***
	(3.45)	(9.38)	(3.60)	(3.63)	(10.01)	(3.67)
$SUE^{Acc} \times Post$ (β_3)			2.003***			2.135***
			(4.08)			(4.47)
SUE^{EW5} (β_2)	2.324***	1.767***	2.279***	2.379***	1.734***	2.284***
	(6.10)	(5.56)	(6.84)	(6.44)	(5.62)	(7.03)
SUE^{EW5} $\times Post$ (β_4)			-0.399			-0.449
			(-0.92)			(-1.06)
Premium (Pre-GPT)	-0.948		-0.972	-0.978		-0.994
	(-1.28)		(-1.47)	(-1.37)		(-1.55)
Premium (Post-GPT)		1.654***	1.429**		1.785***	1.590***
		(2.62)	(2.39)		(2.93)	(2.74)
Δ Premium (Post-Pre)			**2.401*** **			**2.584*** **
			(2.75)			**(3.03)**
Control Variables						
Log Market Equity	-5.348***	-4.964***	-2.848***	-5.585***	-5.125***	-3.238***
	(-5.35)	(-5.35)	(-6.00)	(-5.59)	(-5.74)	(-7.19)
Book-to-Market	0.413	-1.315	0.168	0.849	-1.245	0.309
	(0.36)	(-1.22)	(0.28)	(0.75)	(-1.16)	(0.53)
Leverage	-4.957**	-3.652*	-0.913	-4.561**	-4.067**	-1.312
	(-2.38)	(-1.70)	(-0.76)	(-2.17)	(-1.97)	(-1.13)
Prior 12-Month Return	-0.295	-0.868**	-0.235	-0.024	-1.038***	-0.129
	(-1.27)	(-2.52)	(-1.27)	(-0.11)	(-3.02)	(-0.71)
Log Price	-1.063	-1.047	-0.482	-0.671	-0.735	-0.029
	(-1.13)	(-1.12)	(-1.05)	(-0.71)	(-0.83)	(-0.07)
Return Volatility	-59.587***	27.445*	14.899	-53.652***	19.210	11.467
	(-3.08)	(1.66)	(1.34)	(-2.87)	(1.20)	(1.05)
Analyst Dispersion	-23.636	-40.842	-15.946	-11.993	-17.323	-2.625
	(-0.40)	(-0.71)	(-0.40)	(-0.20)	(-0.30)	(-0.07)
Busyness	-0.104	-0.245*	-0.172*	-0.132	-0.219*	-0.164*
	(-0.73)	(-1.83)	(-1.79)	(-0.93)	(-1.66)	(-1.72)
$\lvert \Delta IB \rvert$	16.729***	-4.012	2.450	17.171***	-4.609	2.027
	(2.98)	(-0.78)	(0.69)	(3.09)	(-0.91)	(0.58)
$\lvert \Delta Shares \rvert$	3.721	-1.964	1.505	3.251	-2.171	1.376
	(0.79)	(-0.51)	(0.52)	(0.72)	(-0.57)	(0.48)
$\lvert SpecItems \rvert$	-18.458	-8.991	-10.912	-21.400*	-7.112	-9.847
	(-1.50)	(-0.79)	(-1.33)	(-1.77)	(-0.64)	(-1.22)
SBC High	-0.149	0.174	0.106	-0.168	0.186	0.095
	(-0.57)	(0.77)	(0.65)	(-0.66)	(0.83)	(0.59)
Unexpected Item	-0.258	-0.036	-0.123	-0.222	-0.023	-0.103
	(-0.97)	(-0.17)	(-0.78)	(-0.89)	(-0.11)	(-0.68)
Stock Split	-2.679**	-1.110	-1.584*	-2.266*	-0.472	-0.977
	(-2.21)	(-0.92)	(-1.90)	(-1.78)	(-0.42)	(-1.18)

Fiscal Q4	0.600***	-0.045	0.276**	0.606***	0.230	0.417***
	(2.83)	(-0.24)	(2.07)	(2.89)	(1.23)	(3.19)
Institutional Ownership	3.234	-0.702**	-0.470	1.626	-0.543	-0.416
	(1.56)	(-2.02)	(-1.41)	(0.83)	(-1.57)	(-1.25)
Log # Analysts	0.508	-0.218	-0.248	0.648	-0.328	-0.371
	(0.79)	(-0.39)	(-0.65)	(1.04)	(-0.60)	(-0.99)
Controls	Yes	Yes	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes	Yes
Observations	9,405	14,402	23,833	9,405	14,402	23,833
Adjusted R^2	0.085	0.106	0.080	0.092	0.110	0.083

Panel B. Horse race: SUE^{Acc} vs SUE^{IBES}.

Dep. Var.:	Mkt CAR$_{[-1,+1]}$			FF4 CAR$_{[-1,+1]}$		
	(1)	(2)	(3)	(4)	(5)	(6)
	Pre	Post	Pooled	Pre	Post	Pooled
SUE^{Acc} (β_1)	1.428***	3.678***	1.430***	1.527***	3.796***	1.494***
	(4.00)	(10.49)	(4.31)	(4.36)	(10.70)	(4.48)
$SUE^{Acc} \times Post$ (β_3)			2.153***			2.204***
			(4.82)			(4.98)
SUE^{IBES} (β_2)	2.176***	1.460***	2.054***	2.149***	1.407***	1.974***
	(7.08)	(4.44)	(7.31)	(7.20)	(4.20)	(7.01)
SUE^{IBES} $\times Post$ (β_4)			-0.519			-0.484
			(-1.34)			(-1.28)
Premium (Pre-GPT)	-0.747		-0.623	-0.622		-0.480
	(-1.20)		(-1.09)	(-1.03)		(-0.83)
Premium (Post-GPT)		2.219***	2.049***		2.389***	2.208***
		(3.56)	(3.60)		(3.75)	(3.82)
Δ Premium (Post-Pre)			**2.672*** **			**2.688*** **
			(3.44)			**(3.50)**
Controls	Yes	Yes	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes	Yes
Observations	9,405	14,402	23,833	9,405	14,402	23,833
Adjusted R^2	0.087	0.106	0.080	0.094	0.110	0.083

Notes. This table reports OLS estimates of the horse race between the accuracy-weighted surprise SUE^{Acc} and a benchmark surprise, including their interactions with *Post*. SUE^{Acc} uses softmax weights based on each provider's within-firm-quarter z-scored four-quarter rolling MAFE. **Panel A** uses the equal-weighted benchmark SUE^{EW5}; **Panel B** the I/B/E/S benchmark and uses the same firm and event controls as Panel A, suppressed in the table for brevity. Columns (1)–(3) use the market-adjusted $CAR_{[-1,+1]}$; columns (4)–(6) the FF4 four-factor $CAR_{[-1,+1]}$. All columns include firm and year fixed effects and the full control set; standard errors are clustered by firm; continuous variables are winsorized at (1, 99). t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$. Variable definitions are in Appendix A.

TABLE 6. The Accuracy Premium Across GPT-Capability Regimes

Panel A. Benchmark = Equal-Weighted Five-FDP Consensus

Regime:	(1) R0 Pre-ChatGPT	(2) R1 text-only	(3) R2 browsing	(4) R3 native search
SUE^{Acc} (β_1)	1.376***	2.597***	3.753***	3.954***
	(3.45)	(3.38)	(6.19)	(5.69)
Benchmark SUE (β_2)	2.324***	2.419***	1.325***	1.404**
	(6.10)	(3.78)	(2.63)	(2.26)
Premium (β_1-β_2)	**-0.948**	**0.178**	**2.429****	**2.550****
	(-1.28)	**(0.13)**	**(2.35)**	**(2.07)**
Controls / Firm FE / Year FE	Yes	Yes	Yes	Yes
Observations	9,405	3,705	5,529	5,084
Adjusted R^2	0.085	0.158	0.151	0.140

Panel B. Benchmark = I/B/E/S Consensus

Regime:	(1) R0 Pre-ChatGPT	(2) R1 text-only	(3) R2 browsing	(4) R3 native search
SUE^{Acc} (β_1)	1.428***	1.918**	3.362***	5.079***
	(4.00)	(2.35)	(6.17)	(7.87)
Benchmark SUE (β_2)	2.176***	3.144***	1.842***	0.169
	(7.08)	(4.70)	(3.74)	(0.30)
Premium (β_1-β_2)	**-0.747**	**-1.226**	**1.519**	**4.909*****
	(-1.20)	**(-0.86)**	**(1.61)**	**(4.41)**
Controls / Firm FE / Year FE	Yes	Yes	Yes	Yes
Observations	9,405	3,705	5,529	5,084
Adjusted R^2	0.087	0.170	0.155	0.138

Notes. This table estimates the accuracy premium separately within each GPT-capability regime. For each regime the specification is $CAR^{Mkt}_{[-1,+1]} = \alpha + \beta_1 SUE^{Acc}_{i,q} + \beta_2 SUE^{X}_{i,q} + \text{controls} + \alpha_i + \tau_t + \varepsilon$, and the *Accuracy Premium* is the contrast β_1-β_2. Regimes are defined by announcement date: **R0** (before 2022-11-30, Pre-ChatGPT), **R1** (2022-11-30 to 2023-09-26, text-only), **R2** (2023-09-27 to 2024-10-30, browsing), and **R3** (from 2024-10-31, native search). The benchmark X is the equal-weighted five-FDP consensus (Panel A) or I/B/E/S (Panel B). All columns include the full control set, firm and year fixed effects, and firm-clustered standard errors; continuous variables winsorized at $(1,99)$. t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$. Variable definitions are in Appendix A.

TABLE 7. Cross-sectional Tests of the "Accuracy Premium"

Panel A. Concurrent Cross-FDP SUE Dispersion.

Benchmark:	Equal-Weighted		IBES	
	(1) Low	(2) High	(3) Low	(4) High
SUE^{Acc} (β_1)	4.138***	1.433***	1.569**	1.748***
	(6.37)	(4.23)	(2.43)	(5.57)
Benchmark SUE (β_2)	3.318***	2.024***	7.496***	1.560***
	(5.86)	(6.41)	(10.03)	(6.10)
$SUE^{Acc} \times Post^{\text{Scr}}$ (β_3)	2.748**	1.610***	1.439	1.746***
	(2.48)	(2.89)	(1.34)	(3.63)
Benchmark $SUE \times Post^{\text{Scr}}$ (β_4)	0.712	-0.602	2.130*	-0.681*
	(0.76)	(-1.20)	(1.86)	(-1.70)
Premium (Pre-Scrape)	0.821	-0.591	-5.927***	0.188
	(0.73)	(-0.96)	(-4.58)	(0.35)
Premium (Post-Scrape)	2.857*	1.620**	-6.619***	2.615***
	(1.71)	(1.98)	(-3.81)	(3.81)
Δ Premium (Post-Pre)	**2.037**	**2.211****	**-0.691**	**2.427*****
	(1.06)	**(2.21)**	**(-0.33)**	**(2.99)**
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	11,832	11,846	11,832	11,846
Adjusted R^2	0.111	0.075	0.127	0.075

Panel B. Special-Items Magnitude.

Benchmark:	Equal-Weighted		IBES	
	(1) Low	(2) High	(3) Low	(4) High
SUE^{Acc} (β_1)	2.345***	1.613***	1.940***	1.969***
	(4.36)	(4.49)	(4.34)	(5.48)
Benchmark SUE (β_2)	2.153***	2.078***	2.465***	1.615***
	(4.21)	(6.41)	(6.46)	(5.25)
$SUE^{Acc} \times Post^{\text{Scr}}$ (β_3)	0.247	2.035***	0.440	2.583***
	(0.31)	(3.39)	(0.58)	(4.69)
Benchmark $SUE \times Post^{\text{Scr}}$ (β_4)	0.799	-0.483	0.659	-1.055**
	(1.14)	(-0.92)	(0.98)	(-2.28)
Premium (Pre-Scrape)	0.192	-0.464	-0.525	0.354
	(0.19)	(-0.74)	(-0.68)	(0.58)
Premium (Post-Scrape)	-0.360	2.054**	-0.743	3.992***
	(-0.28)	(2.48)	(-0.65)	(5.36)
Δ Premium (Post-Pre)	**-0.552**	**2.518****	**-0.219**	**3.638*****
	(-0.39)	**(2.40)**	**(-0.16)**	**(3.92)**
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	11,823	11,828	11,823	11,828
Adjusted R^2	0.089	0.073	0.095	0.071

Notes. This table estimates the accuracy-weighted horse race against the equal-weighted (EW5) and

IBES benchmarks separately for two cross-sectional partitions. **Panel A** splits firm-quarters at the median concurrent cross-FDP SUE inter-quartile range. **Panel B** splits firm-quarters at the median of $|SpecItems|$. Within each panel, columns (1)–(2) use SUE^{EW5} and columns (3)–(4) use SUE^{IBES}. The Accuracy Premium is the contrast SUE^{Acc}-Benchmark SUE; the bolded Δ row reports the post-pre change. Here $Post^{\text{Scr}}$ (post-scraping) equals one for announcements on or after the onset of GPT browsing (2023-09-27), distinct from the post-ChatGPT *Post* used in the other tables. All panels include the full control vector, firm fixed effects, year fixed effects, and firm-clustered standard errors. The dependent variable is the market-adjusted $CAR_{[-1,+1]}$. Continuous variables are winsorized at the (1, 99) level. t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$. Variable definitions are in Appendix A.

TABLE 8. LLM Visibility and the Accuracy Premium by GPT Regimes

Panel A. Equal-Weighted Five-FDP Benchmark.

Regime:	(1) R0 Pre-ChatGPT	(2) R1 text-only	(3) R2 browsing	(4) R3 native search
SUE^{Acc} (β_1)	1.361**	3.503***	3.077***	3.055***
	(2.57)	(5.05)	(2.95)	(4.18)
SUE^{EW5} (β_2)	2.252***	2.185***	1.885***	2.002***
	(4.09)	(3.61)	(3.65)	(3.21)
$SUE^{Acc} \times High\_Vis$ (β_3)	0.020	-2.274	1.550*	1.815*
	(0.05)	(-1.49)	(1.79)	(1.93)
$SUE^{EW5} \times High\_Vis$ (β_4)	0.153	0.889	-1.133**	-1.269
	(0.28)	(0.79)	(-2.11)	(-1.51)
Accuracy premium = loading on SUE^{Acc} - loading on benchmark *SUE*:				
Premium (Low Vis)	-0.892	1.317	1.192	1.053
	(-0.90)	(1.12)	(0.89)	(0.83)
Premium (High Vis)	-1.025	-1.845	**3.875*****	**4.136*****
	(-1.57)	(-0.94)	**(4.60)**	**(4.16)**
Premium (High - Low)	-0.133	-3.163	**2.683****	**3.083***
	(-0.15)	(-1.22)	**(2.12)**	**(1.85)**
Controls / Firm FE / Year-qtr FE	Yes	Yes	Yes	Yes
Observations	9,399	3,705	5,529	5,084
Adjusted R^2	0.086	0.163	0.165	0.143

Panel B. I/B/E/S Benchmark.

Regime:	(1) R0 Pre-ChatGPT	(2) R1 text-only	(3) R2 browsing	(4) R3 native search
SUE^{Acc} (β_1)	1.607***	3.138***	2.737***	4.029***
	(4.07)	(3.15)	(6.05)	(8.81)
SUE^{IBES} (β_2)	1.859***	2.424***	2.403**	0.974**
	(3.60)	(2.96)	(2.28)	(2.13)
$SUE^{Acc} \times High\_Vis$ (β_3)	-0.382	-2.353*	1.667**	2.376***
	(-1.03)	(-1.74)	(2.44)	(3.68)
$SUE^{IBES} \times High\_Vis$ (β_4)	0.669	1.362	-1.380	-1.856***
	(0.97)	(1.45)	(-1.45)	(-3.06)
Accuracy premium = loading on SUE^{Acc} - loading on benchmark *SUE*:				
Premium (Low Vis)	-0.252	0.714	0.335	3.054***
	(-0.31)	(0.41)	(0.26)	(4.00)
Premium (High Vis)	-1.303**	-3.001**	**3.382*****	**7.286*****
	(-2.06)	(-2.12)	**(6.17)**	**(10.56)**
Premium (High - Low)	-1.052	-3.715*	**3.047****	**4.232*****
	(-1.07)	(-1.72)	**(2.05)**	**(3.93)**
Controls / Firm FE / Year-qtr FE	Yes	Yes	Yes	Yes
Observations	9,399	3,705	5,529	5,084
Adjusted R^2	0.088	0.175	0.169	0.142

Notes. This table estimates the accuracy-weighted horse race $CAR^{Mkt}_{[-1,+1]} = \alpha + \beta_1\, SUE^{Acc}_{i,q} + \beta_2\, SUE^{X}_{i,q} + \beta_3\, SUE^{Acc}_{i,q} \times High\_Vis_i + \beta_4\, SUE^{X}_{i,q} \times High\_Vis_i + \text{controls} + \text{FE} + \varepsilon$ *separately within each GPT-capability regime*, where $SUE^{Acc}_{i,q}$ is the accuracy-weighted multi-vendor surprise with softmax weights based on each provider's within-firm-quarter z-scored four-quarter rolling MAFE. Regimes track ChatGPT's information-retrieval capability by announcement date: **R0** (before 2022-11-30, Pre-ChatGPT), **R1** (2022-11-30 to 2023-09-26, text-only), **R2** (2023-09-27 to 2024-10-30, browsing / Bing), and **R3** (from 2024-10-31, native ChatGPT Search). **Panel A** uses the equal-weighted five-FDP benchmark SUE^{EW5}; **Panel B** the I/B/E/S benchmark. The *Accuracy Premium* is the contrast SUE^{Acc}-Benchmark SUE, reported for low- and high-visibility firms (Premium (Low Vis) = β_1-β_2; Premium (High Vis) = $(\beta_1 + \beta_3)$-$(\beta_2 + \beta_4)$); the bolded Premium (High - Low) row is the triple difference β_3-β_4. All specifications include the full control vector, firm fixed effects, and year-quarter fixed effects; the dependent variable is the market-adjusted $CAR_{[-1,+1]}$. Continuous variables are winsorized at the (1, 99) level; standard errors are two-way clustered by firm and calendar year. t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$. Variable definitions are in Appendix A.

TABLE 9. LLM Visibility Is Distinct from Classical Visibility

Panel A. Overlap of high-visibility classifications (%)

	LLM Visibility	Firm Size	Analyst Coverage	Inst Ownership
LLM Visibility	100	50.2	49.7	47.1
Firm Size	51.0	100	79.6	40.7
Analyst Coverage	50.6	79.8	100	44.5
Inst Ownership	47.9	40.7	44.5	100

Panel B. Accuracy-premium concentration: Equal-weighted benchmark

Visibility measure	Univariate ΔPremium		Horse-race
	High group	Low group	triple diff.
LLM Visibility	**3.263*****	**0.300**	**3.172***
	(2.59)	**(0.25)**	**(1.80)**
	[11,219]	[11,825]	
Firm Size	0.723	2.447**	0.111
	(0.55)	(2.18)	(0.05)
	[11,901]	[11,143]	
Analyst Coverage	0.109	3.076**	-3.161
	(0.08)	(2.57)	(-1.38)
	[11,861]	[11,183]	
Inst Ownership	3.224**	0.646	3.222*
	(2.51)	(0.58)	(1.86)
	[11,593]	[11,451]	
Controls / Firm FE / Year FE	Yes	Yes	Yes
Observations	23,044		23,044

Panel C. Accuracy-premium concentration: I/B/E/S benchmark

Visibility measure	Univariate ΔPremium		Horse-race
	High group	Low group	triple diff.
LLM Visibility	**5.743*****	**-0.125**	**5.841*****
	(5.67)	**(-0.12)**	**(4.03)**
	[11,219]	[11,825]	
Firm Size	2.191	2.462***	2.010
	(1.28)	(2.76)	(0.82)
	[11,901]	[11,143]	
Analyst Coverage	0.441	3.198***	-4.099*
	(0.29)	(3.59)	(-1.77)
	[11,861]	[11,183]	
Inst Ownership	2.564**	2.418**	0.625
	(2.46)	(2.08)	(0.41)
	[11,593]	[11,451]	
Controls / Firm FE / Year FE	Yes	Yes	Yes
Observations	23,044		23,044

Notes. Each measure (LLM visibility; firm size, log market equity; analyst coverage, log average analyst count; institutional ownership) is split at its cross-firm median; the visibility score is discrete, so firms tied at the median join the Low group, which is why the split is not exactly even. **Panel A**: cell (i, j) is the percentage of firms in measure i's High group that measure j also classifies as High (diagonal = 100). **Panels B and C**: ΔPremium is the post-pre change in the accuracy premium $\beta(SUE^{Acc})$-$\beta(SUE^{X})$ against benchmark $X =$ the equal-weighted five-FDP consensus (Panel B) or I/B/E/S (Panel C), estimated within each measure's High and Low group; the horse-race triple $(SUE^{Acc}\text{-}SUE^{X}) \times Post^{\text{Scr}} \times High_m$ is from one pooled regression including the triple for every measure jointly. Throughout, $Post^{\text{Scr}}$ is the post-scraping indicator (announcements on or after the onset of browsing, 2023-09-27), as in the cross-sectional tests (Table 7). All specifications include the full control set, firm and year fixed effects, and firm-clustered standard errors; continuous variables winsorized at $(1, 99)$. t-statistics in parentheses and group observation counts in brackets (each measure's High and Low groups partition the same 23,044 firm-quarters); $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$. Variable definitions are in Appendix A.

TABLE 10. The Predictive Value of the I/B/E/S Consensus: Measurement and Post-ChatGPT Pricing

	Single post-ChatGPT indicator		GPT-capability regimes	
Dep. Var.:	Mkt CAR (1)	FF4 CAR (2)	Mkt CAR (3)	FF4 CAR (4)
SUE^{IBES} (β_1)	2.237***	2.241***	2.237***	2.241***
	(8.71)	(9.02)	(8.74)	(9.05)
$SUE^{IBES} \times Post$ (β_2)	0.534	0.565*		
	(1.54)	(1.67)		
$SUE^{IBES} \times$ Regime 1 (text-only)			0.868*	1.016**
			(1.91)	(2.29)
$SUE^{IBES} \times$ Regime 2 (browsing)			0.785	0.591
			(1.62)	(1.18)
$SUE^{IBES} \times$ Regime 3 (native search)			0.162	0.268
			(0.33)	(0.57)
$SUE^{IBES} \times Predictive\ Value^{IBES}$ (β_3)	2.172***	2.119***	2.174***	2.125***
	(4.03)	(4.11)	(4.05)	(4.13)
$SUE^{IBES} \times Predictive\ Value^{IBES} \times Post$ (β_4)	**2.049*** **	**2.098*** **		
	(2.68)	**(2.82)**		
$SUE^{IBES} \times Predictive\ Value^{IBES} \times$ Regime 1			0.919	0.941
			(0.93)	(0.97)
$SUE^{IBES} \times Predictive\ Value^{IBES} \times$ Regime 2			**1.435**	**1.777***
			(1.45)	**(1.80)**
$SUE^{IBES} \times Predictive\ Value^{IBES} \times$ Regime 3			**3.540*** **	**3.365*** **
			(3.08)	**(3.02)**
Lower-order terms / Controls	Yes	Yes	Yes	Yes
Firm FE / Year FE	Yes	Yes	Yes	Yes
Observations	22,356	22,356	22,356	22,356
Adjusted R^2	0.082	0.083	0.082	0.083

Notes. This table measures, for each firm, the predictive value of the I/B/E/S consensus — how closely it has anticipated the firm's future street numbers — and tests whether the post-ChatGPT strengthening of announcement-window reactions concentrates in firms whose consensus predicts well. $Predictive\ Value^{IBES}_{i,q} = 1\text{-}error_{i,q}$, where $error_{i,q}$ is the consensus's trailing relative anticipation error: the mean own-basis distance $|actual_{i,q+4}\text{-}consensus_{i,q}|$ over up to the twenty most recent source quarters whose year-ahead actuals are public at the announcement (at least twelve required; a stock-split guard applies), scaled by the window mean $|actual|$ so the measure is comparable across firms, and winsorized at $(1,99)$. The measure is continuous (mean 0.56, standard deviation 0.30), higher values meaning the consensus has anticipated the firm's year-ahead street numbers more closely, and is implementable in real time. Coverage is 22,364 of the 23,833 firm-quarters (8,730 pre-ChatGPT), with no industry restriction; the uncovered remainder are recent listings with fewer than twelve scoreable history quarters. The same construction at the current-quarter horizon yields the consensus's trailing *accuracy*; across firms the two are highly correlated (Pearson 0.72, Spearman 0.78; 0.82 corrected for split-half measurement error), so the accuracy dimension used throughout the paper carries most of this informativeness content and is preferred there for econometric convenience — the current-quarter error requires no four-quarter maturation and is observable in real time. SUE^{IBES} is the I/B/E/S consensus surprise;

all columns include the Table 5 controls (including dispersion, volatility, and the unexpected-item flag), with firm and year fixed effects and firm-clustered standard errors. The pre-period gradient β_3 reflects the familiar association between predictability and the ERC; the identified claim is the bolded post-ChatGPT increment β_4, comparable in size to the pre-period gradient — the gradient roughly doubles after ChatGPT (about 0.6 per standard deviation of the measure). The level terms are indistinguishable from zero: high-predictive-value firms earn no different announcement return per se, and the effect operates entirely through the *SUE* loading. Columns (3)–(4) replace the single *Post* indicator with the three GPT-capability regimes of Appendix A (text-only from November 30, 2022; browsing from September 27, 2023; native search from October 31, 2024): the increment rises monotonically across regimes and reaches significance under native search — the same capability profile as the accuracy premium of Table 6, and consistent with the retrieval-intensive nature of a forecast track record, which browsing and native search made accessible. The native-search regime spans only the final five sample quarters, so the regime profile is descriptive and the pooled increment in columns (1)–(2) is the formal test. The increment is not announcement composition in disguise: jointly adding interactions of a large-unexpected-item flag (seven Compustat non-recurring items mapped per Bochkay et al., 2022, top-decile magnitude screen) leaves β_4 unchanged (2.04, $t = 2.67$, market-adjusted; 2.09, $t = 2.83$, FF4), and the two measures are essentially uncorrelated (-0.02). Nor is it selection toward I/B/E/S specifically: in an untabulated decomposition of the surprise into the component common to all providers and the I/B/E/S-specific tilt, both the gradient (2.3, $t = 3.7$) and its post-ChatGPT increment (1.9, $t = 2.0$) price the common component, so the evidence concerns the consensus signal of well-predicted firms rather than provider choice. A within-firm-quarter *four-provider rank* version of the same score — each provider's own-basis error against its own year-ahead actual — gives a stacked post-period rank-ERC gradient of 0.320 ($t = 4.14$) and a post-pre shift of 0.242 ($t = 2.56$), concentrated in Capital IQ provider-events. t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 11. Additional Tests: Outage Shock and Information Acquisition

Panel A. ChatGPT outage: accuracy premium on the day after a major outage

Dep. Var.: day-0 AR_0 (market-adj., %)	Benchmark: Equal-Weighted (1)	Benchmark: IBES (2)
SUE^{Acc} (β_1)	1.631***	1.599***
	(6.93)	(7.78)
$SUE^{Acc} \times Outage$ (β_3)	1.178	-1.685*
	(0.80)	(-1.85)
Benchmark SUE (β_2)	0.626***	0.654***
	(3.13)	(4.10)
Benchmark $SUE \times Outage$ (β_4)	-2.608*	1.488**
	(-1.79)	(2.27)
Premium (non-outage)	1.005**	0.945***
	(2.48)	(2.88)
Premium (outage)	4.792*	-2.228*
	(1.77)	(-1.65)
Δ Premium (Outage-Non)	**3.787**	**-3.173****
	(1.37)	**(-2.29)**
Controls / Firm FE / Year FE	Yes	Yes
Observations	14,402	14,402
Adjusted R^2	0.051	0.052

Panel B. Information acquisition: platform brand-search share by regime × vendor group

Dep. Var.: $S_{\text{brand}}/S_{\text{all platforms}}$	(1) Pre/Post-GPT	(2) Pre + 3-Regime
Post-GPT × Free Distributor	0.021***	
	(15.44)	
Post-GPT × Free FDP	-0.075***	
	(-63.80)	
Post-GPT × Paywalled FDP	-0.071***	
	(-64.91)	
Regime 1 × Free Distributor		0.021***
		(9.35)
Regime 1 × Free FDP		-0.075***
		(-35.57)
Regime 1 × Paywalled FDP		-0.070***
		(-32.85)
Regime 2 × Free Distributor		0.020***
		(10.43)
Regime 2 × Free FDP		-0.073***
		(-33.39)
Regime 2 × Paywalled FDP		-0.069***
		(-38.23)

Regime 3 × Free Distributor		0.022***
		(10.96)
Regime 3 × Free FDP		-0.075***
		(-42.82)
Regime 3 × Paywalled FDP		-0.073***
		(-40.93)
Day-level controls (searches, # earnings, trading day)	Yes	Yes
Calendar-month FE	Yes	Yes
Observations (platform-day)	37,247	37,247
Adjusted R^2	0.046	0.046

Notes. **Panel A** tests whether the accuracy premium collapses on earnings announcements that follow a major ChatGPT outage. The sample is post-ChatGPT firm-quarters; $Outage_{i,q} = 1$ if a ChatGPT incident of at least 240 minutes (from the OpenAI incident log) occurred on the calendar day immediately preceding the announcement (84 outage firm-quarters). The dependent variable is the day-0 market-adjusted abnormal return; both columns include firm and year fixed effects, the full control set, and firm-clustered standard errors. The Accuracy Premium is the loading on SUE^{Acc} minus the loading on the benchmark SUE; the bolded row reports its outage-non-outage difference. **Panel B** reports platform-day evidence on investors' information acquisition: the dependent variable is a platform's share of total platform-specific Google search, regressed on regime × vendor-group interactions. The 17 platforms are grouped into Free Distributors (Barchart, CNBC, Finviz, Forbes, Investing.com, MarketWatch, Motley Fool, Reuters, Seeking Alpha, TradingView, WSJ, Yahoo Finance), the Free FDP (Zacks), and Paywalled FDPs (Bloomberg, Capital IQ, FactSet, IBES); Free Distributor in the pre-ChatGPT regime is the omitted baseline. Column (1) uses a single Post-GPT indicator; column (2) splits it into three GenAI regimes. Standard errors are HC1. t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$. Variable definitions are in Appendix A.

Appendix A: Variable Definitions

Variable	Definition	Source
Panel A. Across-FDP firm-quarter outcome and surprise variables (indexed by firm *i*, quarter *q*).		
$CAR^{Mkt}_{i,q,[-1,+1]}$	Three-day market-adjusted cumulative abnormal return around the earnings-announcement date, $\sum_{t=-1}^{+1}(r_{i,t}\text{-vwretd}_t)$, in percent.	CRSP
$CAR^{FFC}_{i,q,[-1,+1]}$	Three-day FFC four-factor (FF3+UMD) abnormal CAR over $[-1,+1]$ with factor loadings estimated on the $[-260,-11]$ window.	CRSP
$BHAR^{Mkt}_{i,q,[-2,+1]}$	Four-day market-adjusted buy-and-hold abnormal return over $[-2,+1]$.	CRSP
$BHAR^{FFC}_{i,q,[-2,+1]}$	Four-day FFC four-factor buy-and-hold abnormal return over $[-2,+1]$.	CRSP
$SUE^{g}_{i,q}$	Per-FDP standardized unexpected earnings for provider $g \in \{\text{FactSet, Bloomberg, Zacks, CapIQ, IBES}\}$: $(\text{actual EPS}^{g}_{i,q}\text{-consensus EPS}^{g}_{i,q})/\text{price}_{i,q}$.	FactSet; Bloomberg; Zacks; S&P CapIQ; I/B/E/S
$SUE^{Acc}_{i,q}$	Accuracy-weighted multi-vendor SUE, $\sum_g w^{g}_{i,q}\, SUE^{g}_{i,q}$, where $w^{g}_{i,q}$ are softmax weights based on each FDP's z-scored four-quarter rolling MAFE.	Calculated
$SUE^{EW5}_{i,q}$	Equal-weighted multi-vendor SUE, $(1/5)\sum_g SUE^{g}_{i,q}$.	Calculated
Standardized Pred. Surp.$^{X}_{i,q}$	Min-max standardized predicted surprise $(Consensus^{Acc}_{i,q}\text{-}Consensus^{X}_{i,q})/\text{price}_{i,q}$ for benchmark $X \in \{EW5, IBES\}$.	Calculated
Panel B. Firm and event controls (indexed by firm *i*, quarter *q*).		
Log Market Equity$_{i,q-1}$	Natural log of firm *i*'s market capitalization at the end of quarter q-1.	CRSP
Book-to-Market$_{i,q-1}$	Book value of equity divided by market value of equity at quarter-end q-1.	Compustat; CRSP
Leverage$_{i,q-1}$	Long-term debt plus debt in current liabilities, scaled by total assets, at q-1.	Compustat
Prior 12-Month Return$_{i,q}$	Buy-and-hold raw return over the twelve months ending one month before firm *i*'s quarter-q announcement.	CRSP
Log Price$_{i,q-1}$	Natural log of share price at quarter-end q-1.	CRSP
Return Volatility$_{i,q}$	Standard deviation of daily returns over the $[-260,-11]$ window before the announcement.	CRSP
Analyst Dispersion$_{i,q}$	Standard deviation of analyst EPS forecasts divided by absolute consensus.	I/B/E/S
Busyness$_{i,q}$	Natural log of the number of S&P 1500 firms announcing earnings on firm *i*'s announcement day.	Compustat

(continued on next page)

Variable	Definition	Source
$\lvert\Delta IB\rvert_{i,q}$	Absolute year-over-year change in income before extraordinary items, scaled by lagged total assets.	Compustat
$\lvert\Delta Shares\rvert_{i,q}$	Absolute year-over-year change in shares outstanding, scaled by lagged shares.	Compustat
$\lvert SpecItems\rvert_{i,q}$	Absolute value of special items, scaled by lagged total assets.	Compustat
$SBC\ High_{i,q}$	Indicator equal to one if stock-based compensation expense exceeds the sample median.	Compustat
$Unexpected\ Item_{i,q}$	Indicator equal to one if the quarter contains a non-recurring item flagged by Compustat.	Compustat
$Stock\ Split_{i,q}$	Indicator equal to one if a stock split was effective within [-30, +30] trading days of the announcement.	CRSP
$Fiscal\ Q4_{i,q}$	Indicator equal to one for fiscal fourth-quarter announcements.	Compustat
$Institutional\ Ownership_{i,q\text{-}1}$	Fraction of shares held by institutional investors at q-1.	Thomson Reuters 13F
$\#\ Analysts_{i,q}$	Number of unique analysts in the IBES detail file forecasting EPS for firm-quarter (i,q).	I/B/E/S
Panel C. Cross-sectional partitioning and identification variables.		
$LLM\ Visibility_i$	EPS-extraction indicator $\times \log(1+\text{unique cited domains})$, averaged across ChatGPT and Perplexity for direct queries about firm i's quarterly EPS.	DataForSEO
$Cross\text{-}FDP\ SUE\ Dispersion_{i,q}$	Inter-quartile range of $SUE^{g}_{i,q}$ across the five FDPs in firm-quarter (i,q).	Calculated
$\lvert SpecItems\rvert_{i,q}\,(\%)$	Absolute special items as a percentage of absolute income before extraordinary items, $\lvert SpecItems\rvert_{i,q}/\lvert IB\rvert_{i,q} \times 100$.	Compustat
$IBES\,Up_i$	Indicator equal to one if IBES's rounded average per-firm-quarter accuracy rank among the five FDPs rose by at least two positions from pre- to post-ChatGPT.	Calculated
$IBES\,Down_i$	Indicator equal to one if IBES's rank fell by at least two positions.	Calculated
$Outage_{i,q}$	Indicator equal to one if a ChatGPT incident of at least 240 minutes occurred on the calendar day immediately preceding firm i's quarter-q announcement date.	OpenAI
$Large\ Unexpected\ Item$ $(LUI_{i,q})$	Indicator equal to one if firm-quarter (i,q) contains a large unexpected item: any of seven Compustat non-recurring items mapped per Bochkay et al. (2022), with restructuring, acquisition, and special items requiring top-decile magnitude. Refines *Unexpected Item* (Panel B) with a magnitude screen.	Compustat
$INFO_{i,q}$	Indicator equal to one if firm-quarter (i,q) contains no large unexpected item, 1-$LUI_{i,q}$.	Calculated
Panel D. Information-acquisition platform-day variables (indexed by platform p, day t).		

(continued on next page)

Variable	Definition	Source
Brand Share$_{p,t}$	Platform p's share of total platform-specific Google search volume on day t, $S_{p,t}/\sum_{p'} S_{p',t}$, summed across the 17 sample platforms.	Google Trends; DataForSEO
Log Total Platform Searches$_t$	Natural log of total Google search volume across the 17 sample platforms on day t.	Google Trends; DataForSEO
Log # Earnings Today$_t$	Natural log of one plus the number of S&P 1500 firms announcing earnings on day t.	Compustat
Trading Day$_t$	Indicator equal to one on NYSE trading days; zero otherwise.	CRSP
Free Distributor$_p$	Indicator equal to one for platforms in the Free Distributors group (Barchart, CNBC, Finviz, Forbes, Investing.com, MarketWatch, Motley Fool, Reuters, Seeking Alpha, TradingView, WSJ, Yahoo Finance).	Classification
Free FDP$_p$	Indicator equal to one for Zacks.	Classification
Paywalled FDP$_p$	Indicator equal to one for Bloomberg, Capital IQ, FactSet, and IBES.	Classification
Post-ChatGPT$_t$	Indicator equal to one for days on or after 2022-11-30 (ChatGPT public release).	OpenAI
Regime k_t ($k = 1,2,3$)	Mutually exclusive indicators for the three post-GenAI regimes (Regime 1: 2022-11-30 to 2023-09-26; Regime 2: 2023-09-27 to 2024-08-31; Regime 3: 2024-09-01 onward).	OpenAI

Appendix B: Worked Examples of GenAI Aggregation of Analyst Consensus

This appendix shows the prompt and ChatGPT (gpt-4o) responses for two illustrative large-cap firms, IBM and Microsoft. Responses were obtained through the DataForSEO LLM Responses API and represent the unedited output of the model. Each firm subsection ends with a summary table mapping the cited aggregators to their upstream forecast data provider (FDP).

B.0 Prompt

What is the analyst consensus EPS forecast for [TICKER] for the next quarter? Please report the figures compiled by multiple third-party aggregators and cite each source.

B.1 International Business Machines Corporation (IBM)

ChatGPT (gpt-4o) response:

Here's the latest analyst consensus EPS forecast for IBM's next quarter, as compiled from several third-party aggregators.

- **MarketBeat (via Fiscal.ai)**: Q2 2026 average EPS estimate is **$2.96**, based on 2 estimates (low $2.92, high $2.99) (**marketbeat.com**).
- **TipRanks**: next quarter's earnings estimate is **$3.02**, with a range from $2.92 to $3.09 (**tipranks.com**).
- **Barchart**: next quarter average earnings estimate is **$2.96**, based on 6 estimates (low $2.73, high $3.50) (**barchart.com**).
- **Yahoo Finance**: Jun 2026 average EPS estimate is **$3.04**, based on 14 analysts (low $2.80, high $3.21) (**finance.yahoo.com**).

Summary table.

Source	Forecast	Upstream FDP
MarketBeat	$2.96	Independent aggregator
TipRanks	$3.02	Independent aggregator
Barchart	$2.96	Independent aggregator
Yahoo Finance	$3.04	S&P Capital IQ

B.2 Microsoft Corporation (MSFT)

ChatGPT (gpt-4o) response:

Here are the analyst consensus EPS forecasts for Microsoft (MSFT) for the **next quarter**, as compiled by multiple third-party aggregators. Each figure is accompanied by its source.

- **MarketBeat**: for Q4 FY2026, the consensus EPS estimate is **$4.04**, based on 5 analyst forecasts (**marketbeat.com**).
- **Zacks Investment Research**: consensus estimate for the next quarter (June 2026) is **$4.13**, derived from 13 analyst estimates (**zacks.com**).
- **TipRanks**: next EPS estimate is **$4.24** (**tipranks.com**).

- **ChartMill**: consensus EPS estimate of **$4.33** for the next earnings (Q4 FY2026), based on 39 analysts (**chartmill.com**).
- **Seeking Alpha**: an earnings preview lists the consensus EPS at $4.06 for Q3 FY2026 (already reported), so not applicable for the upcoming quarter (**seekingalpha.com**).
- **Yahoo Finance**: for the period "Next qtr. (Jun 2026)", average EPS estimate is **$4.26**, based on 29 analysts (low $3.74, high $5.54) (**finance.yahoo.com**).
- **Investing.com**: I was unable to locate a clear next-quarter EPS consensus figure (**investing.com**).

Summary table.

Source	**Forecast**	**Upstream FDP**
MarketBeat	$4.04	Independent aggregator
Zacks	$4.13	Zacks
TipRanks	$4.24	Independent aggregator
Yahoo Finance	$4.26	S&P Capital IQ
ChartMill	$4.33	Independent aggregator

Appendix C: Numerical Illustration of the Accuracy-Weighted SUE

This appendix illustrates how the accuracy-weighted SUE is constructed for a single firm-quarter. Suppose firm i in quarter q has the following backward-looking accuracy histories ($acc^g_{i,q}$ = -$MAFE^g_{i,q}$ over the prior four quarters, in price-scaled units) across the five FDPs, alongside their current-quarter standardized surprises $SUE^g_{i,q}$ (in percent of price).

Panel A. Weight construction from accuracy history.

Provider g	$acc^g_{i,q}$	$acc\_z^g_{i,q}$	$\exp(acc\_z^g_{i,q})$	$w^g_{i,q}$
IBES	-0.0010	$+1.5$	4.48	0.57
FactSet	-0.0020	$+0.5$	1.65	0.21
Capital IQ	-0.0030	0.0	1.00	0.13
Zacks	-0.0040	-0.8	0.45	0.06
Bloomberg	-0.0050	-1.2	0.30	0.04
Sum			7.88	1.00

I/B/E/S has the smallest absolute forecast error (largest acc^g) and therefore the highest within-row z-score, which the softmax transform maps to the largest weight (0.57). Bloomberg has the largest absolute forecast error and receives the smallest weight (0.04). The weights sum to one and reduce to the uniform 0.20 vector when the five accuracy scores are equal.

Panel B. Composite construction.

Provider g	$SUE^g_{i,q}$	$w^g_{i,q} \times SUE^g_{i,q}$	$0.20 \times SUE^g_{i,q}$
IBES	$+1.20$	$+0.68$	$+0.24$
FactSet	$+0.80$	$+0.17$	$+0.16$
Capital IQ	$+1.00$	$+0.13$	$+0.20$
Zacks	$+1.50$	$+0.09$	$+0.30$
Bloomberg	-0.50	-0.02	-0.10
Total		$SUE^{Acc}_{i,q} = +1.05$	$SUE^{EW5}_{i,q} = +0.80$

The accuracy-weighted composite $SUE^{Acc}_{i,q} = +1.05$ exceeds the equal-weighted composite $SUE^{EW5}_{i,q} = +0.80$ because the two most-accurate providers in this row (I/B/E/S and FactSet) reported positive surprises, while the least-accurate provider (Bloomberg) was the lone outlier on the downside. The accuracy tilt $SUE^{Acc}_{i,q}$-$SUE^{EW5}_{i,q} = +0.25$ captures the accuracy-driven correction. In the horse-race regression of Table 6, this tilt is the source of identification between the two composites: when the tilt is informative about announcement-window returns beyond the equal-weighted aggregate, the market is loading on provider-specific accuracy rather than treating all five FDPs symmetrically.

ONLINE APPENDIX

"Synthesizing the Consensus: Generative AI, LLM Visibility, and the Integration of Multi-Provider Consensus Forecasts"

This Online Appendix validates accuracy as the weighting dimension of the multi-vendor synthesis. Table OA1 shows that the accuracy tilt is a weakly dominant design choice: the accuracy-weighted composite is strictly closer to realized earnings than either the equal-weighted or the I/B/E/S-only heuristic in both regimes (Panel A), while the five providers' consensus slopes for future fundamentals are statistically indistinguishable (Panel B) — so accuracy weighting strictly gains precision at no cost in fundamental informativeness.

TABLE OA1. Validating Accuracy as the Weighting Dimension: Strict Precision Gains, No Informativeness Cost

Panel A. Expectational precision: paired differences in $|SUE|$ against the accuracy-weighted composite.

	Pooled	Pre-ChatGPT	Post-ChatGPT
$\|SUE^{EW5}\|$-$\|SUE^{Acc}\|$	0.033^{***}	0.033^{***}	0.033^{***}
	(13.30)	(10.03)	(10.94)
Share with $\|SUE^{Acc}\|$ strictly smaller	52.3%	53.6%	51.5%
$\|SUE^{IBES}\|$-$\|SUE^{Acc}\|$	0.045^{***}	0.049^{***}	0.042^{***}
	(14.44)	(12.10)	(11.99)
Share with $\|SUE^{Acc}\|$ strictly smaller	62.0%	62.7%	61.6%
Observations	23,833	9,426	14,407

Panel B. Fundamental informativeness: consensus slopes for four-quarter-ahead outcomes.

	FactSet	Zacks	CapIQ	IBES
Future cash flows (θ_g)	1.016^{**}	0.807^{**}	0.822^{**}	0.845^{**}
	(2.11)	(2.08)	(1.98)	(2.00)
Wald test of slope equality		$F = 0.80$, $p = 0.493$		
Future street earnings (θ_g)	1.260^{**}	1.098^{***}	1.198^{**}	1.282^{**}
	(2.46)	(2.66)	(2.50)	(2.54)
Wald test of slope equality		$F = 1.75$, $p = 0.155$		

Notes. This table validates the choice of forecast accuracy as the weighting dimension of the manuscript's composite. **Panel A** reports Diebold–Mariano-style paired tests: the mean difference in absolute standardized surprise between each heuristic composite and the accuracy-weighted composite, estimated on the full firm-quarter panel with firm-clustered standard errors (the formal counterpart of Figure 2). The accuracy-weighted composite is strictly closer to the realized actual in both regimes against both benchmarks. **Panel B** asks what accuracy-based weighting could cost along the provider-level informativeness dimension of Bochkay et al. (2022): each provider's consensus slope for four-quarter-ahead fundamentals, estimated on a stacked provider panel (2015–2025, the four providers with native point-in-time history; provider-specific intercepts, lagged-cash-flow loadings, and industry effects; firm-clustered standard errors). The slopes are statistically indistinguishable across providers for both outcomes, so tilting weights toward accurate providers sacrifices no fundamental informativeness. Bloomberg, excluded above because its native point-in-time history begins in 2020Q2, is indistinguishable from FactSet, CapIQ, and I/B/E/S on the common 2020Q2-onward window (cash-flow slopes 2.16 versus 2.17–2.18); the five-provider equality test on that shorter window rejects only marginally ($p = 0.09$ cash flows, $p = 0.03$ earnings), driven by a lower Zacks slope consistent with the attenuation from its thin, penny-rounded consensus. Together with the persistence of accuracy rankings (Table 4, Spearman $\rho = 0.71$) and their firm-level variation, the panels imply that accuracy weighting weakly dominates: it strictly improves expectational precision, costs nothing in informativeness, and is the only quality dimension observable and exploitable in real time. t-statistics in parentheses; $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.