Synthesizing the Consensus: Generative AI, LLM Visibility, and the Pricing of Multi-Provider Earnings Forecasts*

Leonard Yang Liu†

October 2026

Abstract

This paper studies how generative AI (GenAI) changes capital markets' use of consensus earnings forecasts from five forecast data providers (FDPs) that construct distinct consensus numbers for the same firm-quarter. In S&P 1500 firm-quarters over 2021–2025 covered by all five FDPs, I find evidence consistent with GenAI facilitating the market's selective reliance on individual FDPs and an accuracy-weighted synthesis of their forecasts. After ChatGPT, announcement reactions align more closely with provider accuracy ranks and consensus predictive value. Returns per percentage point of surprise rise 2.40 percentage points more under the accuracy-weighted than the equal-weighted consensus and 2.67 more than under I/B/E/S. Across the GPT-capability regimes, this accuracy premium is undetectable before browsing, positive under browsing, and significant against both benchmarks only under native search. The premium also concentrates where special items are large and, against I/B/E/S, where providers disagree. To capture cross-firm variation in retrieval, I construct *LLM Visibility*, which measures the extent to which a firm's consensus forecasts are visible to LLMs. The premium concentrates in high-visibility firms only after browsing arrives. The evidence is consistent with retrieval-enabled GenAI lowering the cost of synthesizing competing consensus forecasts and changing the earnings benchmark that prices reflect.

Keywords: Generative AI, Analyst, Forecast Data Provider, Consensus Forecast, Information Processing, Earnings Response Coefficient.

JEL classification: G11, G14, G41, M41.

*I am deeply grateful to my dissertation committee for their invaluable guidance and support: Musa Subasi (co-chair), Michael Kimbrough (co-chair), Rebecca Hann, Nick Seybert, Jingyi Qian, and Erkut Y. Ozbay. I also thank Matt Ege, Antonis Kartapanis, Kimberlyn Munevar, Brady Twedt, Gus De Franco, Stephanie Cheng, Yuan Zhang, Sophia Hamm, Jasmijn Bol, and seminar and workshop participants at the University of Maryland, Texas A&M University, and Tulane University for helpful comments. I gratefully acknowledge the Texas A&M University AI Dissertation Proposal Award in Accounting. All errors are my own.

†Department of Accounting and Information Assurance, Robert H. Smith School of Business, University of Maryland. Email: yliu337@umd.edu.

1 Introduction

Sell-side analyst consensus forecasts are the market's primary benchmark for evaluating realized earnings (e.g., Bradshaw et al., 2017; Kothari, 2001). The earnings surprise relative to this benchmark helps explain announcement-window stock price reactions and post-announcement drift (e.g., Bartov et al., 2002; Bernard and Thomas, 1989; Livnat and Mendenhall, 2006). The consensus, however, is not a single number. Multiple *forecast data providers* (FDPs), including I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks, construct economically distinct consensus forecasts for the same firm-quarter.[1] Their proprietary aggregation methods differ in screening rules, treatment of nonrecurring items, and analyst-coverage thresholds (Bochkay et al., 2022; Larocque et al., 2026), producing consensuses whose relative accuracy varies across firms and over time. Prior evidence shows that market expectations tend to be closer to the consensuses that turn out to be more accurate, indicating that accuracy is an important dimension in the market's selection among available consensuses.[2]

Yet predicting which consensus will be most accurate entails substantial processing costs.[3] These costs may thus limit the extent to which the market selects consensuses based on accuracy. Consistent with this notion, announcement returns align more closely with the I/B/E/S surprise than with most other providers' surprises (Larocque et al., 2026), suggesting that I/B/E/S is the consensus source on which the market mainly relies. Generative AI (GenAI) can reduce the processing costs that lead the market to default to one FDP. Therefore, this paper examines whether the introduction of GenAI changes the market's reliance on individual providers and the combination of their forecasts reflected in prices.

[1] In a sample of 94,030 firm-quarters covered by all five FDPs over 2002–2020, Larocque et al. (2026) find that the resulting earnings surprises differ beyond rounding in 58 percent of firm-quarters. In 9 percent, at least one FDP signals a miss while another signals a beat.

[2] Market reactions align more closely with higher-quality providers when providers disagree about whether earnings beat or miss expectations (Larocque et al., 2026).

[3] Investors must recognize that the forecasts differ, retrieve the providers' past forecasts and actual earnings, and assess each provider's accuracy for the firm. They must then translate that assessment into weights and repeat the process for each firm-quarter they follow. These tasks span the awareness, acquisition, and integration costs in the disclosure-processing-cost framework of Blankespoor et al. (2020).

Following ChatGPT's release in late 2022, industry surveys document widespread investor adoption (AIMA, 2025; CFA Institute, 2024).[4] A retrieval-enabled GenAI query can bring forecasts and earnings histories from multiple websites into one response and support their combination according to accuracy. By lowering processing costs, GenAI may increase reliance on accurate providers and shift the market's earnings benchmark toward accuracy-weighted synthesis. In the framework I develop, the shift is larger when synthesis is more accessible or offers greater benefits.

These predictions are not obvious *ex ante*, for two reasons. First, although GenAI can retrieve the providers' forecasts and compute combination weights, the market may not use it for this task if the cost of synthesis exceeds its benefit. An investor must still assess each provider's record for the firm and verify what the model retrieves. Second, limited-attention research shows that competing information can weaken investors' responses to relevant news (Hirshleifer et al., 2009). Easier access may therefore increase the information investors must assess without improving their expectations. Beyond these reasons, the size of the cost reduction is unknown. It may grow as the models' retrieval capability expands, and it may differ across firms according to what a browsing model can retrieve. The tests compare regimes of retrieval capability over time and compare firms by what a browsing model retrieves. Whether GenAI strengthens the market's use of relative accuracy remains an empirical question.

Using S&P 1500 firm-quarters from 2021 to 2025, I find evidence consistent with GenAI facilitating the market's selective reliance on individual FDPs and its accuracy-weighted synthesis of their forecasts. First, after ChatGPT's release, the market relies more on the providers with better recent accuracy for the firm. The increment in the earnings response coefficient associated with a one-rank improvement in accuracy roughly doubles. Relative to the firms in which I/B/E/S holds its rank, the market's response to the I/B/E/S surprise weakens where I/B/E/S loses accuracy rank and does not change significantly where it gains rank. This asymmetry is consistent with a pre-ChatGPT market that defaults to I/B/E/S. Second, announcement returns align more closely with a surprise that weights providers by recent accuracy. In horse-race regressions, I refer to the difference between

[4]Recent research connects GenAI use to financial analysis and capital-market outcomes (e.g., Bertomeu et al., 2025; Blankespoor et al., 2026; Chang et al., 2026; Cheng et al., 2025; Liu and Subasi, 2026; Xue et al., 2025).

its return loading and that of a benchmark surprise as the *accuracy premium*. After ChatGPT, this premium widens by 2.40 percentage points against the equal-weighted benchmark and by 2.67 percentage points against the I/B/E/S benchmark. Third, the premium is undetectable before browsing, positive under browsing, and significant against both benchmarks only under native search. This timing is consistent with a role for retrieval. The premium also concentrates where special items are large and, against I/B/E/S, where providers disagree. Finally, the premium is larger in firms with high *LLM Visibility*, a firm-level measure of how visible the firm's consensus forecasts are to large language models (LLMs). The difference appears only once GPT can retrieve.

The empirical analysis first examines how reliance on individual FDPs changes with GenAI's availability, before and after ChatGPT's release. The sample includes firm-quarters with consensus earnings-per-share (EPS) forecasts from all five providers.[5] I rank the five providers by trailing accuracy within each firm-quarter. I then estimate how the earnings response coefficient (ERC; e.g., Collins and Kothari, 1989; Easton and Zmijewski, 1989) on each provider's surprise varies with that rank. The ERC increment per one-rank improvement in accuracy rises from 0.306 before ChatGPT to 0.621 after.[6] The response to each provider's own surprise also rises, significantly so for four of the five providers. Where I/B/E/S loses rank, the response to its surprise declines significantly relative to the firms in which it holds its rank, whereas the change where it gains rank is not significant. Sell-side analysts show the same one-sided pattern. After ChatGPT, they herd toward the prevailing I/B/E/S consensus significantly less in firms in which I/B/E/S loses rank around the release than in firms in which it holds its rank. The difference is 4.8 percentage points relative to a 32-percent base rate.

A stronger response to accurate providers does not establish how the market combines their forecasts. I therefore construct an accuracy-weighted standardized unexpected earnings (SUE) measure using each FDP's trailing forecast accuracy for the firm. Weights based on forecast errors

[5]The sample begins in 2021 to provide about two years of announcements before ChatGPT's release on November 30, 2022, and three years after it. Bloomberg's consensus history is accessible only for five years.

[6]The slope of the ERC in predictive value, measured by how closely the I/B/E/S consensus anticipates the firm's street earnings a year ahead, also more than doubles after ChatGPT. Because the measure is close to a firm's earnings predictability, this increment is also consistent with the 2022 shift in discount rates.

observed before the announcement apply to each provider's own surprise.[7] I compare this measure with the I/B/E/S surprise and the equal-weighted surprise. The I/B/E/S surprise captures reliance on a single provider. The equal-weighted surprise captures combination without differentiating providers by accuracy. Each horse-race regression includes the accuracy-weighted surprise and one benchmark surprise. After ChatGPT, the loading on the accuracy-weighted surprise rises, while the loading on the benchmark surprise stays flat or declines. The premium widens significantly against both benchmarks. Each widening amounts to 79 to 88 percent of the pre-ChatGPT ERC on the I/B/E/S surprise estimated on its own.

To examine when the premium emerges, I distinguish four GPT-capability regimes based on milestones in ChatGPT's publicly available retrieval capability: pre-ChatGPT, text-only, browsing, and native search. The native-search regime begins with the launch of ChatGPT search, which can search the web automatically when appropriate. In the pre-ChatGPT and text-only regimes, the premium is statistically indistinguishable from zero against both benchmarks. Against I/B/E/S, it is an insignificant 1.52 percentage points under browsing and a significant 4.91 percentage points under native search. Against equal weighting, it reaches a significant 2.43 and 2.55 percentage points under browsing and native search, respectively. These patterns are consistent with retrieval facilitating synthesis across providers.

The framework also predicts a larger premium when accuracy weighting offers greater benefits. Against I/B/E/S, the premium concentrates in firm-quarters with high cross-FDP disagreement. Against equal weighting, the change in the premium is similar across the two halves of this partition. Against both benchmarks, the premium concentrates in firm-quarters with large special items, where differences in providers' treatment of nonrecurring items are more consequential. Together, the regime and cross-sectional results support a shift toward accuracy-weighted synthesis that varies with retrieval capability and the benefits of distinguishing among providers.

Timing alone cannot separate retrieval from concurrent developments because every firm enters a regime on the same date. A comparison across firms within a regime addresses this limitation,

[7] Each provider's surprise pairs its consensus with its own reported actual earnings. The market's weights need not equal mine.

because a development unrelated to retrieval would not concentrate in the firms whose forecasts a browsing model retrieves. I measure *LLM Visibility* with queries issued in 2026, after the sample period, to a browsing-enabled model, ChatGPT with GPT-4o.[8] The gap in the premium between high- and low-visibility firms is insignificant at the five-percent level in the pre-ChatGPT and text-only regimes. The gap reaches 2.7 to 3.0 percentage points under browsing ($p < 0.05$) and 3.1 to 4.2 percentage points under native search ($p < 0.10$). By randomization inference, the browsing-regime gap is significant at the five-percent level against both benchmarks when visibility enters as a continuous measure, and the native-search gap is significant against I/B/E/S when visibility enters as a median split. Under either form, the gap reaches the five-percent level by this inference in at most one of the two regimes.

This paper makes three contributions. First, I contribute to the literature on FDPs and consensus construction (Bochkay et al., 2022; Larocque et al., 2026) by documenting changes in the earnings benchmark reflected in prices. The market's relative reliance on providers varies with the information environment. The findings thus extend the older question of which composite forecast proxies for the market's expectation (Brown, 1991; O'Brien, 1988) to the choice among providers. For researchers and practitioners, a widely used provider's consensus may therefore become a better or worse proxy for market expectations as information technology changes.

Second, I contribute to the literature on the capital-market consequences of GenAI (e.g., Bertomeu et al., 2025; Blankespoor et al., 2026; Chang et al., 2026; Cheng et al., 2025; Liu and Subasi, 2026; Lopez-Lira and Tang, 2026; Xue et al., 2025) and to the disclosure-processing-cost literature reviewed by Blankespoor et al. (2020), in which earlier studies examine EDGAR, XBRL, social media, and robo-journalism (Blankespoor, 2019; Blankespoor et al., 2014, 2018; Drake et al., 2015; Loughran and McDonald, 2017). I extend these studies from the processing of disclosures to a recurring task, the synthesis of competing forecasts. The findings are consistent with accuracy-weighted synthesis as a channel through which GenAI influences earnings-announcement prices. The accuracy premium rises in regimes with retrieval, and in those regimes it is larger in firms

[8] *LLM Visibility* is the number of source citations in the model's response, averaged across three queries for the firm's next-quarter consensus EPS. The measure is empirically distinct from conventional proxies for firm prominence.

with high *LLM Visibility* and in firm-quarters with large special items. These patterns support a processing-cost explanation in which the market's use of provider quality depends on the cost of acquiring the providers' forecasts and the benefit of distinguishing among providers.

Third, I contribute to the literature on the information environment and capital markets (e.g., Beyer et al., 2010; Bushee et al., 2010) through the construction of *LLM Visibility*, a firm-level measure of how visible a firm's consensus forecasts are to LLMs. Research measures a firm's information environment by the information available to investors. In the GenAI era, that environment also includes the information available to the models that investors query. *LLM Visibility* proxies for that availability at the firm level and can be adapted as models evolve. The premium's concentration in high-visibility firms once GPT can retrieve suggests that this information enters prices.

The paper proceeds as follows. Section 2 reviews the literature and develops the framework and its predictions. Section 3 describes the data, the GPT-capability regimes, and the accuracy-weighted surprise. Section 4 tests selective reliance on individual FDPs and estimates the accuracy premium as GenAI's capability changes and where synthesis offers greater benefits. Section 5 constructs *LLM Visibility* and tests whether the premium concentrates in firms whose forecasts a browsing model retrieves, and whether it does so only in the browsing and native-search regimes. Section 6 presents additional tests. Section 7 concludes.

2 Related Literature and Theoretical Framework

2.1 Related Literature

2.1.1 Consensus Forecasts and Forecast Data Providers

Research on analyst forecasts links the quality of earnings expectations to investors' interpretation of earnings news. Consensus forecasts provide a benchmark for evaluating realized earnings in both research and practice (Bartov et al., 2002; Beyer et al., 2010; Bradshaw et al., 2017; Kothari, 2001). At the individual-analyst level, prices respond more strongly to revisions from more accurate or higher-status analysts (Clement and Tse, 2003; Park and Stice, 2000; Stickel, 1992). This evidence raises the question of whether the market also distinguishes among the providers that aggregate analysts' forecasts.

The quality of a consensus depends partly on how its provider constructs it. FDPs screen analysts' submissions and determine which earnings components enter their reported forecasts and actuals. Bochkay et al. (2022) show that changes in I/B/E/S's treatment of unexpected charges and gains alter the predictive properties of street earnings and price discovery around announcements.[9] Their evidence establishes that providers' construction choices shape the information investors receive. It also motivates separating changes in the forecasts providers supply from changes in how investors use competing forecasts.

Among these construction choices, the earnings definition has the longest research record. Bradshaw and Sloan (2002) document a growing gap between GAAP and street earnings. They also show that the market response to street earnings has displaced GAAP earnings as a primary determinant of stock prices. The reported actual is itself a construction, since I/B/E/S actuals differ systematically from the non-GAAP earnings that managers disclose (Bentley et al., 2018) and often from the actual implied by analysts' own forecasts (Brown and Larocque, 2013). A provider's surprise therefore depends on the basis of its actual as well as on its forecast.

Cross-provider evidence links these differences to market responses. Larocque et al. (2026) document economically meaningful differences among five FDPs' consensus forecasts and earnings surprises for the same firm-quarter. When providers disagree about whether a firm misses or beats expectations, market reactions are more likely to align with higher-quality providers. Their composite quality measure includes historical forecast accuracy and the predictive value of street earnings, among other attributes. On average, market reactions align more closely with I/B/E/S and Capital IQ than with the other providers. This evidence leaves open whether lower processing costs shift market expectations toward a combination that uses accuracy information across providers.

The providers' differences revive the older question of which composite forecast proxies for the market's expectation. Brown et al. (1987) treat the choice among proxies for the market's assessment of unexpected earnings as a measurement-error problem in return regressions. The most

[9] *Street earnings* reflect adjustments to GAAP earnings for items such as restructuring charges, impairments, and gains on asset sales. Pairing each forecast with its provider's reported actual avoids mixing sources, but does not guarantee a common accounting basis.

current forecast is at least as accurate as the mean or median forecast (O'Brien, 1988). Timely composites of recent forecasts are more accurate than the mean of all forecasts (Brown, 1991). Consensus forecasts outside the five providers extend the same question. Jame et al. (2016) find that the Estimize consensus, built from forecasts contributed through an open platform, adds information about earnings and about the market's expectation beyond the sell-side consensus. These studies take the forecasts at hand as given and ask how to combine them. The providers' differences add a prior choice among sources that aggregate overlapping analysts under different rules.

Evidence on analyst quality shows that identifying more accurate forecasts need not benefit investors. Michaely et al. (2024) find limited economic benefits from using the average forecast of high-quality analysts in place of the broader consensus, although dispersion among those analysts' forecasts contains economically useful information. They discuss error cancellation, imperfect persistence in analyst quality, and the costs of identifying high-quality analysts as reasons why quality-based selection need not improve investment outcomes. Their analysis concerns aggregation across analysts within a provider. Combining FDP consensus forecasts introduces differences in providers' screening and earnings-adjustment rules, in addition to differences in their underlying analyst forecasts. The within-provider evidence thus calls for a test of the benefits of cross-provider synthesis that does not presume that recent accuracy yields useful combination weights.

2.1.2 Disclosure Processing Costs and Limited Attention

The disclosure-processing-cost literature explains why investors may rely on a familiar consensus even when alternatives are publicly available. Blankespoor et al. (2020) distinguish the costs of becoming aware of information, acquiring it, and integrating it into decisions. Research on Twitter and EDGAR examines how dissemination and access affect investors' use of disclosures (Blankespoor et al., 2014; Drake et al., 2015; Loughran and McDonald, 2017). Investors' searches for a firm rise about two weeks before its earnings announcement (Drake et al., 2012). Where investors search more beforehand, the price response at the announcement is smaller. Studies of XBRL and robo-journalism examine the effects of making disclosed information easier to process (Blankespoor, 2019; Blankespoor et al., 2018). Open conference calls and the business press also

broaden access to firm information (Bushee et al., 2003, 2010; Kimbrough, 2005). Integration is costly for professionals and individuals alike. Less readable annual reports are associated with less accurate and more dispersed analyst forecasts (Lehavy et al., 2011). Individual investors invest more in firms with clear and concise disclosures (Lawrence, 2013). Blankespoor et al. (2020) note that making a disclosure more available lowers the costs of awareness and acquisition and leaves integration costs largely in place. Limited-attention research shows why access alone need not ensure that investors incorporate available information into prices (Cohen and Lou, 2012; deHaan et al., 2015; DellaVigna and Pollet, 2009; Hirshleifer et al., 2009). The same research implies that adding sources can lower the quality of the earnings expectation when investors cannot assess each source's reliability cheaply (Blankespoor et al., 2020; Hirshleifer et al., 2009). With five providers of unequal accuracy, the value of the additional sources depends on the investor's ability to assess them.

The FDP setting applies this framework to the choice among competing earnings benchmarks. Investors face several forecasts of the same firm's earnings, each constructed under different rules and accompanied by a different accuracy record. Access to these forecasts still leaves investors with the task of assessing their relative reliability. Even at the analyst level, investors extract only part of the information that observable characteristics carry about future forecast accuracy (Clement and Tse, 2003). Because that reliability can vary across firms and over time, a provider's overall reputation may be an imperfect guide to the forecast most relevant for a particular announcement. The relevant processing task is to assess how much each forecast should contribute to the earnings expectation against which the announcement is evaluated.

2.1.3 Artificial Intelligence in Capital Markets

Earlier applications of AI in finance often use models developed for a particular prediction or classification task (Goldstein et al., 2021). For example, Gu et al. (2020) show that machine learning improves return prediction by capturing nonlinear interactions among predictors. Adapting such a model to a new investment problem can require preparing inputs, estimating and validating it, and maintaining the resulting process. These requirements create costs that can limit investors' ability to apply the technology.

GenAI offers a more flexible way to obtain analytical assistance. A pretrained LLM can perform different tasks specified through natural-language instructions and examples, without being retrained for each task (Brown et al., 2020). Investors can use the same model to request an interpretation, compare alternatives, and refine an analysis through follow-up questions. This flexibility may reduce the expertise and effort required to adapt an analytical tool to a particular investment question. Evidence from other settings supports this reduction in the expertise required. In a field study of customer-support agents and in a preregistered experiment with professionals, GenAI assistance raises productivity most for less experienced and lower-ability workers (Brynjolfsson et al., 2025; Noy and Zhang, 2023). The same models also carry information of their own. ChatGPT-derived assessments of news headlines predict stock returns (Lopez-Lira and Tang, 2026). Li et al. (2025) find that GPT's own earnings forecasts vary in quality across contexts. They conclude that the friction can shift from processing information to evaluating the model's output.

Retrieval adds a separate capability that is particularly relevant to consensus forecasts. When connected to browsing tools, an LLM can obtain current external information to supplement its training data and the material supplied by users (OpenAI, 2023). In the FDP setting, this capability can connect a request to compare providers with the acquisition of their forecasts and earnings histories. Conventional software can also retrieve forecasts and calculate combination weights. The proposed benefit of retrieval-enabled GenAI is to make these steps accessible through one natural-language interface. That access can lower the cost of comparing sources repeatedly across firms and quarters.

A browsing model answers a question from the pages it retrieves (Nakano et al., 2021; OpenAI, 2023). When the documents available for retrieval change, the answers change with them (Lewis et al., 2020). For a question about a firm, the answer therefore depends on which pages the model reaches and on how much it draws on each. How much an answer draws on a retrieved page depends on the page's content (Aggarwal et al., 2024). In systems that answer from web search, the source citations support only part of what the answers state (Liu et al., 2023). A model can also answer

from memory in place of retrieval (Lewis et al., 2020; Lopez-Lira et al., 2025). The same question can also draw different answers from successive versions of one model (Chen et al., 2024).

AI also changes the production of financial information. Analysts at brokerages with greater AI adoption issue more accurate forecasts (Kimbrough et al., 2026). GenAI narrows the forecast-accuracy gap between nonprofessional and professional contributors (Liu and Subasi, 2026). On Seeking Alpha, AI adoption increases authors' output and coverage, although AI-generated articles elicit weaker market responses than human-written articles (Bradshaw et al., 2026). These findings show that AI can change both the supply and the quality of information available to investors.

Studies of investor use connect these analytical capabilities to financial decisions. Brokerage queries and survey responses show that investors use GenAI to interpret and contextualize financial information (Blankespoor et al., 2026). Retail trading also becomes more closely aligned with AI-derived earnings-call sentiment after ChatGPT's release (Chang et al., 2026). Trading declines during ChatGPT outages (Cheng et al., 2025). Bid–ask spreads widen during Italy's ChatGPT ban (Bertomeu et al., 2025). Neither design answers the question of this paper. In an event study of eight outages, Cheng et al. (2025) measure trading in five-minute intervals during interruptions that last from 18 minutes to four hours, 119 minutes on average. In an archival study of the ban, Bertomeu et al. (2025) observe one country without the tool for four weeks. The synthesis this paper studies is a standing practice that announcement-window prices reflect across firm-quarters. An interruption that brief, or a ban confined to one country, would dilute its effect below detection.

In a review of the 2025 *Journal of Accounting Research* conference papers, Cao et al. (2026) conclude that GenAI can reduce investors' processing costs. They add that its effect on the information environment depends on the cost of verifying its output. Building on this view, I examine whether, as GenAI diffuses, the market places greater weight on more accurate providers and on an accuracy-weighted synthesis of their forecasts.

2.2 Theoretical Framework and Predictions

Figure 1 summarizes a framework that weighs the benefits of synthesis against its processing costs. I define *synthesis* as retrieving forecasts and their histories, assessing source quality, and combining the forecasts using that assessment. GenAI's retrieval capability varies along two dimensions. First, it grows over time across the GPT-capability regimes, four periods marked by ChatGPT's retrieval capability. Second, across firms, some firms' consensus forecasts are more visible to LLMs than others'. Greater capability and greater visibility lower the cost of synthesis (the awareness, acquisition, and integration costs of assessing competing providers). A lower cost enables the market to price an accuracy-weighted synthesis across FDPs in place of a single-FDP or equal-weighted benchmark. The benefit of synthesis moderates the shift, which should be largest where providers disagree and where special items are large. Combining forecasts can reduce error by diversifying provider-specific noise (Ahlers and Lakonishok, 1983; Bates and Granger, 1969; Clemen, 1989; Granger and Ramanathan, 1984; Stock and Watson, 2004; Timmermann, 2006). The two benchmarks pool different inputs. I/B/E/S combines individual analysts' forecasts, whereas the equal-weighted benchmark combines FDP consensus forecasts. Neither adjusts weights across FDPs to reflect their recent accuracy for a particular firm.

Predictable differences in provider accuracy create scope for improving on these benchmarks. FDP accuracy varies across firms and over time (Larocque et al., 2026). Providers' methodologies also affect the information content of their reported earnings (Bochkay et al., 2022). Assigning greater weight to more accurate providers can improve on equal weighting (Bates and Granger, 1969). The gain depends on whether recent accuracy predicts subsequent accuracy. Noisy or unstable estimates can offset the benefit of assigning unequal weights.

Processing costs can prevent investors from realizing this benefit. Investors must first recognize that FDP forecasts and their relative accuracy differ. They must then retrieve each provider's current forecast and historical consensus–actual pairs, assess its accuracy for the firm, and use that assessment to assign weights. These tasks span the awareness, acquisition, and integration costs in Blankespoor et al. (2020). Repeating them across firm-quarters requires time, and automating

them at scale requires programming and data-engineering expertise. An investor may therefore prefer a familiar provider or an equal-weighted combination when the expected gain from accuracy weighting is smaller than its cost. Blankespoor et al. (2020) state the same condition for disclosures in general. Where the cost of processing a disclosure exceeds the expected trading gain, investors rationally leave it unprocessed, and its information can stay out of prices.

Retrieval-enabled GenAI can lower both the recurring time cost and the expertise required for synthesis. An investor can ask an LLM to retrieve forecasts and earnings histories from public websites, assess accuracy, and compute a weighted consensus.[10] The task requires consensus and actual earnings histories, without reverse-engineering providers' proprietary methodologies. Both an analyst with terminal access and an individual investor can issue such a query. The framework therefore requires synthesis by investors whose expectations enter prices, without specifying an investor group.

The accuracy-weighted surprise proxies for the synthesis that this process supports. The GPT-capability regimes provide the most direct evidence on retrieval over time. I test the role of accuracy in integration indirectly, through the selective-reliance tests and the special-items and dispersion partitions.

The framework predicts greater use of provider quality as the cost of synthesis falls. I measure quality primarily by trailing forecast accuracy for the firm and also examine predictive value for future performance. The framework holds providers' construction rules and the resulting street earnings fixed.[11] Changes in awareness, acquisition, and integration costs yield two predictions, the first in two parts.

> **Prediction 1a.** After ChatGPT's release, the market places greater relative weight on FDP forecasts with higher recent accuracy or greater predictive value for future performance. The

[10] Online Appendix A.2 presents a pilot that asks a browsing model to complete these steps. Table OA1 reports the October 2026 pilot results. ChatGPT retrieves per-site histories whose reported EPS figures match one of the five providers' actuals within two cents at a mean rate of 98 percent per response. It computes site-level accuracy and an accuracy-weighted consensus in 28 to 62 percent of responses, depending on the prompt.

[11] Section 4.2.2 examines this assumption at ChatGPT's release and the browsing boundary. A supply-side change, such as the 2009 I/B/E/S methodology change studied by Bochkay et al. (2022), provides a separate explanation for changes in the market's use of forecasts.

alignment of announcement returns with an accuracy-weighted synthesis strengthens relative to their alignment with a single-FDP consensus or an equal-weighted combination.

The shift should also be larger in the firm-quarters where the benefit of synthesis is larger. When providers' surprises disagree more, the choice of weights moves the earnings expectation more. Larocque et al. (2026) find that, when providers disagree about a beat or a miss, the market's reaction is more likely to align with the higher-quality providers. Because providers differ in their treatment of unexpected charges and gains (Bochkay et al., 2022; Larocque et al., 2026), disagreement over what belongs in street earnings should be largest where special items are economically material. Special items also carry information about future earnings that prices do not fully reflect (Burgstahler et al., 2002). The choice among providers' surprises therefore has its largest consequences where such items are large. Section 4.2.3 partitions the sample on disagreement and on special items.

Prediction 1b. The post-ChatGPT shift is stronger in firm-quarters with greater cross-FDP disagreement and in those with larger special items.

The shift in Prediction 1a should differ across firms. The awareness and acquisition costs of synthesis fall only for the forecasts that a browsing model retrieves. The reason is that the model answers from the pages it reaches and otherwise from memory (Lopez-Lira et al., 2025; Nakano et al., 2021). The shift toward accuracy-weighted synthesis should be larger for the firms whose forecasts are more visible to the model. Before ChatGPT can browse, the model answers from its training data, apart from a temporary browsing beta. The difference between firms should therefore appear only once the model can retrieve.

LLM Visibility measures the second dimension, visibility across firms. I observe how much the model draws on the web for a firm when asked for its consensus, through the source citations in its answer. The citation count proxies for this visibility and is not a record of which providers' forecasts the model weighs. Where the model draws more, the fall in acquisition costs is larger. The browsing model retrieves forecasts from free websites, without access to FDP subscription products. Among the five providers, those websites name only Zacks and S&P Global Market Intelligence, the publisher of the Capital IQ consensus, as their current sources (see Table OA2 for details). The correspondence between the accuracy-weighted surprise and the forecasts retrieved

by the model is therefore an assumption, not an observed match. *LLM Visibility* provides the most direct evidence on retrieval across firms within a regime, whereas the GPT-capability regimes vary retrieval capability over time, on the same dates for every firm. A firm characteristic that proxies for prominence would predict a difference in every regime. Section 5 constructs the measure and compares firms within each regime, and Online Appendix A.1 describes the query.

Prediction 2. The post-ChatGPT shift is stronger for firms whose consensus forecasts are more visible to LLMs, and this difference appears only once GPT can retrieve.

3 Data, Sample, and the Accuracy-Weighted Surprise

3.1 Heterogeneity in the FDPs' Consensus Forecasts

Differences in analyst coverage and earnings definitions produce distinct consensus forecasts across FDPs. I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks each construct and sell analyst-consensus EPS information. Each provider includes a different subset of analysts and applies its own recency cutoffs and outlier screens. Providers also differ in their treatment of charges, gains, and other adjustments to reported earnings (Bochkay et al., 2022; Larocque et al., 2026). These choices enter both their consensus forecasts and their reported actuals.[12]

The providers also differ in accuracy, and no single provider consistently supplies the most accurate forecast. Bloomberg has the largest trailing absolute error in most firm-quarters, so it generally receives the smallest accuracy weight. I/B/E/S ranks highest on average across composite measures of FDP quality (Larocque et al., 2026). I/B/E/S is nonetheless the most accurate provider in only 16.9 percent of firm-quarters in my S&P 1500 sample (untabulated), below the 20 percent implied by a uniform draw across five providers. These differences create scope for weighting providers according to their recent accuracy for each firm.

3.2 Data and Sample

I collect quarterly consensus EPS forecasts and reported actuals separately from all five FDPs. I obtain I/B/E/S and Zacks data through WRDS, using the I/B/E/S Summary History file and the

[12]Table OA3 reports the share of firm-quarters in which reported actuals differ for each of the ten provider pairs, and Online Appendix C.1 describes the providers' origins.

Zacks Estimates database, respectively. I download FactSet, Bloomberg, and S&P Capital IQ histories for each firm through proprietary terminals or institutional portals.[13]

Table 1 summarizes the sample selection. The initial sample contains 29,395 firm-quarters across 1,502 S&P 1500 firms with earnings announcements during 2021–2025. I apply the S&P 1500 constituent list at the data-collection date to every quarter in the sample, so membership is not point in time. I use I/B/E/S announcement dates to center all event windows. Requiring a consensus from every provider in the announcement quarter leaves 25,027 firm-quarters. I then exclude observations with missing market outcomes, accuracy-weighted surprises, controls, constructed returns, or predicted surprises, and remove singleton firms. The final sample contains 23,833 firm-quarters across 1,334 firms. The underlying consensus and actual histories extend into 2020 for the earliest sample announcements.

Table 2 reports descriptive statistics for the main variables, with all continuous variables winsorized at the 1st and 99th percentiles. Firm characteristics come from CRSP and Compustat. Quarter and announcement characteristics come from Compustat, CRSP, and I/B/E/S. Measures of the information environment come from I/B/E/S and the Thomson Reuters 13F Holdings database. Appendix A defines the controls. DataForSEO provides the *LLM Visibility* queries, issued through its LLM Responses API in 2026, after the sample period (Section 5.1), and the estimated monthly visits from Google organic search to the GenAI tools (Section 3.3 and Online Appendix C.5). The retail share of trading comes from the WRDS TAQ Intraday Indicators (Section 6 and Online Appendix C.4).

3.3 The GPT-Capability Regimes

Because variation in GenAI use across investors cannot be measured, I rely primarily on time-series variation in ChatGPT's retrieval capability, which is public and dated. I therefore divide the sample into four regimes using milestones in ChatGPT's publicly available retrieval capability,

[13] Larocque et al. (2026) follow a similar collection procedure for their 2002–2020 study. Bloomberg histories require retrieval by ticker through Bloomberg Query Language functions. The S&P 1500 provides a feasible universe for that retrieval, because coverage declines outside its large-, mid-, and small-cap indices and the design requires a consensus from all five providers. The five-year limit on accessible consensus history determines the sample start.

so that I can follow the premium as the capability evolves.[14] The regimes are pre-ChatGPT (R0), text-only (R1, beginning November 30, 2022), browsing (R2, beginning September 27, 2023), and native search (R3, beginning October 31, 2024).[15] If the premium reflects retrieval, it should not appear in the text-only months. Figure 2 depicts the four regimes, the milestones that separate them, and the number of firm-quarters in each regime's horse-race regression. I assign each announcement to the regime in effect on its announcement date, with R3 ending on December 31, 2025. R1 also includes a temporary browsing beta, so it is not a period uniformly without retrieval. Browsing allows the model to retrieve current forecasts from multiple websites in response to a query. R3 marks the launch of ChatGPT search, which can search automatically when appropriate or let the user select search manually. The post-ChatGPT indicator $Post_q$ equals one for announcements on or after November 30, 2022, and the post-browsing indicator $Post_q^{\text{Browse}}$ equals one for announcements on or after September 27, 2023.

3.4 Dependent Variables

I infer the market's reliance on a consensus from the association between its implied surprise and announcement returns. Other things equal, returns should respond more strongly to surprises measured against a consensus that more closely reflects market expectations. The immediate-reaction tests use the three-day cumulative abnormal return $CAR_{[-1,1]}$. The

[14] When browsing is restored in September 2023, ChatGPT receives 334 times as many visits from Google organic search as Perplexity. At the launch of ChatGPT search, it receives 12.8 times as many as the next-largest tracked tool. In December 2025, it receives 5.9 times as many as Gemini. These comparisons establish ChatGPT's prominence among the tools that DataForSEO covers. They do not measure total usage or investor adoption. The milestones also occur during a broader expansion of GenAI web retrieval. Perplexity launches in December 2022, Bing Chat on February 7, 2023, and public access to Bard begins on March 21, 2023. These tools offer web retrieval before ChatGPT's September 2023 restoration of browsing. Figure OA1 in Online Appendix C.5 plots the estimated monthly visits from Google organic search to the GenAI tools covered by DataForSEO. Bing Chat and Bard are absent from DataForSEO's traffic estimates. I therefore use ChatGPT's milestones to organize the timing tests, without attributing changes across regimes exclusively to ChatGPT.

[15] A browsing beta begins rolling out to all Plus users in May 2023, is suspended on July 3, and returns on September 27, initially for paid subscribers. The R2 boundary marks this restoration, not browsing's first availability. Browsing is already integrated with other Plus tools in November 2023, without requiring users to switch models. The R3 boundary marks the launch of ChatGPT search, not the first integration of browsing into the main product. Search initially reaches Plus and Team users and the SearchGPT waitlist. Access expands to all logged-in users in December 2024.

trading-strategy test uses the four-day buy-and-hold abnormal return $BHAR_{[-2,1]}$. For each window, I construct Fama–French–Carhart four-factor (FFC) and market-adjusted abnormal returns.[16]

3.5 The Accuracy-Weighted SUE

I construct the accuracy-weighted surprise by combining the five providers' surprises using their recent forecast accuracy for the firm. Let g index I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks. Provider g's standardized unexpected earnings for firm i in quarter q is $SUE^g_{i,q} = (Actual^g_{i,q} - Consensus^g_{i,q})/p_{i,q}$, where $p_{i,q}$ is the closing share price two trading days before the announcement. The consensus is the provider's last forecast for quarter q before the announcement. The actual is the earnings figure that the same provider reports. The accuracy-weighted surprise, with the weights $w^g_{i,q}$ defined in Equation (3), is

$$SUE^{Acc}_{i,q} = \sum_g w^g_{i,q} SUE^g_{i,q}. \tag{1}$$

The equal-weighted benchmark, $SUE^{EW5}_{i,q}$, assigns each provider a weight of $1/5$. The I/B/E/S benchmark, $SUE^{IBES}_{i,q}$, uses that provider's surprise alone.

Combining surprises accommodates the providers' differing earnings definitions without selecting one provider's actual as the common target. Pairing each consensus with its own reported actual avoids the mismatch that subtracting one provider's forecast from another's actual would introduce. An accounting adjustment cancels within a surprise only if it enters the forecast and actual by the same amount. The accuracy-weighted surprise can be written as

$$SUE^{Acc}_{i,q} = \frac{\sum_g w^g_{i,q} Actual^g_{i,q} - \sum_g w^g_{i,q} Consensus^g_{i,q}}{p_{i,q}}. \tag{2}$$

Thus, the accuracy-weighted surprise applies the same weights to forecasts and reported actuals. This identity does not harmonize providers' earnings definitions. The difference between the accuracy-weighted surprise and a benchmark surprise therefore contains both consensus and actual-earnings differences. The horse race in Section 4.2.1 does not separate these components. The

[16]The BHAR window begins at day -2 to include any pre-announcement leakage. Both windows end at day $+1$ to capture the next trading session's response to announcements after the market close. All returns are expressed in percent. Inferences are unchanged across the two adjustments. Tables presenting one adjustment report market-adjusted returns. The FFC counterparts are available on request.

predicted-surprise test in Online Appendix B.4 isolates the consensus component and shows that it carries economic value on its own.

To construct the weights, I first calculate each provider's price-scaled absolute forecast error, $Err^g_{i,q} = |Actual^g_{i,q} - Consensus^g_{i,q}|/p_{i,q}$. I then average these errors over the four quarters preceding quarter q to obtain each provider's mean absolute forecast error, $MAFE^g_{i,q}$.[17] I define accuracy as $Acc^g_{i,q} = -MAFE^g_{i,q}$, so higher values indicate greater accuracy. The rolling window excludes the announcement quarter's forecast error.

I next standardize accuracy by its mean $\overline{Acc}_{i,q}$ and standard deviation $\text{sd}(Acc_{i,q})$ across the five providers within each firm-quarter:

$$Acc_z^g_{i,q} = \frac{Acc^g_{i,q} - \overline{Acc}_{i,q}}{\text{sd}(Acc_{i,q})}.$$

I convert these standardized accuracies into weights using the softmax transform:

$$w^g_{i,q} = \frac{\exp(Acc_z^g_{i,q})}{\sum_h \exp(Acc_z^h_{i,q})}. \tag{3}$$

The weights are positive, sum to one, and increase with a provider's trailing accuracy relative to the others.

This construction implements the forecast-combination intuition that provider accuracy can guide aggregation (Bates and Granger, 1969; Timmermann, 2006). It does not estimate the population-optimal weights, which also depend on forecast-error covariances.[18] The four-quarter window balances recency against estimation noise, and updating it each quarter allows providers' weights to change. The average within-firm-quarter minimum and maximum weights are 0.029 and 0.420, respectively (untabulated), compared with the equal weight of 0.200. The correlation between SUE^{Acc} and SUE^{EW5} is approximately 0.82 (untabulated).

[17] When fewer than four errors are available, I average the observed errors within this window, dividing their sum by the number observed. Table OA3, Panel C, reports the number of firm-quarters with complete histories in each regime. All five providers have complete four-quarter histories in 44 percent of pre-ChatGPT firm-quarters, compared with 88 to 92 percent in each post-ChatGPT regime.

[18] Standardization makes the weights insensitive to the absolute size of the accuracy differences. Online Appendix B.1 illustrates the construction for a hypothetical firm-quarter, and Online Appendix B.2 derives the optimal weights when the provider-specific errors are uncorrelated.

4 GenAI Availability, Selective Reliance, and Accuracy-Weighted Synthesis Across FDPs

This section provides a chain of evidence on how the market's selection and synthesis of the five FDPs' consensus forecasts change after ChatGPT's release and as its retrieval capability grows. Section 4.1 asks whether the market relies on each FDP's consensus selectively by its accuracy, and whether that selectivity strengthens after ChatGPT. Under Prediction 1a, once GenAI makes each provider's track record readily retrievable, the market should weight a provider's surprise by how accurate that provider has recently been for the firm. Section 4.2 asks whether the selectivity carries into the aggregation that prices reflect, first after the release and then as retrieval capability grows, and where the benefits of synthesis are largest. A market can favor one accurate provider without combining across providers at all. The second question therefore tests the stronger claim.

4.1 Selective Reliance on Individual FDPs

4.1.1 Selective Reliance After ChatGPT

The first test asks whether, after ChatGPT, the ERC on an individual FDP's surprise rises more steeply with that FDP's accuracy rank. I rank the five FDPs within each firm-quarter by *ex ante* trailing MAFE (rank 1 = least, 5 = most accurate), stack the five provider observations for each firm-quarter, and estimate

$$CAR_{i,q} = \alpha + \beta\, SUE^{g}_{i,q} + \delta\, Rank^{g}_{i,q} + \gamma\left(SUE^{g}_{i,q} \times Rank^{g}_{i,q}\right) + \varepsilon_{i,q,g}, \tag{4}$$

where γ is the ERC increment per one-rank accuracy gain.[19] Because rank is assigned within each firm-quarter, identification comes from the same announcement return's differential sensitivity to the five providers' surprises. A firm- or time-level shock, including methodology drift common to all providers, shifts all five observations alike and cannot by itself create differences in the ERC across accuracy ranks. The first test leaves open, however, whether the market's reliance on a given provider rises and falls with that provider's record.

The second test addresses this question directly by fixing the provider and exploiting cross-firm variation in how its accuracy rank changes around ChatGPT. I classify firms by the change in

[19] The earnings response coefficient is the slope of announcement returns on the earnings surprise. A larger ERC means the market treats the surprise as more informative.

I/B/E/S's mean rank between the pre- and post-ChatGPT windows. The indicators $IBES\,Up_i$ and $IBES\,Down_i$ mark improvements and deteriorations of at least two positions. I then estimate the triple-interaction specification

$$\begin{aligned} CAR_{i,q} = \alpha + \beta_1 SUE^{IBES}_{i,q} + \beta_2 SUE^{IBES}_{i,q} \times Post_q \times IBES\,Up_i \\ + \beta_3 SUE^{IBES}_{i,q} \times Post_q \times IBES\,Down_i + Controls + \text{FE} + \varepsilon_{i,q}, \end{aligned} \tag{5}$$

with the lower-order terms included but not displayed.[20] Selective reliance predicts $\beta_3 < 0$. If the pre-ChatGPT market already defaults to I/B/E/S, a gain in its rank adds little, so $\beta_2 \approx 0$.

Table 3 reports both tests. After ChatGPT, the price reaction to a given surprise depends increasingly on which provider produces it.[21] In Panel A, the per-rank ERC increment γ roughly doubles, from 0.306 before to 0.621 after, and the difference is significant ($p < 0.01$). The provider rows show positive and significant post-ChatGPT per-rank increments for every provider except Zacks.[22]

Panel B shows that the response to the I/B/E/S surprise weakens significantly where I/B/E/S loses rank, relative to the firms in which it holds its rank, and does not change significantly where it gains rank. For the firms in which I/B/E/S loses rank, the post-ChatGPT triple interaction ranges from -2.3 to -2.8 across specifications ($p < 0.05$). The point estimates for the firms in which it gains rank are about half of those for the firms in which it loses rank. This asymmetry is consistent with a pre-ChatGPT market that already defaults to I/B/E/S. The baseline I/B/E/S ERC rises after ChatGPT, consistent with the decline in surprise magnitudes in Panel (b) of Figure OA2. The triple interaction therefore measures a relative decline.

4.1.2 Analyst Herding Toward the I/B/E/S Consensus

The partition in Table 3, Panel B, rests on the change in I/B/E/S's rank, so attenuation alone could lower the I/B/E/S ERC where I/B/E/S loses rank. I therefore need a measure of reliance that

[20] The lower-order terms are the two-way interactions among SUE^{IBES}, $Post_q$, and the two rank-change indicators. The indicators themselves are firm-level and are absorbed by the firm fixed effects. $Post_q$ itself is not included, and the year fixed effects absorb its variation except within fiscal years 2022 and 2023.

[21] The rank selectivity could come from a weaker reaction after ChatGPT to the less accurate providers, or from a weaker reaction to every provider. Table OA4, Panel B, reports the ERC on each provider's own surprise. The ERC rises after ChatGPT for every provider. Among the four providers whose increase is significant, it ranges from 0.68 for FactSet to 1.45 for Zacks. Bloomberg's increment of 0.18 is insignificant. The selectivity is therefore a relative reweighting.

[22] Versions of Panel A estimated with one regression per provider–rank pair, in both absolute-value and signed form, yield the same inferences (untabulated).

does not rest on an ERC. Analysts' forecast revisions provide such a measure from the forecasting process itself. Each time a sell-side analyst revises a forecast, the analyst chooses how far to move toward the prevailing consensus (Clement and Tse, 2005). The strength of that tendency indicates how much weight the analyst assigns to the consensus as a signal. If the post-ChatGPT information environment devalues a provider whose accuracy has declined, analysts should herd less toward the I/B/E/S consensus in the firms in which I/B/E/S loses accuracy rank than in the firms in which it holds its rank. The regressions use 122,001 quarterly EPS forecast revisions for the sample firms over 2019–2025. The indicator *Herding* equals one when the revised forecast falls between the analyst's own prior and the prevailing I/B/E/S consensus. I regress the indicator on $Post_q$ and its interactions with the $IBES\,Up_i$ and $IBES\,Down_i$ indicators of Table 3, with the firm and event controls of Equation (5), the analyst controls of Clement (1999), either analyst and firm fixed effects or analyst×firm fixed effects, revision-year fixed effects, and fixed effects for twenty ventiles of the distance between the analyst's prior and the consensus.[23]

Table 4 reports the estimates. After ChatGPT, herding toward the consensus falls in the firms in which I/B/E/S loses rank, relative to the firms in which I/B/E/S holds its rank. The $Post \times IBES\,Down$ coefficient is -4.8 percentage points (t = -2.62), about 15 percent of the 32-percent base rate, while the $Post \times IBES\,Up$ coefficient is positive and insignificant. This decline mirrors the one-sided pattern of the market's reweighting in Table 3.

Table 4 also reports the estimates by GPT-capability regime. Because the partition rests on the rank change over the whole post-ChatGPT period, selective reliance predicts a relative decline in every post-ChatGPT regime. The regime split is therefore descriptive, and columns (1) and (2) provide the formal test. The decline is already present in the text-only regime. This evidence supports the market-level results but does not identify the mechanism on its own.

[23] I/B/E/S is the only provider for which forecast-level data are available to me. The test therefore asks whether analysts' reliance on the I/B/E/S consensus changes after ChatGPT, not whether analysts substitute toward the consensus forecasts of the other FDPs. Online Appendix C.2 describes the revision sample, the prevailing consensus, and the specification in full.

4.1.3 An Alternative Proxy for the Quality of the Consensus Forecast: Predictive Value

Trailing accuracy is one proxy for the quality of a consensus forecast. *Predictive Value*, how closely the consensus anticipates the firm's future street earnings, is an alternative proxy. If the post-ChatGPT selectivity reflects a general preference for consensus quality, and is not an artifact of the trailing-accuracy construction, the same pattern should also appear on this proxy, as Prediction 1a anticipates. The weighting analyses of Section 4.2 operate on accuracy, not on predictive value, because accuracy is realized at every announcement and varies by firm-quarter. The test proceeds in two stages.

The first stage measures predictive value for each firm-quarter from the firm's own history of I/B/E/S consensus forecasts and street actuals. The measure scores each firm-quarter by how closely a linear function of the I/B/E/S consensus tracks the firm's street earnings a year ahead over a window of up to twenty quarters.[24]

The second stage tests whether the ERC on the I/B/E/S surprise rises with predictive value and whether it rises more steeply after ChatGPT. I estimate

$$\begin{aligned} CAR_{i,q} = \alpha + \beta_1 SUE^{IBES}_{i,q} + \beta_2 SUE^{IBES}_{i,q} \times Post_q + \beta_3 SUE^{IBES}_{i,q} \times Predictive\ Value^{IBES}_{i,q} \\ + \beta_4 SUE^{IBES}_{i,q} \times Predictive\ Value^{IBES}_{i,q} \times Post_q + Controls + \text{FE} + \varepsilon_{i,q}, \end{aligned} \tag{6}$$

with the lower-order terms included but not displayed. The regression includes the firm and event controls of Equation (5) and firm and year fixed effects. The coefficient β_1 is the pre-ChatGPT ERC at zero predictive value, and β_2 is its post-ChatGPT change. The coefficient β_3 is the pre-ChatGPT slope of the ERC in predictive value, and β_4, the coefficient of interest, is its post-ChatGPT increment. A second version replaces $Post_q$ in every term with the indicators of the three GPT-capability regimes after ChatGPT's release, so that each regime carries its own increment.

Table 5 reports the estimates. Before ChatGPT, the slope of the ERC in predictive value is already strongly positive ($\beta_3 = 2.76$, $t = 4.70$). Firms whose consensus predicts well have higher ERCs (Collins and Kothari, 1989), so the identified claim is the post-ChatGPT increment,

[24] Predictive value and accuracy largely measure the same underlying quality at different horizons, with a cross-firm correlation of 0.72. Accuracy therefore captures most of what predictive value measures. Online Appendix C.3 gives the two-stage construction, the measure's coverage, and two alternative explanations for the increment.

$\beta_4 = 3.11$ ($t = 3.53$). The increment exceeds β_3, so the slope more than doubles after ChatGPT. Because the measure is close to a firm's earnings predictability, the increment is also consistent with an explanation that does not involve GenAI, such as the 2022 shift in discount rates. Across the GPT-capability regimes, the increment is insignificant in the text-only regime, significant under browsing, and largest under native search. The regime profile is descriptive, and the pooled increment is the formal test.

4.2 GPT-Capability Regimes and the Market's Reaction to the Accuracy-Weighted Surprise

The expectation against which the market judges an announcement is an aggregation of the consensus forecasts across FDPs, and the two benchmarks either use the I/B/E/S consensus alone or weight the five providers equally. The accuracy-weighted surprise weights each provider's surprise by its recent accuracy for the firm. The test is whether announcement returns respond to the accuracy-weighted surprise beyond their response to the two benchmarks. Section 4.2.2 follows the premium across the GPT-capability regimes. Section 4.2.3 tests Prediction 1b, that the premium concentrates where accuracy weighting offers greater benefits.

4.2.1 The Accuracy Premium: Horse-Race Tests

I test the accuracy-weighted surprise SUE^{Acc} against the two benchmarks. The first is the equal-weighted surprise SUE^{EW5} and the second the I/B/E/S surprise SUE^{IBES}. The pooled specification is

$$\begin{aligned} CAR_{i,q} \;=\; \alpha \;+\; \beta_1 SUE^{Acc}_{i,q} \;+\; \beta_2 SUE^{X}_{i,q} \;+\; \beta_3 SUE^{Acc}_{i,q} \times Post_q \\ +\; \beta_4 SUE^{X}_{i,q} \times Post_q \;+\; Controls \;+\; \text{FE} \;+\; \varepsilon_{i,q}, \end{aligned} \tag{7}$$

for $X \in \{EW5, IBES\}$. The pre-period accuracy premium is the contrast $\beta_1 - \beta_2$. Its post-minus-pre change, $\beta_3 - \beta_4$, is the central test of Prediction 1a. The dependent variable is the three-day CAR under both risk adjustments, market-adjusted and FFC four-factor, and year fixed effects are for the fiscal year of quarter q.[25] Online Appendix B.4 validates the construct (see Table OA5 and Figure OA2 for details). Provider accuracy ranks persist across quarters, partly by construction.

[25] Although SUE^{Acc} and the benchmark SUE^{X} are correlated by construction ($\rho \approx 0.82$ for $X = EW5$), the identifying quantity is the contrast $\beta_1 - \beta_2$, whose sampling variance is a property of the estimand, not of the horse-race parameterization. Because every test in this section uses the contrast, the inflated standard errors of the individual slopes do not affect the inference.

The predicted surprise, the pre-announcement difference between the accuracy-weighted and a benchmark consensus, forecasts announcement returns.

Table 6 reports the horse-race estimates. After ChatGPT, the market tilts toward the accuracy-weighted surprise relative to both benchmarks. In Panel A (equal-weighted benchmark), the pre-ChatGPT accuracy premium is statistically indistinguishable from zero, consistent with the equal-weighted surprise already capturing the pre-period information set. After ChatGPT, the premium turns positive, as the accuracy-weighted slope rises while the equal-weighted slope falls. The falling benchmark slope does not mean the equal-weighted surprise loses information, because in a horse race each slope is conditional on the other measure. The formal test is the pooled interaction contrast, which estimates the premium's post-minus-pre change at 2.401 ($p < 0.01$) under the market-adjusted CAR and 2.584 under the FFC four-factor CAR.

Panel B repeats the test against the I/B/E/S benchmark. There the pre-ChatGPT accuracy premium is negative but statistically insignificant. The post-ChatGPT premium rises to 2.049 ($p < 0.01$), a post-minus-pre change of 2.672 ($p < 0.01$) under the market-adjusted CAR and 2.688 under the FFC four-factor CAR.

The comparison with I/B/E/S mixes two gains, because a composite of several providers gains over I/B/E/S from pooling them as well as from weighting them by accuracy. The equal-weighted benchmark isolates the second only when its surprises share one earnings basis. Because Bloomberg's reported actuals differ from the other providers' far more often than theirs differ from one another (see Table OA3, Panel A, for details), Table OA4, Panel D, repeats the pooled horse race with the four providers other than Bloomberg. Without Bloomberg, the accuracy-weighted surprise is priced above I/B/E/S before ChatGPT and below the equal-weighted surprise of the same four providers. The pooled tests of Table 6 therefore rest on the post-ChatGPT change in the five-FDP premium, not on its level. The post-ChatGPT change in the premium is then 3.24 against I/B/E/S and 4.56 against the equal-weighted surprise of the other four providers, each significant at the five-percent level despite larger standard errors than those of the five-FDP estimates. The rise in the premium therefore does not rest on Bloomberg's surprise.

A more accurate composite proxies better for the market's expectation and would earn a higher slope for that reason alone. Its accuracy advantage over each benchmark, however, changes little from the pre- to the post-ChatGPT period, so a better proxy cannot explain a premium that widens (see Panel (b) of Figure OA2 for details). The two periods differ in one input to the accuracy-weighted surprise. Complete four-quarter windows for all five providers are rarer before ChatGPT (Section 3.5), so the pre-period weights rest on fewer quarters. A noisier pre-period accuracy-weighted surprise would lower its pre-period loading and raise the measured change.

The magnitudes are economically considerable. The post-ChatGPT increments of 2.4 and 2.7 percentage points are 79 to 88 percent of the pre-ChatGPT ERC on the I/B/E/S surprise estimated on its own (see Table OA4, Panel B, for details). Scaled by a typical accuracy tilt, $SUE^{Acc} - SUE^{EW5}$, instead of a full percentage point of surprise, the magnitudes are more modest. The counterfactual is a market that relies equally on the accuracy-weighted and the equal-weighted consensus, under which the premium is zero (Online Appendix B.3). Relative to it, a one-standard-deviation tilt of roughly 0.29 percent of price (implied by the standard deviations of Table 2 and the correlation of Section 3.5) moves the three-day CAR by about 0.2 percentage points, half the post-ChatGPT premium times the tilt, or roughly 3 percent of the CAR's standard deviation.

4.2.2 The Time Series of the Accuracy Premium

The pooled contrast of Table 6 dates the shift only to the ChatGPT boundary. If the premium reflects the synthesis of Prediction 1a, it should strengthen as the tools' retrieval capabilities expand, because retrieval is the step of synthesis that the tests identify (Section 2.2). I therefore re-estimate the horse race of Equation (7) within the four GPT-capability regimes of Section 3.3.

If the premium reflects attention to AI, sentiment, or any other shock coincident with ChatGPT's release, it should already appear in the text-only months. If it reflects retrieval, it should not. This test has three limitations. First, it tests for a discrete shock at the release date, not for a gradual process that happens to accelerate in late 2023. Second, it does not rule out shocks that coincide with the browsing boundary itself. Third, it is partial, because Perplexity, Bing Chat, and Bard

answer with web retrieval during the text-only regime and a ChatGPT browsing beta runs from May to early July 2023 before it is suspended.

Table 7 reports the horse race by regime, against the equal-weighted five-FDP benchmark in Panel A and against the I/B/E/S benchmark in Panel B. In the pre-ChatGPT and text-only regimes, the accuracy premium is statistically indistinguishable from zero on both comparisons. Under browsing it turns positive on both, at 2.43 ($p < 0.05$) in Panel A and 1.52, which remains insignificant, in Panel B. Under native search it is significant on both, at 2.55 ($p < 0.05$) in Panel A and 4.91 ($p < 0.01$) in Panel B. The premium is thus detectable only once the model can retrieve forecasts (see Panel (a) of Figure 3).

A regression with both boundaries does not separate the change at ChatGPT's release from the change at the browsing boundary.[26] The regime estimates of Table 7, not the regression, place the premium in the browsing and native-search regimes.

Other developments coincide with the regime boundaries, and timing alone cannot separate them from retrieval.[27] The visibility gap of Section 5, which appears only when retrieval arrives, weighs against the SEC guidance, the providers' AI tools, and the rise in analyst coverage as explanations of the premium.

4.2.3 Cross-Sectional Evidence: Variation in the Benefits of Synthesis at the Firm-Quarter Level

Under Prediction 1b, the market has more reason to synthesize where the gain from weighting the providers by accuracy is largest, that is, where the providers disagree about the surprise

[26] Starting the post-period at the browsing boundary lowers the change in the premium to 2.48 ($p < 0.01$) against I/B/E/S and 1.77 ($p < 0.05$) against the equal-weighted benchmark. A regression with both boundaries separates neither change from zero at the five-percent level, whereas their sum is significant (see Table OA4, Panel A, for details).

[27] Five developments could coincide with the boundaries. First, the SEC's December 2022 update to its guidance on non-GAAP measures cannot be separated from the release. At the five-percent level, four of the ten provider pairs change how often their reported actuals disagree at the release, and none at the browsing boundary (see Table OA3, Panel B, for details). Second, the providers' own AI tools arrive inside the post-period, Capital IQ's ChatIQ twelve days into native search, so part of the largest step against I/B/E/S may come from that tool. Third, analyst coverage rises at the browsing and native-search boundaries, by 0.12 and 0.33 analysts within firms (untabulated), and cannot be separated from the regime timing. Fourth, the other GenAI tools add web retrieval on the dates that Section 3.3 reports, and, taken together, those that DataForSEO tracks draw fewer visits from Google organic search than ChatGPT in every month from October 2022 to December 2025. Figure OA1 in Online Appendix C.5 plots the monthly visits of each tool. Fifth, the FDPs' consensus forecasts do not converge on I/B/E/S, with median absolute gaps to the I/B/E/S consensus between 1.08 and 1.58 basis points of price for FactSet, Bloomberg, and Capital IQ in every regime (untabulated). The gap is the absolute difference between a provider's consensus and the I/B/E/S consensus, scaled by the price that scales SUE, in basis points.

and where their treatment of nonrecurring items differs most. A premium that concentrates in those firm-quarters reflects synthesis, not a uniform shift in the earnings response. The regime estimates of Table 7 place the accuracy premium in the browsing and native-search regimes, so the cross-sectional tests replace the post-ChatGPT indicator $Post_q$ with the post-browsing indicator $Post_q^{\text{Browse}}$ of Section 3.3. I re-estimate the horse-race specification in Equation (7) within the below- and above-median halves of two partitions. First, *Concurrent Cross-FDP SUE Dispersion* is the interquartile range of the five FDPs' SUE for the same firm-quarter, and accuracy weighting has the greatest scope to be selective among providers when this dispersion is high. Second, *Special-Items Magnitude* is the absolute value of special items scaled by market capitalization two trading days before the announcement. Because FDPs differ in their treatment of unexpected charges and gains, cross-FDP disagreement over what belongs in street earnings should be largest where such items are economically material.

Table 8 reports the estimates by partition.[28] The post-browsing change in the premium is significant only in the predicted half of each partition (see Figure 4). In Panel A (*Concurrent Cross-FDP SUE Dispersion*), the change is 2.211 ($p < 0.05$) against the equal-weighted benchmark and 2.427 ($p < 0.01$) against I/B/E/S in the high-dispersion subsample, versus insignificant estimates in the low-dispersion subsample. The equal-weighted comparison requires a caveat, because the low-dispersion point estimate is close in magnitude to its high-dispersion counterpart and differs mainly in precision. Against I/B/E/S, the low-dispersion estimate turns negative. In that subsample the level of the premium is negative and precisely estimated, because, before the browsing boundary, its announcement returns load about 4.8 times as heavily on the I/B/E/S surprise as on the accuracy-weighted surprise. This loading has two possible sources. First, the I/B/E/S surprise may be the better proxy for the market's expectation when the providers agree. Second, a few firm-quarters with a large gap between the two surprises may identify the contrast.

[28] A variant of Table 8 that uses the post-ChatGPT boundary in place of the post-browsing indicator and adds a median split on *LLM Visibility* yields similar inferences, with the premium concentrating in the high half of each partition (untabulated).

Panel B (*Special-Items Magnitude*) requires no qualification. The change is 2.518 ($p < 0.05$) against the equal-weighted benchmark and 3.638 ($p < 0.01$) against I/B/E/S in the high-special-items subsample, versus point estimates of the opposite sign in the low-special-items subsample.

5 Cross-Firm Variation: LLM Visibility and the Accuracy Premium

The regime estimates of Section 4.2.2 are consistent with retrieval capability as the step on which the market's synthesis rests. This section uses variation in *LLM Visibility* across firms.

5.1 The Measure

I construct *LLM Visibility*, a firm-level measure of how visible a firm's consensus forecasts are to LLMs. For each firm, I query a browsing-enabled model, ChatGPT (GPT-4o with web search), three times for the next-quarter consensus EPS and receive each answer with links to the sources it cites. The measure is the number of these source citations, repeats included, averaged over the three queries. High-visibility firms are those above the cross-firm median, and firms at the median join the low group.[29]

The count indexes how much the model draws on the web for the firm when asked, not which providers' forecasts the model weighs. The model's citations concentrate on 17 sites, and the sites that name none of the five providers hold most of the citations. Across firm-quarters the count has a mean of 13.3 and a standard deviation of 3.0 (see Table 2 for details). I do not observe why one firm draws more citations than another. Section 5.2 tests whether the count is another name for firm prominence.

The count proxies for a firm's visibility to the model in the browsing and native-search regimes under the assumption that the set of websites carrying a firm's forecasts is stable over the sample. Drift in that set, where it is unrelated to announcement returns, adds classification error. The error attenuates the differences between high- and low-visibility firms and does not produce them. Across the three queries the firm ranking is only weakly stable. The Spearman–Brown reliability of the

[29] Online Appendix A.1 describes the query, and, with Table OA2, reports the sites the model cites and the providers they name as the sources of their estimates. The queries run on May 28, 2026, for 1,304 firms (3,912 responses) and on June 1 for 56 firms (168 responses), three per firm, after the end of the sample period. Eleven of the June firms are also among the May firms.

three-query mean is 0.11, and a median split on two of the three queries agrees with a split on the third in 52 percent of firms. This instability is a second source of classification error, with the same attenuating direction.

5.2 Distinctness from Firm Prominence

The comparison of high- and low-visibility firms is informative only if *LLM Visibility* is not another name for firm prominence. This subsection tests that condition against firm size, analyst coverage, and institutional ownership.

Table 9 reports the overlap of the high-visibility half with the top halves of firm size, analyst coverage, and institutional ownership (Panel A) and a horse race of the four measures against each benchmark (Panels B and C). *LLM Visibility* overlaps little with the conventional proxies for firm prominence. The overlap between the high-visibility half and the top half of each proxy is 0.47 to 0.51 in either direction, near the 0.50 that independent median splits imply, whereas the overlap between the top halves of firm size and analyst coverage is 0.80. In the horse race, which includes every measure's high-group triple interaction in one regression, *LLM Visibility*'s triple interaction remains, at 3.17 against the equal-weighted benchmark ($p < 0.10$) and 5.84 against I/B/E/S ($p < 0.01$). Firm size has no significant triple interaction against either benchmark. Analyst coverage has none at the five-percent level. Institutional ownership has one only against the equal-weighted benchmark, and only at the ten-percent level.

Table 9, Panel B, also splits the change against the equal-weighted benchmark, with the post-period starting at the browsing boundary, at the median of firm size and of analyst coverage. The change is 2.45 ($t = 2.18$) in firms below the median size and 0.72 above it, and 3.08 ($t = 2.57$) in firms below the median analyst coverage and 0.11 above it. The change is therefore significant only in the smaller and less-followed halves of the sample. The horse race does not separate the two halves of either split.

5.3 LLM Visibility and the Accuracy Premium by Regime

Every firm enters a GPT-capability regime on the same date, so the regime estimates of Section 4.2.2 cannot separate retrieval from the developments that coincide with the regime

boundaries. With *LLM Visibility* (Section 5.1), I compare firms within each regime. A model synthesizes only the consensus forecasts that it retrieves, so under Prediction 2 the premium should be larger in the firms whose forecasts are more visible to it. A development unrelated to retrieval would not concentrate in those firms. I estimate, separately within each GPT-capability regime r and for each benchmark X of Equation (7),

$$\begin{aligned} CAR_{i,q} = \alpha + \beta_1 SUE^{Acc}_{i,q} + \beta_2 SUE^{X}_{i,q} + \beta_3 SUE^{Acc}_{i,q} \times High\,Visibility_i \\ + \beta_4 SUE^{X}_{i,q} \times High\,Visibility_i + Controls + \text{FE} + \varepsilon_{i,q}, \end{aligned} \tag{8}$$

with the three-day market-adjusted CAR as the dependent variable, the firm and event controls of Equation (7), firm and year-quarter fixed effects for quarter q, and standard errors two-way clustered by firm and fiscal year. The low-visibility premium is $\beta_1 - \beta_2$, the high-visibility premium is $(\beta_1 + \beta_3) - (\beta_2 + \beta_4)$, and the high-minus-low gap within the regime is $Premium_r^{\text{High}-\text{Low}} = \beta_{3,r} - \beta_{4,r}$. If acquisition is a precondition for synthesis, the gap should be absent before the model can retrieve and should emerge once it can.

Table 10 reports the estimates of Equation (8) by regime. The gap $Premium_r^{\text{High}-\text{Low}}$ is statistically insignificant at the five-percent level in the pre-ChatGPT and text-only regimes (R0, R1), the difference-in-differences analog of flat pre-trends.[30] In the browsing and native-search regimes (R2, R3) the gap is positive against both benchmarks. Against the equal-weighted benchmark it is 2.7 in the browsing regime ($p < 0.05$) and 3.1 under native search ($p < 0.10$). Against I/B/E/S it is 3.0 ($p < 0.05$) and 4.2 ($p < 0.01$). The premium thus concentrates in firms whose forecasts are visible to the model, and it does so only in the regimes in which the model can retrieve them (see Panels (b) and (c) of Figure 3). This timing is consistent with acquisition as a precondition for synthesis. Once native search arrives, the premium is not confined to high-visibility firms, because the low-visibility premium against I/B/E/S in R3 is 3.05 ($p < 0.01$).

Because the year clusters within a regime are few, I also judge each gap by a randomization p-value. The randomization p-value, reported in Tables 11 and OA6, is the share of the 5,000

[30] A formal pre-trend test that stacks the two pre-retrieval regimes finds the change in the gap from R0 to R1 statistically indistinguishable from zero ($p = 0.22$ against I/B/E/S and 0.16 against the equal-weighted benchmark; untabulated).

random splits of the firms (654 high-visibility firms each) whose gap is at least as large as the estimated gap, in its direction. For the continuous measure and the pooled regressions of Section 5.4, it is the same share of 2,000 permutations of the measure across firms. By randomization inference the median split's gap reaches the five-percent level only under native search against I/B/E/S. Against the equal-weighted benchmark the randomization p-value is 0.097 in the browsing regime and 0.108 under native search. Against I/B/E/S it is 0.050 (0.0502 before rounding) and 0.031.

The regime timing also weighs against the residual concern that *LLM Visibility* merely proxies for firm prominence. Because I apply the visibility split to every regime, a prominence account resting on a static firm characteristic would predict a premium gap in the pre-ChatGPT and text-only regimes as well, whereas the gap appears only when retrieval arrives.

5.4 The Continuous Measure and the Change with Retrieval

Table 11, Panel B, tests whether the gap depends on the median split, with *High Visibility* in Equation (8) replaced by the standardized log of one plus the firm's three-query mean citation count. The gap for the continuous measure is expressed per standard deviation. In the browsing regime, the gap is significant by randomization inference against both benchmarks, at 2.90 against the equal-weighted benchmark (randomization p-value 0.003) and 1.74 against I/B/E/S (0.037). Under native search, neither gap reaches the five-percent level. The decision rule for the continuous measure's regime gaps, fixed before the run, requires such a gap in both retrieval regimes against at least one benchmark, so the rule is not met.

Table 11, Panel C, pools the four regimes, with controls and fixed effects common to all of them, and estimates the change in each gap with retrieval (see Table OA6 for details). The gap rises significantly with retrieval under both forms of the measure and against both benchmarks.[31] Against the equal-weighted benchmark, the change is 2.92 under the median split (randomization p-value 0.040) and 1.72 per standard deviation under the continuous measure (0.020). Against I/B/E/S, it is 5.85 under the median split, which none of the 2,000 permutations reaches, and 2.52 under

[31]The decision rule for the change with retrieval, fixed before the run, requires a positive change with a randomization p-value of at most 0.05 against at least one benchmark and no significantly negative change against the other. Both forms of the measure meet this rule.

the continuous measure (0.001). Part of the median split's rise against I/B/E/S reverses a negative pre-retrieval gap. The continuous measure's pre-retrieval gaps are insignificant by randomization inference, so its gap emerges with retrieval.

6 Additional Tests

This section presents two sets of additional tests. The first set asks whether the premium reflects access to more forecasts instead of accuracy-weighted synthesis, in a race against the Estimize consensus. The second set asks which investors perform the synthesis, through partitions by retail share of trading, institutional ownership, and analyst coverage. A GenAI query retrieves forecasts beyond the five providers, such as the Estimize consensus, so the market after ChatGPT may use more sources instead of weighting the five providers by accuracy. The two explanations make opposite predictions in a race against the Estimize consensus itself, a free consensus of many contributors outside the five providers. Table OA4, Panel C, reports the race under native search and its change from the pre-ChatGPT regime. Under native search, the premium over the Estimize consensus is 5.07 ($t = 2.38$), an increase of 5.16 ($t = 2.81$) from the pre-ChatGPT regime. Cheaper access to more forecasts does not explain the premium against I/B/E/S.[32]

The second set partitions the firms by investor group. If one investor group performs the synthesis, the premium should concentrate in the firms that this group holds or trades. Table OA4, Panel E, reports the high-minus-low difference in the post-ChatGPT change in the premium in median splits of the retail share of trading, institutional ownership, and analyst coverage. Against the equal-weighted benchmark, the change in the premium is larger in firms that fewer analysts follow. It does not differ significantly between high- and low-retail-share firms or between high- and low-institutional-ownership firms. The premium is therefore consistent with synthesis by whichever investors set announcement prices. The field evidence on professionals' use of AI tools also predicts these two nulls. In interviews and surveys, sell-side analysts describe adopting AI mainly to improve

[32] Jame et al. (2016) find Estimize forecasts informative about earnings. The test has two limits. First, the Estimize firm-quarters are 58 percent of the sample and belong to larger, more closely followed firms. Second, on that subsample only the premium against I/B/E/S rises significantly from the pre-ChatGPT regime to native search, so the test bears on the I/B/E/S comparison.

their efficiency, with the tools absorbing routine processing (Christ et al., 2025; Kimbrough et al., 2026). On the buy side, nearly all fund managers in an industry survey report using GenAI in their work, mostly for summarizing documents, general research, and writing (AIMA, 2025). In a survey of hedge fund managers, most funds use GenAI tools for investment purposes, and almost all of the adopters use them to process data and text such as earnings calls (Sheng et al., 2026).

7 Conclusion

How the market combines forecasts from competing providers determines the consensus against which it evaluates earnings. I find that, after ChatGPT's release, the market places greater relative weight on more accurate providers while continuing to respond to the other providers' forecasts. Where the relative accuracy of I/B/E/S declines, both the market's response to its surprise and sell-side analysts' herding toward the I/B/E/S consensus fall relative to firms in which I/B/E/S holds its rank. Announcement returns also become more closely aligned with an accuracy-weighted synthesis relative to the I/B/E/S and equal-weighted benchmarks.

The timing of the accuracy premium suggests that retrieval contributes to this shift. The premium increases in the browsing and native-search regimes, with the largest point estimates under native search. It is also larger in firm-quarters with large special items and, against I/B/E/S, in firm-quarters with greater disagreement across providers. These patterns are consistent with lower retrieval costs making synthesis more accessible, particularly where distinguishing among providers offers greater benefits.

Furthermore, the shift varies with *LLM Visibility*. The accuracy premium concentrates in the firms for which a browsing model draws more on the web when asked for their consensus forecasts, and it does so only once the model can retrieve. I measure visibility with one model at one point in time. The evidence does not isolate retrieval's causal effect. Instead, it locates the market's synthesis in the firms whose forecasts a browsing model retrieves and in the browsing and native-search regimes.

The paper contributes in three ways. First, I extend research on FDPs and consensus construction (Bochkay et al., 2022; Larocque et al., 2026) by showing that the market's reliance on

competing providers changes with the information environment. The question of which forecast best represents the market's expectation (Brown, 1991; O'Brien, 1988) therefore also involves a choice among providers. How closely a widely used provider's consensus tracks that expectation may change with information technology. Second, for research on the capital-market consequences of GenAI (e.g., Bertomeu et al., 2025; Cheng et al., 2025; Liu and Subasi, 2026), I provide evidence consistent with accuracy-weighted synthesis across providers as a channel linking GenAI to prices. I also extend the disclosure-processing-cost framework of Blankespoor et al. (2020) to the recurring task of retrieving competing forecasts, assessing their accuracy, and combining them. Third, for research on the information environment and capital markets (e.g., Beyer et al., 2010; Bushee et al., 2010), I introduce *LLM Visibility*, a firm-level measure of how visible a firm's consensus forecasts are to LLMs. Such studies typically focus on what investors themselves can access. *LLM Visibility* extends this view to what the models that investors query can retrieve. The concentration of the premium in high-visibility firms suggests that prices also reflect what these models retrieve.

Three limitations warrant consideration. First, I measure *LLM Visibility* after the sample period, and changes in model behavior over time can limit its correspondence with each firm's visibility to the model during that period (Chen et al., 2024). Second, the GPT-capability regimes coincide with changes in other GenAI tools and the broader information environment, and variation in GenAI use across investors cannot be measured, because usage data at the individual level are unavailable. The investor-group partitions therefore cannot identify which investors use GenAI for the synthesis. Together, these limitations prevent me from claiming that GenAI causes the shift. Finally, the evidence does not establish that investors themselves perform the synthesis. Any participant in the dissemination chain who combines providers' forecasts, such as the financial media, could produce the same pricing pattern. Where the free websites that the browsing model cites disclose a source, they name Zacks or S&P Global Market Intelligence, the publisher of the Capital IQ consensus. Most citations, however, go to sites that disclose no source, so the participants who supply their figures cannot be identified.

References

Aggarwal, P., V. Murahari, T. Rajpurohit, A. Kalyan, K. Narasimhan, and A. Deshpande. 2024. GEO: Generative engine optimization. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 5–16. Association for Computing Machinery.

Ahlers, D. and J. Lakonishok. 1983. A study of economists' consensus forecasts. *Management Science* 29 (10): 1113–1125.

AIMA. 2025. Charting the course: Lessons from AI leaders in alternative investment. Alternative Investment Management Association (AIMA) report, September 2025. Available at: https://www.aima.org/compass/insights/artificial-intelligence-ai/lessons-from-ai-leaders-in-alternative-investment.html.

Bartov, E., D. Givoly, and C. Hayn. 2002. The rewards to meeting or beating earnings expectations. *Journal of Accounting and Economics* 33 (2): 173–204.

Bates, J. M. and C. W. J. Granger. 1969. The combination of forecasts. *Operational Research Quarterly* 20 (4): 451–468.

Battalio, R. H. and R. R. Mendenhall. 2005. Earnings expectations, investor trade size, and anomalous returns around earnings announcements. *Journal of Financial Economics* 77 (2): 289–319.

Bentley, J. W., T. E. Christensen, K. H. Gee, and B. C. Whipple. 2018. Disentangling managers' and analysts' non-GAAP reporting. *Journal of Accounting Research* 56 (4): 1039–1081.

Bernard, V. L. and J. K. Thomas. 1989. Post-earnings-announcement drift: Delayed price response or risk premium? *Journal of Accounting Research* 27 (Supplement): 1–36.

Bertomeu, J., Y. Lin, Y. Liu, and Z. Ni. 2025. The impact of generative AI on information processing: Evidence from the ban of ChatGPT in Italy. *Journal of Accounting and Economics* 80 (1): 101782.

Beyer, A., D. A. Cohen, T. Z. Lys, and B. R. Walther. 2010. The financial reporting environment: Review of the recent literature. *Journal of Accounting and Economics* 50 (2–3): 296–343.

Blankespoor, E. 2019. The impact of information processing costs on firm disclosure choice: Evidence from the XBRL mandate. *Journal of Accounting Research* 57 (4): 919–967.

Blankespoor, E., G. S. Miller, and H. D. White. 2014. The role of dissemination in market liquidity: Evidence from firms' use of Twitter. *The Accounting Review* 89 (1): 79–112.

Blankespoor, E., E. deHaan, and C. Zhu. 2018. Capital market effects of media synthesis and dissemination: Evidence from robo-journalism. *Review of Accounting Studies* 23 (1): 1–36.

Blankespoor, E., E. deHaan, and I. Marinovic. 2020. Disclosure processing costs, investors' information choice, and equity market outcomes: A review. *Journal of Accounting and Economics* 70 (2–3): 101344.

Blankespoor, E., J. Croom, and S. M. Grant. 2026. Generative AI and investor processing of financial information. *Journal of Accounting and Economics* 82 (2): 101908.

Bochkay, K., S. Markov, M. Subasi, and E. H. Weisbrod. 2022. The roles of data providers and analysts in the production, dissemination, and pricing of street earnings. *Journal of Accounting Research* 60 (5): 1695–1740.

Boehmer, E., C. M. Jones, X. Zhang, and X. Zhang. 2021. Tracking retail investor activity. *Journal of Finance* 76 (5): 2249–2305.

Bradshaw, M. T. and R. G. Sloan. 2002. GAAP versus the street: An empirical assessment of two alternative definitions of earnings. *Journal of Accounting Research* 40 (1): 41–66.

Bradshaw, M. T., Y. Ertimur, and P. C. O'Brien. 2017. Financial analysts and their contribution to well-functioning capital markets. *Foundations and Trends in Accounting* 11 (3): 119–191.

Bradshaw, M. T., C. Ma, B. P. Yost, and Y. Zou. 2026. Generative AI use by capital market information intermediaries: Evidence from Seeking Alpha. *Journal of Accounting Research* 64 (3): 1233–1286.

Brown, L. D. 1991. Forecast selection when all forecasts are not equally recent. *International Journal of Forecasting* 7 (3): 349–356.

Brown, L. D. and S. A. Larocque. 2013. I/B/E/S reported actual EPS and analysts' inferred actual EPS. *The Accounting Review* 88 (3): 853–880.

Brown, L. D., P. A. Griffin, R. L. Hagerman, and M. E. Zmijewski. 1987. An evaluation of alternative proxies for the market's assessment of unexpected earnings. *Journal of Accounting and Economics* 9 (2): 159–193.

Brown, T., B. Mann, N. Ryder, M. Subbiah, et al. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc.

Brynjolfsson, E., D. Li, and L. R. Raymond. 2025. Generative AI at work. *Quarterly Journal of Economics* 140 (2): 889–942.

Burgstahler, D., J. Jiambalvo, and T. Shevlin. 2002. Do stock prices fully reflect the implications of special items for future earnings? *Journal of Accounting Research* 40 (3): 585–612.

Bushee, B. J., D. A. Matsumoto, and G. S. Miller. 2003. Open versus closed conference calls: The determinants and effects of broadening access to disclosure. *Journal of Accounting and Economics* 34 (1–3): 149–180.

Bushee, B. J., J. E. Core, W. Guay, and S. J. W. Hamm. 2010. The role of the business press as an information intermediary. *Journal of Accounting Research* 48 (1): 1–19.

Cao, S. S., W. X. Chen, G. Ma, and S. Srinivasan. 2026. Generative AI in capital markets: Information production, dissemination, and processing. *Journal of Accounting Research* 64 (3): 1427–1450.

CFA Institute. 2024. Employer survey results. CFA Institute survey report on AI and generative AI in the investment industry, August 2024. Available at: https://www.cfainstitute.org/sites/default/files/-/media/documents/survey/AI-disruption-employer-survey-2024.pdf.

Chang, A. Y., X. Dong, X. Martin, and C. Zhou. 2026. AI democratization and trading inequality. *Journal of Accounting Research* 64 (3): 1287–1331.

Chen, L., M. Zaharia, and J. Zou. 2024. How is ChatGPT's behavior changing over time? *Harvard Data Science Review* 6 (2).

Cheng, Q., P. Lin, and Y. Zhao. 2025. Does generative AI facilitate investor trading? Early evidence from ChatGPT outages. *Journal of Accounting and Economics* 80 (2–3): 101821.

Christ, M. H., M. Kim, and M. A. Yip. 2025. Word on the street: Field evidence of artificial intelligence use in sell-side equity research. SSRN Working Paper No. 5071836. Available at: https://ssrn.com/abstract=5071836.

Clemen, R. T. 1989. Combining forecasts: A review and annotated bibliography. *International Journal of Forecasting* 5 (4): 559–583.

Clement, M. B. 1999. Analyst forecast accuracy: Do ability, resources, and portfolio complexity matter? *Journal of Accounting and Economics* 27 (3): 285–303.

Clement, M. B. and S. Y. Tse. 2003. Do investors respond to analysts' forecast revisions as if forecast accuracy is all that matters? *The Accounting Review* 78 (1): 227–249.

Clement, M. B. and S. Y. Tse. 2005. Financial analyst characteristics and herding behavior in forecasting. *Journal of Finance* 60 (1): 307–341.

Cohen, L. and D. Lou. 2012. Complicated firms. *Journal of Financial Economics* 104 (2): 383–400.

Collins, D. W. and S. P. Kothari. 1989. An analysis of intertemporal and cross-sectional determinants of earnings response coefficients. *Journal of Accounting and Economics* 11 (2–3): 143–181.

deHaan, E., T. Shevlin, and J. Thornock. 2015. Market (in)attention and the strategic scheduling and timing of earnings announcements. *Journal of Accounting and Economics* 60 (1): 36–55.

DellaVigna, S. and J. M. Pollet. 2009. Investor inattention and Friday earnings announcements. *Journal of Finance* 64 (2): 709–749.

Drake, M. S., D. T. Roulstone, and J. R. Thornock. 2012. Investor information demand: Evidence from Google searches around earnings announcements. *Journal of Accounting Research* 50 (4): 1001–1040.

Drake, M. S., D. T. Roulstone, and J. R. Thornock. 2015. The determinants and consequences of information acquisition via EDGAR. *Contemporary Accounting Research* 32 (3): 1128–1161.

Easton, P. D. and M. E. Zmijewski. 1989. Cross-sectional variation in the stock market response to accounting earnings announcements. *Journal of Accounting and Economics* 11 (2–3): 117–141.

Goldstein, I., C. S. Spatt, and M. Ye. 2021. Big data in finance. *Review of Financial Studies* 34 (7): 3213–3225.

Granger, C. W. J. and R. Ramanathan. 1984. Improved methods of combining forecasts. *Journal of Forecasting* 3 (2): 197–204.

Gu, S., B. Kelly, and D. Xiu. 2020. Empirical asset pricing via machine learning. *Review of Financial Studies* 33 (5): 2223–2273.

Hirshleifer, D., S. S. Lim, and S. H. Teoh. 2009. Driven to distraction: Extraneous events and underreaction to earnings news. *Journal of Finance* 64 (5): 2289–2325.

Holthausen, R. W. and R. E. Verrecchia. 1988. The effect of sequential information releases on the variance of price changes in an intertemporal multi-asset market. *Journal of Accounting Research* 26 (1): 82–106.

Jame, R., R. Johnston, S. Markov, and M. C. Wolfe. 2016. The value of crowdsourced earnings forecasts. *Journal of Accounting Research* 54 (4): 1077–1110.

Kim, O. and R. E. Verrecchia. 1991. Trading volume and price reactions to public announcements. *Journal of Accounting Research* 29 (2): 302–321.

Kimbrough, M. D. 2005. The effect of conference calls on analyst and market underreaction to earnings announcements. *The Accounting Review* 80 (1): 189–219.

Kimbrough, M. D., L. Y. Liu, and M. Subasi. 2026. AI-powered analysts. SSRN Working Paper No. 7176461. Available at: https://ssrn.com/abstract=7176461.

Kothari, S. P. 2001. Capital markets research in accounting. *Journal of Accounting and Economics* 31 (1–3): 105–231.

Larocque, S. A., J. Watkins, and E. H. Weisbrod. 2026. Consensus? An examination of differences in earnings information across forecast data providers. *Journal of Accounting Research* (forthcoming).

Lawrence, A. 2013. Individual investors and financial disclosure. *Journal of Accounting and Economics* 56 (1): 130–147.

Lehavy, R., F. Li, and K. Merkley. 2011. The effect of annual report readability on analyst following and the properties of their earnings forecasts. *The Accounting Review* 86 (3): 1087–1115.

Lewis, P., E. Perez, A. Piktus, F. Petroni, et al. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474. Curran Associates, Inc.

Li, E. X., M. Shen, Z. Tu, and D. Zhou. 2025. The promise and peril of generative AI: Evidence from GPT as sell-side analysts. arXiv preprint arXiv:2412.01069v2. Available at: https://arxiv.org/abs/2412.01069.

Liu, L. Y. and M. Subasi. 2026. When does generative AI level the playing field? Evidence from crowdsourced earnings forecasts. SSRN Working Paper No. 7253378. Available at: https://ssrn.com/abstract=7253378.

Liu, N. F., T. Zhang, and P. Liang. 2023. Evaluating verifiability in generative search engines. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 7001–7025. Association for Computational Linguistics.

Livnat, J. and R. R. Mendenhall. 2006. Comparing the post-earnings announcement drift for surprises calculated from analyst and time series forecasts. *Journal of Accounting Research* 44 (1): 177–205.

Lopez-Lira, A. and Y. Tang. 2026. Can ChatGPT forecast stock price movements? Return predictability and large language models. *Journal of Financial Economics* 184: 104335.

Lopez-Lira, A., Y. Tang, and M. Zhu. 2025. The memorization problem: Can we trust LLMs' economic forecasts? SSRN Working Paper No. 5217505. Available at: https://ssrn.com/abstract=5217505.

Loughran, T. and B. McDonald. 2017. The use of EDGAR filings by investors. *Journal of Behavioral Finance* 18 (2): 231–248.

Michaely, R., A. Rubin, D. Segal, and A. Vedrashko. 2024. Do differences in analyst quality matter for investors relying on consensus information? *Management Science* 70 (2): 751–772.

Nakano, R., J. Hilton, S. Balaji, J. Wu, et al. 2021. WebGPT: Browser-assisted question-answering with human feedback. arXiv preprint arXiv:2112.09332. Available at: https://arxiv.org/abs/2112.09332.

Noy, S. and W. Zhang. 2023. Experimental evidence on the productivity effects of generative artificial intelligence. *Science* 381 (6654): 187–192.

O'Brien, P. C. 1988. Analysts' forecasts as earnings expectations. *Journal of Accounting and Economics* 10 (1): 53–83.

OpenAI. 2023. ChatGPT plugins. OpenAI blog post, March 23, 2023. Available at: https://openai.com/blog/chatgpt-plugins.

Park, C. W. and E. K. Stice. 2000. Analyst forecasting ability and the stock price reaction to forecast revisions. *Review of Accounting Studies* 5 (3): 259–272.

Sheng, J., Z. Sun, B. Yang, and A. L. Zhang. 2026. Generative AI and asset management. *Review of Financial Studies* (forthcoming).

Stickel, S. E. 1992. Reputation and performance among security analysts. *Journal of Finance* 47 (5): 1811–1836.

Stock, J. H. and M. W. Watson. 2004. Combination forecasts of output growth in a seven-country data set. *Journal of Forecasting* 23 (6): 405–430.

Timmermann, A. 2006. Forecast combinations. In Elliott, G., C. W. J. Granger, and A. Timmermann, editors, *Handbook of Economic Forecasting*, volume 1, chapter 4, pages 135–196. Elsevier.

Walther, B. R. 1997. Investor sophistication and market earnings expectations. *Journal of Accounting Research* 35 (2): 157–179.

Xue, J., Q. Zhang, and W. Zhu. 2025. Generative AI for analysts. SSRN Working Paper No. 5821142. Available at: https://ssrn.com/abstract=5821142.

Appendix A: Variable Definitions

Variable	Definition	Source
Panel A. Across-FDP firm-quarter outcome and surprise variables		
$CAR^{m}_{i,q,[-1,1]}$, $BHAR^{m}_{i,q,[-2,1]}$	For the adjustment m, market (*Mkt*: the daily return $r_{i,t}$ minus r^{M}_{t}, the CRSP value-weighted index return on day t) or Fama–French–Carhart four-factor (*FFC*: the market, size, value, and momentum factors, with loadings estimated over $[-280,-32]$ from at least 60 daily returns), the three-day cumulative abnormal return over $[-1,1]$ (*CAR*) and the four-day buy-and-hold abnormal return over $[-2,1]$ (*BHAR*) around firm i's quarter-q announcement, in percent.	CRSP
$SUE^{g}_{i,q}$	For provider g, one of I/B/E/S (*IBES*), FactSet, Bloomberg, S&P Capital IQ (*CapIQ*), and Zacks, the standardized unexpected earnings $(Actual^{g}_{i,q} - Consensus^{g}_{i,q})/p_{i,q}$, in percent of price, where $Consensus^{g}_{i,q}$ is g's last consensus EPS for quarter q before the announcement, $Actual^{g}_{i,q}$ its own reported actual EPS for the same quarter, and $p_{i,q}$ the closing share price two trading days before the announcement.	I/B/E/S; FactSet; Bloomberg; S&P Capital IQ; Zacks; CRSP
$SUE^{Acc}_{i,q}$	Accuracy-weighted SUE, $\sum_g w^{g}_{i,q} SUE^{g}_{i,q}$, with the weights $w^{g}_{i,q}$ of Panel C.	Calculated
$SUE^{X}_{i,q}$	For the benchmark X, equal weighting of the five providers (*EW5*) or I/B/E/S (*IBES*), the benchmark surprise, $SUE^{EW5}_{i,q} = (1/5)\sum_g SUE^{g}_{i,q}$ or I/B/E/S's own $SUE^{IBES}_{i,q}$.	Calculated
Standardized Pred. Surp.$^{X}_{i,q}$	Min-max standardized predicted surprise $(Consensus^{Acc}_{i,q} - Consensus^{X}_{i,q})/p_{i,q}$, where $Consensus^{Acc}_{i,q} = \sum_g w^{g}_{i,q} Consensus^{g}_{i,q}$ and $Consensus^{X}_{i,q}$ is the benchmark's consensus, the equal-weighted five-FDP consensus or I/B/E/S's own.	Calculated
Panel B. Firm and event controls		
Log Market Equity$_{i,q}$	Natural log of firm i's market capitalization, in millions of dollars, two trading days before the announcement.	CRSP
Book-to-Market$_{i,q}$	Book value of equity for the fiscal quarter that ends in calendar quarter q, divided by market capitalization two trading days before the announcement.	Compustat; CRSP
Leverage$_{i,q}$	Long-term debt plus debt in current liabilities, scaled by total assets, for the fiscal quarter that ends in calendar quarter q.	Compustat
Prior 12-Month Return$_{i,q}$	Buy-and-hold raw return over the twelve months ending one month before firm i's quarter-q announcement.	CRSP
Log Price$_{i,q}$	Natural log of share price two trading days before the announcement.	CRSP
Return Volatility$_{i,q}$	Standard deviation of daily returns over the 60 trading days before the announcement.	CRSP
Analyst Dispersion$_{i,q}$	Standard deviation of analyst EPS forecasts in the I/B/E/S summary file, divided by the share price two trading days before the announcement.	I/B/E/S; CRSP

(continued on next page)

Variable	Definition	Source
$Busyness_{i,q}$	Natural log of the number of firms, among those with a consensus from at least one of the five FDPs, that announce earnings on firm i's announcement day.	I/B/E/S; Compustat
$\lvert\Delta IB\rvert_{i,q}$	Absolute change in income before extraordinary items from the previous fiscal quarter, for the fiscal quarter that ends in calendar quarter q, scaled by market capitalization two trading days before the announcement.	Compustat; CRSP
$\lvert\Delta Shares\rvert_{i,q}$	Absolute change in shares outstanding from the previous fiscal quarter, scaled by the previous quarter's shares outstanding, for the fiscal quarter that ends in calendar quarter q.	Compustat
$\lvert SpecItems\rvert_{i,q}$	Absolute value of special items for the fiscal quarter that ends in calendar quarter q, scaled by market capitalization two trading days before the announcement.	Compustat; CRSP
$SBC\ High_{i,q}$	Indicator equal to one if stock-based compensation (SBC) expense for calendar quarter q, scaled by market capitalization two trading days before the announcement, exceeds the sample median. The expense is the quarterly figure when reported and one quarter of the annual figure otherwise. A missing expense takes the firm's last reported value, or its first later value when none precedes it.	Compustat; CRSP
$Unexpected\ Item_{i,q}$	Indicator equal to one if Compustat special items for the fiscal quarter that ends in calendar quarter q are nonzero.	Compustat
$Stock\ Split_{i,q}$	Indicator equal to one if a stock split's ex-date falls in calendar quarter q.	CRSP
$Fiscal\ Q4_{i,q}$	Indicator equal to one if the fiscal quarter that ends in calendar quarter q is the firm's fourth fiscal quarter.	Compustat
$Institutional\ Ownership_{i,q}$	Shares held by 13F institutions at the end of calendar quarter q, divided by shares outstanding at the end of that quarter and capped at one. The last available value carries forward within the firm.	Thomson Reuters 13F (WRDS); CRSP
$\#\ Analysts_{i,q}$	Number of EPS estimates in the I/B/E/S summary consensus for quarter q at the last statistical period on or before the announcement. It enters the regressions as the log of one plus the count.	I/B/E/S
Panel C. Cross-sectional partitioning and identification variables		
$LLM\ Visibility_i$	Number of source citations, repeats included, in the response of ChatGPT (GPT-4o with web search) to a query for firm i's next-quarter consensus EPS, averaged over three queries issued on May 28 and June 1, 2026.	DataForSEO
$Concurrent\ Cross\text{-}FDP\ SUE\ Dispersion_{i,q}$	Interquartile range of $SUE^{g}_{i,q}$ across the five FDPs in firm-quarter (i,q).	Calculated
$IBES\,Up_i$, $IBES\,Down_i$	Indicators equal to one if I/B/E/S's rounded average per-firm-quarter accuracy rank among the five FDPs rises (falls) by at least two positions from pre- to post-ChatGPT. Firms observed only before or only after the release are in neither group.	Calculated
Large Unexpected Item ($LUI_{i,q}$)	Indicator equal to one if firm-quarter (i,q) contains a large unexpected item: any of seven Compustat nonrecurring items mapped per Bochkay et al. (2022), with restructuring, acquisition, and special items requiring top-decile magnitude.	Compustat

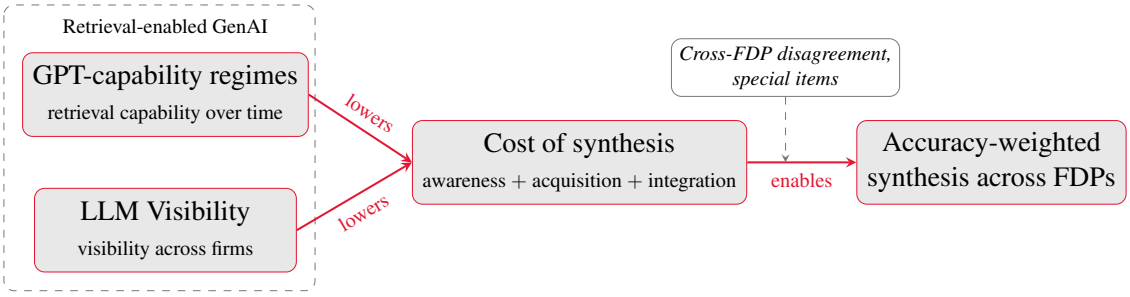
(continued on next page)

Variable	Definition	Source
$Post_q$, $Post_q^{\text{Browse}}$, regimes R0–R3	Announcement-date indicators: $Post_q = 1$ on or after November 30, 2022 (ChatGPT's public release); $Post_q^{\text{Browse}} = 1$ on or after September 27, 2023 (the restoration of browsing). The GPT-capability regimes partition the same dates: R0, Pre-ChatGPT (before November 30, 2022); R1, Text-only (to September 26, 2023); R2, Browsing (to October 30, 2024); R3, Native search (October 31, 2024, to December 31, 2025, the sample end).	OpenAI
$MAFE_{i,q}^{g}$	Provider g's trailing mean absolute forecast error, the average of $\lvert Actual_{i,q-k}^{g} - Consensus_{i,q-k}^{g}\rvert / p_{i,q-k}$ over those of the four preceding quarters ($k = 1,\ldots,4$) with an observed error.	Calculated
$w_{i,q}^{g}$	Accuracy weight on provider g: the softmax of the within-firm-quarter z-score of $-MAFE_{i,q}^{g}$ across the five FDPs (Section 3.5).	Calculated
$Rank_{i,q}^{g}$	Provider g's within-firm-quarter accuracy rank by $MAFE_{i,q}^{g}$ (1 = least accurate, 5 = most accurate).	Calculated
$Predictive\ Value_{i,q}^{IBES}$	One minus the ratio of the summed absolute residuals of the backward regression of the year-ahead I/B/E/S street actual on the I/B/E/S consensus over up to the twenty most recent source quarters to the summed absolute year-ahead actuals over the same quarters, winsorized at the 1st and 99th percentiles (Online Appendix C.3); defined for 22,364 firm-quarters.	I/B/E/S
$High\ c_{i,q}$, $High\ c_i$	For a characteristic c, an indicator equal to one when c exceeds its median, with ties assigned to Low. For the firm-quarter characteristics *Concurrent Cross-FDP SUE Dispersion* and *Special-Items Magnitude* ($\lvert SpecItems\rvert$, *High Special Items*), the median is taken across firm-quarters (i,q). For the firm characteristics, it is taken across firms (i): *LLM Visibility* (*High Visibility*); in Table 9, firm size (*Log Market Equity*), analyst coverage (*# Analysts*), and *Institutional Ownership*; and in Section 6, *Retail Share of Trading*, the firm's mean *Institutional Ownership*, and its mean *# Analysts* over its announcements before November 30, 2022, firms without such an announcement left unclassified.	Calculated
$Retail\ Share\ of\ Trading_i$	Retail dollar volume (buys plus sells) divided by total dollar volume, each summed over firm i's trading days from January 4, 2021, to November 29, 2022, before ChatGPT's release. Retail trades are those that the WRDS TAQ Intraday Indicators classify as retail by the sub-penny rule of Boehmer et al. (2021).	WRDS TAQ Intraday Indicators; CRSP
$Reported\text{-}Actual\ Disagreement_{i,q}^{gh}$	For a pair of providers g and h, an indicator equal to one if their reported actuals for firm-quarter (i,q) are not identical (Table OA3, Panels A and B).	FactSet; Bloomberg; Zacks; S&P Capital IQ; I/B/E/S
Panel D. Analyst revision-level variables (Section 4.1.2)		
$Herding_{a,i,q,t}$	Indicator equal to one when the revised forecast falls between the analyst's own prior forecast and the prevailing I/B/E/S consensus.	I/B/E/S

(continued on next page)

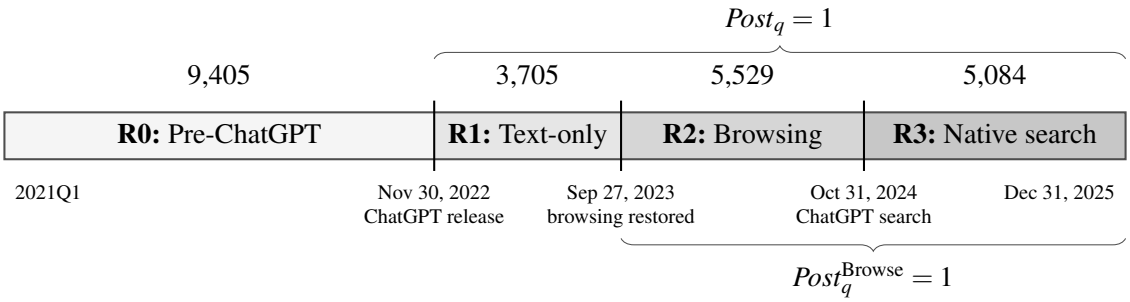
Variable	Definition	Source
Prevailing I/B/E/S Consensus$_{i,q,t}$	Mean estimate from the latest I/B/E/S statistical-period summary strictly preceding the revision date, at most 45 days old and based on at least three estimates.	I/B/E/S
Consensus Distance$_{a,i,q,t}$	Absolute distance between the analyst's prior forecast and the prevailing consensus, scaled by the closing share price on the last trading day before the announcement, in percent; enters the regressions as twenty ventile fixed effects.	I/B/E/S; CRSP
General (Firm) Experience$_{a,t}$	Number of years since the analyst's first I/B/E/S forecast (first forecast for firm *i*).	I/B/E/S
Firms (Industries) Covered$_{a,t}$	Number of distinct firms (two-digit SIC industries) for which the analyst issues forecasts in the revision year.	I/B/E/S; CRSP
Brokerage Size$_{a,t}$	Number of analysts issuing forecasts at the analyst's brokerage in the revision year.	I/B/E/S
Forecast Horizon$_{a,i,q,t}$	Number of days from the revision date to the earnings announcement. Experience, coverage, brokerage size, and horizon enter the regressions as the log of one plus the variable.	I/B/E/S

FIGURE 1. Theoretical Framework



Notes. This figure summarizes the theoretical framework. Retrieval-enabled generative AI (GenAI) varies along two dimensions. Over time, its retrieval capability grows across the GPT-capability regimes, from text-only to browsing and native search. Across firms, some firms' consensus forecasts are more visible to LLMs than others'. *LLM Visibility* proxies for this visibility with the source citations in GPT's answers, measured after the sample period with one model. Greater retrieval capability and greater visibility lower the cost of synthesizing analyst consensus forecasts across the five forecast data providers (FDPs), the joint cost of awareness, acquisition, and integration in the Blankespoor et al. (2020) decomposition. Lower costs enable the market to price an accuracy-weighted synthesis across providers in place of a single-FDP or equal-weighted benchmark. The strength of the channel is moderated by how valuable synthesis is for a given firm-quarter (cross-FDP disagreement and special-items magnitude). These mechanisms map to Predictions 1a, 1b, and 2, developed in Section 2. Prediction 1a's quality dimension covers both trailing accuracy and the consensus's predictive value for the firm's future results.

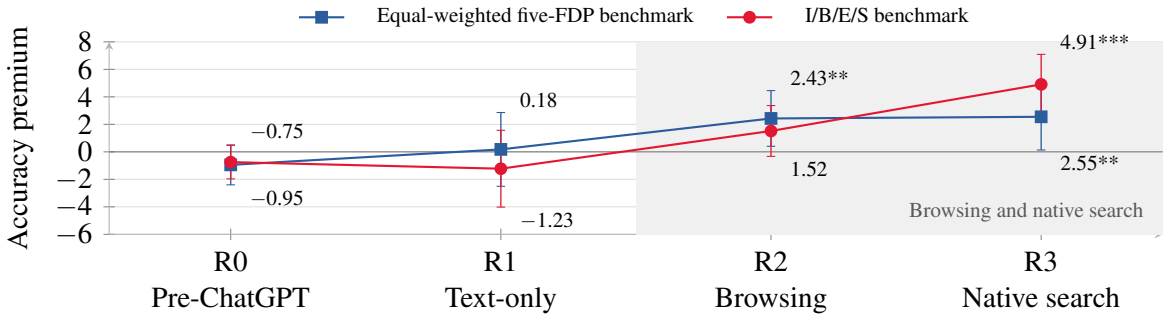
FIGURE 2. The GPT-Capability Regimes



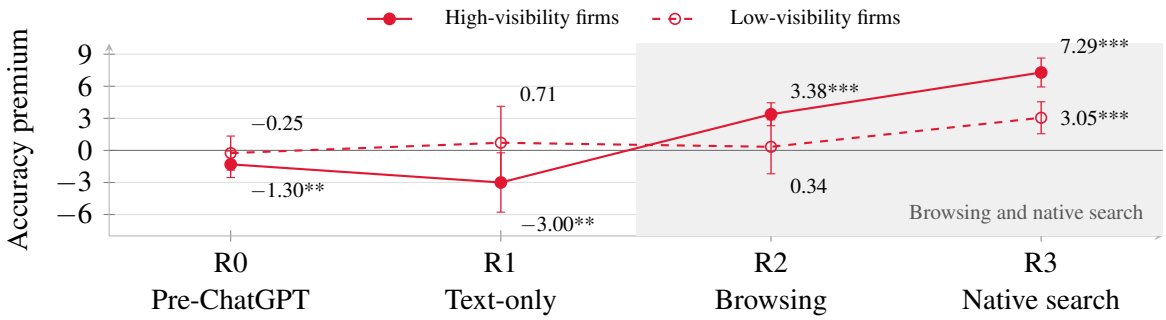
Notes. The bar divides the sample period into the four GPT-capability regimes at milestones in ChatGPT's publicly available retrieval capability. The number above each regime is the number of firm-quarters in that regime's horse-race regression (Table 7). The post-ChatGPT indicator $Post_q$ equals one on or after November 30, 2022, and the post-browsing indicator $Post_q^{\text{Browse}}$ equals one on or after September 27, 2023. R1 also contains a temporary ChatGPT browsing beta, from May to July 3, 2023.

FIGURE 3. The Accuracy Premium Across GPT-Capability Regimes, All Firms and by LLM Visibility

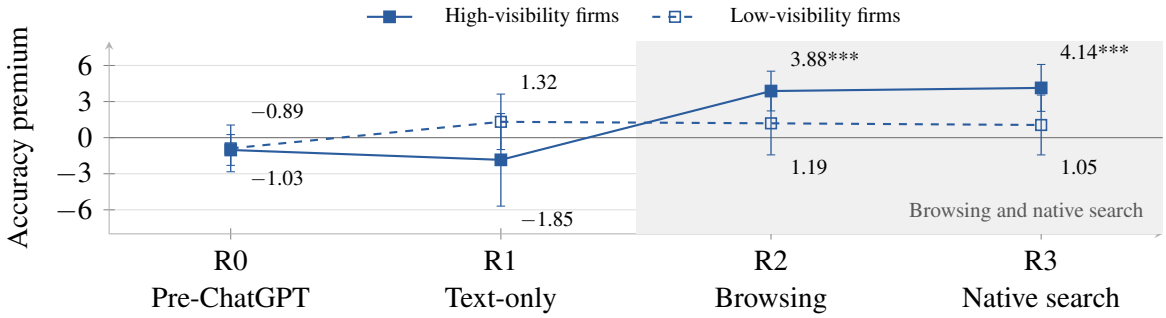
(a) All firms, by benchmark



(b) High- against low-visibility firms, I/B/E/S benchmark

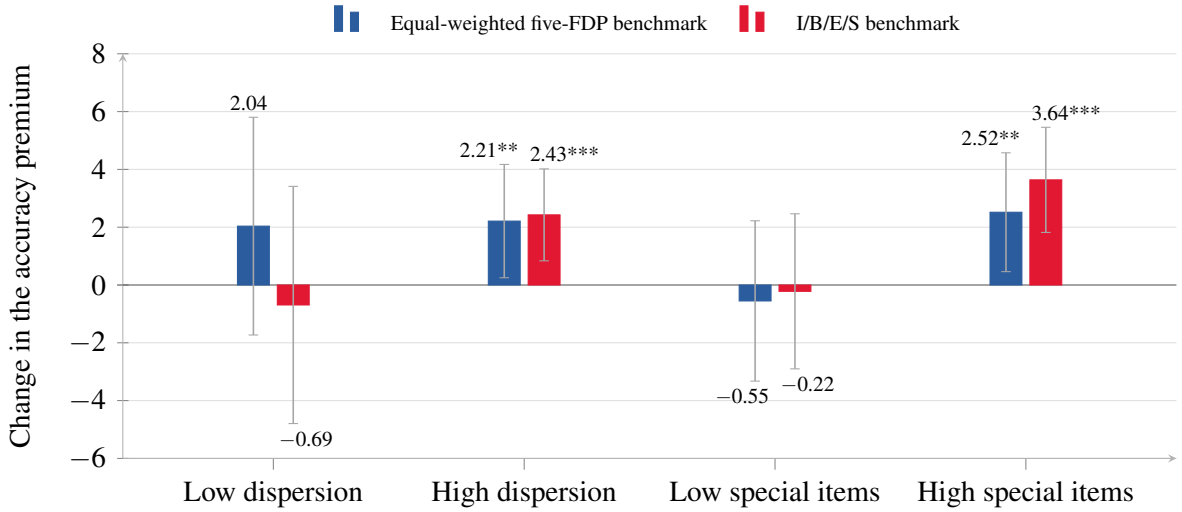


(c) High- against low-visibility firms, equal-weighted five-FDP benchmark



Notes. This figure plots the accuracy premium, the contrast $\beta_1 - \beta_2$ between the loading of the three-day market-adjusted $CAR_{[-1,1]}$ (in percent) on the accuracy-weighted surprise SUE^{Acc} and its loading on the benchmark surprise (both in percent of price), estimated separately within each GPT-capability regime (R0 to R3, defined in Appendix A). Panel (a) plots the estimates of Table 7 for all firms, against the I/B/E/S benchmark (red circles) and the equal-weighted five-FDP benchmark (blue squares). Panels (b) and (c) plot the estimates of Table 10 for high-visibility firms (solid lines, filled markers) and low-visibility firms (dashed lines, open markers), against the I/B/E/S benchmark in Panel (b) and the equal-weighted five-FDP benchmark in Panel (c). In Panels (b) and (c) the vertical distance between the two lines is the high-minus-low gap $\beta_3 - \beta_4$ of Table 10. The shaded band marks the regimes in which the model can retrieve forecasts from the web. Whiskers are 95 percent confidence intervals computed from the t-statistics of the two tables (firm-clustered in Table 7, two-way clustered by firm and fiscal year in Table 10), and each label reports the estimate with its significance ($^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$).

FIGURE 4. The Change in the Accuracy Premium by Concurrent Cross-FDP SUE Dispersion and by Special-Items Magnitude



Notes. This figure plots the post-browsing change in the accuracy premium, $\beta_3 - \beta_4$ of Equation (7) with the post-browsing indicator in place of $Post_q$, estimated within the below- and above-median halves of the two partitions of Table 8: Panel A on Concurrent Cross-FDP SUE Dispersion (the first two groups) and Panel B on Special-Items Magnitude (the last two). The blue bars use the equal-weighted five-FDP benchmark and the red bars the I/B/E/S benchmark. The bar labels report the change, in percentage points of announcement return per percentage point of surprise, with its significance ($^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$), and the whiskers are 95 percent confidence intervals computed from the firm-clustered t-statistics of Table 8.

TABLE 1. Sample Selection

Sample construction (firm × quarter)	**Observations**	**# Firms**
S&P 1500 firm-quarters with an earnings announcement (announcement date 2021–2025)	29,395	1,502
Retain firm-quarters with non-missing consensus from all five FDPs (FactSet, Bloomberg, Zacks, Capital IQ, and I/B/E/S)	25,027	1,353
Drop firm-quarters missing market-outcome variables, the accuracy-weighted surprise, or required controls	(1,179)	(13)
Drop firm-quarters missing any constructed return or predicted-surprise variable	(10)	(1)
Drop firms with only one firm-quarter (no within-firm variation for firm fixed effects)	(5)	(5)
Final Sample	**23,833**	**1,334**

Notes. This table reports the selection of the firm-quarter sample. Counts in parentheses are the firm-quarters (firms) dropped at each step. The other rows report the firm-quarters (firms) that remain.

TABLE 2. Descriptive Statistics

	N	Mean	SD	Min	p25	Median	p75	Max
1. Dependent Variables: Abnormal Returns (%)								
$CAR^{Mkt}_{[-1,1]}$	23,833	0.204	7.739	-22.532	-3.933	0.168	4.255	23.535
$CAR^{FFC}_{[-1,1]}$	23,833	0.196	7.635	-22.779	-3.776	0.186	4.135	23.606
$BHAR^{Mkt}_{[-2,1]}$	23,833	0.228	7.978	-22.803	-4.165	0.177	4.418	24.992
$BHAR^{FFC}_{[-2,1]}$	23,833	0.193	7.847	-22.935	-3.966	0.111	4.255	24.819
2. Per-FDP Standardized Earnings Surprise (SUE^g, %)								
$SUE^{FactSet}$	23,833	0.132	0.534	-2.190	0.000	0.074	0.238	2.484
$SUE^{Bloomberg}$	23,833	-0.222	1.659	-10.049	-0.362	-0.027	0.162	5.780
SUE^{Zacks}	23,833	0.149	0.467	-1.606	0.000	0.075	0.240	2.291
SUE^{CapIQ}	23,833	0.137	0.530	-2.185	0.001	0.076	0.243	2.477
SUE^{IBES}	23,833	0.130	0.548	-2.271	0.000	0.075	0.241	2.502
3. Cross-FDP Composites and Predicted Surprise								
SUE^{Acc}	23,833	0.110	0.450	-1.656	-0.015	0.059	0.212	1.972
SUE^{EW5}	23,833	0.066	0.510	-1.936	-0.064	0.041	0.201	1.942
SUE^{Acc4}	23,832	0.134	0.429	-1.448	0.001	0.073	0.231	1.960
SUE^{EW4}	23,833	0.137	0.441	-1.461	0.001	0.076	0.239	1.948
$Standardized\ Pred.Surp.^{EW5}$	23,833	0.606	0.116	0.000	0.590	0.619	0.631	1.000
$Standardized\ Pred.Surp.^{IBES}$	23,833	0.483	0.109	0.000	0.469	0.478	0.493	1.000
4. Firm and Event Controls								
Log Market Equity	23,833	8.992	1.448	6.390	7.923	8.769	9.942	12.883
Book-to-Market	23,833	0.440	0.361	-0.159	0.170	0.362	0.634	1.670
Leverage	23,833	0.314	0.218	0.001	0.133	0.301	0.444	1.008
Prior 12-Month Return	23,833	0.201	0.487	-0.589	-0.101	0.112	0.379	2.400
Log Price	23,833	4.166	0.980	1.880	3.491	4.178	4.831	6.580
Return Volatility	23,833	0.021	0.009	0.009	0.015	0.019	0.025	0.054
Analyst Dispersion	23,833	0.002	0.003	0.000	0.000	0.001	0.002	0.018
Busyness	23,833	4.944	0.973	1.946	4.431	5.124	5.684	6.299
$\lvert \Delta IB \rvert$	23,833	0.012	0.026	0.000	0.001	0.004	0.011	0.178
$\lvert \Delta Shares \rvert$	23,833	0.010	0.021	0.000	0.001	0.003	0.009	0.150
$\lvert SpecItems \rvert$	23,833	0.003	0.010	0.000	0.000	0.000	0.002	0.074
SBC High	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
Unexpected Item	23,833	0.686	0.464	0.000	0.000	1.000	1.000	1.000
Stock Split	23,833	0.003	0.051	0.000	0.000	0.000	0.000	1.000
Fiscal Q4	23,833	0.250	0.433	0.000	0.000	0.000	0.000	1.000
Institutional Ownership	23,833	0.850	0.238	0.011	0.813	0.940	1.000	1.000
# Analysts	23,833	10.914	6.536	2.000	6.000	10.000	15.000	31.000
5. Identification, Weighting, and Partitioning Variables								
$Predictive\ Value^{IBES}$	22,364	0.666	0.252	-0.036	0.558	0.756	0.851	0.947
$Rank^g$	119,165	3.000	1.414	1.000	2.000	3.000	4.000	5.000

(continued on next page)

	N	Mean	SD	Min	p25	Median	p75	Max
Concurrent Cross-FDP SUE Dispersion	23,833	0.108	0.261	0.000	0.007	0.024	0.081	1.804
Log # Analysts	23,833	2.322	0.573	1.099	1.946	2.398	2.773	3.466
LLM Visibility	23,827	13.257	2.985	7.667	11.000	13.000	15.000	21.333
Post-ChatGPT (Post)	23,833	0.604	0.489	0.000	0.000	1.000	1.000	1.000
Post-Browsing ($Post^{Browse}$)	23,833	0.448	0.497	0.000	0.000	0.000	1.000	1.000
R1 (Text-only)	23,833	0.157	0.364	0.000	0.000	0.000	0.000	1.000
R2 (Browsing)	23,833	0.233	0.423	0.000	0.000	0.000	0.000	1.000
R3 (Native search)	23,833	0.215	0.411	0.000	0.000	0.000	0.000	1.000
6. Cross-Sectional Partitioning Indicators								
High Visibility	23,827	0.489	0.500	0.000	0.000	0.000	1.000	1.000
High Concurrent Cross-FDP SUE Dispersion	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
High Special Items	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
High Firm Size (Table 9)	23,044	0.516	0.500	0.000	0.000	1.000	1.000	1.000
High Analyst Coverage (Table 9)	23,044	0.515	0.500	0.000	0.000	1.000	1.000	1.000
High Inst. Ownership (Table 9)	23,044	0.503	0.500	0.000	0.000	1.000	1.000	1.000
IBES Accuracy-Rank Up (IBES Up)	23,833	0.096	0.295	0.000	0.000	0.000	0.000	1.000
IBES Accuracy-Rank Down (IBES Down)	23,833	0.066	0.248	0.000	0.000	0.000	0.000	1.000
7. Analyst Revision Panel (Section 4.1.2)								
Herding	122,001	0.323	0.468	0.000	0.000	0.000	1.000	1.000
Consensus Distance (% of price)	122,001	0.227	0.379	0.004	0.031	0.084	0.238	2.321
General Experience (years)	122,001	16.025	11.341	0.416	6.940	13.317	22.563	41.462
Firm Experience (years)	122,001	7.281	6.822	0.079	1.996	5.183	10.691	33.435
# Firms Covered	122,001	23.255	8.999	4.000	17.000	22.000	28.000	51.000
# Industries Covered	122,001	9.357	4.455	2.000	6.000	9.000	12.000	21.000
Brokerage Size	122,001	81.793	65.585	2.000	31.000	65.000	114.000	251.000
Forecast Horizon (days)	122,001	69.957	38.847	2.000	29.000	85.000	91.000	162.000
8. Additional Tests (Section 6)								
SUE^{Est}	13,830	0.072	0.425	-1.659	-0.033	0.030	0.152	1.975
Retail Share of Trading	23,708	0.051	0.027	0.021	0.035	0.043	0.059	0.174
High Retail Share of Trading	23,476	0.497	0.500	0.000	0.000	0.000	1.000	1.000
High Inst. Ownership (pre-ChatGPT)	23,515	0.499	0.500	0.000	0.000	0.000	1.000	1.000
High # Analysts (pre-ChatGPT)	23,515	0.510	0.500	0.000	0.000	1.000	1.000	1.000

Notes. This table reports descriptive statistics for the sample of S&P 1500 firm-quarters with announcement dates in 2021–2025. *LLM Visibility* and *High Visibility* are defined for 23,827 of the 23,833 firm-quarters. *# Analysts* enters the regressions as *Log # Analysts*, the log of one plus the count. $Rank^g$ is summarized over the provider-firm-quarters of Table 3, Panel A, and the high-group indicators of Table 9 over that table's horse-race sample. Category 8 reports each variable on the firm-quarters for which it is defined. Continuous variables are winsorized at the 1st and 99th percentiles. Appendix A defines the variables, and the notes to Table OA4 define SUE^{Acc4}, SUE^{EW4}, and SUE^{Est}.

TABLE 3. Earnings Response Coefficients by Provider Accuracy Rank and by Change in I/B/E/S's Accuracy Rank

Panel A. Earnings response coefficients by provider accuracy rank

	Pre-ChatGPT		Post-ChatGPT	
Provider	β	γ	β	γ
I/B/E/S	2.591***	0.022	2.611***	0.295**
	(6.60)	(0.17)	(5.49)	(2.09)
FactSet	2.316***	0.050	2.059***	0.380**
	(4.78)	(0.36)	(3.80)	(2.42)
Bloomberg	0.985***	-0.058	0.789***	0.162**
	(5.82)	(-0.80)	(5.01)	(2.10)
Zacks	3.603***	-0.139	4.757***	-0.020
	(8.35)	(-1.04)	(8.57)	(-0.13)
Capital IQ	1.107**	0.543***	2.483***	0.349**
	(2.40)	(3.94)	(4.94)	(2.40)
Stacked	**1.206*** **	**0.306*** **	**0.973*** **	**0.621*** **
	(9.21)	**(6.93)**	**(8.09)**	**(11.65)**
Observations	47,130		72,035	
Stacked Δ (Post−Pre)			**-0.232**	**0.316*** **
			(-1.38)	**(4.91)**

Panel B. Heterogeneous post-ChatGPT effects by firm-level I/B/E/S accuracy-rank change

Dep. Var.:	Market-adjusted $CAR_{[-1,1]}$		FFC $CAR_{[-1,1]}$	
	(1)	(2)	(3)	(4)
SUE^{IBES}	2.859***	2.845***	2.775***	2.776***
	(13.91)	(14.03)	(13.67)	(13.84)
$SUE^{IBES} \times Post$	1.029***	0.978***	1.136***	1.065***
	(3.56)	(3.32)	(4.11)	(3.76)
$SUE^{IBES} \times IBES\,Up$	0.693	0.628	0.803	0.737
	(1.24)	(1.16)	(1.46)	(1.39)
$SUE^{IBES} \times IBES\,Down$	1.397*	1.582**	1.913***	2.080***
	(1.81)	(2.00)	(2.61)	(2.74)
$\boldsymbol{SUE^{IBES} \times Post \times IBES\,Up}$	**-1.248**	**-1.135**	**-1.291**	**-1.185**
	(-1.44)	**(-1.30)**	**(-1.40)**	**(-1.26)**
$\boldsymbol{SUE^{IBES} \times Post \times IBES\,Down}$	**-2.278****	**-2.344****	**-2.758*****	**-2.791****
	(-2.27)	**(-2.15)**	**(-2.70)**	**(-2.51)**
$Post \times IBES\,Up$	0.353	0.232	0.437	0.297
	(1.13)	(0.68)	(1.37)	(0.85)
$Post \times IBES\,Down$	0.544	0.713	0.678*	0.859*
	(1.34)	(1.61)	(1.68)	(1.95)
Controls	No	Yes	No	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	23,833	23,833	23,833	23,833
Adjusted R^2	0.052	0.070	0.054	0.072

Notes. This table examines whether, after ChatGPT, the market relies more heavily on the consensus forecasts of the more accurate FDPs. Panel A reports the estimates of Equation (4) without controls or fixed effects, where γ is the ERC increment per one-rank accuracy improvement. Each provider row is a separate regression, and the Stacked row pools the five FDPs' surprises. The Stacked Δ row reports the $SUE \times Post$ and $SUE \times Rank \times Post$ coefficients of the pooled regression that interacts every term with *Post*. Panel B reports the estimates of Equation (5) at the firm-quarter level. The dependent variable in Panel A is the market-adjusted $CAR_{[-1,1]}$. Continuous variables are winsorized at the 1st and 99th percentiles. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 4. Analyst Herding Toward the I/B/E/S Consensus and the Accuracy-Rank Partition

	Single post-ChatGPT indicator		GPT-capability regimes	
	(1)	(2)	(3)	(4)
Post	0.044**	0.049**		
	(2.01)	(2.23)		
Post × *IBES Up* (β_1)	0.012	0.016		
	(0.71)	(0.90)		
Post × *IBES Down* (β_2)	-0.048***	-0.056***		
	(-2.62)	(-2.63)		
R1 (Text-only) × *IBES Up* ($\beta_{1,1}$)			0.002	0.006
			(0.11)	(0.30)
R1 (Text-only) × *IBES Down* ($\beta_{2,1}$)			-0.049**	-0.053**
			(-2.07)	(-1.98)
R2 (Browsing) × *IBES Up* ($\beta_{1,2}$)			0.016	0.021
			(0.84)	(1.09)
R2 (Browsing) × *IBES Down* ($\beta_{2,2}$)			-0.059**	-0.065**
			(-2.45)	(-2.41)
R3 (Native search) × *IBES Up* ($\beta_{1,3}$)			0.019	0.023
			(0.83)	(0.96)
R3 (Native search) × *IBES Down* ($\beta_{2,3}$)			-0.034	-0.049**
			(-1.59)	(-2.19)
Spread ($\beta_1 - \beta_2$)	**0.060****	**0.072*****		
	(2.54)	**(2.79)**		
Spread, R1 ($\beta_{1,1} - \beta_{2,1}$)			**0.052***	**0.059***
			(1.78)	**(1.88)**
Spread, R2 ($\beta_{1,2} - \beta_{2,2}$)			**0.075****	**0.086*****
			(2.57)	**(2.75)**
Spread, R3 ($\beta_{1,3} - \beta_{2,3}$)			**0.052***	**0.072****
			(1.79)	**(2.32)**
Lower-order terms	Yes	Yes	Yes	Yes
Firm and event controls	Yes	Yes	Yes	Yes
Analyst controls	Yes	Yes	Yes	Yes
Analyst FE	Yes	No	Yes	No
Firm FE	Yes	No	Yes	No
Analyst×Firm FE	No	Yes	No	Yes
Revision-year FE	Yes	Yes	Yes	Yes
Distance-ventile FE	Yes	Yes	Yes	Yes
Observations	122,001	120,868	122,001	120,868
Adjusted R^2	0.107	0.107	0.107	0.107

Notes. This table examines whether analysts' herding toward the prevailing I/B/E/S consensus falls after ChatGPT in the firms in which I/B/E/S loses accuracy rank, relative to the firms in which it holds its rank. The dependent variable is *Herding*, with a base rate of 0.32. The sample comprises 122,001 quarterly EPS forecast revisions over 2019–2025. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 5. The Predictive Value of the I/B/E/S Consensus and Post-ChatGPT Pricing

	Single post-ChatGPT indicator		GPT-capability regimes	
Dep. Var.:	Mkt CAR	FFC CAR	Mkt CAR	FFC CAR
	(1)	(2)	(3)	(4)
SUE^{IBES} (β_1)	1.768***	1.766***	1.765***	1.765***
	(5.77)	(5.87)	(5.79)	(5.89)
$SUE^{IBES} \times Post$ (β_2)	-0.194	-0.151		
	(-0.46)	(-0.35)		
$SUE^{IBES} \times$ R1 (Text-only)			0.395	0.580
			(0.76)	(1.09)
$SUE^{IBES} \times$ R2 (Browsing)			0.151	-0.118
			(0.25)	(-0.18)
$SUE^{IBES} \times$ R3 (Native search)			-0.936	-0.755
			(-1.39)	(-1.15)
$SUE^{IBES} \times Predictive\ Value^{IBES}$ (β_3)	2.764***	2.736***	2.773***	2.745***
	(4.70)	(4.78)	(4.72)	(4.81)
$SUE^{IBES} \times Predictive\ Value^{IBES} \times Post$ (β_4)	**3.110*** **	**3.123*** **		
	(3.53)	**(3.51)**		
$SUE^{IBES} \times Predictive\ Value^{IBES} \times$ R1			**1.653**	**1.589**
			(1.60)	**(1.52)**
$SUE^{IBES} \times Predictive\ Value^{IBES} \times$ R2			**2.454****	**2.882****
			(2.13)	**(2.41)**
$SUE^{IBES} \times Predictive\ Value^{IBES} \times$ R3			**5.101*** **	**4.790*** **
			(3.44)	**(3.26)**
Lower-order terms	Yes	Yes	Yes	Yes
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	22,356	22,356	22,356	22,356
Adjusted R^2	0.084	0.085	0.084	0.086

Notes. This table examines whether the earnings response coefficient (ERC) on the I/B/E/S surprise rises more steeply with the predictive value of the I/B/E/S consensus after ChatGPT. The dependent variable is the three-day $CAR_{[-1,1]}$, market-adjusted (Mkt) or Fama–French–Carhart four-factor (FFC). Columns (1)–(2) estimate Equation (6), and columns (3)–(4) replace $Post_q$ with the GPT-capability regimes R1 to R3. The coefficient β_3 is the pre-ChatGPT slope of the ERC in predictive value, and β_4 is its post-ChatGPT increment. $Predictive\ Value^{IBES}$ covers 22,364 of the 23,833 firm-quarters (8,730 pre-ChatGPT), of which 22,356 enter the regressions. The uncovered firm-quarters, mostly of recent listings, have fewer than twelve source quarters. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 6. The Accuracy Premium: Horse Races Between the Accuracy-Weighted and Benchmark Surprises

Panel A. Horse race: SUE^{Acc} vs. SUE^{EW5}

Dep. Var.:	Mkt CAR$_{[-1,1]}$			FFC CAR$_{[-1,1]}$		
	(1) Pre	(2) Post	(3) Pooled	(4) Pre	(5) Post	(6) Pooled
SUE^{Acc} (β_1)	1.376***	3.420***	1.307***	1.401***	3.519***	1.290***
	(3.45)	(9.38)	(3.60)	(3.63)	(10.01)	(3.67)
$SUE^{Acc} \times Post$ (β_3)			2.003***			2.135***
			(4.08)			(4.47)
SUE^{EW5} (β_2)	2.324***	1.767***	2.279***	2.379***	1.734***	2.284***
	(6.10)	(5.56)	(6.84)	(6.44)	(5.62)	(7.03)
$SUE^{EW5} \times Post$ (β_4)			-0.399			-0.449
			(-0.92)			(-1.06)
Premium (Pre-ChatGPT)	-0.948		-0.972	-0.978		-0.994
	(-1.28)		(-1.47)	(-1.37)		(-1.55)
Premium (Post-ChatGPT)		1.654***	1.429**		1.785***	1.590***
		(2.62)	(2.39)		(2.93)	(2.74)
Δ Premium (Post−Pre)			**2.401*** **			**2.584*** **
			(2.75)			**(3.03)**
Control Variables						
Log Market Equity	-5.348***	-4.964***	-2.848***	-5.585***	-5.125***	-3.238***
	(-5.35)	(-5.35)	(-6.00)	(-5.59)	(-5.74)	(-7.19)
Book-to-Market	0.413	-1.315	0.168	0.849	-1.245	0.309
	(0.36)	(-1.22)	(0.28)	(0.75)	(-1.16)	(0.53)
Leverage	-4.957**	-3.652*	-0.913	-4.561**	-4.067**	-1.312
	(-2.38)	(-1.70)	(-0.76)	(-2.17)	(-1.97)	(-1.13)
Prior 12-Month Return	-0.295	-0.868**	-0.235	-0.024	-1.038***	-0.129
	(-1.27)	(-2.52)	(-1.27)	(-0.11)	(-3.02)	(-0.71)
Log Price	-1.063	-1.047	-0.482	-0.671	-0.735	-0.029
	(-1.13)	(-1.12)	(-1.05)	(-0.71)	(-0.83)	(-0.07)
Return Volatility	-59.587***	27.445*	14.899	-53.652***	19.210	11.467
	(-3.08)	(1.66)	(1.34)	(-2.87)	(1.20)	(1.05)
Analyst Dispersion	-23.636	-40.842	-15.946	-11.993	-17.323	-2.625
	(-0.40)	(-0.71)	(-0.40)	(-0.20)	(-0.30)	(-0.07)
Busyness	-0.104	-0.245*	-0.172*	-0.132	-0.219*	-0.164*
	(-0.73)	(-1.83)	(-1.79)	(-0.93)	(-1.66)	(-1.72)
$\lvert\Delta IB\rvert$	16.729***	-4.012	2.450	17.171***	-4.609	2.027
	(2.98)	(-0.78)	(0.69)	(3.09)	(-0.91)	(0.58)
$\lvert\Delta Shares\rvert$	3.721	-1.964	1.505	3.251	-2.171	1.376
	(0.79)	(-0.51)	(0.52)	(0.72)	(-0.57)	(0.48)
$\lvert SpecItems\rvert$	-18.458	-8.991	-10.912	-21.400*	-7.112	-9.847
	(-1.50)	(-0.79)	(-1.33)	(-1.77)	(-0.64)	(-1.22)
SBC High	-0.149	0.174	0.106	-0.168	0.186	0.095
	(-0.57)	(0.77)	(0.65)	(-0.66)	(0.83)	(0.59)
Unexpected Item	-0.258	-0.036	-0.123	-0.222	-0.023	-0.103
	(-0.97)	(-0.17)	(-0.78)	(-0.89)	(-0.11)	(-0.68)
Stock Split	-2.679**	-1.110	-1.584*	-2.266*	-0.472	-0.977
	(-2.21)	(-0.92)	(-1.90)	(-1.78)	(-0.42)	(-1.18)
Fiscal Q4	0.600***	-0.045	0.276**	0.606***	0.230	0.417***
	(2.83)	(-0.24)	(2.07)	(2.89)	(1.23)	(3.19)
Institutional Ownership	3.234	-0.702**	-0.470	1.626	-0.543	-0.416
	(1.56)	(-2.02)	(-1.41)	(0.83)	(-1.57)	(-1.25)

(continued on next page)

Table 6 continued from previous page

Dep. Var.:	Mkt CAR$_{[-1,1]}$			FFC CAR$_{[-1,1]}$		
	(1) Pre	(2) Post	(3) Pooled	(4) Pre	(5) Post	(6) Pooled
Log # Analysts	0.508	-0.218	-0.248	0.648	-0.328	-0.371
	(0.79)	(-0.39)	(-0.65)	(1.04)	(-0.60)	(-0.99)
Controls	Yes	Yes	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes	Yes
Observations	9,405	14,402	23,833	9,405	14,402	23,833
Adjusted R^2	0.085	0.106	0.080	0.092	0.110	0.083

Panel B. Horse race: SUE^{Acc} vs. SUE^{IBES}

Dep. Var.:	Mkt CAR$_{[-1,1]}$			FFC CAR$_{[-1,1]}$		
	(1) Pre	(2) Post	(3) Pooled	(4) Pre	(5) Post	(6) Pooled
SUE^{Acc} (β_1)	1.428***	3.678***	1.430***	1.527***	3.796***	1.494***
	(4.00)	(10.49)	(4.31)	(4.36)	(10.70)	(4.48)
$SUE^{Acc} \times Post$ (β_3)			2.153***			2.204***
			(4.82)			(4.98)
SUE^{IBES} (β_2)	2.176***	1.460***	2.054***	2.149***	1.407***	1.974***
	(7.08)	(4.44)	(7.31)	(7.20)	(4.20)	(7.01)
$SUE^{IBES} \times Post$ (β_4)			-0.519			-0.484
			(-1.34)			(-1.28)
Premium (Pre-ChatGPT)	-0.747		-0.623	-0.622		-0.480
	(-1.20)		(-1.09)	(-1.03)		(-0.83)
Premium (Post-ChatGPT)		2.219***	2.049***		2.389***	2.208***
		(3.56)	(3.60)		(3.75)	(3.82)
Δ Premium (Post−Pre)			**2.672*** **			**2.688*** **
			(3.44)			**(3.50)**
Controls	Yes	Yes	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes	Yes
Observations	9,405	14,402	23,833	9,405	14,402	23,833
Adjusted R^2	0.087	0.106	0.080	0.094	0.110	0.083

Notes. This table examines whether announcement returns respond to the accuracy-weighted surprise beyond their response to a benchmark surprise, before and after ChatGPT. Each column estimates the horse race of Equation (7), without the $Post_q$ interactions in the single-period columns. The accuracy premium is $\beta_1 - \beta_2$, except in the post-ChatGPT row of the pooled columns, where it is $(\beta_1 + \beta_3) - (\beta_2 + \beta_4)$. The Δ Premium row reports the post-minus-pre change, $\beta_3 - \beta_4$. Continuous variables are winsorized at the 1st and 99th percentiles. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 7. The Accuracy Premium Across GPT-Capability Regimes

Panel A. Benchmark: equal-weighted five-FDP consensus

Regime:	(1) R0 Pre-ChatGPT	(2) R1 Text-only	(3) R2 Browsing	(4) R3 Native search
SUE^{Acc} (β_1)	1.376***	2.597***	3.753***	3.954***
	(3.45)	(3.38)	(6.19)	(5.69)
Benchmark SUE (β_2)	2.324***	2.419***	1.325***	1.404**
	(6.10)	(3.78)	(2.63)	(2.26)
Premium ($\beta_1 - \beta_2$)	**-0.948**	**0.178**	**2.429****	**2.550****
	(-1.28)	**(0.13)**	**(2.35)**	**(2.07)**
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	9,405	3,705	5,529	5,084
Adjusted R^2	0.085	0.158	0.151	0.140

Panel B. Benchmark: I/B/E/S consensus

Regime:	(1) R0 Pre-ChatGPT	(2) R1 Text-only	(3) R2 Browsing	(4) R3 Native search
SUE^{Acc} (β_1)	1.428***	1.918**	3.362***	5.079***
	(4.00)	(2.35)	(6.17)	(7.87)
Benchmark SUE (β_2)	2.176***	3.144***	1.842***	0.169
	(7.08)	(4.70)	(3.74)	(0.30)
Premium ($\beta_1 - \beta_2$)	**-0.747**	**-1.226**	**1.519**	**4.909*****
	(-1.20)	**(-0.86)**	**(1.61)**	**(4.41)**
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	9,405	3,705	5,529	5,084
Adjusted R^2	0.087	0.170	0.155	0.138

Notes. This table examines whether the accuracy premium differs across the GPT-capability regimes, R0 to R3. The dependent variable is the market-adjusted $CAR_{[-1,1]}$. Each column estimates Equation (7) within one regime, without the $Post_q$ interactions. Continuous variables are winsorized at the 1st and 99th percentiles. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 8. Cross-Sectional Tests of the Accuracy Premium

Panel A. Concurrent Cross-FDP SUE Dispersion

Benchmark:	Equal-weighted		I/B/E/S	
	(1) Low	(2) High	(3) Low	(4) High
SUE^{Acc} (β_1)	4.138***	1.433***	1.569**	1.748***
	(6.37)	(4.23)	(2.43)	(5.57)
Benchmark SUE (β_2)	3.318***	2.024***	7.496***	1.560***
	(5.86)	(6.41)	(10.03)	(6.10)
$SUE^{Acc} \times Post^{\text{Browse}}$ (β_3)	2.748**	1.610***	1.439	1.746***
	(2.48)	(2.89)	(1.34)	(3.63)
Benchmark $SUE \times Post^{\text{Browse}}$ (β_4)	0.712	-0.602	2.130*	-0.681*
	(0.76)	(-1.20)	(1.86)	(-1.70)
Premium (Pre-Browsing)	0.821	-0.591	-5.927***	0.188
	(0.73)	(-0.96)	(-4.58)	(0.35)
Premium (Post-Browsing)	2.857*	1.620**	-6.619***	2.615***
	(1.71)	(1.98)	(-3.81)	(3.81)
Δ Premium (Post−Pre)	**2.037**	**2.211****	**-0.691**	**2.427*****
	(1.06)	**(2.21)**	**(-0.33)**	**(2.99)**
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	11,832	11,846	11,832	11,846
Adjusted R^2	0.111	0.075	0.127	0.075

Panel B. Special-Items Magnitude

Benchmark:	Equal-weighted		I/B/E/S	
	(1) Low	(2) High	(3) Low	(4) High
SUE^{Acc} (β_1)	2.345***	1.613***	1.940***	1.969***
	(4.36)	(4.49)	(4.34)	(5.48)
Benchmark SUE (β_2)	2.153***	2.078***	2.465***	1.615***
	(4.21)	(6.41)	(6.46)	(5.25)
$SUE^{Acc} \times Post^{\text{Browse}}$ (β_3)	0.247	2.035***	0.440	2.583***
	(0.31)	(3.39)	(0.58)	(4.69)
Benchmark $SUE \times Post^{\text{Browse}}$ (β_4)	0.799	-0.483	0.659	-1.055**
	(1.14)	(-0.92)	(0.98)	(-2.28)
Premium (Pre-Browsing)	0.192	-0.464	-0.525	0.354
	(0.19)	(-0.74)	(-0.68)	(0.58)
Premium (Post-Browsing)	-0.360	2.054**	-0.743	3.992***
	(-0.28)	(2.48)	(-0.65)	(5.36)
Δ Premium (Post−Pre)	**-0.552**	**2.518****	**-0.219**	**3.638*****
	(-0.39)	**(2.40)**	**(-0.16)**	**(3.92)**
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	11,823	11,828	11,823	11,828
Adjusted R^2	0.089	0.073	0.095	0.071

Notes. This table examines whether the change in the accuracy premium at the browsing boundary is larger in firm-quarters with greater cross-FDP disagreement and in those with larger special items. The dependent variable is the market-adjusted $CAR_{[-1,1]}$. Each column estimates Equation (7), with $Post_q^{\text{Browse}}$ in place of $Post_q$, within the below- or above-median half of the panel's variable. The accuracy premium is $\beta_1 - \beta_2$ before the browsing boundary and $(\beta_1 + \beta_3) - (\beta_2 + \beta_4)$ after it. The Δ Premium row reports the change, $\beta_3 - \beta_4$. Each panel's sample comprises the firm-quarters with non-missing values for all regression variables. Continuous variables are winsorized at the 1st and 99th percentiles. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 9. LLM Visibility vs. Conventional Proxies for Firm Prominence

Panel A. Overlap of the High groups

	LLM Visibility	Firm Size	Analyst Coverage	Inst. Ownership
LLM Visibility	**1.000**	**0.502**	**0.497**	**0.471**
Firm Size	0.510	1.000	0.796	0.407
Analyst Coverage	0.506	0.798	1.000	0.445
Inst. Ownership	0.479	0.407	0.445	1.000

Panel B. Accuracy-premium concentration: equal-weighted benchmark

Measure	Univariate ΔPremium		Horse-race
	High group	Low group	triple interaction
LLM Visibility	**3.263*****	**0.300**	**3.172***
	(2.59)	**(0.25)**	**(1.80)**
Firm Size	0.723	2.447**	0.111
	(0.55)	(2.18)	(0.05)
Analyst Coverage	0.109	3.076**	-3.161
	(0.08)	(2.57)	(-1.38)
Inst. Ownership	3.224**	0.646	3.222*
	(2.51)	(0.58)	(1.86)
Controls	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes
Year FE	Yes	Yes	Yes
Observations	23,044		23,044

Panel C. Accuracy-premium concentration: I/B/E/S benchmark

Measure	Univariate ΔPremium		Horse-race
	High group	Low group	triple interaction
LLM Visibility	**5.743*****	**-0.125**	**5.841*****
	(5.67)	**(-0.12)**	**(4.03)**
Firm Size	2.191	2.462***	2.010
	(1.28)	(2.76)	(0.82)
Analyst Coverage	0.441	3.198***	-4.099*
	(0.29)	(3.59)	(-1.77)
Inst. Ownership	2.564**	2.418**	0.625
	(2.46)	(2.08)	(0.41)
Controls	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes
Year FE	Yes	Yes	Yes
Observations	23,044		23,044

Notes. This table examines whether *LLM Visibility*, how visible a firm's consensus forecasts are to LLMs, is distinct from conventional proxies for firm prominence. In Panel A, the cell in row m and column n is the share of firms in measure m's High group (above its cross-firm median) that measure n also classifies as High (diagonal $= 1$). In Panels B and C, ΔPremium is the change in the accuracy premium, $\beta(SUE^{Acc}) - \beta(SUE^{X})$, at the browsing boundary, estimated within each measure's High and Low groups. The horse-race triple interaction, $(SUE^{Acc} - SUE^{X}) \times Post^{\text{Browse}} \times High_m$, comes from one pooled regression that includes the triple interaction of every measure. The dependent variable is the market-adjusted $CAR_{[-1,1]}$. The univariate estimates use each measure's own sample, 23,833 firm-quarters for *LLM Visibility*, 23,063 for firm size and institutional ownership, and 23,044 for analyst coverage. The six firm-quarters without *LLM Visibility* fall in its Low group. The horse race uses the 23,044 firm-quarters with non-missing values for all four measures, which the Observations rows report. In that sample, the High groups hold 11,219 to 11,901 firm-quarters and the Low groups 11,143 to 11,825. Continuous variables are winsorized at the 1st and 99th percentiles. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 10. LLM Visibility and the Accuracy Premium by GPT-Capability Regime

Panel A. Equal-weighted five-FDP benchmark

Regime:	(1) R0 Pre-ChatGPT	(2) R1 Text-only	(3) R2 Browsing	(4) R3 Native search
SUE^{Acc} (β_1)	1.361**	3.503***	3.077***	3.055***
	(2.57)	(5.05)	(2.95)	(4.18)
SUE^{EW5} (β_2)	2.252***	2.185***	1.885***	2.002***
	(4.09)	(3.61)	(3.65)	(3.21)
$SUE^{Acc} \times High\,Visibility$ (β_3)	0.020	-2.274	1.550*	1.815*
	(0.05)	(-1.49)	(1.79)	(1.93)
$SUE^{EW5} \times High\,Visibility$ (β_4)	0.153	0.889	-1.133**	-1.269
	(0.28)	(0.79)	(-2.11)	(-1.51)
Accuracy premium = loading on SUE^{Acc} − loading on benchmark *SUE*:				
Premium (Low Vis)	-0.892	1.317	1.192	1.053
	(-0.90)	(1.12)	(0.89)	(0.83)
Premium (High Vis)	-1.025	-1.845	3.875***	4.136***
	(-1.57)	(-0.94)	(4.60)	(4.16)
Gap (High Vis − Low Vis)	**-0.133**	**-3.163**	**2.683****	**3.083***
	(-0.15)	**(-1.22)**	**(2.12)**	**(1.85)**
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year-quarter FE	Yes	Yes	Yes	Yes
Observations	9,399	3,705	5,529	5,084
Adjusted R^2	0.086	0.163	0.165	0.143

Panel B. I/B/E/S benchmark

Regime:	(1) R0 Pre-ChatGPT	(2) R1 Text-only	(3) R2 Browsing	(4) R3 Native search
SUE^{Acc} (β_1)	1.607***	3.138***	2.737***	4.029***
	(4.07)	(3.15)	(6.05)	(8.81)
SUE^{IBES} (β_2)	1.859***	2.424***	2.403**	0.974**
	(3.60)	(2.96)	(2.28)	(2.13)
$SUE^{Acc} \times High\,Visibility$ (β_3)	-0.382	-2.353*	1.667**	2.376***
	(-1.03)	(-1.74)	(2.44)	(3.68)
$SUE^{IBES} \times High\,Visibility$ (β_4)	0.669	1.362	-1.380	-1.856***
	(0.97)	(1.45)	(-1.45)	(-3.06)
Accuracy premium = loading on SUE^{Acc} − loading on benchmark *SUE*:				
Premium (Low Vis)	-0.252	0.714	0.335	3.054***
	(-0.31)	(0.41)	(0.26)	(4.00)
Premium (High Vis)	-1.303**	-3.001**	3.382***	7.286***
	(-2.06)	(-2.12)	(6.17)	(10.56)
Gap (High Vis − Low Vis)	**-1.052**	**-3.715***	**3.047****	**4.232*****
	(-1.07)	**(-1.72)**	**(2.05)**	**(3.93)**
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year-quarter FE	Yes	Yes	Yes	Yes
Observations	9,399	3,705	5,529	5,084
Adjusted R^2	0.088	0.175	0.169	0.142

Notes. This table examines whether the accuracy premium is larger in firms whose consensus forecasts are more visible to LLMs, within each GPT-capability regime. The dependent variable is the market-adjusted $CAR_{[-1,1]}$. Each column estimates Equation (8) within one regime. The accuracy premium is $\beta_1 - \beta_2$ for low-visibility firms and $(\beta_1 + \beta_3) - (\beta_2 + \beta_4)$ for high-visibility firms, and the gap is $\beta_3 - \beta_4$. Each regime's sample comprises the firm-quarters with non-missing values for all regression variables. Continuous variables are winsorized at the 1st and 99th percentiles. Appendix A defines the variables. Standard errors are two-way clustered by firm and fiscal year, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE 11. The Accuracy Premium Gap by LLM Visibility: Median Split, Continuous Measure, and Change with Retrieval

Panel A. The median split

Regime:	(1) R0 Pre-ChatGPT	(2) R1 Text-only	(3) R2 Browsing	(4) R3 Native search
Equal-weighted five-FDP benchmark:				
Gap (High Vis − Low Vis)	-0.133	-3.163	2.683*	3.083
t-statistic	(-0.15)	(-1.22)	(2.12)	(1.85)
Randomization *p*-value	[0.462]	[0.126]	[0.097]	[0.108]
I/B/E/S benchmark:				
Gap (High Vis − Low Vis)	-1.052	-3.715	3.047*	4.232**
t-statistic	(-1.07)	(-1.72)	(2.05)	(3.93)
Randomization *p*-value	[0.197]	[0.114]	[0.050]	[0.031]
Observations	9,399	3,705	5,529	5,084

Panel B. The continuous measure

Regime:	(1) R0 Pre-ChatGPT	(2) R1 Text-only	(3) R2 Browsing	(4) R3 Native search
Equal-weighted five-FDP benchmark:				
Gap (per SD)	0.014	-0.506	2.897***	1.252
t-statistic	(0.02)	(-0.46)	(4.87)	(1.49)
Randomization *p*-value	[0.492]	[0.364]	[0.003]	[0.167]
I/B/E/S benchmark:				
Gap (per SD)	-0.442	-0.278	1.741**	1.802*
t-statistic	(-0.52)	(-0.33)	(3.63)	(3.46)
Randomization *p*-value	[0.245]	[0.438]	[0.037]	[0.066]
Observations	9,399	3,705	5,529	5,084

Panel C. The change with retrieval

	Pre-retrieval (R0–R1)	Retrieval (R2–R3)	Change
Equal-weighted five-FDP benchmark:			
Median split (High Vis − Low Vis)	-0.814	2.109*	2.923**
t-statistic	(-0.90)	(3.35)	(2.64)
Randomization *p*-value	[0.233]	[0.064]	[0.040]
Continuous measure (per SD)	-0.169	1.555**	1.723**
t-statistic	(-0.32)	(2.95)	(2.44)
Randomization *p*-value	[0.396]	[0.016]	[0.020]
I/B/E/S benchmark:			
Median split (High Vis − Low Vis)	-1.969**	3.879***	5.847***
t-statistic	(-3.57)	(4.96)	(6.00)
Randomization *p*-value	[0.022]	[0.003]	[0.000]
Continuous measure (per SD)	-0.508	2.012***	2.520***
t-statistic	(-1.48)	(6.84)	(5.78)
Randomization *p*-value	[0.161]	[0.001]	[0.001]
Observations	23,827	23,827	23,827

Notes. This table examines whether the accuracy premium gap between high- and low-visibility firms changes with retrieval. Each cell reports the gap, $\beta_3 - \beta_4$, from the specification of Table 10, or its change. Panel A repeats the median-split gaps of Table 10. Panel B replaces *High Visibility* with the standardized log of one plus the firm's three-query mean citation count. The gap for this continuous measure is expressed per standard deviation. Panel C pools the four regimes and interacts each surprise with $Post_q^{\text{Browse}}$, the indicator of R2 and R3. In Panel A the reference distribution is 5,000 random splits of the firms, each with 654 High firms. In Panels B and C it is 2,000 permutations of the measure across firms. Appendix A defines the variables. Standard errors are two-way clustered by firm and fiscal year, with *t*-statistics in parentheses. The randomization *p*-value is the one-sided share of the reference distribution with an estimate at least as large as the cell's, in its direction. *, **, and *** denote one-sided randomization-inference *p*-values below 0.10, 0.05, and 0.01.

Online Appendix

Synthesizing the Consensus: Generative AI, LLM Visibility, and the Pricing of Multi-Provider Earnings Forecasts

Leonard Yang Liu
University of Maryland
yliu337@umd.edu

This Online Appendix contains supplementary material for the paper:

- **Online Appendix A: The LLM Visibility Query and an Integration-Prompt Pilot**
  - A.1 The Query Behind LLM Visibility and the Providers Behind the Cited Sites
  - A.2 An Integration-Prompt Pilot
- **Online Appendix B: The Accuracy-Weighted SUE, Optimal Synthesis, and the Pricing Model**
  - B.1 Numerical Illustration of the Accuracy-Weighted SUE
  - B.2 Optimal Synthesis Across Providers
  - B.3 The Pricing Model and the Stability of Cross-Provider Dynamics
  - B.4 Validating the Accuracy-Weighted Construct
- **Online Appendix C: Supplementary Details**
  - C.1 Institutional Background of the Five Providers
  - C.2 The Analyst Revision Sample and the Herding Specification
  - C.3 The Predictive-Value Measure and Two Alternative Explanations
  - C.4 The Change in the Premium by Firm Characteristic
  - C.5 Search Traffic to the GenAI Tools

Online Appendix A: The LLM Visibility Query and an Integration-Prompt Pilot

This appendix first describes the query behind the *LLM Visibility* measure of Section 5.1 and reports the providers that the cited sites name as the sources of their estimates (A.1). It then presents a pilot in which browsing models are asked to rank and weight the cited sites by accuracy (A.2).

A.1 The Query Behind LLM Visibility and the Providers Behind the Cited Sites

The query asks ChatGPT (`gpt-4o-2024-08-06` with web search) for a firm's next-quarter consensus EPS and for every source that publishes one, with each source's name, estimate, number of contributing analysts, and URL, together with the firm's next earnings date and its most recent reported EPS. Each response returns the sources it cites, and the measure counts these source citations, repeats included. I issue the query three times per firm on May 28 and June 1, 2026 (Section 5.1). Table OA2 reports, for the 17 sites that receive 88.7 percent of the citations in the May 28 run, the provider that each site names as the source of its estimates. Of the five providers, the sites name only Zacks and S&P Global Market Intelligence, the publisher of the Capital IQ consensus, as their current sources.

A.2 An Integration-Prompt Pilot

I define synthesis as three steps (Section 2.2). The first retrieves each provider's consensus and actual. The second evaluates each provider's quality from its past performance for the firm. The third weights the forecasts by that assessment. The GPT-capability regimes after ChatGPT's release and the *LLM Visibility* measure vary only the first step, so this subsection asks whether a browsing model performs the other two when asked. If models could only retrieve, the investor would have to weight the providers by accuracy by hand, and the cost reduction on which the framework rests would be partial. On October 9, 2026, I send two prompts, P1 in free text and P2 in a fixed JSON format, to ChatGPT (`gpt-6.1-sol` with web search) and to Perplexity (`sonar`) for 50 sample firms (IBM, Microsoft, and 48 drawn from the grid of visibility halves and size terciles). The prompts ask for each site's consensus and reported EPS for the four quarters reported in 2025, its mean absolute error, a ranking and weights of the sites by accuracy, and an accuracy-weighted consensus for the next quarter. Table OA1 reports the results. The histories that ChatGPT retrieves

match one of the five providers' own actuals within two cents at a mean rate of 98 percent per response. ChatGPT completes the chain to an accuracy-weighted consensus in 28 to 62 percent of its responses, depending on the prompt, and Perplexity in 24 percent of its P2 responses. The pilot covers 50 firms with a model released after the sample period, so it shows capability at the time of the pilot, not use during the sample. The completion rate also depends on the prompt's wording. Under the P2 system message, which requires every number to come from a cited page, 35 percent of the 100 P2 responses (two per firm) complete the chain. Once an amended message lets the model compute the errors, ranking, weights, and weighted consensus (prompt P2b), 62 percent do.

Online Appendix B: The Accuracy-Weighted SUE, Optimal Synthesis, and the Pricing Model

This appendix first illustrates how the accuracy-weighted SUE is constructed for a single firm-quarter (B.1). It then derives the expectation that an investor forms from the five providers' consensus forecasts and the gain from synthesis over a single provider and over equal weighting (B.2). It then gives a pricing model under which the accuracy premium is a demand-side object and shows that the cross-provider structure on which the weights operate is stable (B.3). Finally, it validates the accuracy-weighted construct through the persistence of the providers' accuracy ranks, the accuracy of the three surprise measures, and the predicted surprise's trading value (B.4).

B.1 Numerical Illustration of the Accuracy-Weighted SUE

Suppose firm i in quarter q has the following backward-looking accuracy histories across the five FDPs, alongside their current-quarter standardized surprises. The accuracy $Acc^g_{i,q} = -MAFE^g_{i,q}$ is measured over the prior four quarters in price-scaled units, and the surprises $SUE^g_{i,q}$ are in percent of price. The z-scores use the population standard deviation across the five providers (divisor five).

Provider g	$Acc^g_{i,q}$	$Acc_z^g_{i,q}$	$\exp(Acc_z^g_{i,q})$	$w^g_{i,q}$
I/B/E/S	-0.0010	1.41	4.11	0.52
FactSet	-0.0020	0.71	2.03	0.26
Capital IQ	-0.0030	0.00	1.00	0.13
Zacks	-0.0040	-0.71	0.49	0.06
Bloomberg	-0.0050	-1.41	0.24	0.03
Sum			7.88	1.00

I/B/E/S has the smallest absolute forecast error (largest Acc^g) and therefore the highest z-score, which the softmax transform maps to the largest weight (0.52). Bloomberg has the largest absolute forecast error and receives the smallest weight (0.03). The weights sum to one.

Provider g	$SUE^{g}_{i,q}$	$w^{g}_{i,q} \times SUE^{g}_{i,q}$	$0.20 \times SUE^{g}_{i,q}$
I/B/E/S	1.20	0.63	0.24
FactSet	0.80	0.21	0.16
Capital IQ	1.00	0.13	0.20
Zacks	1.50	0.09	0.30
Bloomberg	-0.50	-0.02	-0.10
Total		$SUE^{Acc}_{i,q} = 1.04$	$SUE^{EW5}_{i,q} = 0.80$

Entries are rounded to two decimals, while the totals are computed from unrounded weights, so the displayed cells need not sum exactly to the reported totals. The accuracy-weighted surprise $SUE^{Acc}_{i,q} = 1.04$ exceeds the equal-weighted surprise $SUE^{EW5}_{i,q} = 0.80$ for two reasons. First, I/B/E/S, with a weight of 0.52 against 0.20, reports a surprise of 1.20. Second, the low weight on Bloomberg removes most of its surprise of -0.50. These two contributions outweigh the low weight on Zacks's surprise of 1.50, the largest of the five. The accuracy tilt $SUE^{Acc}_{i,q} - SUE^{EW5}_{i,q} = 0.24$ captures the accuracy-driven correction. In the horse-race regression of Table 6, this tilt is the source of identification between the two surprises. When the tilt is informative about announcement-window returns beyond the equal-weighted surprise, the market loads on provider-specific accuracy instead of treating all five FDPs alike.

B.2 Optimal Synthesis Across Providers

A simple model illustrates why synthesizing the providers' consensus forecasts can benefit investors. Following the forecast-combination logic of Bates and Granger (1969), consider an investor who uses the forecasts to estimate the earnings revealed at the announcement. The investor seeks to minimize expected squared forecast error.

Fix a firm-quarter and suppress its subscripts. Let provider g's consensus forecast be

$$s^g = x + u + v^g, \qquad g = 1, \ldots, K,$$

where x is realized earnings, u is a common forecast error, and v^g is a provider-specific forecast error.[1]

[1] I include the common error u because providers may share analyst inputs or forecasting mistakes. The same optimal weights and comparisons with the most accurate single provider and equal weighting obtain when $u = 0$. The common error adds the same σ_u^2 to each forecast's expected squared error and leaves the optimal weights and absolute accuracy gains unchanged.

Assume that all errors have mean zero, that the provider-specific errors are mutually uncorrelated and uncorrelated with u, and that $\mathrm{Var}(u) = \sigma_u^2$ and $\mathrm{Var}(v^g) = a_g > 0$. I treat these variances as known.

The investor forms a combined forecast $c(w) = \sum_g w_g s^g$, with $\sum_g w_g = 1$. Its expected squared error is

$$\mathrm{E}\left[(c(w) - x)^2\right] = \sigma_u^2 + \sum_g w_g^2 a_g.$$

Combining forecasts cannot remove the common error, but it can reduce provider-specific noise. Minimizing this expression gives

$$w_g^* = \frac{a_g^{-1}}{\sum_h a_h^{-1}}, \qquad \mathrm{E}\left[(c(w^*) - x)^2\right] = \sigma_u^2 + \frac{1}{\sum_g a_g^{-1}}. \tag{B1}$$

More reliable providers receive greater weight, while the investor retains information from the other providers.

This combination improves on the most accurate single provider:

$$\sigma_u^2 + \frac{1}{\sum_g a_g^{-1}} < \sigma_u^2 + \min_g a_g,$$

provided that at least two providers have finite, strictly positive provider-specific error variances. For example, if all five providers have the same provider-specific error variance a, averaging reduces that component of forecast error from a to $a/5$. The common error variance remains unchanged.

Accuracy weighting also weakly improves on equal weighting:

$$\sigma_u^2 + \frac{1}{\sum_g a_g^{-1}} \leq \sigma_u^2 + \frac{1}{K^2} \sum_g a_g,$$

by the Cauchy–Schwarz inequality, with equality only when all a_g are equal. Thus, combining providers can improve accuracy even when they are equally reliable. Assigning different weights provides an additional benefit when their reliabilities differ.

The model gives investors a mathematical reason to synthesize consensus forecasts. Different providers contain some distinct information, and combining them reduces the influence of provider-specific errors. The empirical procedure in Section 3.5 implements this intuition by assigning greater weight to providers with better trailing accuracy. Its estimated weights

need not equal the population-optimal weights above, so Figure OA2 assesses the resulting accuracy gains empirically.

B.3 The Pricing Model and the Stability of Cross-Provider Dynamics

Consensus forecasts become more accurate over 2021–2025, whether from secular improvement in analyst research, from GenAI assistance reaching the analysts whose forecasts feed all five FDPs, or from shifting sample composition. A premium that follows this trend does not show how the market synthesizes the providers' forecasts. This subsection gives a pricing model under which the premium is a demand-side object and shows that the cross-provider structure on which the weights operate is stable.

Write each provider's surprise as $SUE^g_{i,q} = N_{i,q} + \eta^g_{i,q}$, where $N_{i,q}$ is the value-relevant earnings news of announcement (i,q) and the measurement components η^g may have any joint distribution and any drift over time. The set $\{\eta^g\}$ captures the supply side (forecast accuracy and its evolution). Two properties separate demand from supply.

First, the construction is supply-invariant. Because Equation (3) standardizes accuracy within firm-quarter before the softmax, replacing $MAFE^g_{i,q}$ by $a_{i,q}MAFE^g_{i,q} + b_{i,q}$ $(a_{i,q} > 0)$ uniformly across the five providers leaves every weight $w^g_{i,q}$ unchanged.

Second, let the market's pre-announcement expectation of street earnings be a convex combination of the accuracy-weighted consensus and a benchmark consensus, $E^m_{i,q} = \omega\, Consensus^{Acc}_{i,q} + (1-\omega)\, Consensus^X_{i,q}$, where $Consensus^{Acc}_{i,q} = \sum_g w^g_{i,q}\, Consensus^g_{i,q}$ is the accuracy-weighted consensus level, $Consensus^X_{i,q}$ is the benchmark's, $X \in \{EW5, IBES\}$, and ω is the market's relative reliance on the accuracy-weighted synthesis, the direct analog of the expectation mixture in Walther (1997). Each provider's consensus pairs with its own reported actual (Section 3.5), so the surprise inherits the same mixture and the return equation is

$$CAR_{i,q} \;=\; \theta\left[\,\omega\, SUE^{Acc}_{i,q} \;+\; (1-\omega)\, SUE^X_{i,q}\right] \;+\; \varepsilon_{i,q}, \tag{B2}$$

where θ is the return response per unit of perceived surprise (Holthausen and Verrecchia, 1988; Kim and Verrecchia, 1991) and ε collects pricing components outside the span of the two surprises. Equation (B2) reparameterizes the horse-race slopes for each period ($\theta = \beta_1 + \beta_2$ and

$\omega = \beta_1/(\beta_1+\beta_2)$ before ChatGPT, with $\beta_1+\beta_3$ and $\beta_2+\beta_4$ in their place after) and imposes no restriction. The accuracy premium in each period, the loading on SUE^{Acc} minus the loading on SUE^{X}, is therefore $\theta(2\omega-1)$, linear in ω given θ. Neither the level of forecast accuracy nor its improvement appears in it, so the supply side can reach the premium only through ω, the market's choice. In these units, the slopes in columns (1) and (2) of Table 6, Panel A, imply $\omega = 0.37$ before ChatGPT and 0.66 after.

The improvement in accuracy is common to the providers, not a convergence among them. With firm fixed effects, the log realized error is 10.1 log points lower under browsing ($t = 4.0$) and 10.6 lower under native search ($t = 3.8$), with no change in the text-only regime (untabulated). Table OA3, Panel C, reports the distribution of the *ex ante* most accurate provider by regime. Across the three post-ChatGPT regimes, the period over which the premium emerges, the distribution shows no significant change ($\chi^2(8) = 10.8$, $p = 0.21$). A one-time step down in Bloomberg's share at the ChatGPT boundary (17.2 to 14.5 percent in Panel C) predates the premium's emergence at the browsing boundary and is the kind of reordering that the trailing-MAFE weights track.

The relative structure keeps its predictive validity. The within-firm-quarter rank correlation between the trailing-MAFE ordering and the realized-error ordering of the announcement quarter is 0.217 in R0 and 0.215 in R3, with the R3-minus-R0 difference statistically indistinguishable from zero ($t = 0.0$; untabulated). A transient dip in the text-only regime (-0.034, t = -2.7; untabulated) has recovered by the regimes in which the premium appears.

Two further tests support the construct. First, in untabulated paired tests, the absolute value of the accuracy-weighted surprise is smaller than that of either benchmark in both periods ($t >$ 10). Second, an untabulated stacked panel of the four providers with native point-in-time history (Bloomberg's begins only in 2020Q2) regresses four-quarter-ahead earnings and cash flows on each provider's consensus, with provider intercepts, industry effects, and firm-clustered standard errors. The consensus slopes are statistically indistinguishable across providers. Tilting weights toward accurate providers therefore loses no detectable predictive content.

B.4 Validating the Accuracy-Weighted Construct

This subsection validates the accuracy-weighted construct of Section 3.5. It reports the persistence of the providers' accuracy ranks, compares the accuracy of the three surprise measures, and tests whether the predicted surprise earns a return after the announcement.

Weighting by trailing accuracy is informative only if a provider's accuracy persists. Table OA5, Panel A, reports the persistence of the providers' accuracy ranks. Provider ranks persist (Spearman $\rho = 0.71$), with the least and most accurate providers retaining their ranks 76 and 65 percent of the time, roughly three to four times the uniform baseline.[2] Figure OA2 compares the accuracy of the three surprise measures. Out of sample, SUE^{Acc} has the smallest mean absolute error of the three in both the pre- and post-ChatGPT periods.

If accuracy persists, the difference between the accuracy-weighted and a benchmark consensus, known before the announcement, should also earn a return after it. I measure this difference as the *predicted surprise*

$$Pred.Surp._{i,q}^{X} = \frac{Consensus_{i,q}^{Acc} - Consensus_{i,q}^{X}}{p_{i,q}}, \qquad X \in \{EW5, IBES\}, \tag{B3}$$

where $Consensus_{i,q}^{Acc} = \sum_g w_{i,q}^{g} Consensus_{i,q}^{g}$ is the consensus EPS built with the weights of Equation (3), $Consensus_{i,q}^{X}$ the corresponding benchmark consensus, and $p_{i,q}$ the closing share price two trading days before the announcement. Table OA5, Panel B, reports its trading value. The standardized predicted surprise forecasts the four-day BHAR in all four columns ($p < 0.01$).[3] As in the trading-strategy literature (e.g., Battalio and Mendenhall, 2005; Michaely et al., 2024), a long–short portfolio on the quintiles of the predicted surprise earns 0.63 to 0.73 percent per announcement against the equal-weighted benchmark and 1.35 percent against I/B/E/S ($p < 0.01$).

[2] Part of this persistence is mechanical, because adjacent four-quarter windows share three quarters.

[3] The tests min-max standardize the predicted surprise to the unit interval. Table 2 reports the standard deviations of the two standardized predicted surprises, 0.116 and 0.109. Per standard deviation, the market-adjusted slopes of Panel B imply BHAR increments of about 0.29 and 0.37 percentage points against the equal-weighted and I/B/E/S benchmarks.

Online Appendix C: Supplementary Details

This appendix gives the detail behind five parts of the paper. It describes the five providers' institutional background (C.1), the analyst revision sample and herding specification (C.2), the predictive-value measure and two alternative explanations (C.3), the change in the premium by firm characteristic (C.4), and the search traffic to the GenAI tools behind the regime timing (C.5).

C.1 Institutional Background of the Five Providers

This subsection gives the providers' origins and the extent to which their reported actuals agree. I/B/E/S is credited with creating the consensus-earnings-estimate industry in the early 1970s, and Zacks claims to have first distributed quarterly consensus forecasts in the early 1980s. Bloomberg, Capital IQ, and FactSet come to the segment from adjacent product lines (Larocque et al., 2026). Beyond the pairwise disagreement of Table OA3, all five providers report the same actual in only 6.3 percent of firm-quarters, and the largest cross-provider gap reaches at least five cents per share in 78 percent (untabulated).

C.2 The Analyst Revision Sample and the Herding Specification

This subsection describes the analyst revision sample and the herding specification of Section 4.1.2, whose estimates are in Table 4. A revision pairs consecutive forecasts by the same analyst for the same firm-quarter, issued within the 180 days preceding the announcement and no more than 120 days apart. The sample contains 170,956 revisions by 2,283 analysts covering 759 firms, of which 122,001 with the full firm and analyst control sets enter the regressions (120,868 in the analyst$\times$firm pair specifications). The consensus toward which herding is measured is the prevailing I/B/E/S summary consensus that analysts observe, the mean estimate from the most recent statistical-period summary strictly preceding the revision date (at most 45 days old, at least three estimates). The regressions of Section 4.1.2 interact $Post_q$ with the $IBES\,Up_i$ and $IBES\,Down_i$ indicators of Table 3 (β_1 and β_2) or, in columns (3)–(4) of Table 4, with the three post-ChatGPT GPT-capability regimes ($\beta_{1,r}$ and $\beta_{2,r}$). Appendix A defines the analyst controls of Clement (1999). The distance fixed effects are twenty ventiles of the price-scaled distance between the analyst's prior and the consensus, and standard errors cluster by firm.

C.3 The Predictive-Value Measure and Two Alternative Explanations

This subsection describes the predictive-value measure of Section 4.1.3, the relation between predictive value and trailing accuracy, and two alternative explanations for the post-ChatGPT increment. The first stage of the test measures predictive value for each firm-quarter from the firm's own history of I/B/E/S consensus forecasts and street actuals. For a source quarter s, the year-ahead pair is the I/B/E/S consensus for quarter s and the I/B/E/S street actual for quarter $s+4$, the same fiscal quarter one year later. The window $S_{i,q}$ holds up to the twenty most recent source quarters whose year-ahead actuals are announced before the start of the calendar quarter in which firm i announces quarter q. I require at least twelve such quarters. For each firm-quarter, I regress the year-ahead actual on the consensus over the window by ordinary least squares.[4] The intercept $a_{i,q}$ and the slope $b_{i,q}$ are fitted on the window, and $e_{i,s}$ is the error of source quarter s,

$$Actual^{IBES}_{i,s+4} = a_{i,q} + b_{i,q}\, Consensus^{IBES}_{i,s} + e_{i,s}, \qquad s \in S_{i,q}. \tag{C1}$$

The residual $\hat{e}_{i,s}$ is the year-ahead actual net of its fitted value. Predictive value is one minus the ratio of the window's summed absolute residuals to its summed absolute year-ahead actuals,

$$Predictive\ Value^{IBES}_{i,q} = 1 - \frac{\sum_{s\in S_{i,q}} \left|\hat{e}_{i,s}\right|}{\sum_{s\in S_{i,q}} \left|Actual^{IBES}_{i,s+4}\right|}, \tag{C2}$$

with the ratio winsorized at the 1st and 99th percentiles.

The regression pairs the I/B/E/S consensus with the same provider's street actual because the providers' reported actuals disagree pervasively (see Table OA3 for details). A pair of quarters s and $s+4$ enters only if no stock split falls between them. The fitted intercept and slope absorb a level bias and a scale error, so the measure scores how closely a linear function of the consensus tracks the year-ahead actual. Predictive value is a continuous cross-firm variable (mean 0.67, standard deviation 0.25). It is negative for the 2 percent of firm-quarters whose summed absolute residuals exceed their summed absolute actuals. It covers 22,364 of the 23,833 firm-quarters (8,730 pre-ChatGPT), of which 22,356 enter the regressions. The window holds the full twenty source quarters for 21,140 of the 22,364. The test covers the I/B/E/S consensus alone, because

[4] The fitted slope is negative in 20.8 percent of the windows, and the R-squared is below 0.05 in 31.2 percent.

the measure requires up to twenty quarters of realized year-ahead outcomes and the five-FDP panel is too short to supply them for every FDP. Applied at the current-quarter horizon, the same procedure yields a like-for-like measure of the consensus's trailing accuracy, and across firms the two are highly correlated (Pearson 0.72, Spearman 0.75; 0.79 correcting for split-half measurement error; untabulated).

Two alternative explanations for the increment of Table 5 receive little support. The first is announcement composition. Adding the large-unexpected-item flag of Bochkay et al. (2022) with its full set of interactions leaves the increment at 3.14 ($t = 3.60$). The flag is nearly uncorrelated with predictive value (0.02; both untabulated). The second is selection toward I/B/E/S specifically. An untabulated decomposition splits the I/B/E/S surprise into a common component (the average of the I/B/E/S surprise and the mean surprise of FactSet, Zacks, and Capital IQ) and an I/B/E/S-specific component (half their difference). The common component carries both the slope of the ERC in predictive value (3.0, $t = 4.1$) and its post-ChatGPT increment (2.7, $t = 2.5$) on the market-adjusted $CAR_{[-1,1]}$.

C.4 The Change in the Premium by Firm Characteristic

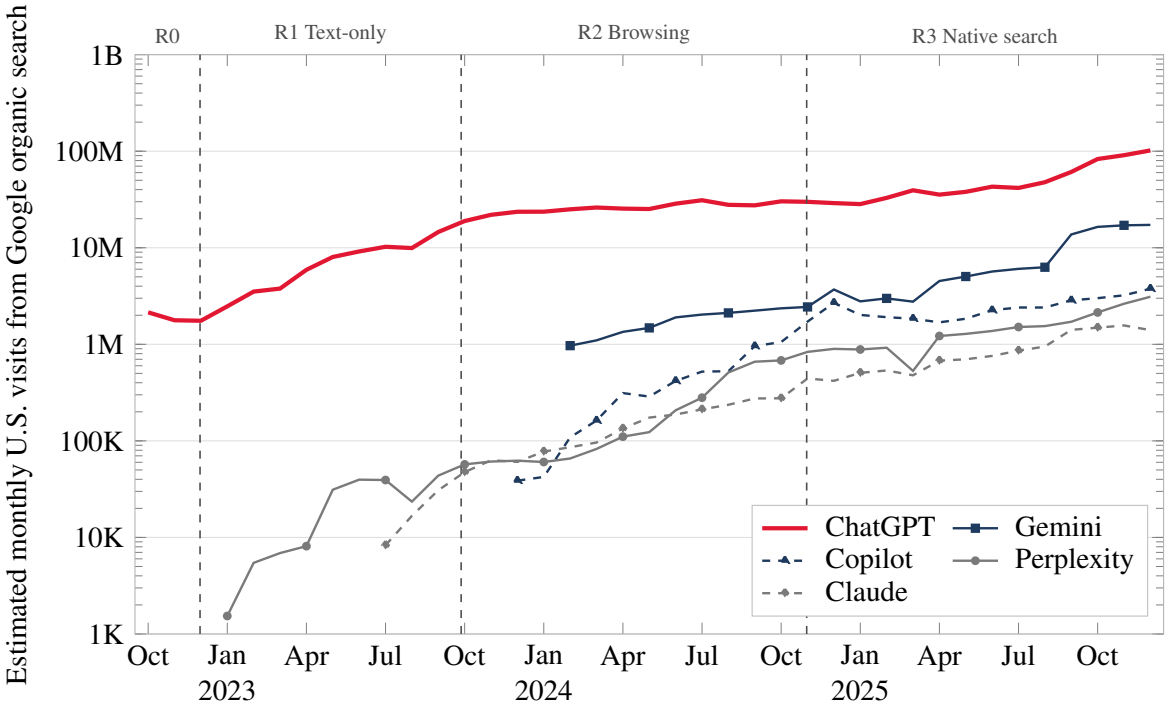
This subsection describes the median splits of Section 6, whose estimates are in Table OA4, Panel E. The splits divide the firms at the pre-ChatGPT median of the retail share of trading, of institutional ownership, and of analyst coverage, and the panel reports the high-minus-low difference in the post-ChatGPT change in the premium against each benchmark. Appendix A defines the three characteristics and the split.

C.5 Search Traffic to the GenAI Tools

This subsection reports the search-traffic evidence behind the regime timing of Section 3.3. Figure OA1 plots the estimated monthly visits that each tool receives from Google organic search in the United States, October 2022 to December 2025. When browsing is restored in September 2023, ChatGPT receives 334 times as many of these visits as Perplexity. At the launch of ChatGPT search, in October 2024, it receives 12.8 times as many as the next-largest tracked tool. In December 2025, it receives 5.9 times as many as Gemini. In every month of the window, ChatGPT draws

more of these visits than the other tracked tools together. The measure counts visits that arrive through Google search results, not total usage or investor adoption, and Bing Chat and Bard are absent from the source.

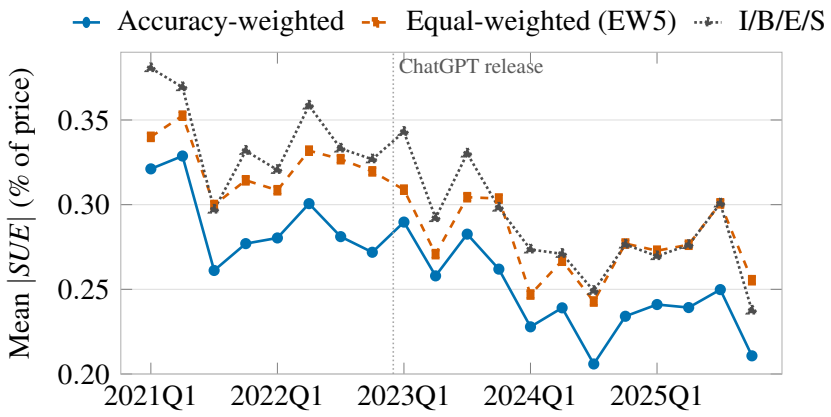
FIGURE OA1. Monthly Organic Search Traffic to GenAI Tools



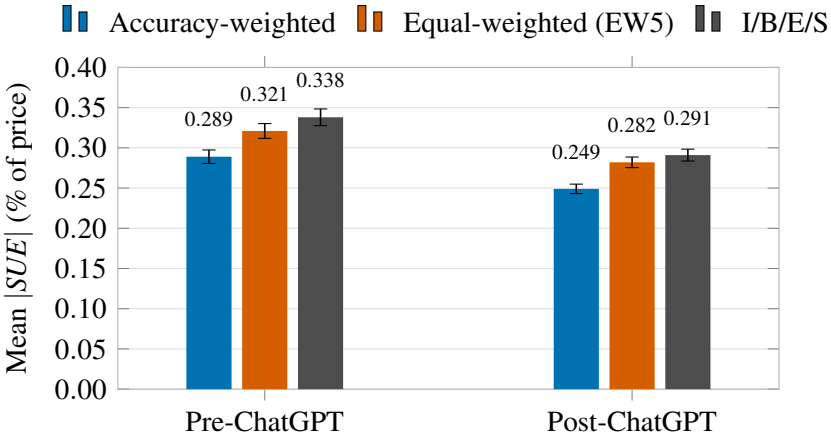
Notes. This figure compares the estimated monthly visits from Google organic search to ChatGPT with those to the other GenAI tools that launch near the regime boundaries. Each line plots a tool's estimated monthly visits from Google organic search results in the United States (DataForSEO), October 2022 to December 2025, on a logarithmic scale. The measure counts visits that arrive through Google search results, not total visits. ChatGPT is the sum of `openai.com` and `chatgpt.com`, because the product moves from `chat.openai.com` to `chatgpt.com` in mid-2024. The ChatGPT line therefore includes visits to `openai.com` in October and November 2022, before ChatGPT's release. Before June 2024, `chatgpt.com` registers fewer than 34,000 visits a month. Gemini is `gemini.google.com`, the name Bard takes in February 2024, and Copilot is `copilot.microsoft.com`, live since November 2023. Each line starts in the first month with at least 1,000 estimated visits. Bing Chat (February 2023) and Bard (March 2023) are not in the source, which has no `bing.com` or `bard.google.com` series. The dashed vertical lines mark the starts of the GPT-capability regimes R1, R2, and R3 of Appendix A (November 30, 2022; September 27, 2023; October 31, 2024).

FIGURE OA2. Mean Magnitude of *SUE*: Accuracy-Weighted vs. the Two Benchmarks

(a) Quarterly mean $|SUE|$ by surprise measure



(b) Mean $|SUE|$ by surprise measure



Notes. This figure compares the accuracy of the accuracy-weighted surprise SUE^{Acc} against the equal-weighted surprise SUE^{EW5} and the I/B/E/S surprise SUE^{IBES} across the five-FDP sample (FactSet, Bloomberg, Zacks, Capital IQ, and I/B/E/S) of S&P 1500 firm-quarters for fiscal quarters 2021Q1–2025Q4. $|SUE|$ is the absolute standardized unexpected earnings, scaled by the closing share price two trading days before the announcement and expressed in percent, winsorized at the 1st and 99th percentiles. Panel (a) plots the quarterly mean $|SUE|$ for each of the three surprise measures over twenty fiscal quarters. The vertical reference line marks the public release of ChatGPT (November 30, 2022). Panel (b) shows the pooled mean $|SUE|$ within each of the pre- and post-ChatGPT periods, with 95 percent confidence intervals from the cell-level standard error of the mean. The *y*-axis starts at zero so that bar height is proportional to the mean.

TABLE OA1. Integration-Prompt Pilot: What the Models Do When Asked

Model	Prompt	N	Share of responses with (a)	(b)	(c)	(d)	(a)–(d)	Actual match
ChatGPT `gpt-6.1-sol`	P1	100	0.77	0.67	0.63	0.34	0.28	0.97
ChatGPT `gpt-6.1-sol`	P2	100	0.81	0.64	0.52	0.40	0.35	0.98
ChatGPT `gpt-6.1-sol`	P2b	50	0.82	0.74	0.96	0.78	0.62	0.98
Perplexity `sonar`	P1	50	0.30	0.34	0.14	0.20	0.06	0.77
Perplexity `sonar`	P2	50	0.72	0.42	0.40	0.36	0.24	0.90

Notes. This table examines whether browsing models perform the steps of synthesis when asked. N is the number of responses per model and prompt. P2b is the P2 user prompt with the amended system message of Online Appendix A.2, run on ChatGPT only. Of its 50 responses, one per firm, 17 come from a rerun on October 11, 2026, of the calls that failed. In 16 of these 17, the model takes October 9 as the reference date. Columns (a) to (d) give the share of responses that (a) report a consensus and a reported EPS for at least two quarters for each of at least two sites, (b) rank at least two sites by accuracy, (c) assign at least two numeric weights that sum to one (0.9 to 1.1, or 90 to 110 percent), and (d) report a numeric accuracy-weighted consensus. Column (a)–(d) gives the share of responses that meet all four. Actual match is the mean, over responses with at least one claimed reported EPS in the firm's 2025 announcement window, of the share of those claims within $0.02 of the reported actual of any of the five providers for that firm and quarter. The grade A or B rows of Table OA2 valid on October 9, 2026, map eight domains, all to S&P Capital IQ or Zacks. Because no response ranks two mapped sites with different providers, the table omits the agreement between the models' site rankings and the providers' accuracy ranks.

TABLE OA2. Upstream Providers Disclosed by the Sites That the Model Cites

Site	Share of citations (%)	Disclosed provider	Grade
marketbeat.com	21.5	Fiscal.ai	
chartmill.com	20.0	None named	
alphaquery.com	6.1	Zacks Investment Research, Inc.	A
businessquant.com	6.0	Business Quant (self-attribution, no vendor named)	
marketxls.com	5.8	MarketXLS (self-attribution only, no estimate vendor named)	
barchart.com	4.8	Zacks (spelled "Zack's" on the page). The page footer says "Fundamental data provided by Zacks and Morningstar"	A
benzinga.com	4.3	None named	
fxempire.com	4.1	None named	
tipranks.com	4.1	None named for EPS or revenue estimates. The FAQ describes TipRanks' own collection of analyst ratings	
coincodex.com	3.9	None named	
zacks.com	3.3	Zacks (its own consensus database, compiled from brokerage feeds)	B
stockanalysis.com	3.0	S&P Global Market Intelligence (consensus and forecasts) and TipRanks (individual analyst data)	A
investing.com	1.0	None named	
finance.yahoo.com	0.4	S&P Global Market Intelligence, from January 29, 2025	B
finance.yahoo.com	0.4	LSEG Data and Analytics (called Refinitiv in the December 16, 2022, and October 1, 2023, copies), until January 20, 2025	B
nasdaq.com	0.1	Zacks Investment Research	A
seekingalpha.com	0.1	S&P Global Market Intelligence	B
fiscal.ai	0.0	S&P Market Intelligence LLC and its suppliers, for "certain financials and estimates." The API terms say only "Third-Party suppliers"	B

Notes. This table reports, for the 17 sites that receive 88.7 percent of the citations in the May 28 run of the queries behind *LLM Visibility*, the provider that each site names as the source of its estimates. Share of citations is the site's share of the 51,933 citations in the 3,912 responses of that run. Grade A means that the site's estimates page names the provider, grade B that a help, methodology, data-source, or terms page does, and a blank grade that the site names none of the five forecast data providers. Yahoo Finance has two rows because its help page names LSEG Data and Analytics in copies up to January 20, 2025, and S&P Global Market Intelligence from January 29, 2025. MarketBeat credits its estimates to Fiscal.ai, whose terms name S&P Global Market Intelligence only for certain financials and estimates. The table therefore leaves MarketBeat undisclosed. Counting MarketBeat as an S&P site would raise the share of citations that go to sites carrying S&P's estimates from 3.5 to 25.1 percent.

TABLE OA3. Reported-Actual Disagreement and Provider Accuracy Shares Across FDPs and GPT-Capability Regimes

Panel A. Pairwise disagreement over the 23,833 sample firm-quarters (percent)

	I/B/E/S	FactSet	Bloomberg	Zacks
FactSet	22.27			
Bloomberg	88.50	88.66		
Zacks	33.63	38.30	90.84	
Capital IQ	16.31	21.06	88.95	32.41

Panel B. Pairwise disagreement by GPT-capability regime (percent)

Regime:	(1) R0 Pre-ChatGPT	(2) R1 Text-only	(3) R2 Browsing	(4) R3 Native search	t-statistic R1−R0	t-statistic R2−R1	Joint p
I/B/E/S–FactSet	24.62	23.11	22.15	17.47	-4.09	0.03	0.000
I/B/E/S–Bloomberg	88.49	87.95	88.88	88.52	-1.03	1.64	0.357
I/B/E/S–Zacks	32.89	34.27	34.21	33.92	2.27	0.28	0.066
I/B/E/S–Capital IQ	17.02	16.36	16.10	15.22	-1.04	-1.28	0.021
FactSet–Bloomberg	88.86	87.84	88.64	88.91	-2.08	1.83	0.022
FactSet–Zacks	38.63	39.25	38.86	36.42	-0.37	1.44	0.502
FactSet–Capital IQ	23.17	22.12	20.89	16.61	-3.97	-1.11	0.000
Bloomberg–Zacks	90.55	90.58	91.17	91.24	0.37	1.10	0.195
Bloomberg–Capital IQ	88.75	88.59	89.19	89.31	-0.56	0.79	0.291
Zacks–Capital IQ	31.81	32.80	33.07	32.53	1.28	0.64	0.182

Panel C. Provider accuracy shares by GPT-capability regime (percent)

	R0 Pre-ChatGPT	R1 Text-only	R2 Browsing	R3 Native search
I/B/E/S	16.9	18.4	17.6	19.7
FactSet	20.9	19.8	22.4	21.7
Bloomberg	17.2	14.5	14.4	13.6
Zacks	29.5	30.9	29.1	28.6
Capital IQ	16.1	16.9	17.0	16.8
Firm-quarters	4,167	3,287	5,077	4,631

Joint test, all providers: distribution equal across the post-ChatGPT regimes, $\chi^2(8) = 10.8$, $p = 0.214$.

Notes. This table reports how often the FDPs' reported actuals differ and how often each FDP's consensus is the most accurate. Panel A reports the percentage of the 23,833 sample firm-quarters in which a pair's reported actuals are not identical. All five actuals are non-missing for every sample firm-quarter. In Panel B, the t-statistics test the within-firm differences between adjacent regimes in a regression of the pair's disagreement indicator on regime indicators with firm fixed effects. Joint p is the p-value of the Wald test that the three regime coefficients are zero. Panel C reports the share of the firm-quarters with complete trailing-error windows for all five providers in which each provider's consensus is the *ex ante* most accurate (lowest trailing four-quarter $MAFE^g$). A firm-quarter with tied lowest $MAFE^g$ is credited to each tied provider, so a column can sum to slightly above 100 percent. The joint Wald test stacks the five provider-share equations. Appendix A defines the variables. Standard errors are clustered by firm.

TABLE OA4. Additional Tests of the Accuracy Premium

Panel A. The timing of the premium within the post-period

	Equal-weighted five-FDP		I/B/E/S	
Post-period from browsing: Δ Premium	**1.766****	**(2.04)**	**2.476*****	**(3.21)**
Both boundaries: Δ Premium at release	**2.256***	**(1.87)**	**1.784**	**(1.61)**
Both boundaries: Δ Premium at browsing	**0.216**	**(0.18)**	**1.236**	**(1.13)**
Both boundaries: cumulative Δ Premium	**2.472*****	**(2.62)**	**3.020*****	**(3.62)**

Panel B. The response to each provider's own surprise

	ERC before ChatGPT		Post-ChatGPT increment	
I/B/E/S	3.022***	(16.05)	0.694**	(2.57)
FactSet	2.783***	(13.73)	0.676**	(2.45)
Bloomberg	1.102***	(9.48)	0.179	(1.24)
Capital IQ	3.247***	(15.86)	0.687**	(2.47)
Zacks	3.618***	(16.94)	1.452***	(4.72)

Panel C. The premium over the Estimize consensus

	Market-adjusted CAR		FFC CAR	
Premium over the Estimize consensus, R3	**5.067****	**(2.38)**	**5.448****	**(2.57)**
Change in that premium, R0 to R3	**5.164*****	**(2.81)**	**5.195*****	**(2.81)**

Panel D. The results without Bloomberg (four providers)

	Equal-weighted four-FDP		I/B/E/S	
Premium (Pre-ChatGPT)	-3.573***	(-2.62)	1.565**	(2.11)
Premium (Post-ChatGPT)	0.982	(0.68)	4.804***	(6.14)
Δ Premium (Post−Pre)	**4.555****	**(2.41)**	**3.239*****	**(3.22)**

Panel E. The change in the premium by firm characteristic

	Equal-weighted five-FDP		I/B/E/S	
High-minus-low Δ Premium at the pre-ChatGPT median of:				
Retail Share of Trading	**1.284**	**(0.72)**	**1.453**	**(0.95)**
Institutional Ownership	**1.757**	**(1.00)**	**0.482**	**(0.31)**
Analyst Coverage	**-7.281*****	**(-4.28)**	**-1.849**	**(-1.14)**

Notes. This table reports additional tests of the accuracy premium, the loading on SUE^{Acc} minus that on the benchmark surprise, and of each provider's earnings response coefficient (ERC). Panel A starts the post-period at the browsing boundary (September 27, 2023) or uses both boundaries. Panel B estimates each provider's ERC in one regression per provider. Panel C reports the premium over the Estimize consensus on the firm-quarters with an Estimize consensus (2,564 under native search, 13,791 in all). Its benchmark surprise SUE^{Est} is the Estimize reported EPS minus the consensus of the day before the announcement (at least three estimates, at most 30 days old), scaled by the price that scales SUE. Panel D repeats the pooled horse race of Table 6 with the four providers other than Bloomberg (23,832 firm-quarters), SUE^{Acc4} with the weights of Section 3.5 computed across the four and SUE^{EW4} their mean. Panel E estimates each high-minus-low difference in a pooled regression per characteristic and benchmark (23,476 firm-quarters for the retail share, 23,515 for the other two). The dependent variable is the market-adjusted $CAR_{[-1,1]}$ unless a column head names another. The sample is the 23,833 firm-quarters unless noted. All regressions include the firm and event controls and firm and year fixed effects. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

TABLE OA5. Accuracy Persistence and the Predicted-Surprise Trading Strategy

Panel A. Persistence of FDP accuracy ranks (trailing-MAFE rank, quarter t to $t+1$)

	$\Pr(Rank_{t+1} = j \mid Rank_t = i)$				
$Rank_t$	$j=1$ *(least)*	$j=2$	$j=3$	$j=4$	$j=5$ *(most)*
1 *(least)*	0.761	0.121	0.049	0.034	0.036
2	0.121	0.541	0.197	0.089	0.052
3	0.047	0.195	0.467	0.211	0.080
4	0.034	0.091	0.208	0.483	0.185
5 *(most)*	0.038	0.052	0.079	0.184	0.647

Spearman $\rho = 0.71$; overall same-rank rate $= 0.58$; 109,915 quarter-to-quarter pairs.

Panel B. Trading strategy from predicted surprise

Dep. Var.:	Market-adj. $BHAR_{[-2,1]}$		FFC $BHAR_{[-2,1]}$	
Benchmark:	Equal-weighted	I/B/E/S	Equal-weighted	I/B/E/S
	(1)	(2)	(3)	(4)
Standardized Pred. Surp.	**2.493*****	**3.433*****	**2.440*****	**3.357*****
	(4.10)	**(5.35)**	**(4.08)**	**(5.25)**
Long−Short Return (Q5−Q1)	**0.73%*****	**1.35%*****	**0.63%*****	**1.35%*****
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations (firm-quarters)	23,833	23,833	23,833	23,833
Adjusted R^2	0.016	0.017	0.018	0.019

Notes. This table examines whether synthesizing the five FDPs' forecasts by recent accuracy brings tradable benefits to investors. Panel A reports the share of provider-quarters at Rank i that move to Rank j the following quarter (rows sum to one up to rounding). Panel B regresses the four-day announcement-window abnormal return on the standardized predicted surprise against each benchmark. The Long−Short row reports the abnormal return, in percent per announcement, from going long the top quintile and short the bottom quintile of the predicted surprise over the $[-2,1]$ window. Continuous variables are winsorized at the 1st and 99th percentiles. Appendix A defines the variables. Standard errors are clustered by firm, with t-statistics in parentheses. $^{*}p<0.10$, $^{**}p<0.05$, $^{***}p<0.01$.

TABLE OA6. Randomization Inference for the Change in the Accuracy Premium Gap with Retrieval

	Gap				
	Pre-retrieval (R0–R1)	Retrieval (R2–R3)	Change	Randomization p-value	2.5th–97.5th percentiles of the permuted changes
Equal-weighted five-FDP benchmark:					
Median split (High Vis − Low Vis)	-0.814	2.109	2.923**	0.040	[-3.35, 3.27]
Continuous measure (per SD)	-0.169	1.555	1.723**	0.020	[-1.77, 1.64]
I/B/E/S benchmark:					
Median split (High Vis − Low Vis)	-1.969	3.879	5.847***	0.000	[-3.15, 3.09]
Continuous measure (per SD)	-0.508	2.012	2.520***	0.001	[-1.53, 1.58]

Notes. This table examines whether the accuracy premium gap between high- and low-visibility firms widens from the pre-retrieval regimes, R0 and R1, to the retrieval regimes, R2 and R3. It compares each change with the changes obtained when the measure is reassigned across firms at random. Each row reports one test from the pooled regression of Table 11, Panel C, with the gap before retrieval, the gap with retrieval, and the change between them. The randomization p-value is the share of 2,000 permutations of the measure across firms with a change at least as large as the estimated change, in its direction. Each firm keeps one value in all its quarters, and for the median split each permutation is a random split with 654 High firms. The last column gives the 2.5th and 97.5th percentiles of the permuted changes. Table 11, Panel C, reports the t-statistics, with standard errors two-way clustered by firm and fiscal year. Each regression uses 23,827 firm-quarters. Appendix A defines the variables. *, **, and *** denote one-sided randomization-inference p-values below 0.10, 0.05, and 0.01.