

Synthesizing the Consensus: Generative AI, LLM Visibility, and the Integration of Multi-Provider Consensus Forecasts*

Leonard Yang Liu[†]

September 2026

[\[Newest Version\]](#)

Abstract

The consensus earnings forecast is not one number: forecast data providers (FDPs) construct economically distinct consensus numbers for the same firm-quarter. This paper asks whether generative AI (GenAI) changes how equity markets synthesize consensus forecasts across providers. I study a panel of S&P 1500 firm-quarters from 2021 to 2025 with consensus forecasts from five FDPs—FactSet, Bloomberg, Zacks, S&P Capital IQ, and I/B/E/S. After ChatGPT’s release, the market reacts more strongly to forecasts from more-accurate FDPs: provider-specific earnings response coefficients align more closely with each FDP’s accuracy rank. Post-ChatGPT, market reactions also align increasingly with an accuracy-weighted consensus across the five FDPs rather than with the I/B/E/S-only or equal-weighted heuristic alternatives. The shift concentrates where GenAI-enabled synthesis is most valuable: in firm-quarters with high cross-FDP disagreement and larger nonrecurring items, and—once GenAI tools can retrieve forecasts from the web—in firms with high large language model (LLM) visibility. A long–short strategy trading the gap between the accuracy-weighted and heuristic consensus measures earns economically meaningful returns. Google search data support a post-ChatGPT shift in investor search toward free forecast distributors, and tests using major ChatGPT outages as exogenous shocks point to GenAI as the driver of the shift. The evidence suggests that GenAI helps investors both gather and integrate forecast information—and changes how that information is impounded into prices.

Keywords: Generative AI, Analyst, Forecast Data Provider, Consensus Forecast, Earnings Surprise, Information Processing, Earnings Response Coefficient.

JEL classification: G11, G14, G41, M41.

*I am deeply grateful to my dissertation committee for their invaluable guidance and support: Musa Subasi (chair), Michael Kimbrough (co-chair), Rebecca Hann, Nick Seybert, Jingyi Qian, and Erkut Y. Ozbay. I gratefully acknowledge the Texas A&M University AI Dissertation Award in Accounting. All errors are my own.

[†]Department of Accounting and Information Assurance, Robert H. Smith School of Business, University of Maryland. Email: yliu337@umd.edu.

1 Introduction

“The consensus” earnings forecast is not one number. Sell-side analyst forecasts serve as the primary market benchmark against which realized earnings are evaluated (e.g., [Bradshaw et al., 2017](#); [Kothari, 2001](#)). The earnings surprise relative to this consensus is a key determinant of announcement-window stock price reactions and post-announcement drift (e.g., [Bartov et al., 2002](#); [Bernard and Thomas, 1989](#); [Livnat and Mendenhall, 2006](#)), with measurable consequences for price discovery, trading volume, and firms’ own disclosure choices. Yet multiple *forecast data providers* (FDPs) construct economically distinct consensus numbers for the same firm-quarter through proprietary aggregation rules. [Larocque et al. \(2025\)](#) document these differences across the five leading FDPs: I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks.¹ Although these FDPs are widely used and cited for investment decisions ([Coleman et al., 2025](#); [Xue et al., 2025](#)), little is known about whether the market reacts differently across FDPs and, more importantly, how it selects and integrates across them.

The question matters because acquiring and integrating data across multiple providers is costly.² Since the release of ChatGPT, generative AI (GenAI) has diffused rapidly into investor research and analysis workflows, reshaping the economics of how financial information is gathered and combined. Industry surveys document widespread adoption ([CFA Institute, 2024](#)), and a growing body of research shows that these tools substantially lower the costs of acquiring, integrating, and synthesizing financial information (e.g., [Bertomeu et al., 2025](#); [Blankespoor et al., 2026](#); [Chang et al., 2026](#); [Cheng et al., 2025](#); [Liu and Subasi, 2026](#); [Xue et al., 2025](#)). Prior technological advances such as EDGAR, XBRL, and social-media disclosure each lowered a specific component of investors’ information-processing costs (e.g., [Blankespoor, 2019](#); [Blankespoor et al., 2014](#); [Drake et al., 2015](#)), but none automated what GenAI uniquely enables: cross-FDP synthesis of analyst consensus forecasts. This paper asks whether the post-ChatGPT market values an accuracy-weighted

¹In a sample of 94,030 firm-quarters covered by all five FDPs over 2002–2020, [Larocque et al. \(2025\)](#) find that the FDPs report different earnings surprises (beyond rounding) for the same firm-quarter 58 percent of the time, and that at least one FDP signals a miss while another signals a beat in 9 percent of firm-quarters.

²[Larocque et al. \(2025\)](#) show that announcement-window returns align more with I/B/E/S and Capital IQ signals than with the other FDPs, consistent with acquisition cost being a binding factor.

multi-FDP consensus more than the conventional single-FDP consensus or the equal-weighted average across providers.³

The case for cross-FDP synthesis rests on the substantive disagreement among providers, who differ in their screening rules, treatment of unusual items, and analyst-coverage thresholds (Bochkay et al., 2022). Which FDP yields the most accurate consensus varies across firms and over time, with I/B/E/S ranking highest on average, consistent with its prevalence as the academic-research standard (Larocque et al., 2025). This variation implies meaningful gains from synthesizing across providers. Historically, such synthesis was impractical: investors had to recognize that the FDPs disagree, subscribe to each consensus separately, and reconcile differing methodologies.

Recognize, subscribe, reconcile—these three frictions correspond to the awareness, acquisition, and integration costs in the information-processing-cost framework of Blankespoor et al. (2020). GenAI attenuates all three, but the binding friction for accuracy-weighted multi-FDP aggregation has historically been the third: cross-source integration of heterogeneous consensus measures, which prior single-source technologies left intact.⁴ The framework yields two empirical predictions: (i) the post-ChatGPT market should shift its implicit consensus toward an accuracy-weighted synthesis across multiple FDPs, and (ii) this shift should concentrate where GenAI synthesis is most accessible and most valuable.

Neither prediction is obvious ex ante. The strongest reason for doubt is theoretical: in the forecast-combination literature, combining heterogeneous, partially independent forecasts almost always dominates sub-selecting on accuracy (e.g., Ahlers and Lakonishok, 1983; Bates and Granger, 1969). Consistent with this logic, Michaely et al. (2024) show that constructing a consensus from high-quality analysts within a single provider does not improve investor outcomes; the frictions they identify carry over with even greater force to the cross-FDP setting. A separate concern is

³The accuracy-weighted construction is not a claim that the market literally applies this weighting rule. It serves as a tractable proxy for the broader hypothesis that the post-ChatGPT market draws on multiple FDPs and discriminates by source accuracy; if so, market reactions should align more closely with this benchmark than with the single-FDP or equal-weighted alternatives, even if the specific weights the market applies differ from mine.

⁴A single LLM query of the form “What is the consensus EPS forecast for [ticker] next quarter?” typically returns a synthesized response citing several free third-party distributors (e.g., Yahoo Finance, MarketBeat, TipRanks). Appendix B reports two worked examples (IBM and Microsoft) showing the cited distributors, their reported figures, and the disclosed upstream FDP where available.

that lower acquisition costs do not imply more selective integration: a richer information set need not improve forecast accuracy when the synthesis itself is unselective, and LLMs in particular can rely heavily on memorization rather than reasoning (e.g., [Lopez-Lira et al., 2025](#)). Whether GenAI nonetheless shifts the market toward a more selective synthesis is therefore an empirical question.

I test these predictions using a panel of S&P 1500 firm-quarters from 2021 to 2025 with consensus EPS forecasts for the same firm-quarter from all five FDPs.⁵ After ChatGPT’s release, the market reacts more strongly to forecasts from accurate FDPs and discounts those from less accurate ones. In a within-firm-quarter rank-based test that estimates provider-specific earnings response coefficients (ERCs) by accuracy rank (e.g., [Collins and Kothari, 1989](#); [Easton and Zmijewski, 1989](#)), the ERC increment per one-rank accuracy gain roughly doubles after ChatGPT, from 0.31 to 0.62 (post-pre differential 0.32, $p < 0.01$). A complementary firm-level design partitions firms by whether I/B/E/S’s accuracy rank rose or fell from pre- to post-ChatGPT. The result is asymmetric: the market response to I/B/E/S-based earnings surprises weakens significantly where I/B/E/S lost rank but not where it gained, consistent with pre-GenAI overweighting of the I/B/E/S consensus.

Given this evidence of selective use, I next test whether the market’s aggregation rule itself has shifted toward an accuracy-weighted synthesis. I construct an accuracy-weighted standardized unexpected earnings (SUE) in which each FDP’s SUE is weighted by its trailing forecast accuracy, and benchmark it against two heuristic alternatives: the single-FDP I/B/E/S consensus and the equal-weighted average across the five FDPs.⁶ Post-ChatGPT, ERCs on the accuracy-weighted surprise rise sharply, while ERCs on the I/B/E/S and equal-weighted surprises remain flat or decline. The resulting accuracy premium widens significantly against both heuristic benchmarks. In economic terms, from pre- to post-ChatGPT, the three-day announcement CAR generated by a one-percentage-point standardized surprise rises 2.40 percentage points more under the

⁵The sample begins in 2021 for two reasons: (i) it yields a relatively balanced panel of pre- and post-ChatGPT firm-quarters around the November 30, 2022 cut, and (ii) Bloomberg’s historical consensus data are accessible only for the past five years. Given the different methodologies used by FDPs, I compute each FDP’s forecast accuracy from its own end-of-quarter consensus EPS and reported actual EPS.

⁶Specifically, the weight on FDP g is the softmax of the negative of its z -scored four-quarter rolling mean absolute forecast error, so providers with better recent track records receive larger weights. [Figure 2](#) confirms that the accuracy-weighted SUE has the smallest mean absolute forecast error of the three benchmarks, with the difference statistically significant. All variable definitions are in [Appendix A](#).

accuracy-weighted SUE than under the equal-weighted benchmark, and 2.67 percentage points more than under I/B/E/S.

The timing and the mechanism of this shift both point to GenAI. Re-estimated within four GPT-capability regimes, the premium against I/B/E/S climbs from negative in the pre-ChatGPT and text-only regimes to 1.52 under browsing and 4.91 under native search—a dose–response staircase that tracks the tools’ retrieval capabilities. The gating condition is acquisition: a firm-level *LLM visibility*⁷—how many forecast sources a browsing-enabled model actually retrieves for the firm—is empirically distinct from classical prominence proxies, and the premium concentrates in high-visibility firms only once the model can browse. Together, these shifts imply that the consensus reflected in market prices has moved from heuristic alternatives toward an accuracy-weighted multi-FDP synthesis.

Two cross-sectional tests show that the accuracy premium concentrates in settings where GenAI-enabled synthesis is most valuable. Because the regime evidence localizes the premium to the browsing and native-search regimes, these tests use a post-scraping indicator—equal to one for announcements on or after the onset of GPT browsing (September 27, 2023), when web retrieval first made cross-provider synthesis feasible—in place of the post-ChatGPT indicator. First, the premium concentrates in firm-quarters with high cross-FDP disagreement, where accuracy-weighting offers the largest improvement over heuristic benchmarks. Second, it concentrates in firm-quarters with large nonrecurring items: because FDPs differ most in their treatment of unusual items (Bochkay et al., 2022), this is where multi-FDP synthesis most attenuates single-provider methodology noise.

I next examine whether the accuracy-weighted consensus is economically valuable to investors.⁸ I construct a predicted earnings surprise—the price-scaled difference between the accuracy-weighted and each heuristic consensus—as an ex ante predictor of post-announcement returns. Regressing the four-day $[-2, +1]$ announcement-window buy-and-hold abnormal return (BHAR) on the standardized predicted surprise yields a positive coefficient ($p < 0.01$ for both

⁷Operationally, LLM visibility is an EPS-extraction indicator times $\log(1 + \text{unique cited domains})$, averaged across ChatGPT and Perplexity responses to firm-specific EPS queries (from the DataForSEO LLM Responses API).

⁸This is not a foregone conclusion: Michaely et al. (2024) show that historically a consensus from high-quality analysts within a single provider did not generate significantly better investor outcomes.

heuristic benchmarks), indicating that the predicted surprise carries return-relevant information. Following the earnings-surprise trading-strategy literature (e.g., [Battalio and Mendenhall, 2005](#); [Michaely et al., 2024](#)), a quintile long–short portfolio on the predicted surprise earns 0.63–0.73 percent per announcement against the equal-weighted benchmark and 1.35 percent against I/B/E/S (all $p < 0.01$).⁹

Three additional tests probe the GenAI channel. First, a forecast-based informativeness measure captures, firm by firm and in real time, how closely the consensus has anticipated the firm’s year-ahead street numbers over the preceding twenty quarters. The post-ChatGPT rise in announcement-window reactions concentrates in firms with an informative consensus: the ERC–informativeness gradient roughly doubles, significant under both return adjustments; the increment is orthogonal to announcement composition; and it prices the common cross-provider consensus rather than any single provider. Second, GenAI might merely substitute for information investors previously acquired directly from free distributors. Google search-volume data, a standard proxy for investor information demand (e.g., [Da et al., 2011](#); [Drake et al., 2012](#)), suggest otherwise: post-ChatGPT search share for free forecast-data distributors rises relative to paywalled FDPs, consistent with GenAI acting as a gateway to, rather than a substitute for, direct investor acquisition. Third, following [Cheng et al. \(2025\)](#) and [Liu and Subasi \(2026\)](#), I use major ChatGPT outages as an exogenous shock to GenAI availability. On announcement days following a major outage, the accuracy premium against I/B/E/S collapses while the equal-weighted contrast shows no significant change—an asymmetry consistent with GenAI as the integration technology underlying the premium: when multi-FDP synthesis is unavailable, investors revert specifically to the I/B/E/S consensus.¹⁰

This paper makes three contributions. First, I show that the market’s reliance on FDPs is not static but technology-dependent, extending recent evidence on cross-FDP differences in analyst

⁹Specifically, the portfolio goes long firm-quarters with positive predicted surprise and short those with negative predicted surprise, holding each position over the four-day $[-2, +1]$ announcement window. The per-trade alpha is estimated by panel OLS with firm-clustered standard errors, and annualized by scaling by 252 trading days divided by the four-day holding window. Results are robust to alternative event windows, including $[-2, +2]$ and $[-1, +1]$.

¹⁰Outages are sourced from OpenAI’s public status-page history; I require an incident of at least 240 minutes, yielding four incidents across five unique calendar days through 2024.

consensus information ([Larocque et al., 2025](#)), which documents substantial methodological and accuracy differences across providers. When information-processing costs fall, the market reweights its implicit consensus toward the more accurate provider's information, with measurable effects on earnings-announcement returns. For researchers, the implication is methodological: the FDP whose consensus best reflects market expectations shifts with information technology.

Second, I add a distinct mechanism to the literature on the capital-market consequences of generative AI (e.g., [Bertomeu et al., 2025](#); [Blankespoor et al., 2026](#); [Chang et al., 2026](#); [Cheng et al., 2025](#); [Liu and Subasi, 2026](#); [Lopez-Lira and Tang, 2023](#); [Xue et al., 2025](#)). Prior work documents that GenAI lowers investor processing costs and shifts prices around specific events such as outages or regulation. The mechanism here is synthesis: GenAI makes it cheap for the market to combine forecasts across multiple FDPs, and this combined consensus is now priced into earnings-announcement reactions.

Third, I sharpen the disclosure-processing-cost framework of [Blankespoor et al. \(2020\)](#) by isolating which cost component drives the market's accuracy-weighted use of consensus information. A long line of work shows how prior single-source technologies reduce specific cost components: EDGAR for filing acquisition ([Drake et al., 2015](#); [Loughran and McDonald, 2017](#)), XBRL for structured-data processing ([Blankespoor, 2019](#)), social media for awareness latency ([Blankespoor et al., 2014](#)), and automated text processing for retail-investor processing ([Blankespoor et al., 2018](#)). The evidence here shows that the binding friction for accuracy-weighted multi-FDP aggregation was cross-source integration of heterogeneous consensus measures. The earlier technologies left that friction intact; GenAI relaxes it, with measurable consequences for earnings-announcement pricing.

The remainder of the paper is organized as follows. Section 2 reviews the related literature and presents the theoretical framework. Section 3 describes the empirical setting. Section 4 presents initial evidence that the market begins to rely selectively on more-accurate FDPs after ChatGPT. Section 5 tests whether the post-ChatGPT market prices the accuracy-weighted synthesis across FDPs, documents its tradeable returns, traces the premium across GPT-capability regimes, and

identifies the cross-sectional drivers of the resulting accuracy premium. Section 6 shows that LLM visibility—a direct measure of GenAI forecast acquisition—is distinct from classical prominence proxies and gates the premium. Section 7 constructs a forecast-based informativeness measure and tests its post-ChatGPT pricing, examines investor information acquisition, and exploits major ChatGPT outages as an exogenous shock to GenAI availability. Section 8 concludes.

2 Related Literature and Theoretical Framework

2.1 Forecast Data Providers and the Construction of Analyst Consensus

The market discriminates on forecast quality at the level of the individual analyst; no comparable discrimination has been documented at the level of the consensus. Sell-side analyst earnings forecasts have long served as the benchmark proxy for market expectations in research and practice (Bartov et al., 2002; Beyer et al., 2010; Bradshaw et al., 2017; Kothari, 2001), in part because analysts integrate firm-specific guidance, industry knowledge, and macroeconomic conditions in ways that simple time-series extrapolation does not (Bradshaw et al., 2012; Schipper, 1991). Individual forecast accuracy varies systematically with an analyst’s experience, employer resources, and portfolio complexity (Clement, 1999; Mikhail et al., 1997), and revisions from more accurate or higher-status analysts are met with larger price reactions (Clement and Tse, 2003; Park and Stice, 2000; Stickel, 1992). At the consensus level, however, the analogous accuracy-weighted sub-selection has not been shown to dominate the pooled benchmark, even within a single provider (Michaely et al., 2024). Understanding why discrimination stops at the analyst level begins with how the consensus is made.

The consensus number itself is not a primitive of nature: third-party forecast data providers (FDPs) construct it by pulling each underlying analyst forecast through proprietary screening, exclusion, and adjustment rules.¹¹ Bochkay et al. (2022) document the first-order capital-market consequences of these construction choices using Thomson Reuters’ (now Refinitiv’s) 2009

¹¹FDPs differ, among other dimensions, in (i) which broker estimates they include (some exclude forecasts deemed stale or non-comparable; others apply tighter recency cutoffs), (ii) how they treat nonrecurring items (some report GAAP-based earnings, others a “street” figure adjusted for one-time charges), and (iii) rounding and staleness conventions.

methodology change for I/B/E/S.¹² An FDP's treatment of unexpected charges and gains, they show, feeds directly into the time-series predictive properties of street earnings and into price discovery around announcements. The FDP layer, on this evidence, acts as a producer of information rather than a conduit for it.

The same construction discretion breeds pervasive disagreement across providers. Comparing the five FDPs in a large sample, [Larocque et al. \(2025\)](#) find that providers report different earnings surprises (beyond rounding) for the same firm-quarter 58 percent of the time, and that in 9 percent of firm-quarters at least one FDP signals a beat while at least one other signals a miss. These disagreements are associated with differences in price responsiveness, abnormal volatility, and liquidity around earnings announcements. The accuracy ranking across FDPs is itself unstable across firms and over time, so no fixed choice of provider reliably delivers the most accurate signal.

Disagreement of this magnitude leaves substantial room, in principle, for investors to synthesize across providers rather than anchor on a single one.¹³ If the market could anticipate each FDP's accuracy ex ante, accuracy-weighted aggregation across providers could outperform any single-provider benchmark. Historically, the market appears not to have exploited this room: [Larocque et al. \(2025\)](#) document that I/B/E/S ranks highest on their composite measures of FDP quality and that market reactions align most closely with the I/B/E/S surprise, consistent with I/B/E/S serving as the historical anchor.

The closest evidence on why the room went unexploited comes from [Michaely et al. \(2024\)](#), who examine the within-provider analog. A consensus built from high-quality analysts inside I/B/E/S yields no meaningful improvement in investor outcomes over the full-analyst consensus—a null they trace to three frictions: error cancellation in the broader pool, imperfect persistence in

¹²In September 2009, Thomson Reuters stopped relying on analysts' ex post consensus to determine the I/B/E/S street-earnings actual when an earnings press release reported unexpected charges or gains, adopting instead a uniform rule that immediately excludes such items from GAAP earnings. Because the change rolled out gradually, [Bochkay et al. \(2022\)](#) exploit the staggered variation in a difference-in-differences design.

¹³Classical forecast-combination theory ([Bates and Granger, 1969](#); [Clemen, 1989](#); [Timmermann, 2006](#)) establishes that the accuracy gain from combining forecasts rises with the heterogeneity of their errors and falls with their pairwise correlation; whether the gain is realizable in practice depends on the cost of locating and combining the underlying sources, a point developed below.

analyst quality, and the cognitive cost of identifying high-quality forecasters.¹⁴ Each of these frictions operates with greater force across providers, where methodologically distinct consensus numbers must also be reconciled. The open question is whether a large enough reduction in the cost of identifying, acquiring, and integrating accurate sources can unlock the consensus-level discrimination that the historical environment failed to produce. The next subsection turns to the technologies that determine this cost.

2.2 Technological Advancements and Information Processing Costs in Capital Markets

Whether investors use available information depends on what it costs them to do so. [Blankespoor et al. \(2020\)](#) provide the canonical economic frame, decomposing the cost of using any piece of information into three components: *awareness* of its existence, *acquisition* of the information itself, and *integration* of that information into the investor’s decision. Each component is sensitive to technology, and an exogenous reduction in any one of them—holding the underlying information fixed—can measurably change investor behavior, with downstream consequences for asset prices and for firms’ own reporting choices.

An extensive empirical literature documents these effects component by component, with much of the identifying variation coming from single-source technological shocks introduced by regulators or by changes in market infrastructure. For awareness, [Blankespoor et al. \(2014\)](#) find that firm-initiated social-media dissemination broadens the investor base that responds to earnings news. On the acquisition side, [Drake et al. \(2015\)](#) show that EDGAR search activity tracks firm-level information demand and that surges in retrieval precede price adjustments, while [Loughran and McDonald \(2017\)](#) use granular access logs to show that few investors view most filings but that the filings investors do open are associated with substantive price reactions. Integration evidence begins with the SEC’s XBRL mandate, which required firms to tag financial-statement items in machine-readable form: [Blankespoor \(2019\)](#) documents that the resulting reduction in investors’ processing costs shifted firms’ own disclosure choices. On the textual margin,

¹⁴This pattern echoes a broader theme in the analyst and disclosure literature: forecast dispersion is itself a priced characteristic ([Diether et al., 2002](#)), and disclosure clarity is linked to investor behavior at the individual-disclosure level ([Lawrence, 2013](#)).

Tetlock (2007) and Loughran and McDonald (2011) establish that machine-readable sentiment and finance-specific dictionaries extract economically meaningful information from narrative disclosure, and Blankespoor et al. (2018) demonstrate that robo-journalism (automated article generation) lowers the processing cost retail investors face when consuming financial information.

A parallel literature approaches these costs from the demand side. Da et al. (2011) establish Google search as a measure of retail attention that predicts short-horizon returns, and Drake et al. (2012) document that earnings-window search proxies information demand more directly than trading-based measures. Attention is itself a binding constraint, with documented inattention along multiple margins: complicated firms (Cohen and Lou, 2012), competing simultaneous announcements (Driskill et al., 2020; Hirshleifer et al., 2009), Friday timing (Della Vigna and Pollet, 2009), strategic scheduling (deHaan et al., 2015), and high-fear states (Da et al., 2015).

Closest to the FDP-aggregation question, a disseminator-side literature ties broader access to firm disclosures to changes in price formation: conference calls open to all investors coincide with different retail-trading and volatility patterns (Bushee et al., 2003) and with reduced post-announcement underreaction (Kimbrough, 2005), and business-press coverage is associated with narrower spreads and deeper markets around earnings releases (Bushee et al., 2010).

A common limit runs through this work: each technological advance lowered the cost of one component within a single information source. None enabled the low-cost cross-source integration of heterogeneous numerical outputs from multiple distinct producers—combining, say, I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks consensus estimates in real time. That friction survived every one of these technologies.

2.3 Generative AI in Capital Markets

GenAI is the first technology positioned to relax this cross-source friction. It arrives against the backdrop of a machine-learning literature in capital markets whose central finding is well established: nonlinear, high-dimensional methods such as neural networks and tree-based models extract incremental predictive signal from financial data that traditional linear asset-pricing regressions miss (Goldstein et al., 2021; Gu et al., 2020). What distinguishes the current generation of LLMs is

the conjunction of three features prior technologies lacked: they process unstructured information from heterogeneous sources, they synthesize it through natural-language queries without specialized retrieval infrastructure or programming expertise, and they operate at low marginal cost per query. Cross-source synthesis that once demanded the time and dedicated infrastructure of an institutional team collapses to the marginal cost of a single natural-language query—placing the capability within reach of retail investors for the first time.

Early evidence indicates that this shift is already price-relevant. [Lopez-Lira and Tang \(2023\)](#) show that ChatGPT-derived sentiment signals predict cross-sectional returns; [Cheng et al. \(2025\)](#) use ChatGPT service outages as exogenous shocks and find that investor responses to information degrade when the tool is unavailable; [Bertomeu et al. \(2025\)](#) exploit Italy’s 2023 ChatGPT ban for country-level identification; and [Chang et al. \(2026\)](#) show that retail traders’ alignment with AI-derived earnings-call sentiment rose sharply after ChatGPT’s release, narrowing the information-processing gap between less-sophisticated retail traders and more-sophisticated short sellers.

Most directly relevant, a separate strand examines the production side—AI’s effects on the analyst research function itself. [Kimbrough et al. \(2026\)](#) find that analysts at brokerages with greater AI integration produce more accurate earnings forecasts and that their revisions are more informative to capital markets. AI-generated content, [Bradshaw et al. \(2026\)](#) document, reached 13.5 percent of Seeking Alpha after ChatGPT’s release, with AI-using authors covering more firms but producing less-informative pieces per article. [Liu and Subasi \(2026\)](#) find an analogous democratization on the crowdsourced side, with GenAI closing the accuracy gap between non-professional and professional forecasters. Together, these papers indicate that AI is reshaping the production of analyst research. How GenAI reshapes the consumption of that research—in particular, how investors aggregate across the heterogeneous consensus numbers that competing FDPs construct—is the gap this paper fills.

2.4 Theoretical Framework and Predictions

Two opposing forces shape the framework, and the balance between them determines whether consensus-level discrimination can emerge. The first comes from the forecast-combination literature: combining heterogeneous, partially independent forecasts almost always dominates relying on any single source ([Bates and Granger, 1969](#); [Clemen, 1989](#); [Granger and Ramanathan, 1984](#); [Stock and Watson, 2004](#); [Timmermann, 2006](#)). The intuition is the wisdom-of-crowds principle popularized by [Surowiecki \(2004\)](#): averaging diversifies away idiosyncratic forecaster errors, so sub-selecting on individual accuracy can lose more from forgone diversification than it gains in precision. Applied to the analyst-consensus setting, this force rationalizes the historical equilibrium in which the market anchors on a single pooled benchmark (typically the I/B/E/S consensus)—itself already an average over many analysts—and treats cross-FDP differences as noise. It is also consistent with the within-FDP evidence in [Ahlers and Lakonishok \(1983\)](#) and [Michaely et al. \(2024\)](#), in which more sophisticated within-source sub-selection fails to deliver consistent gains over the broader pool. Taken alone, this force implies that the historical anchor should persist.

The opposing force is heterogeneity in what the providers produce: the documented cross-sectional and time-series variation in FDP accuracy ([Larocque et al., 2025](#)), together with the methodology-driven information content in [Bochkay et al. \(2022\)](#). When cross-provider dispersion is large, carries a non-trivial accuracy-relevant component, and correlates with exploitable provider-specific characteristics, a sufficiently low-cost synthesis across providers can dominate the pooled benchmark. The model-based earnings-forecast literature makes a parallel point: cross-sectional models deliver more accurate forecasts than any single analyst for specific sub-samples ([Hou et al., 2012](#)). The cost of identifying, acquiring, and integrating accurate sources is the state variable that resolves the tension: when this cost is high, the wisdom-of-crowds force prevails and the market anchors on a pooled single-FDP consensus; when it is low, accuracy-weighted synthesis becomes feasible and the ordering can reverse.

Historically, that cost was high: for cross-FDP synthesis, all three components of the [Blankespoor et al. \(2020\)](#) decomposition bind at once. Even when consensus data are freely

redistributed through retail-facing platforms, an investor seeking the cross-FDP picture must know which platform carries which FDP’s consensus (awareness), visit each individually to retrieve the number (acquisition), and reconcile inconsistent presentations of methodology-driven differences (integration). These steps recur for every firm-quarter an investor follows, and synthesis at scale further demands programming and data-engineering expertise. This combination confined the capability to specialists and anchored the market on a single source.

GenAI relaxes both barriers—the per-firm time cost and the expertise requirement—at once. A single LLM query identifies candidate FDP sources, retrieves their numerical content, and integrates them into a synthesized response that surfaces the implied consensus and the cited distributors, all without specialized infrastructure or programming skill. The effective cost of cross-FDP synthesis collapses to that of issuing one query.¹⁵ Figure 1 summarizes the framework, which yields three empirical predictions:

Prediction 1. After ChatGPT’s release, the consensus reflected in earnings-announcement prices shifts away from heuristic aggregation (single-FDP anchoring or equal-weighting across FDPs) toward a more sophisticated synthesis that draws on multiple FDPs and places greater weight on the more accurate ones.

Prediction 2. The post-ChatGPT shift is stronger when and where GenAI-mediated synthesis is more accessible (LLMs effectively retrieve firm-level analyst information) and more valuable (cross-FDP disagreement and its methodology-driven component are large).

The framework operates entirely on the demand side—a scope condition that frames the final prediction. Falling awareness, acquisition, and integration costs change how investors use the FDPs’ numbers, while the numbers themselves—each provider’s construction rules and the resulting street earnings—are held fixed. The restriction yields a falsifiable corollary that separates the framework from a supply-side alternative in the spirit of the 2009 I/B/E/S methodology change studied by [Bochkay et al. \(2022\)](#): if the post-ChatGPT patterns instead reflected a change in how the street

¹⁵[Appendix B](#) reports two worked examples (IBM and Microsoft) showing the actual LLM responses, the cited distributors, and the upstream FDPs where disclosed.

numbers are constructed, the predictive quality of street earnings for future performance should shift at the ChatGPT boundary, and differentially so in announcements containing large unexpected items, where construction discretion binds most.

Prediction 3. The post-ChatGPT shift extends beyond trailing accuracy: announcement reactions load increasingly on providers whose consensus is more informative about the firm’s future performance.

3 Institutional Background, Data, Sample, and Key Measures

The paper’s tests demand a specific empirical setting: providers that genuinely disagree, a panel that observes every provider’s consensus for the same firm-quarter, outcomes that register the market’s reaction, and a surprise benchmark that embeds provider accuracy. This section supplies each in turn.

3.1 Institutional Background

Cross-FDP disagreement begins with industry structure. Five major data providers—I/B/E/S, FactSet, Bloomberg, S&P Capital IQ, and Zacks—construct and sell analyst-consensus EPS information to sell-side brokerages, buy-side investors, financial media outlets, and academic researchers. I/B/E/S is credited with creating the consensus-earnings-estimate industry in the early 1970s, and Zacks claims to have first distributed quarterly consensus forecasts in the early 1980s. Bloomberg, Capital IQ, and FactSet entered the segment from adjacent product lines (terminal data, investment-banking software, and printed “fact sets” delivered to Wall Street clients, respectively).¹⁶

Two sources of variation drive cross-FDP differences in consensus numbers. First, the underlying analyst set differs: each FDP includes a different subset of contributing sell-side analysts, applies its own recency cutoffs (e.g., Bloomberg drops estimates more than 30 days stale), and runs proprietary statistical checks to filter outliers. Second, the methodology for constructing “street” earnings varies: each FDP decides for itself how to treat unexpected charges, gains, and other

¹⁶Each FDP today operates within a broader financial-data platform: I/B/E/S under LSEG Data & Analytics (formerly Refinitiv); Capital IQ under S&P Global; FactSet and Bloomberg as standalone parents. Bloomberg bundles earnings estimates with its core terminal offering, while the other four sell estimates as standalone or modular products. See [Larocque et al. \(2025\)](#) for a detailed institutional review.

adjustments to reported earnings (Bochkay et al., 2022). Zacks, for example, counts stock-based compensation as a recurring expense in street earnings; I/B/E/S has historically excluded it. Such divergence is common: 46.7 percent of firm-quarters in the 2002–2020 sample of Larocque et al. (2025) contain a large unexpected item that FDPs treat inconsistently. Together, these coverage and methodology differences underlie the cross-FDP disagreement documented in Section 2; Larocque et al. (2025) provide a comprehensive FDP-by-FDP review.

No single FDP, however, strictly dominates the others at the firm-quarter level. I/B/E/S ranks highest on average across composite measures of FDP quality (Larocque et al., 2025) and has long served as the de facto anchor for research and trading alike, yet it is the most accurate provider in only 16.9 percent of firm-quarters in my S&P 1500 sample and the least accurate in 12.1 percent. In the remaining 71 percent of firm-quarters it is neither, and the extreme accuracy ranks are dispersed across the other four FDPs. This churn in accuracy leadership across firms and over time gives the cross-FDP synthesis question its economic content: a static anchoring choice cannot exploit it; an adaptive multi-FDP rule can.

3.2 Data and Sample

Exploiting this firm-quarter variation requires parallel consensus histories from all five FDPs, which no single database supplies. I therefore collect quarterly consensus EPS forecasts and reported actuals from each provider separately. I/B/E/S and Zacks are available through WRDS (the I/B/E/S Summary History file and the Zacks Estimates database, respectively). The remaining three FDPs (FactSet, Bloomberg, and S&P Capital IQ) sit behind proprietary terminals or institutional portals, from which I download quarterly consensus histories firm by firm.¹⁷

The sample begins with all S&P 1500 firm-quarters that announce earnings over 2021–2025: 29,395 firm-quarters across 1,502 firms.¹⁸ I then require consistent coverage across all five FDPs,

¹⁷Larocque et al. (2025) follow a similar collection procedure for their 2002–2020 study, downloading or hand-collecting data from each FDP’s proprietary system. Bloomberg consensus data, in particular, must be retrieved via the Bloomberg Query Language (BQL) function ticker by ticker, and Bloomberg’s historical consensus reaches back only five years, which constrains the sample start date.

¹⁸The November 30, 2022 ChatGPT release falls roughly in the middle of this window, yielding a relatively balanced panel of pre- and post-ChatGPT firm-quarters. The five-year window is also the maximum permitted by Bloomberg’s historical-retrieval constraint described above.

which removes 4,368 firm-quarters lacking the consensus from at least one provider. Dropping firm-quarters with missing market-outcome variables or controls (1,179 observations), missing constructed return or predicted-surprise variables (10), and singleton firms (5) brings the final sample to 23,833 firm-quarters across 1,334 firms. Table 1 details the full selection procedure; Table 2 reports descriptive statistics for the main variables.¹⁹

Firm controls and market-outcome variables are drawn from WRDS-hosted databases: accounting fundamentals from Compustat, returns and trading data from CRSP, institutional ownership from the Thomson Reuters 13F Holdings database, and analyst-coverage characteristics from I/B/E/S. Appendix A defines all variables.

3.3 Dependent Variables

The market’s reliance on any given consensus is not directly observable, so I infer it from the price response at the earnings announcement: the surprise implied by a consensus on which the market relies more heavily should move prices more. The main outcome variables are short-window abnormal returns: a three-day cumulative abnormal return $CAR_{[-1,+1]}$ for the immediate-reaction tests and a four-day buy-and-hold abnormal return $BHAR_{[-2,+1]}$ for the trading-strategy test.²⁰ For each window I use two risk adjustments: a Fama-French-Carhart four-factor (FFC) abnormal return and a market-adjusted return.²¹

3.4 Constructing the Accuracy-Weighted SUE

The paper’s central construct is an accuracy-weighted aggregation of the five FDP consensus signals. Let $g \in \{\text{FactSet, Bloomberg, Zacks, CapIQ, IBES}\}$ index providers and let $SUE_{i,q}^g$ denote provider g ’s price-scaled standardized unexpected earnings for firm i in quarter q . The

¹⁹All continuous variables in the analyses that follow are winsorized at the (1,99) level.

²⁰The four-day BHAR window adds a pre-announcement day to capture any leakage and a post-announcement day to capture overnight drift for announcements made after the market close.

²¹FFC factor loadings are estimated on the [-260,-11] pre-announcement window; the market-adjusted return subtracts the contemporaneous CRSP value-weighted index from each daily return. All returns are expressed in percent. All inferences are robust to both adjustments: Tables 4, 5, and 11 report results under both adjustments side by side (Table 3 Panel B likewise), while Tables 3 (Panel A), 6, 7, 8, 9, 10 (Panel B), and 12 (Panel A) report only the market-adjusted return for brevity (FFC counterparts deliver identical inferences and are available on request).

accuracy-weighted composite is

$$SUE_{i,q}^{Acc} = \sum_g w_{i,q}^g SUE_{i,q}^g, \quad (1)$$

with weights $w_{i,q}^g$ derived from each provider’s recent forecast accuracy. The equal-weighted benchmark sets $w_{i,q}^g = 1/5$, yielding $SUE_{i,q}^{EW5}$; the single-FDP I/B/E/S benchmark requires no aggregation. Entering both composites in the same regression tests whether the market loads more on the accuracy-tilted aggregate than on its naive equal-weighted counterpart.

Weight construction. The weights map each provider’s recent, firm-specific track record into its share of the composite. I first compute the price-scaled absolute forecast error, $err_{i,q}^g = |actual_{i,q}^g - consensus_{i,q}^g|/p_{i,q}$.²² A backward-looking four-quarter rolling mean then yields a recent mean absolute forecast error, $MAFE_{i,q}^g = \frac{1}{4} \sum_{k=1}^4 err_{i,q-k}^g$, whose sign I flip so that larger values denote greater accuracy: $acc_{i,q}^g = -MAFE_{i,q}^g$. Because the rolling window excludes quarter q itself, the weights use only information available before the quarter- q announcement and cannot induce look-ahead bias. I standardize across the five providers within firm-quarter, $acc_z_{i,q}^g = (acc_{i,q}^g - \overline{acc}_{i,q})/sd(acc_{i,q})$, and convert the z-scores into convex weights via the softmax (multinomial-logit) transform,

$$w_{i,q}^g = \frac{\exp(acc_z_{i,q}^g)}{\sum_h \exp(acc_z_{i,q}^h)}. \quad (2)$$

The resulting weights are strictly positive, sum to one within each firm-quarter, preserve the accuracy ordering, and reduce to the uniform 0.2 vector when all five major FDPs are equally accurate. [Appendix C](#) works through this construction for a single hypothetical firm-quarter, tracing how raw accuracy histories pass through the four layers to produce both the weights $w_{i,q}^g$ and the composite $SUE_{i,q}^{Acc}$.

Design choices. Three choices in the construction merit comment. First, the within-firm-quarter standardization (rather than cross-sectional) is essential because earnings noise differs

²²The reported actual EPS is itself FDP-specific: each provider computes its own “street” earnings actual under its own screening and adjustment rules. Pairing each FDP’s consensus with its own actual ensures that the resulting forecast error reflects only ex ante predictive accuracy, not methodology mismatch between the consensus and the realized value. This is the approach taken by [Larocque et al. \(2025\)](#) and is consistent with how each FDP reports its own pre-announcement consensus relative to its own post-announcement actual.

across firms and across periods: what matters economically is relative accuracy for this firm this quarter, not absolute accuracy compared with other firms or other quarters. Second, the softmax transform, equivalent to the exponential weighting scheme in the forecast-combination literature (e.g., [Bates and Granger, 1969](#); [Timmermann, 2006](#)), occupies a middle ground between equal weighting (which ignores accuracy) and winner-take-all selection (which would amplify the noise in the four-quarter MAFE estimate). It also admits a maximum-entropy interpretation: among weight vectors delivering the same average standardized accuracy, the softmax imposes the least additional structure. Third, the backward-looking four-quarter window balances recency against statistical noise: shorter windows overweight transient accuracy shocks; longer windows blur structural changes in provider quality.

Empirical weight distribution. In practice, the weights tilt meaningfully away from equal weighting while stopping well short of winner-take-all. The average within-firm-quarter minimum weight is 0.028 and the corresponding maximum is 0.414, against a uniform benchmark of 0.200: a typical firm-quarter cuts its worst provider to roughly three percent of the composite and concentrates roughly forty-one percent on its best. SUE^{Acc} and SUE^{EW5} are highly correlated ($\rho \approx 0.82$) but not collinear—leaving distinct variation for the pricing tests to exploit.

4 GenAI and Selective Reliance Across FDPs: Initial Evidence

This section presents the first evidence on [Prediction 1](#): if GenAI lets the market discriminate across FDPs, price reactions should rise with provider accuracy, and rise more steeply after ChatGPT. Two tests implement this logic—a within-firm-quarter test of whether announcement returns track provider accuracy ranks, and a firm-level test of whether reliance on the historical anchor, I/B/E/S, tracks its own accuracy.

Accuracy is the natural dimension along which to rank providers. A validation exercise in the Online Appendix (Table [OA1](#)) supports the choice: the accuracy-weighted consensus is strictly closer to realized earnings than either heuristic in both regimes (paired tests, $t > 10$), while providers' consensus slopes for future fundamentals are statistically indistinguishable, suggesting that the accuracy tilt sacrifices little fundamental informativeness. At the firm level, moreover, the accuracy

and informativeness of the consensus—constructed identically at the current-quarter and year-ahead horizons—are themselves highly correlated (Pearson 0.72; Section 7). Ranking on accuracy rather than informativeness is therefore largely a matter of econometric convenience.

Design. The first test asks whether the ERC rises with provider accuracy rank. Ranking the five FDPs each firm-quarter by ex ante trailing MAFE (rank 1 = least, 5 = most accurate) and pooling in long form, I estimate

$$CAR_{i,q,[-1,+1]}^{Mkt} = \alpha + \beta SUE_{i,q}^g + \delta Rank_{i,q}^g + \gamma (SUE_{i,q}^g \times Rank_{i,q}^g) + \varepsilon_{i,q,g}, \quad (3)$$

where γ is the ERC increment per one-rank accuracy gain.²³ Because rank is assigned within firm-quarter, identification comes from differential reactions to the same announcement: firm- and time-level confounders, including methodology drift common to all providers, difference out. The design cannot show, however, whether reliance moves with a provider's own record.

The second test asks exactly that: it fixes the provider and exploits cross-firm variation in its accuracy. Classifying firms by the change in I/B/E/S's mean rank between the pre- and post-ChatGPT windows ($IBESUp_i/IBESDown_i$ for moves of ≥ 2 positions), I estimate the triple-interaction specification

$$\begin{aligned} CAR_{i,q} = & \alpha + \beta_1 SUE_{i,q}^{IBES} + \beta_2 SUE_{i,q}^{IBES} \times Post_q \times IBESUp_i \\ & + \beta_3 SUE_{i,q}^{IBES} \times Post_q \times IBESDown_i + \text{controls} + \text{FE} + \varepsilon_{i,q}, \end{aligned} \quad (4)$$

with the lower-order terms included but not displayed. The triple interactions carry the test: selective reliance predicts $\beta_3 < 0$ and $\beta_2 \approx 0$.

Results. Post-ChatGPT, the price reaction to a given surprise depends increasingly on which provider produced it. Table 3 reports both tests. In Panel A, the per-rank ERC increment γ roughly doubles, from 0.306 before to 0.621 after; the post-pre differential is 0.316 (both $p < 0.01$).²⁴ Panel B shows that the adjustment is accuracy-conditioned: the response to I/B/E/S

²³The earnings response coefficient (ERC) is the slope of announcement returns on the earnings surprise—the abnormal return the market delivers per unit of surprise. A larger ERC means the market treats the surprise as more informative; γ measures how much that slope rises when the surprise comes from a one-rank-more-accurate provider.

²⁴Provider-by-rank cell-level versions of Panel A, estimated separately in absolute-value and signed form, yield the same inferences (untabulated).

weakens significantly where it *lost* rank—the post-ChatGPT triple interaction for I/B/E/S-down firms ranges from -2.3 to -2.8 across specifications ($p < 0.05$)—but not where it gained or held.²⁵ The reweighting is thus asymmetric, unwinding reliance where the historical anchor’s accuracy advantage has eroded. The next section asks whether the market goes further and prices an accuracy-weighted synthesis across providers.

5 GenAI and the Accuracy-Weighted Synthesis Across FDPs

This section tests the paper’s central claim: after ChatGPT, the market prices the accuracy-weighted synthesis of all five FDPs, as [Prediction 1](#) asserts, beyond what the heuristic alternatives capture. Section 4 establishes the precondition—the post-ChatGPT market discriminates across FDPs by accuracy—but discrimination is a weaker claim than synthesis, because a market can favor a single accurate provider without ever combining across providers. The analysis proceeds in three steps. I first validate the construct itself: accuracy ranks persist, the composite is the most accurate of the three, and its ex ante difference from the heuristics earns an ex post trading return. A horse race then pits the accuracy-weighted SUE against the two heuristic benchmarks; the resulting accuracy premium is statistically and economically significant after ChatGPT. Finally, testing [Prediction 2](#), I trace the premium across the four GPT-capability regimes and locate its cross-sectional concentration where the framework predicts: in firm-quarters with greater cross-FDP disagreement and greater methodological divergence across providers.

5.1 Validating the Accuracy-Weighted Construct

Premise and benefits. The design rests on one premise: weighting by *trailing* accuracy is informative only if accuracy *persists*. Given persistence, two benefits should follow: SUE^{Acc} should be more accurate than the heuristics, and it should carry trading value—its difference from a heuristic, known ex ante, earning an ex post return. For the latter I form the *predicted surprise*

$$Pred.Surp_{i,q}^X = (Consensus_{i,q}^{Acc} - Consensus_{i,q}^X) / p_{i,q}, \quad X \in \{EW5, IBES\}, \quad (5)$$

and test whether it forecasts the four-day announcement BHAR and a quintile long–short return.

²⁵A generalized post-ChatGPT shift in risk premia or attention would move all providers’ ERCs together; neither the within-announcement rank gradient nor the cross-firm asymmetry fits that account.

Results. All three conditions hold. Provider ranks are highly persistent (Spearman $\rho = 0.71$): in Table 4, Panel A, the least- and most-accurate providers retain their ranks 76 and 65 percent of the time, roughly three times the uniform baseline. In Figure 2, SUE^{Acc} has the smallest mean absolute error of the three composites in both regimes. Trading value materializes in Panel B: the standardized predicted surprise forecasts the four-day BHAR in all four columns ($p < 0.01$), and a long–short portfolio on its quintiles earns 0.63–0.73 percent per announcement against equal-weighting and 1.35 percent against I/B/E/S ($p < 0.01$).

5.2 The Accuracy Premium: Horse-Race Tests

Design. With the construct validated, the horse race asks whether the market prices it. I race the accuracy-weighted composite SUE^{Acc} against two natural benchmarks: the equal-weighted composite SUE^{EW5} , which treats each FDP as equally informative, and the I/B/E/S-only consensus SUE^{IBES} , the historical anchor. The pooled specification is

$$\begin{aligned} CAR_{i,q} = & \alpha + \beta_1 SUE_{i,q}^{Acc} + \beta_2 SUE_{i,q}^X + \beta_3 SUE_{i,q}^{Acc} \times Post_q \\ & + \beta_4 SUE_{i,q}^X \times Post_q + \text{controls} + \text{FE} + \varepsilon_{i,q}, \end{aligned} \quad (6)$$

for $X \in \{EW5, IBES\}$. The pre-period accuracy premium is the contrast $\beta_1 - \beta_2$; its post-pre change, $\beta_3 - \beta_4$, is the central test of Prediction 1. A positive estimate of $\beta_3 - \beta_4$ indicates that after ChatGPT the market’s loading shifts toward the accuracy-weighted consensus relative to the heuristic benchmark. I report both the three-day market-adjusted abnormal return and the Fama-French-Carhart four-factor abnormal return. The same contrast is the estimand for every test in this section: the regime and cross-sectional analyses re-estimate it within subperiods and subsamples.²⁶

Results. After ChatGPT, the market tilts toward the accuracy-weighted composite at the expense of both heuristics (Table 5). In Panel A (equal-weighted benchmark), the pre-ChatGPT accuracy premium is statistically indistinguishable from zero (-0.972), consistent with the equal-weighted

²⁶Although SUE^{Acc} and the benchmark SUE^X are weighted averages of the same five per-FDP surprises and are therefore correlated by construction ($\rho \approx 0.82$ for $X = EW5$), they are not collinear: with $R_j^2 \approx 0.67$, the variance inflation factor $VIF_j = 1/(1-R_j^2) \approx 3.0$, so each individual slope’s standard error is only $\sqrt{3.0} \approx 1.7$ times its single-regressor counterpart. The identifying quantity, however, is the contrast $\beta_1 - \beta_2$, whose variance $\text{Var}(\hat{\beta}_1) + \text{Var}(\hat{\beta}_2) - 2\text{Cov}(\hat{\beta}_1, \hat{\beta}_2)$ is materially less inflated, because positively correlated regressors yield positively correlated slope estimates. The horse race is therefore the conservative specification: a reparameterization using the accuracy tilt $SUE^{Acc} - SUE^X$ directly would recover the same contrast with cleaner individual slopes.

composite already capturing the pre-period information set. The premium then widens sharply: the accuracy-weighted slope rises from 1.376 to 3.420 while the equal-weighted slope falls from 2.324 to 1.767; the formal test is the contrast, whose post-pre change is 2.401 ($p < 0.01$) under the market-adjusted CAR and 2.584 ($p < 0.01$) under the Fama-French-Carhart four-factor abnormal return. Panel B (I/B/E/S benchmark) adds the complementary result: the pre-ChatGPT accuracy premium is negative (-0.623) but statistically insignificant—a sign directionally consistent with the over-reliance on I/B/E/S documented in Section 4. Post-ChatGPT, the premium climbs to 2.049 ($p < 0.01$), a post-pre change of 2.672 ($p < 0.01$) under the market-adjusted CAR and 2.688 ($p < 0.01$) under the Fama-French-Carhart adjustment.

The magnitudes are economically large. After ChatGPT, a one-percentage-point standardized surprise generates a three-day announcement CAR roughly 2.4 percentage points larger under the accuracy-weighted consensus than under the equal-weighted benchmark, and 2.7 percentage points larger than under I/B/E/S—contrasts comparable to the entire pre-ChatGPT ERC on the I/B/E/S consensus itself, 2.2 in Panel B. The post-ChatGPT market loads on the accuracy-weighted synthesis with a weight that rivals its historical reliance on the I/B/E/S consensus. The shift tracks real informational content: Figure 2 shows that the accuracy-weighted consensus also delivers the smallest mean absolute forecast error of the three composites, and the difference is statistically significant.²⁷ Yet the pooled contrast dates the shift only to the ChatGPT boundary; attribution to GenAI demands a finer clock.

5.3 The Accuracy Premium Across GPT-Capability Regimes

Motivation. Prediction 2 supplies that clock: if the premium reflects GenAI-mediated synthesis, it should strengthen as the tools’ retrieval capabilities expand—a dose–response pattern rather than a single jump at the ChatGPT release. I therefore re-estimate the horse race of Equation (6) within four GPT-capability regimes: pre-ChatGPT (R0), text-only (R1, from November 30, 2022),

²⁷As emphasized in the introduction, the accuracy-weighted construction serves as a tractable proxy for the broader hypothesis that the post-ChatGPT market weights providers by accuracy; it is not a claim that the market literally applies softmax weights. Formally, the accuracy tilt $SUE_{i,q}^{Acc} - SUE_{i,q}^{EWS} = \sum_g (w_{i,q}^g - 1/5) SUE_{i,q}^g$ is the inner product of each provider’s deviation from the uniform weight and its standardized surprise; this is the object the horse-race contrast identifies.

browsing (R2, from September 27, 2023), and native search (R3, from October 31, 2024).

Results. Table 6 and Figure 3 sharpen the timing. Against I/B/E/S, the accuracy premium’s point estimates are negative in the pre-ChatGPT (-0.75) and text-only (-1.23) regimes and turn positive under browsing (1.52); only under native search is the estimate statistically significant (4.91, $p < 0.01$). Against the equal-weighted benchmark, the path is -0.95, 0.18, 2.43 ($p < 0.05$), and 2.55 ($p < 0.05$), so significance arrives at the browsing step and persists thereafter. The point estimates trace the staircase that Prediction 2 anticipates, and the text-only regime supplies the telling null: a chatbot that cannot yet retrieve forecasts leaves the premium statistically undetectable against either benchmark. Detection accompanies retrieval.

5.4 Cross-Sectional Tests of the Accuracy Premium

Design. Prediction 2 specifies where as well as when: the premium should concentrate in firm-quarters where GenAI-mediated synthesis is more valuable. Because the regime evidence localizes the accuracy premium to the browsing and native-search regimes, the cross-sectional tests replace the post-ChatGPT $Post_q$ with the post-scraping indicator $Post_q^{Scr}$ —equal to one for announcements on or after the onset of GPT browsing (September 27, 2023), when web retrieval first made cross-provider synthesis feasible. I re-estimate the horse-race specification in Equation (6) within the low and high subsamples of two partitions, each probing a distinct channel. (i) *Concurrent cross-FDP SUE dispersion* is the interquartile range of the five FDPs’ SUE for the same firm-quarter—the setting in which provider weighting has the most discriminating work to do. (ii) *Special-items magnitude* is the absolute size of special items relative to the absolute size of income before extraordinary items; because FDPs differ in their treatment of unexpected charges and gains, cross-FDP disagreement over what belongs in street earnings should be largest where such items are economically material.²⁸

Results. Table 7 reports the partition-by-partition estimates; the browsing-era premium change is significant only in the predicted half of each partition. In Panel A (cross-FDP dispersion),

²⁸A third moderator the framework predicts—whether GenAI actually retrieves the firm’s forecasts—requires its own construct and is developed in Section 6. The plain post-ChatGPT-boundary variant of this table, which adds a median split on that LLM-visibility score, yields similar inferences, with the premium concentrating in the high half of each partition (untabulated).

the post-pre change in the accuracy premium is 2.211 ($p < 0.05$) against the equal-weighted benchmark and 2.427 ($p < 0.01$) against I/B/E/S in the high-disagreement subsample, versus insignificant estimates (2.037 and -0.691) in the low-disagreement subsample. A caveat attends the equal-weighted comparison: the low-disagreement point estimate is close in magnitude to its high-disagreement counterpart and differs mainly in precision; the partition separates cleanly only against I/B/E/S, where the low-disagreement estimate turns negative. Panel B (special items) requires no such qualification: the premium change is 2.518 ($p < 0.05$) and 3.638 ($p < 0.01$) under high special items, versus point estimates of the opposite sign (-0.552 and -0.219) under low. Two readings of the mechanism emerge. The concentration in high-dispersion firm-quarters supports the discrimination claim: the premium is statistically detectable only where the five FDPs disagree most about whether the firm beat or missed. The concentration in high special-items firm-quarters supports the integration claim: the premium operates where the methodological choices distinguishing the FDPs most affect the reported surprise.

Together, Tables 4–7 document an accuracy premium that is statistically significant, economically meaningful, tradeable in real time, and concentrated where the framework predicts. The shift from heuristic to accuracy-weighted aggregation is the central empirical signature of [Prediction 1](#). Prices, however, record only the outcome of the market’s synthesis. Section 6 turns to the inputs: did this shift in market behavior arise from a change in how investors acquire and integrate forecast information?

6 LLM Visibility and Effective Synthesis Across FDPs

The preceding section establishes that the post-ChatGPT market prices the accuracy-weighted synthesis; the accessibility clause of [Prediction 2](#) asks what makes that synthesis possible. Integration requires acquisition: the model cannot synthesize forecasts it cannot retrieve. This section develops the pillar named in the paper’s title: *LLM visibility*, a direct measure of that acquisition—an EPS-extraction indicator scaled by the logged count of unique domains a browsing-enabled model cites when asked for the firm’s forward consensus EPS, averaged across

ChatGPT and Perplexity responses.²⁹ It captures what GenAI is *doing*, not what it *can* do. Because the measure is informative only if it amounts to more than repackaged prominence, I show it is distinct from the classical proxies before testing whether acquisition gates the accuracy premium.

A distinct construct. Table 9 shows how little LLM visibility shares with the classical prominence proxies: its high-visibility set overlaps those of firm size, analyst coverage, and institutional ownership only about 47–50 percent of the time—near the 50 percent coin-flip benchmark for independent median splits—whereas the classical proxies overlap one another up to 80 percent. A pooled horse race enters each measure’s premium-concentration triple interaction at once: against I/B/E/S, LLM visibility is the only measure that remains significant, at 5.84 ($p < 0.01$); against the equal-weighted benchmark, its coefficient of 3.17 attains only marginal significance ($p < 0.10$).

Acquisition as a precondition. Distinctness permits the sharper test: if acquisition gates synthesis, the premium should appear in high-visibility firms, and only once the model can browse. I re-estimate the horse race of Equation (6) within each regime, separately for above- and below-median visibility firms, and track the high-minus-low gap in the accuracy premium,

$$\Delta Premium_r = Premium_r^{\text{High vis}} - Premium_r^{\text{Low vis}}, \quad r \in \{R0, R1, R2, R3\}. \quad (7)$$

In Table 8, $\Delta Premium_r$ is statistically insignificant at the five-percent level in the pre-ChatGPT and text-only regimes (R0, R1). Once the model can browse, the gap emerges: 2.7 to 3.0 in the browsing regime (R2, $p < 0.05$), then 3.1 to 4.2 under native search (R3)—significant against I/B/E/S ($p < 0.01$) but only marginally so against equal-weighting ($p < 0.10$). The gap opens only where forecasts are visible to the model *and* only once it can retrieve them—a pattern consistent with acquisition as a precondition for synthesis rather than an incidental correlate of it.

Prominence as an alternative. The regime timing also weighs against the residual concern that LLM visibility merely proxies for firm prominence. A firm’s visibility is essentially time-invariant

²⁹When the browsing-enabled model answers the query, it lists the web pages it consulted (e.g., Yahoo Finance, MarketBeat, TipRanks). Counting the unique domains among these citations measures how many distinct forecast sources the model actually retrieves for a firm—the breadth of its information acquisition—and says nothing about the content of any one estimate. Appendix B reports worked examples.

over this window, so a prominence account would predict a premium gap in the pre-ChatGPT and text-only regimes as well. The gap instead opens only when retrieval arrives, which is difficult to reconcile with a static firm characteristic. The visibility evidence thus supplies what the accessibility clause demands—an accuracy premium that travels with GenAI’s actual retrieval of forecasts—and narrows what remains to the GenAI channel itself, the subject of the tests that follow.

7 Additional Tests

Sections 4 and 5 establish that the post-ChatGPT market prices an accuracy-weighted synthesis of consensus EPS forecasts. This section subjects that conclusion to three further tests, each aimed at a gap the earlier evidence leaves open. The first asks whether the market’s selectivity is specific to trailing accuracy or extends to a second quality dimension—a forecast-based informativeness measure built from the consensus’s own predictive record. The last two probe the GenAI channel itself, first by tracing how the multi-FDP signal reaches investors, then by observing how the accuracy premium responds when major ChatGPT outages briefly withdraw GenAI.

7.1 A Forecast-Based Informativeness Measure

Motivation. Trailing accuracy is one notion of consensus quality; a complementary notion asks how *informative* the consensus is about the firm’s future performance. The distinction disciplines the interpretation of Section 5: if the market’s post-ChatGPT selectivity reflects a general taste for consensus quality, it should price this second dimension as well; if it reflects a peculiarity of the trailing-accuracy construction, it should not.

Design. For each firm, I score the consensus’s predictive record in real time. The building block is the own-basis anticipation error: the absolute per-share distance between the quarter- q I/B/E/S consensus and the seasonally matched quarter- $(q+4)$ street actual, with a stock-split guard—“own basis” because the consensus is paired with the same provider’s street actual. I average this error over up to the twenty most recent source quarters whose year-ahead actuals are already public at the announcement (at least twelve), scale it by the window mean $|actual|$ so it is comparable across firms, and winsorize. The measure is then $Predictive\ Value_{i,q}^{IBES} = 1 - error_{i,q}$, a continuous cross-firm variable (mean 0.56, standard deviation 0.30): higher values mean the

firm’s consensus has anticipated its year-ahead street numbers more closely. It covers 22,364 of the 23,833 firm-quarters (8,730 pre-ChatGPT); the remainder are recent listings with fewer than twelve scoreable history quarters. If post-ChatGPT investors read earnings through tools that weight the consensus by its demonstrated predictive value, the announcement-window response to the surprise should steepen in *Predictive Value*^{IBES} after ChatGPT.

Results. The market prices the informativeness dimension, and does so more steeply after ChatGPT. In Table 10, the ERC gradient on predictive value is already strongly positive before ChatGPT (2.17, $t = 4.0$)—firms with a predictable consensus carry higher ERCs, a familiar association—so the identified claim is the post-ChatGPT increment: +2.05 ($t = 2.68$) market-adjusted and +2.10 ($t = 2.82$) under the Fama-French-Carhart adjustment. The increment matches the pre-period gradient in magnitude, so the slope roughly doubles after ChatGPT; dispersion, volatility, and the unexpected-item flag are among the controls, and the *Predictive Value*^{IBES} level terms are indistinguishable from zero. Columns (3)–(4) trace the increment across the three post-ChatGPT capability regimes: 0.92 ($t = 0.9$) under text-only, 1.43 ($t = 1.5$) under browsing, and 3.54 ($t = 3.1$) under native search. The point estimates rise in the same order as the accuracy premium of Table 6, though the native-search step alone is statistically significant—fitting for a forecast track record that only browsing and native search made retrievable.

Two alternative readings of the increment fail to absorb it. Announcement composition does not account for it: jointly adding interactions of the large-unexpected-item flag of [Bochkay et al. \(2022\)](#) leaves the increment unchanged (2.04, $t = 2.67$), and the two measures are essentially uncorrelated (-0.02). Nor does selection toward I/B/E/S specifically: in an untabulated decomposition of the surprise into the component common to all providers and the I/B/E/S-specific tilt, both the gradient (2.3, $t = 3.7$) and its post-ChatGPT increment (1.9, $t = 2.0$) price the common component, though the increment clears the conventional threshold only narrowly. The post-ChatGPT market rewards the consensus signal of well-predicted firms, wherever it is sourced.

The provider level tells the same story, with one qualification. A within-firm-quarter rank version of the score across four providers (each provider’s own-basis error against its own year-ahead

actual; Bloomberg excluded because its native history begins in 2020Q2) shows no pre-period rank gradient, a post-period gradient of 0.320 ($t = 4.14$), and a shift of 0.242 ($t = 2.56$). But the shift's significance concentrates in Capital IQ provider-events, and the informativeness and accuracy rankings are statistically indistinguishable once corrected for their low rank reliabilities, so accuracy weighting already embeds much of the provider-level informativeness signal.

Beyond the pricing test, the measure rationalizes the paper's operating choice. Constructed identically at the current-quarter horizon, the same procedure yields the consensus's trailing *accuracy*, and across firms the two are highly correlated (Pearson 0.72, Spearman 0.78; 0.82 correcting for split-half measurement error; stable pre- and post-ChatGPT and within announcement quarters). Informativeness and accuracy are thus two horizons of one underlying quality—how predictable the firm's street numbers are by its consensus. That the paper's analyses operate on the accuracy dimension is therefore a matter of econometric convenience rather than substance: the current-quarter error is realized within a quarter rather than after a four-quarter maturation, so accuracy is observable in real time at every announcement while carrying most of the informativeness content.

Two boundaries temper the reading. First, the own-basis pairing rewards internal consistency of street-earnings rules alongside genuine foresight about the firm, and the level gradient partly reflects earnings predictability—which is why the identified claim is confined to the post-ChatGPT increment. Second, the evidence concerns the consensus signal of well-predicted firms rather than provider selection: the untabulated I/B/E/S tilt interactions are at best marginal, and the shift in the rank version is, as noted, concentrated in Capital IQ provider-events (excluding Capital IQ, 0.099, $t = 0.75$; excluding any other provider, 0.26–0.31) rather than uniform across the four.

7.2 Additional Tests on Information Acquisition

Motivation. The remaining two tests turn from what the market prices to how the signal reaches it. The accuracy-weighted synthesis documented in Section 5 requires that investors actually access consensus information from multiple FDPs, and [Appendix B](#) shows where GenAI points them: ChatGPT and Perplexity responses to direct EPS queries overwhelmingly cite free distributor

platforms—Yahoo Finance, MarketWatch, Seeking Alpha, and similar aggregators that publish forecasts from several FDPs without paywalls. Whether GenAI substitutes for direct investor access to these distributors or instead acts as a gateway that broadens it cannot be read off prices.

The two interpretations make opposite predictions for direct Google search of free distributors. Under substitution, investors query ChatGPT instead of typing “Yahoo Finance AAPL earnings” into Google, so search for free distributors should fall post-ChatGPT, narrowing the cross-FDP access channel. The gateway interpretation runs the other way: GenAI surfaces these distributors to investors who would not otherwise visit them, so post-ChatGPT search for free distributors should rise even as ChatGPT itself becomes available. The sign of the post-ChatGPT shift in platform search therefore separates the two interpretations.

Design. I obtain daily Google search volume for 17 platform brands, classified into Free Distributors (twelve aggregator sites including Yahoo Finance, MarketWatch, and Seeking Alpha), Paywalled FDPs (Bloomberg, Capital IQ, FactSet, I/B/E/S), and the Free FDP (Zacks); the full classification is in [Appendix A](#). The dependent variable is each brand’s share of total platform-related search on the day.³⁰ I estimate

$$S_{p,t} = \alpha + \sum_g \delta_g Post_t \times Group_{g,p} + \lambda Post_t + \gamma X_t + \text{calendar-month FE} + \varepsilon_{p,t}, \quad (8)$$

where $Group_{g,p}$ indexes the three platform groups (Free Distributor in the pre-ChatGPT regime is the omitted baseline) and X_t contains controls for total platform-related search, earnings-announcement intensity, and trading-day status. Under the gateway interpretation, the post-ChatGPT shift should be positive for Free Distributors and negative for the two FDP groups.

Results. Post-ChatGPT, platform search reallocates from FDP brands to the free aggregators GenAI cites. Table 11, Panel B, reports the platform-day estimates: column (1) uses a single post-ChatGPT indicator, and column (2) splits the post-ChatGPT period into three GenAI regimes.

³⁰Google Trends reports a relative search index normalized within a query window rather than an absolute volume (see, e.g., [Da et al., 2011](#); [Drake et al., 2012](#)), so I use the DataForSEO Search Volume API, which derives absolute search counts from Google’s keyword-level impression and click data. Scaling by total platform-related search captures general demand for financial-information platforms, the relevant baseline for the substitution-versus-gateway comparison; scaling instead by platform-plus-earnings combinations (e.g., “Yahoo Finance EPS forecast”) yields volumes below DataForSEO’s minimum detectability threshold for most platforms.

In column (1), the post-ChatGPT share of total platform-related search rises by 2.1 percentage points for Free Distributors and falls by 7.5 and 7.1 percentage points for the Free FDP and Paywalled FDPs respectively, all significant at the 1% level. Because shares sum to one across platforms each day, these estimates describe reallocation within the platform universe. Column (2) shows the pattern is stable across the three regimes, with the share gain for Free Distributors holding at roughly 2.0 to 2.2 percentage points throughout.

The pattern fits the gateway interpretation. Post-ChatGPT, investor demand for FDP-branded search falls and demand for free aggregators rises, even though ChatGPT itself can answer the underlying queries directly; investors appear to learn about free distributors through ChatGPT and return to them through direct search, a sequence difficult to reconcile with the alternative in which GenAI merely intermediates existing investor demand. The result broadens the empirical foundation of the accuracy premium: the multi-FDP consensus information the premium requires is in fact flowing to investors. The evidence is aggregate and associational, and platform demand may shift for reasons that happen to coincide with ChatGPT’s arrival; the sign nonetheless does the discriminating work—substitution predicts falling distributor search, and distributor search rose. The final test therefore turns to variation in GenAI availability itself.

7.3 Major ChatGPT Outages as an Exogenous Shock

Motivation. A channel that can be interrupted can be identified. To distinguish the GenAI channel from alternative explanations for the post-ChatGPT accuracy premium, I exploit major ChatGPT outages as exogenous, short-duration shocks to GenAI availability, following [Cheng et al. \(2025\)](#) and [Liu and Subasi \(2026\)](#). On announcement dates immediately following a major outage, investors temporarily lose the integration capability that GenAI provides and should revert to the only heuristic that requires no multi-FDP synthesis: the single-FDP I/B/E/S consensus. The prediction is an asymmetric outage effect across the two horse-race benchmarks: against I/B/E/S, the accuracy premium should weaken on outage days as investors substitute toward that fallback; against the equal-weighted benchmark, it should be unchanged, because the accuracy-weighted and equal-weighted aggregates both require multi-FDP synthesis and become equally infeasible when

GenAI is down.

Design. I collect ChatGPT outage incidents from OpenAI’s public status-page history and define a major outage as any incident lasting at least 240 minutes; this yields four incidents across five unique calendar days through 2024. The treatment indicator $Outage_{i,q}$ equals one for firm-quarter announcement dates falling on the calendar day immediately following a major outage and zero otherwise. I re-estimate the horse-race specification of Equation (6) on the post-ChatGPT subsample, interacting both surprise measures with $Outage_{i,q}$, separately for the equal-weighted and I/B/E/S benchmarks. The small number of treated days limits power on the level of the outage effect; identification rests instead on the asymmetric sign predicted across the two benchmarks.

Results. The estimates in Table 11, Panel A, match the predicted asymmetry. Against the equal-weighted benchmark (column 1), the outage-day change in the accuracy premium is statistically insignificant (3.787, $t = 1.37$): the accuracy-weighted-versus-equal-weighted contrast shows no detectable movement, though with so few treated days the null is imprecise. Against the I/B/E/S benchmark (column 2), the outage-day change is -3.173 ($p < 0.05$), combining a marginally significant decline in the accuracy-weighted slope ($\beta_3 = -1.685$, $p < 0.10$) with a rise in the I/B/E/S slope ($\beta_4 = 1.488$, $p < 0.05$). Investors reweight toward I/B/E/S in the immediate wake of an outage.

The asymmetry across the two benchmarks is the identifying signature of the GenAI channel. Both the accuracy-weighted and equal-weighted aggregates require investors to draw consensus information from all five FDPs; when GenAI access is severed, neither weighting is operative, and their relative loading has no reason to move. The I/B/E/S consensus, by contrast, is the single-FDP fallback that requires no multi-source integration, so investors revert specifically to it: the accuracy premium collapses against the I/B/E/S benchmark while the equal-weighted contrast shows no significant change. This sign pattern is difficult to generate without GenAI-mediated multi-FDP integration. Together with the cross-sectional concentration in Section 5 and the gateway evidence above, the outage test points to GenAI as the channel through which the post-ChatGPT accuracy-weighted synthesis reaches market prices.

8 Conclusion

This paper shows that after the release of ChatGPT, market reactions to earnings announcements increasingly reflect an accuracy-weighted synthesis across the five major forecast data providers, departing from the equal-weighted or single-FDP heuristics that dominated prior decades. Four pieces of evidence support this conclusion. First, the post-ChatGPT market discriminates across FDPs by accuracy: the within-firm-quarter ERC steepens with accuracy rank between regimes, and the firm-level response to the I/B/E/S consensus surprise weakens where I/B/E/S has lost relative accuracy but not where it has gained. Second, the accuracy-weighted multi-provider consensus commands a statistically and economically significant accuracy premium over both the equal-weighted alternative and the I/B/E/S-only consensus. The premium concentrates where GenAI synthesis has the most discriminating work to do (high cross-FDP disagreement) and where the integration burden is heaviest (large special items); it translates into tradeable quintile long–short returns on the predicted surprise. The premium also tracks the enabling technology: it climbs across GPT-capability regimes to 4.91 under native search, and it emerges in firms whose forecasts the model demonstrably retrieves (LLM visibility)—but only once retrieval is possible.

The remaining two pieces tie this shift to GenAI itself. Third, the market’s selectivity extends beyond provider accuracy to forecast-based informativeness: when each firm is scored in real time by how closely the consensus has anticipated its year-ahead street numbers over the preceding twenty quarters, the post-ChatGPT rise in announcement reactions concentrates in firms with an informative consensus. The increment is significant under both return adjustments and orthogonal to announcement composition, and it loads on the common cross-provider consensus rather than an I/B/E/S-specific tilt. Across GPT-capability regimes, its point estimates rise stepwise but reach conventional significance only under native search. Fourth, major ChatGPT outages produce the asymmetry a GenAI-mediated multi-FDP channel predicts: on outage days the accuracy premium against I/B/E/S collapses, while the equal-weighted contrast shows no significant change. A corroborative test of post-ChatGPT search behavior suggests that GenAI broadens, rather than substitutes for, direct investor access to free distributors of cross-FDP consensus.

These findings make three contributions. First, they reposition the choice of forecast data provider from a researcher’s measurement decision to an endogenous market outcome that responds to the information environment, extending recent work on cross-FDP differences in consensus information ([Larocque et al., 2025](#)). The historical reliance on I/B/E/S reflects an equilibrium that held only while cross-provider integration was costly; the implication for measurement practice is that the benchmark against which earnings surprises should be gauged shifts with the technology investors use to read the consensus. Second, they identify a specific channel through which GenAI reaches capital markets: low-cost multi-source synthesis of analyst forecasts, which the market now prices into earnings-announcement reactions. That specificity isolates the analyst-information channel from the broader set of GenAI effects on investor behavior and asset prices documented in recent work (e.g., [Bertomeu et al., 2025](#); [Cheng et al., 2025](#); [Liu and Subasi, 2026](#)). Third, they sharpen the disclosure-processing-cost framework of [Blankespoor et al. \(2020\)](#) by pinpointing cross-source integration—distinct from awareness or acquisition within a single source (e.g., [Blankespoor et al., 2014](#); [Blankespoor, 2019](#); [Drake et al., 2015](#))—as the cost component that prior single-source technologies left intact and that GenAI uniquely relaxes.

Four caveats bound these claims. First, the accuracy-weighted construction is a tractable proxy: the evidence establishes that market reactions align more closely with an accuracy-sensitive synthesis than with the heuristic benchmarks, not that the market applies these particular weights. Second, as a point-in-time snapshot of inherently variable LLM outputs, the LLM visibility measure has limited reliability (e.g., [Chen et al., 2023](#); [Lopez-Lira et al., 2025](#)). Third, the forecast-based informativeness measure of Section 7 rewards internal consistency of street-earnings rules alongside genuine foresight, and its level gradient partly reflects earnings predictability; the identified claim is therefore confined to the post-ChatGPT increment, which prices the common consensus and does not speak to provider selection. The provider-level rank version of the measure is noisy (split-half reliability below 0.35), and its shift concentrates in Capital IQ provider-events rather than spreading uniformly across the four providers. Fourth, the outage test rests on few treated days, none spanning a full calendar day; identification accordingly comes from the asymmetric sign across benchmarks

rather than from level estimates, and the documented magnitude is likely a lower bound.

Several questions remain open. First, only a longer post-ChatGPT panel can settle whether the accuracy premium widens or contracts as GenAI tools evolve and as institutional investors embed them in their workflows. Second, whether the cross-FDP synthesis advantage extends to other forecasted quantities, such as revenue and cash flow, or to industry-specific metrics, is unexamined. Third, the analysis treats FDPs as fixed information producers, but the market's new discrimination across them gives providers an incentive to revisit their own methodologies; this supply-side response bears on both the disclosure-processing-cost framework and the structure of the FDP industry.

The broader lesson is that the market's earnings benchmark is technology-bound. For decades, a single provider's consensus stood in for the market's expectation because integrating across providers was costly; GenAI has lifted that constraint, and the benchmark has moved. The market now weights providers by the accuracy they have earned, not the convenience of reaching them—what investors can cheaply integrate shapes what prices reflect.

References

- Ahlers, D. and J. Lakonishok. 1983. A study of economists' consensus forecasts. *Management Science* 29 (10): 1113–1125. doi:[10.1287/mnsc.29.10.1113](https://doi.org/10.1287/mnsc.29.10.1113).
- Bartov, E., D. Givoly, and C. Hayn. 2002. The rewards to meeting or beating earnings expectations. *Journal of Accounting and Economics* 33 (2): 173–204. [https://doi.org/10.1016/S0165-4101\(02\)00045-9](https://doi.org/10.1016/S0165-4101(02)00045-9).
- Bates, J. M. and C. W. J. Granger. 1969. The combination of forecasts. *Operational Research Quarterly* 20 (4): 451–468. doi:[10.1057/jors.1969.103](https://doi.org/10.1057/jors.1969.103).
- Battalio, R. H. and R. R. Mendenhall. 2005. Earnings expectations, investor trade size, and anomalous returns around earnings announcements. *Journal of Financial Economics* 77 (2): 289–319. <https://doi.org/10.1016/j.jfineco.2004.08.002>.
- Bernard, V. L. and J. K. Thomas. 1989. Post-earnings-announcement drift: Delayed price response or risk premium? *Journal of Accounting Research* 27: 1–36. <https://doi.org/10.2307/2491062>. Supplement.
- Bertomeu, J., Y. Lin, Y. Liu, and Z. Ni. 2025. The impact of generative AI on information processing: Evidence from the ban of ChatGPT in Italy. *Journal of Accounting and Economics* 80 (1): 101782. <https://doi.org/10.1016/j.jacceco.2025.101782>.
- Beyer, A., D. A. Cohen, T. Z. Lys, and B. R. Walther. 2010. The financial reporting environment: Review of the recent literature. *Journal of Accounting and Economics* 50 (2–3): 296–343. <https://doi.org/10.1016/j.jacceco.2010.10.003>.
- Blankespoor, E. 2019. The impact of information processing costs on firm disclosure choice: Evidence from the XBRL mandate. *Journal of Accounting Research* 57 (1): 75–120. doi:[10.1111/1475-679X.12237](https://doi.org/10.1111/1475-679X.12237).
- Blankespoor, E., G. S. Miller, and H. D. White. 2014. The role of dissemination in market liquidity: Evidence from firms' use of Twitter. *The Accounting Review* 89 (1): 79–112. <https://doi.org/10.2308/accr-50576>.
- Blankespoor, E., E. deHaan, and C. Zhu. 2018. Capital market effects of media synthesis and dissemination: Evidence from robo-journalism. *Review of Accounting Studies* 23 (1): 1–36. <https://doi.org/10.1007/s11142-017-9422-2>.
- Blankespoor, E., E. de Haan, and I. Marinovic. 2020. Disclosure processing costs, investors' information choice, and equity market outcomes: A review. *Journal of Accounting and Economics* 70 (2–3): 101344. <https://doi.org/10.1016/j.jacceco.2020.101344>.
- Blankespoor, E., J. Croom, and S. M. Grant. 2026. Generative AI and investor processing of financial information. SSRN Working Paper No. 5053905; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5053905. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5053905.
- Bochkay, K., S. Markov, M. Subasi, and E. H. Weisbrod. 2022. The roles of data providers and analysts in the production, dissemination, and pricing of street earnings. *Journal of Accounting Research* 60 (5): 1695–1740. <https://doi.org/10.1111/1475-679X.12457>.


- Bradshaw, M. T., M. S. Drake, J. N. Myers, and L. A. Myers. 2012. A re-examination of analysts' superiority over time-series forecasts of annual earnings. *Review of Accounting Studies* 17 (4): 944–968. <https://doi.org/10.1007/s11142-012-9185-8>.
- Bradshaw, M. T., Y. Ertimur, and P. C. O'Brien. 2017. Financial analysts and their contribution to well-functioning capital markets. *Foundations and Trends in Accounting* 11 (3): 119–191. <https://doi.org/10.1561/14000000042>.
- Bradshaw, M. T., C. Ma, B. P. Yost, and Y. Zou. 2026. Generative AI use by capital market information intermediaries: Evidence from Seeking Alpha. *Journal of Accounting Research* . doi:10.1111/1475-679X.70053; forthcoming.
- Bushee, B. J., D. A. Matsumoto, and G. S. Miller. 2003. Open versus closed conference calls: The determinants and effects of broadening access to disclosure. *Journal of Accounting and Economics* 34 (1–3): 149–180. [https://doi.org/10.1016/S0165-4101\(02\)00073-3](https://doi.org/10.1016/S0165-4101(02)00073-3).
- Bushee, B. J., J. E. Core, W. Guay, and S. J. W. Hamm. 2010. The role of the business press as an information intermediary. *Journal of Accounting Research* 48 (1): 1–19. <https://doi.org/10.1111/j.1475-679X.2009.00357.x>.
- CFA Institute. 2024. AI industry survey: Investment industry employers' views on the responsible and ethical use of AI. CFA Institute Research and Policy Center Survey Report. <https://www.cfainstitute.org/about/press-room/2024/ai-in-investment-sector-survey>. <https://www.cfainstitute.org/about/press-room/2024/ai-in-investment-sector-survey>.
- Chang, A. Y., X. Dong, X. Martin, and C. Zhou. 2026. AI democratization and trading inequality. *Journal of Accounting Research* . <https://doi.org/10.1111/1475-679X.70063>.
- Chen, L., M. Zaharia, and J. Zou. 2023. How is ChatGPT's behavior changing over time? arXiv preprint arXiv:2307.09009; <https://arxiv.org/abs/2307.09009>.
- Cheng, Q., P. Lin, and Y. Zhao. 2025. Does generative AI facilitate investor trading? early evidence from ChatGPT outages. *Journal of Accounting and Economics* 80 (2): 101807. <https://doi.org/10.1016/j.jacceco.2025.101807>.
- Clemen, R. T. 1989. Combining forecasts: A review and annotated bibliography. *International Journal of Forecasting* 5 (4): 559–583. [https://doi.org/10.1016/0169-2070\(89\)90012-5](https://doi.org/10.1016/0169-2070(89)90012-5).
- Clement, M. B. 1999. Analyst forecast accuracy: Do ability, resources, and portfolio complexity matter? *Journal of Accounting and Economics* 27 (3): 285–303. [https://doi.org/10.1016/S0165-4101\(99\)00013-0](https://doi.org/10.1016/S0165-4101(99)00013-0).
- Clement, M. B. and S. Y. Tse. 2003. Do investors respond to analysts' forecast revisions as if forecast accuracy is all that matters? *The Accounting Review* 78 (1): 227–249. <https://doi.org/10.2308/accr.2003.78.1.227>.
- Cohen, L. and D. Lou. 2012. Complicated firms. *Journal of Financial Economics* 104 (2): 383–400. <https://doi.org/10.1016/j.jfineco.2011.08.006>.
- Coleman, B., T. Dyer, and M. Lang. 2025. The influence of financial data subscriptions on analyst research. Working paper, University of Georgia, Brigham Young University, and University of North Carolina at Chapel Hill.

- Collins, D. W. and S. P. Kothari. 1989. An analysis of intertemporal and cross-sectional determinants of earnings response coefficients. *Journal of Accounting and Economics* 11 (2–3): 143–181. [https://doi.org/10.1016/0165-4101\(89\)90004-9](https://doi.org/10.1016/0165-4101(89)90004-9).
- Da, Z., J. Engelberg, and P. Gao. 2011. In search of attention. *Journal of Finance* 66 (5): 1461–1499. <https://doi.org/10.1111/j.1540-6261.2011.01679.x>.
- Da, Z., J. Engelberg, and P. Gao. 2015. The sum of all FEARS: Investor sentiment and asset prices. *Review of Financial Studies* 28 (1): 1–32. <https://doi.org/10.1093/rfs/hhu072>.
- deHaan, E., T. Shevlin, and J. Thornock. 2015. Market (in)attention and the strategic scheduling and timing of earnings announcements. *Journal of Accounting and Economics* 60 (1): 36–55. <https://doi.org/10.1016/j.jacceco.2015.03.003>.
- DellaVigna, S. and J. M. Pollet. 2009. Investor inattention and Friday earnings announcements. *Journal of Finance* 64 (2): 709–749. <https://doi.org/10.1111/j.1540-6261.2009.01447.x>.
- Diether, K. B., C. J. Malloy, and A. Scherbina. 2002. Differences of opinion and the cross section of stock returns. *Journal of Finance* 57 (5): 2113–2141. <https://doi.org/10.1111/0022-1082.00490>.
- Drake, M. S., D. T. Roulstone, and J. R. Thornock. 2012. Investor information demand: Evidence from Google searches around earnings announcements. *Journal of Accounting Research* 50 (4): 1001–1040. <https://doi.org/10.1111/j.1475-679X.2012.00443.x>.
- Drake, M. S., D. T. Roulstone, and J. R. Thornock. 2015. The determinants and consequences of information acquisition via EDGAR. *Contemporary Accounting Research* 32 (3): 1128–1161. <https://doi.org/10.1111/1911-3846.12119>.
- Driskill, M., M. P. Kirk, and J. W. Tucker. 2020. Concurrent earnings announcements and analysts' information production. *The Accounting Review* 95 (1): 165–189. <https://doi.org/10.2308/accr-52489>.
- Easton, P. D. and M. E. Zmijewski. 1989. Cross-sectional variation in the stock market response to accounting earnings announcements. *Journal of Accounting and Economics* 11 (2–3): 117–141. [https://doi.org/10.1016/0165-4101\(89\)90003-7](https://doi.org/10.1016/0165-4101(89)90003-7).
- Goldstein, I., C. S. Spatt, and M. Ye. 2021. Big data in finance. *The Review of Financial Studies* 34 (7): 3213–3225. <https://doi.org/10.1093/rfs/hhab038>.
- Granger, C. W. J. and R. Ramanathan. 1984. Improved methods of combining forecasts. *Journal of Forecasting* 3 (2): 197–204. <https://doi.org/10.1002/for.3980030207>.
- Gu, S., B. Kelly, and D. Xiu. 2020. Empirical asset pricing via machine learning. *The Review of Financial Studies* 33 (5): 2223–2273. <https://doi.org/10.1093/rfs/hhaa009>.
- Hirshleifer, D., S. S. Lim, and S. H. Teoh. 2009. Driven to distraction: Extraneous events and underreaction to earnings news. *Journal of Finance* 64 (5): 2289–2325. <https://doi.org/10.1111/j.1540-6261.2009.01501.x>.
- Hou, K., M. A. van Dijk, and Y. Zhang. 2012. The implied cost of capital: A new approach. *Journal of Accounting and Economics* 53 (3): 504–526. <https://doi.org/10.1016/j.jacceco.2011.12.001>.

- Kimbrough, M. D. 2005. The effect of conference calls on analyst and market underreaction to earnings announcements. *The Accounting Review* 80 (1): 189–219. <https://doi.org/10.2308/accr.2005.80.1.189>.
- Kimbrough, M. D., M. Subasi, and L. Yang Liu. 2026. AI-powered analysts. Working paper, University of Maryland.
- Kothari, S. P. 2001. Capital markets research in accounting. *Journal of Accounting and Economics* 31 (1–3): 105–231. [https://doi.org/10.1016/S0165-4101\(01\)00030-1](https://doi.org/10.1016/S0165-4101(01)00030-1).
- Larocque, S. A., J. Watkins, and E. H. Weisbrod. 2025. Consensus? an examination of differences in earnings information across forecast data providers. SSRN Working Paper No. 4640370; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4640370. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4640370.
- Lawrence, A. 2013. Individual investors and financial disclosure. *Journal of Accounting and Economics* 56 (1): 130–147. <https://doi.org/10.1016/j.jacceco.2013.05.001>.
- Liu, L. Y. and M. Subasi. 2026. The role of generative ai in democratizing investment research: Evidence from crowdsourced forecasts. Working paper, University of Maryland.
- Livnat, J. and R. R. Mendenhall. 2006. Comparing the post-earnings announcement drift for surprises calculated from analyst and time series forecasts. *Journal of Accounting Research* 44 (1): 177–205. <https://doi.org/10.1111/j.1475-679X.2006.00196.x>.
- Lopez-Lira, A. and Y. Tang. 2023. Can ChatGPT forecast stock price movements? return predictability and large language models. SSRN Working Paper No. 4412788; arXiv:2304.07619; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4412788. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4412788.
- Lopez-Lira, A., Y. Tang, and M. Zhu. 2025. The memorization problem: Can we trust LLMs' economic forecasts? SSRN Working Paper No. 5217505; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5217505. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5217505.
- Loughran, T. and B. McDonald. 2011. When is a liability not a liability? Textual analysis, dictionaries, and 10-Ks. *The Journal of Finance* 66 (1): 35–65. <https://doi.org/10.1111/j.1540-6261.2010.01625.x>.
- Loughran, T. and B. McDonald. 2017. The use of EDGAR filings by investors. *Journal of Behavioral Finance* 18 (2): 231–248. <https://doi.org/10.1080/15427560.2017.1308945>.
- Michael, R., A. Rubin, D. Segal, and A. Vadrashko. 2024. Do differences in analyst quality matter for investors relying on consensus information? *Management Science* 70 (2): 751–772. doi:10.1287/mnsc.2023.4699.
- Mikhail, M. B., B. R. Walther, and R. H. Willis. 1997. Do security analysts improve their performance with experience? *Journal of Accounting Research* 35: 131–157. <https://doi.org/10.2307/2491456>. Supplement.
- Park, C. W. and E. K. Stice. 2000. Analyst forecasting ability and the stock price reaction to

- forecast revisions. *Review of Accounting Studies* 5 (3): 259–272.
<https://doi.org/10.1023/A:1009668711298>.
- Schipper, K. 1991. Analysts' forecasts. *Accounting Horizons* 5 (4): 105–121. December issue.
- Stickel, S. E. 1992. Reputation and performance among security analysts. *Journal of Finance* 47 (5): 1811–1836. doi:[10.1111/j.1540-6261.1992.tb04672.x](https://doi.org/10.1111/j.1540-6261.1992.tb04672.x).
- Stock, J. H. and M. W. Watson. 2004. Combination forecasts of output growth in a seven-country data set. *Journal of Forecasting* 23 (6): 405–430. <https://doi.org/10.1002/for.928>.
- Surowiecki, J. 2004. *The Wisdom of Crowds: Why the Many Are Smarter Than the Few and How Collective Wisdom Shapes Business, Economies, Societies and Nations*. New York, Doubleday. ISBN 0-385-50386-5.
- Tetlock, P. C. 2007. Giving content to investor sentiment: The role of media in the stock market. *The Journal of Finance* 62 (3): 1139–1168. <https://doi.org/10.1111/j.1540-6261.2007.01232.x>.
- Timmermann, A. 2006. Forecast combinations. In Elliott, G., C. W. J. Granger, and A. Timmermann, editors, *Handbook of Economic Forecasting*, volume 1, chapter 4, pages 135–196. Elsevier. [https://doi.org/10.1016/S1574-0706\(05\)01004-9](https://doi.org/10.1016/S1574-0706(05)01004-9).
- Xue, J., Q. Zhang, and W. Zhu. 2025. Generative AI for analysts. SSRN Working Paper No. 5821142; Available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5821142.


FIGURE 1. Theoretical Framework




Notes. This figure summarizes the theoretical framework. Generative AI (GenAI) lowers the cost of synthesizing analyst consensus forecasts across the five forecast data providers (FDPs), the joint cost of *awareness*, *acquisition*, and *integration* in the Blankespoor et al. (2020) decomposition, which in turn enables the market to price a sophisticated, accuracy-weighted synthesis across providers in place of a single-FDP or equal-weighted heuristic. The strength of the channel is moderated by how *accessible* synthesis is for a given firm (LLM visibility, the information-acquisition channel) and by how *valuable* it is (cross-FDP disagreement and special-items magnitude). These mechanisms map to Predictions 1 and 2, developed in Section 2; Prediction 3 adds the demand-side corollary that the numbers themselves are unchanged.

FIGURE 2. Mean Magnitude of SUE : Accuracy-Weighted vs. Heuristic Models

(a) Quarterly mean SUE by consensus method




(b) Mean SUE by consensus method



Notes. This figure compares the accuracy of the accuracy-weighted multi-vendor consensus SUE^{Acc} against the equal-weighted five-FDP consensus SUE^{EW5} and the I/B/E/S consensus SUE^{IBES} across the five-FDP sample (FactSet, Bloomberg, Zacks, Capital IQ, I/B/E/S) of S&P 1500 firm-quarters with announcement dates in 2021Q1–2025Q4. $|SUE|$ is the absolute standardized unexpected earnings scaled as a percentage of pre-announcement price, winsorized at the (1, 99) level. **Panel (a)** plots the quarterly mean $|SUE|$ for each of the three methods over 20 announcement quarters. The vertical reference line marks the public release of ChatGPT (November 30, 2022). **Panel (b)** shows the pooled mean $|SUE|$ within each Pre- vs. Post-ChatGPT regime, with 95% confidence intervals from the cell-level standard error of the mean; the y-axis is truncated below 0.20 (break mark on axis) to expand the visible range.

FIGURE 3. The Accuracy Premium Rises with GPT Retrieval Capability



Notes. This figure plots the *accuracy premium*, the contrast $\beta_1 - \beta_2$ between the announcement-return loading on the accuracy-weighted consensus surprise SUE^{Acc} and the loading on the benchmark surprise, estimated separately within each GPT-capability regime. The estimates are taken directly from Table 6: the red line uses the I/B/E/S benchmark (Panel B) and the navy line the equal-weighted five-FDP benchmark (Panel A). Regimes are ordered by ChatGPT’s information-retrieval capability: **R0** (before 2022-11-30, pre-ChatGPT), **R1** (text-only), **R2** (browsing, from 2023-09-27), and **R3** (native search, from 2024-10-31); the shaded region marks the regimes in which the model can retrieve forecasts from the web. Whiskers are 95% confidence intervals; point labels report the premium and its significance (* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$) from firm-clustered standard errors. The premium is statistically indistinguishable from zero before retrieval is available, turns positive once the model can browse, and grows monotonically thereafter, the dose-response Prediction 2 predicts.

TABLE 1. Sample Selection

Panel A. Main Sample (firm $\times$ quarter)	Observations	# Firms
S&P 1500 firm-quarters with an earnings announcement (announcement date 2021–2025)	29,395	1,502
Retain firm-quarters with non-missing consensus from all five FDPs (FactSet, Bloomberg, Zacks, Capital IQ, and I/B/E/S)	25,027	1,353
Drop firm-quarters missing market-outcome variables or required controls	(1,179)	(13)
Drop firm-quarters missing any constructed return or predicted-surprise variable	(10)	(1)
Drop firms with only one firm-quarter (no within-firm variation for firm fixed effects)	(5)	(5)
Final Sample	23,833	1,334
Panel B. Information Acquisition Tests (platform $\times$ day)	Observations	# Platforms
Identified platforms (5 FDPs + 15 distributors)	20	20
Platform-day observations, 2020–2025	43,820	20
Drop platform-days missing Ads scaling	(2,191)	(1)
Require consistent Trends coverage 2020–2025	(4,382)	(2)
Final Sample	37,247	17

Notes. This table reports the sample selection procedure for the two analytic samples in the paper. **Panel A** reports the firm-quarter sample used in the main across-FDP tests, built from CRSP, Compustat, I/B/E/S, and the five-FDP consensus panel. **Panel B** reports the platform-day sample used in the information-acquisition test.

TABLE 2. Descriptive Statistics

	N	Mean	SD	Min	p25	Median	p75	Max
Panel A. Main Sample (Firm $\times$ quarter).								
Dependent Variables: Abnormal Returns (%)								
$CAR_{[-1,+1]}^{Mkt}$	23,833	0.204	7.739	-22.532	-3.933	0.168	4.255	23.535
$CAR_{[-1,+1]}^{FF4}$	23,833	0.196	7.635	-22.779	-3.776	0.186	4.135	23.606
$BHAR_{[-2,+1]}^{Mkt}$	23,833	0.228	7.978	-22.803	-4.165	0.177	4.418	24.992
$BHAR_{[-2,+1]}^{FF4}$	23,833	0.193	7.847	-22.935	-3.966	0.111	4.255	24.819
Per-FDP Standardized Earnings Surprise (SUE^g, %)								
$SUE^{FactSet}$	23,833	0.132	0.534	-2.190	0.000	0.074	0.238	2.484
$SUE^{Bloomberg}$	23,833	-0.222	1.659	-10.049	-0.362	-0.027	0.162	5.780
SUE^{Zacks}	23,833	0.149	0.467	-1.606	0.000	0.075	0.240	2.291
SUE^{CapIQ}	23,833	0.137	0.530	-2.185	0.001	0.076	0.243	2.477
SUE^{IBES}	23,833	0.130	0.548	-2.271	0.000	0.075	0.241	2.502
Cross-FDP Composites and Predicted Surprise								
SUE^{Acc}	23,833	0.110	0.450	-1.656	-0.015	0.059	0.212	1.972
SUE^{EW5}	23,833	0.066	0.510	-1.936	-0.064	0.041	0.201	1.942
<i>Standardized Pred. Surp.</i> ^{EW5}	23,833	0.606	0.116	0.000	0.590	0.619	0.631	1.000
<i>Standardized Pred. Surp.</i> ^{IBES}	23,833	0.483	0.109	0.000	0.469	0.478	0.493	1.000
Firm and Event Controls								
<i>Log Market Equity</i>	23,833	8.992	1.448	6.390	7.923	8.769	9.942	12.883
<i>Book-to-Market</i>	23,833	0.440	0.361	-0.159	0.170	0.362	0.634	1.670
<i>Leverage</i>	23,833	0.314	0.218	0.001	0.133	0.301	0.444	1.008
<i>Prior 12-Month Return</i>	23,833	0.201	0.487	-0.589	-0.101	0.112	0.379	2.400
<i>Log Price</i>	23,833	4.166	0.980	1.880	3.491	4.178	4.831	6.580
<i>Return Volatility</i>	23,833	0.021	0.009	0.009	0.015	0.019	0.025	0.054
<i>Analyst Dispersion</i>	23,833	0.002	0.003	0.000	0.000	0.001	0.002	0.018
<i>Busyness</i>	23,833	4.944	0.973	1.946	4.431	5.124	5.684	6.299
$ \Delta IB $	23,833	0.012	0.026	0.000	0.001	0.004	0.011	0.178
$ \Delta Shares $	23,833	0.010	0.021	0.000	0.001	0.003	0.009	0.150
$ SpecItems $	23,833	0.003	0.010	0.000	0.000	0.000	0.002	0.074
<i>SBC High</i>	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
<i>Unexpected Item</i>	23,833	0.686	0.464	0.000	0.000	1.000	1.000	1.000
<i>Stock Split</i>	23,833	0.003	0.051	0.000	0.000	0.000	0.000	1.000
<i>Fiscal Q4</i>	23,833	0.250	0.433	0.000	0.000	0.000	0.000	1.000
<i>Institutional Ownership</i>	23,833	0.850	0.238	0.011	0.813	0.940	1.000	1.000
<i># Analysts</i>	23,833	10.914	6.536	2.000	6.000	10.000	15.000	31.000
Cross-Sectional Partitioning Indicators								
<i>High Visibility</i>	23,827	0.489	0.500	0.000	0.000	0.000	1.000	1.000

(continued on next page)

	N	Mean	SD	Min	p25	Median	p75	Max
<i>High Cross-FDP Dispersion</i>	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
<i>High Special Items</i>	23,833	0.500	0.500	0.000	0.000	0.000	1.000	1.000
<i>IBES Accuracy-Rank Up</i>	23,833	0.099	0.299	0.000	0.000	0.000	0.000	1.000
<i>IBES Accuracy-Rank Down</i>	23,833	0.059	0.236	0.000	0.000	0.000	0.000	1.000
<i>ChatGPT Outage (prior day)</i>	23,833	0.004	0.059	0.000	0.000	0.000	0.000	1.000
Panel B. Information Acquisition Sample (Platform $\times$ day).								
Dependent Variable								
<i>Brand Share ($S_{\text{brand}}/S_{\text{allplatforms}}$)</i>	37,247	0.059	0.101	0.000	0.005	0.022	0.071	0.832
Macro and Calendar Controls								
<i>Log Total Platform Searches</i>	37,247	15.089	1.148	12.462	14.239	15.100	15.891	17.732
<i>Log # Earnings Today</i>	37,247	2.996	2.017	0.000	0.693	3.434	4.605	6.410
<i>Trading Day</i>	37,247	0.714	0.452	0.000	0.000	1.000	1.000	1.000
Vendor-Group Indicators								
<i>Free Distributor</i>	37,247	0.706	0.456	0.000	0.000	1.000	1.000	1.000
<i>Free FDP</i>	37,247	0.059	0.235	0.000	0.000	0.000	0.000	1.000
<i>Paywalled FDP</i>	37,247	0.235	0.424	0.000	0.000	0.000	0.000	1.000
Regime Indicators								
<i>Post-ChatGPT</i>	37,247	0.514	0.500	0.000	0.000	1.000	1.000	1.000
<i>Regime 1</i>	37,247	0.137	0.344	0.000	0.000	0.000	0.000	1.000
<i>Regime 2</i>	37,247	0.183	0.386	0.000	0.000	0.000	0.000	1.000
<i>Regime 3</i>	37,247	0.194	0.396	0.000	0.000	0.000	0.000	1.000

Notes. This table reports descriptive statistics for the two analytic samples. **Panel A** reports the across-FDP firm-quarter sample used in the main tests; **Panel B** reports the platform-day sample used in the information-acquisition test. All continuous variables are winsorized at the (1, 99) level. Variable definitions are in [Appendix A](#).

TABLE 3. Initial Evidence on Selective Reliance

Panel A. Earnings-response coefficients by forecast-provider accuracy rank

Provider	Pre-GPT		Post-GPT	
	β	γ	β	γ
IBES	2.591*** (6.60)	+0.022 (0.17)	2.611*** (5.49)	+0.295** (2.09)
FactSet	2.316*** (4.78)	+0.050 (0.36)	2.059*** (3.80)	+0.380** (2.42)
Bloomberg	0.985*** (5.82)	-0.058 (-0.80)	0.789*** (5.01)	+0.162** (2.10)
Zacks	3.603*** (8.35)	-0.139 (-1.04)	4.757*** (8.57)	-0.020 (-0.13)
CapIQ	1.107** (2.40)	+0.543*** (3.94)	2.483*** (4.94)	+0.349** (2.40)
Stacked	1.206*** (9.21)	+0.306*** (6.93)	0.973*** (8.09)	+0.621*** (11.65)
<i>Observations</i>	47,130		72,035	
Stacked Δ (Post-Pre)			-0.232 (-1.38)	+0.316*** (4.91)

Panel B. Heterogeneous post-ChatGPT effects by firm-level IBES accuracy-rank change

Dep. Var.:	Market-adjusted $CAR_{[-1,+1]}$		FF4 $CAR_{[-1,+1]}$	
	(1)	(2)	(3)	(4)
SUE^{IBES}	2.859*** (13.91)	2.845*** (14.03)	2.775*** (13.67)	2.776*** (13.84)
$SUE^{IBES} \times Post$	1.029*** (3.56)	0.978*** (3.32)	1.136*** (4.11)	1.065*** (3.76)
$SUE^{IBES} \times IBES\ Up$	0.693 (1.24)	0.628 (1.16)	0.803 (1.46)	0.737 (1.39)
$SUE^{IBES} \times IBES\ Down$	1.397* (1.81)	1.582** (2.00)	1.913*** (2.61)	2.080*** (2.74)
$SUE^{IBES} \times Post \times IBES\ Up$	-1.248 (-1.44)	-1.135 (-1.30)	-1.291 (-1.40)	-1.185 (-1.26)
$SUE^{IBES} \times Post \times IBES\ Down$	-2.278** (-2.27)	-2.344** (-2.15)	-2.758*** (-2.70)	-2.791** (-2.51)
$Post \times IBES\ Up$	0.353 (1.13)	0.232 (0.68)	0.437 (1.37)	0.297 (0.85)
$Post \times IBES\ Down$	0.544 (1.34)	0.713 (1.61)	0.678* (1.68)	0.859* (1.95)
Controls	No	Yes	No	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	23,833	23,833	23,833	23,833
Adjusted R^2	0.052	0.070	0.054	0.072

Notes. Both panels test whether the market relies more on more accurate forecasts. **Panel A:** within each firm-quarter the five FDPs are ranked by their ex-ante trailing four-quarter MAFE (Rank 1 = least accurate, Rank 5 = most accurate); pooling across firm-quarters, $CAR_{[-1,+1]}^{Mkt} = \alpha + \beta SUE^g + \delta Rank^g +$

$\gamma(SUE^g \times Rank^g)$, where γ is the ERC increment per one-rank accuracy improvement. Each provider row is its own regression; the *Stacked* row pools all five FDPs; the *Stacked Δ* row reports the $SUE \times Post$ and $SUE \times Rank \times Post$ coefficients. **Panel B:** firms are split by the change in IBES's rounded mean rank among the five FDPs between the pre- and post-ChatGPT windows ($IBESUp = 1$ if the rank improved by ≥ 2 positions; $IBESDown = 1$ if it worsened by ≥ 2); the bolded triple interactions are the key coefficients, and Panel B includes firm and year fixed effects. Pre-GPT (Post-GPT) covers announcements before (on or after) 2022-11-30. Standard errors are clustered by firm; continuous variables are winsorized at (1, 99). t -statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Variable definitions are in [Appendix A](#).

TABLE 4. Accuracy Persistence and the Predicted-Surprise Trading Strategy**Panel A. Persistence of FDP accuracy ranks (trailing-MAFE rank, quarter t to $t + 1$)**

$rank_t$	$\Pr(rank_{t+1} = j \mid rank_t = i), \text{ in } \%$				
	$j = 1 \text{ (least)}$	$j = 2$	$j = 3$	$j = 4$	$j = 5 \text{ (most)}$
1 (least)	76.1	12.1	4.9	3.4	3.6
2	12.1	54.1	19.7	8.9	5.2
3	4.7	19.5	46.7	21.1	8.0
4	3.4	9.1	20.8	48.3	18.5
5 (most)	3.8	5.2	7.9	18.4	64.7

Spearman $\rho = 0.71$; overall same-rank rate = 58%; 109,915 quarter-to-quarter pairs.

Panel B. Trading strategy from predicted surprise

Dep. Var.:	Market-adj. $BHAR_{[-2,+1]}$		FF4 $BHAR_{[-2,+1]}$	
Benchmark:	Equal-Weighted	IBES	Equal-Weighted	IBES
	(1)	(2)	(3)	(4)
<i>Standardized Pred. Surp.</i>	2.493*** (4.10)	3.433*** (5.35)	2.440*** (4.08)	3.357*** (5.25)
Long-Short Return (Q5-Q1)	+0.73%***	+1.35%***	+0.63%***	+1.35%***
Controls / Firm FE / Year FE	Yes	Yes	Yes	Yes
Observations (firm-q)	23,833	23,833	23,833	23,833
Adjusted R^2	0.016	0.017	0.018	0.019

Notes. This table motivates the accuracy-weighted consensus. **Panel A** reports the quarter-to-quarter persistence of forecast-provider accuracy ranks. Within each firm-quarter the five FDPs are ranked by their trailing four-quarter mean absolute forecast error (Rank 1 = least accurate, Rank 5 = most accurate, the same ex-ante measure used to build SUE^{Acc}); cell (i, j) is the share of provider-quarters at Rank i that move to Rank j the following quarter (rows sum to 100). **Panel B** horse-races two predicted-surprise signals as predictors of post-announcement abnormal returns: for benchmark $X \in \{EW5, IBES\}$ the predicted surprise is $(Consensus^{Acc} - Consensus^X) / \text{price}$, min-max standardized to $[0, 1]$; columns (1)–(2) use the market-adjusted $BHAR_{[-2,+1]}$ and (3)–(4) the FF4 four-factor $BHAR_{[-2,+1]}$. The Long-Short row reports the abnormal return from going long the top quintile and short the bottom quintile of predicted surprise. All Panel B columns include firm and year fixed effects, the full control set, and firm-clustered standard errors; continuous variables winsorized at (1, 99). t -statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Variable definitions are in [Appendix A](#).

TABLE 5. Accuracy-Weighting: Does the Market Place Higher Weight on More Accurate FDPs?

Panel A. Horse race: SUE^{Acc} vs SUE^{EW5} .

Dep. Var.:	Mkt CAR _[-1,+1]			FF4 CAR _[-1,+1]		
	(1) Pre	(2) Post	(3) Pooled	(4) Pre	(5) Post	(6) Pooled
$SUE^{Acc} (\beta_1)$	1.376*** (3.45)	3.420*** (9.38)	1.307*** (3.60)	1.401*** (3.63)	3.519*** (10.01)	1.290*** (3.67)
$SUE^{Acc} \times Post (\beta_3)$			2.003*** (4.08)			2.135*** (4.47)
$SUE^{EW5} (\beta_2)$	2.324*** (6.10)	1.767*** (5.56)	2.279*** (6.84)	2.379*** (6.44)	1.734*** (5.62)	2.284*** (7.03)
$SUE^{EW5} \times Post (\beta_4)$			-0.399 (-0.92)			-0.449 (-1.06)
Premium (Pre-GPT)	-0.948 (-1.28)		-0.972 (-1.47)	-0.978 (-1.37)		-0.994 (-1.55)
Premium (Post-GPT)		1.654*** (2.62)	1.429** (2.39)		1.785*** (2.93)	1.590*** (2.74)
Δ Premium (Post-Pre)			2.401*** (2.75)			2.584*** (3.03)
Control Variables						
<i>Log Market Equity</i>	-5.348*** (-5.35)	-4.964*** (-5.35)	-2.848*** (-6.00)	-5.585*** (-5.59)	-5.125*** (-5.74)	-3.238*** (-7.19)
<i>Book-to-Market</i>	0.413 (0.36)	-1.315 (-1.22)	0.168 (0.28)	0.849 (0.75)	-1.245 (-1.16)	0.309 (0.53)
<i>Leverage</i>	-4.957** (-2.38)	-3.652* (-1.70)	-0.913 (-0.76)	-4.561** (-2.17)	-4.067** (-1.97)	-1.312 (-1.13)
<i>Prior 12-Month Return</i>	-0.295 (-1.27)	-0.868** (-2.52)	-0.235 (-1.27)	-0.024 (-0.11)	-1.038*** (-3.02)	-0.129 (-0.71)
<i>Log Price</i>	-1.063 (-1.13)	-1.047 (-1.12)	-0.482 (-1.05)	-0.671 (-0.71)	-0.735 (-0.83)	-0.029 (-0.07)
<i>Return Volatility</i>	-59.587*** (-3.08)	27.445* (1.66)	14.899 (1.34)	-53.652*** (-2.87)	19.210 (1.20)	11.467 (1.05)
<i>Analyst Dispersion</i>	-23.636 (-0.40)	-40.842 (-0.71)	-15.946 (-0.40)	-11.993 (-0.20)	-17.323 (-0.30)	-2.625 (-0.07)
<i>Busyness</i>	-0.104 (-0.73)	-0.245* (-1.83)	-0.172* (-1.79)	-0.132 (-0.93)	-0.219* (-1.66)	-0.164* (-1.72)
$ \Delta IB $	16.729*** (2.98)	-4.012 (-0.78)	2.450 (0.69)	17.171*** (3.09)	-4.609 (-0.91)	2.027 (0.58)
$ \Delta Shares $	3.721 (0.79)	-1.964 (-0.51)	1.505 (0.52)	3.251 (0.72)	-2.171 (-0.57)	1.376 (0.48)
$ SpecItems $	-18.458 (-1.50)	-8.991 (-0.79)	-10.912 (-1.33)	-21.400* (-1.77)	-7.112 (-0.64)	-9.847 (-1.22)
<i>SBC High</i>	-0.149 (-0.57)	0.174 (0.77)	0.106 (0.65)	-0.168 (-0.66)	0.186 (0.83)	0.095 (0.59)
<i>Unexpected Item</i>	-0.258 (-0.97)	-0.036 (-0.17)	-0.123 (-0.78)	-0.222 (-0.89)	-0.023 (-0.11)	-0.103 (-0.68)
<i>Stock Split</i>	-2.679** (-2.21)	-1.110 (-0.92)	-1.584* (-1.90)	-2.266* (-1.78)	-0.472 (-0.42)	-0.977 (-1.18)

<i>Fiscal Q4</i>	0.600*** (2.83)	-0.045 (-0.24)	0.276** (2.07)	0.606*** (2.89)	0.230 (1.23)	0.417*** (3.19)
<i>Institutional Ownership</i>	3.234 (1.56)	-0.702** (-2.02)	-0.470 (-1.41)	1.626 (0.83)	-0.543 (-1.57)	-0.416 (-1.25)
<i>Log # Analysts</i>	0.508 (0.79)	-0.218 (-0.39)	-0.248 (-0.65)	0.648 (1.04)	-0.328 (-0.60)	-0.371 (-0.99)
Controls	Yes	Yes	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes	Yes
Observations	9,405	14,402	23,833	9,405	14,402	23,833
Adjusted R^2	0.085	0.106	0.080	0.092	0.110	0.083

Panel B. Horse race: SUE^{Acc} vs SUE^{IBES} .

Dep. Var.:	Mkt $CAR_{[-1,+1]}$			FF4 $CAR_{[-1,+1]}$		
	(1) Pre	(2) Post	(3) Pooled	(4) Pre	(5) Post	(6) Pooled
$SUE^{Acc} (\beta_1)$	1.428*** (4.00)	3.678*** (10.49)	1.430*** (4.31)	1.527*** (4.36)	3.796*** (10.70)	1.494*** (4.48)
$SUE^{Acc} \times Post (\beta_3)$			2.153*** (4.82)			2.204*** (4.98)
$SUE^{IBES} (\beta_2)$	2.176*** (7.08)	1.460*** (4.44)	2.054*** (7.31)	2.149*** (7.20)	1.407*** (4.20)	1.974*** (7.01)
$SUE^{IBES} \times Post (\beta_4)$			-0.519 (-1.34)			-0.484 (-1.28)
Premium (Pre-GPT)	-0.747 (-1.20)		-0.623 (-1.09)	-0.622 (-1.03)		-0.480 (-0.83)
Premium (Post-GPT)		2.219*** (3.56)	2.049*** (3.60)		2.389*** (3.75)	2.208*** (3.82)
Δ Premium (Post-Pre)			2.672*** (3.44)			2.688*** (3.50)
Controls	Yes	Yes	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes	Yes
Observations	9,405	14,402	23,833	9,405	14,402	23,833
Adjusted R^2	0.087	0.106	0.080	0.094	0.110	0.083

Notes. This table reports OLS estimates of the horse race between the accuracy-weighted surprise SUE^{Acc} and a benchmark surprise, including their interactions with *Post*. SUE^{Acc} uses softmax weights based on each provider's within-firm-quarter z -scored four-quarter rolling MAFE. **Panel A** uses the equal-weighted benchmark SUE^{EW5} ; **Panel B** the I/B/E/S benchmark and uses the same firm and event controls as Panel A, suppressed in the table for brevity. Columns (1)–(3) use the market-adjusted $CAR_{[-1,+1]}$; columns (4)–(6) the FF4 four-factor $CAR_{[-1,+1]}$. All columns include firm and year fixed effects and the full control set; standard errors are clustered by firm; continuous variables are winsorized at (1, 99). t -statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Variable definitions are in [Appendix A](#).

TABLE 6. The Accuracy Premium Across GPT-Capability Regimes

Panel A. Benchmark = Equal-Weighted Five-FDP Consensus

Regime:	(1) R0 Pre-ChatGPT	(2) R1 text-only	(3) R2 browsing	(4) R3 native search
$SUE^{Acc} (\beta_1)$	1.376*** (3.45)	2.597*** (3.38)	3.753*** (6.19)	3.954*** (5.69)
Benchmark $SUE (\beta_2)$	2.324*** (6.10)	2.419*** (3.78)	1.325*** (2.63)	1.404** (2.26)
Premium ($\beta_1 - \beta_2$)	-0.948 (-1.28)	0.178 (0.13)	2.429** (2.35)	2.550** (2.07)
Controls / Firm FE / Year FE	Yes	Yes	Yes	Yes
Observations	9,405	3,705	5,529	5,084
Adjusted R^2	0.085	0.158	0.151	0.140

Panel B. Benchmark = I/B/E/S Consensus

Regime:	(1) R0 Pre-ChatGPT	(2) R1 text-only	(3) R2 browsing	(4) R3 native search
$SUE^{Acc} (\beta_1)$	1.428*** (4.00)	1.918** (2.35)	3.362*** (6.17)	5.079*** (7.87)
Benchmark $SUE (\beta_2)$	2.176*** (7.08)	3.144*** (4.70)	1.842*** (3.74)	0.169 (0.30)
Premium ($\beta_1 - \beta_2$)	-0.747 (-1.20)	-1.226 (-0.86)	1.519 (1.61)	4.909*** (4.41)
Controls / Firm FE / Year FE	Yes	Yes	Yes	Yes
Observations	9,405	3,705	5,529	5,084
Adjusted R^2	0.087	0.170	0.155	0.138

Notes. This table estimates the accuracy premium separately within each GPT-capability regime. For each regime the specification is $CAR_{[-1,+1]}^{Mkt} = \alpha + \beta_1 SUE_{i,q}^{Acc} + \beta_2 SUE_{i,q}^X + \text{controls} + \alpha_i + \tau_t + \varepsilon$, and the *Accuracy Premium* is the contrast $\beta_1 - \beta_2$. Regimes are defined by announcement date: **R0** (before 2022-11-30, Pre-ChatGPT), **R1** (2022-11-30 to 2023-09-26, text-only), **R2** (2023-09-27 to 2024-10-30, browsing), and **R3** (from 2024-10-31, native search). The benchmark X is the equal-weighted five-FDP consensus (Panel A) or I/B/E/S (Panel B). All columns include the full control set, firm and year fixed effects, and firm-clustered standard errors; continuous variables winsorized at (1, 99). t -statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Variable definitions are in [Appendix A](#).

TABLE 7. Cross-sectional Tests of the “Accuracy Premium”

Panel A. Concurrent Cross-FDP SUE Dispersion.

Benchmark:	Equal-Weighted		IBES	
	(1) Low	(2) High	(3) Low	(4) High
$SUE^{Acc} (\beta_1)$	4.138*** (6.37)	1.433*** (4.23)	1.569** (2.43)	1.748*** (5.57)
Benchmark $SUE (\beta_2)$	3.318*** (5.86)	2.024*** (6.41)	7.496*** (10.03)	1.560*** (6.10)
$SUE^{Acc} \times Post^{Scr} (\beta_3)$	2.748** (2.48)	1.610*** (2.89)	1.439 (1.34)	1.746*** (3.63)
Benchmark $SUE \times Post^{Scr} (\beta_4)$	0.712 (0.76)	-0.602 (-1.20)	2.130* (1.86)	-0.681* (-1.70)
Premium (Pre-Scrape)	0.821 (0.73)	-0.591 (-0.96)	-5.927*** (-4.58)	0.188 (0.35)
Premium (Post-Scrape)	2.857* (1.71)	1.620** (1.98)	-6.619*** (-3.81)	2.615*** (3.81)
Δ Premium (Post-Pre)	2.037 (1.06)	2.211** (2.21)	-0.691 (-0.33)	2.427*** (2.99)
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	11,832	11,846	11,832	11,846
Adjusted R^2	0.111	0.075	0.127	0.075

Panel B. Special-Items Magnitude.

Benchmark:	Equal-Weighted		IBES	
	(1) Low	(2) High	(3) Low	(4) High
$SUE^{Acc} (\beta_1)$	2.345*** (4.36)	1.613*** (4.49)	1.940*** (4.34)	1.969*** (5.48)
Benchmark $SUE (\beta_2)$	2.153*** (4.21)	2.078*** (6.41)	2.465*** (6.46)	1.615*** (5.25)
$SUE^{Acc} \times Post^{Scr} (\beta_3)$	0.247 (0.31)	2.035*** (3.39)	0.440 (0.58)	2.583*** (4.69)
Benchmark $SUE \times Post^{Scr} (\beta_4)$	0.799 (1.14)	-0.483 (-0.92)	0.659 (0.98)	-1.055** (-2.28)
Premium (Pre-Scrape)	0.192 (0.19)	-0.464 (-0.74)	-0.525 (-0.68)	0.354 (0.58)
Premium (Post-Scrape)	-0.360 (-0.28)	2.054** (2.48)	-0.743 (-0.65)	3.992*** (5.36)
Δ Premium (Post-Pre)	-0.552 (-0.39)	2.518** (2.40)	-0.219 (-0.16)	3.638*** (3.92)
Controls	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Observations	11,823	11,828	11,823	11,828
Adjusted R^2	0.089	0.073	0.095	0.071

Notes. This table estimates the accuracy-weighted horse race against the equal-weighted (EW5) and

IBES benchmarks separately for two cross-sectional partitions. **Panel A** splits firm-quarters at the median concurrent cross-FDP SUE inter-quartile range. **Panel B** splits firm-quarters at the median of $|SpecItems|$. Within each panel, columns (1)–(2) use SUE^{EWS} and columns (3)–(4) use SUE^{IBES} . The Accuracy Premium is the contrast SUE^{Acc} -Benchmark SUE; the bolded Δ row reports the post-pre change. Here $Post^{Scr}$ (post-scraping) equals one for announcements on or after the onset of GPT browsing (2023-09-27), distinct from the post-ChatGPT $Post$ used in the other tables. All panels include the full control vector, firm fixed effects, year fixed effects, and firm-clustered standard errors. The dependent variable is the market-adjusted $CAR_{[-1,+1]}$. Continuous variables are winsorized at the (1, 99) level. t -statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Variable definitions are in [Appendix A](#).

TABLE 8. LLM Visibility and the Accuracy Premium by GPT Regimes

Panel A. Equal-Weighted Five-FDP Benchmark.

Regime:	(1) R0 Pre-ChatGPT	(2) R1 text-only	(3) R2 browsing	(4) R3 native search
$SUE^{Acc} (\beta_1)$	1.361** (2.57)	3.503*** (5.05)	3.077*** (2.95)	3.055*** (4.18)
$SUE^{EW5} (\beta_2)$	2.252*** (4.09)	2.185*** (3.61)	1.885*** (3.65)	2.002*** (3.21)
$SUE^{Acc} \times High_Vis (\beta_3)$	0.020 (0.05)	-2.274 (-1.49)	1.550* (1.79)	1.815* (1.93)
$SUE^{EW5} \times High_Vis (\beta_4)$	0.153 (0.28)	0.889 (0.79)	-1.133** (-2.11)	-1.269 (-1.51)
<i>Accuracy premium = loading on SUE^{Acc} - loading on benchmark SUE:</i>				
Premium (Low Vis)	-0.892 (-0.90)	1.317 (1.12)	1.192 (0.89)	1.053 (0.83)
Premium (High Vis)	-1.025 (-1.57)	-1.845 (-0.94)	3.875*** (4.60)	4.136*** (4.16)
Premium (High - Low)	-0.133 (-0.15)	-3.163 (-1.22)	2.683** (2.12)	3.083* (1.85)
Controls / Firm FE / Year-qtr FE	Yes	Yes	Yes	Yes
Observations	9,399	3,705	5,529	5,084
Adjusted R^2	0.086	0.163	0.165	0.143

Panel B. I/B/E/S Benchmark.

Regime:	(1) R0 Pre-ChatGPT	(2) R1 text-only	(3) R2 browsing	(4) R3 native search
$SUE^{Acc} (\beta_1)$	1.607*** (4.07)	3.138*** (3.15)	2.737*** (6.05)	4.029*** (8.81)
$SUE^{IBES} (\beta_2)$	1.859*** (3.60)	2.424*** (2.96)	2.403** (2.28)	0.974** (2.13)
$SUE^{Acc} \times High_Vis (\beta_3)$	-0.382 (-1.03)	-2.353* (-1.74)	1.667** (2.44)	2.376*** (3.68)
$SUE^{IBES} \times High_Vis (\beta_4)$	0.669 (0.97)	1.362 (1.45)	-1.380 (-1.45)	-1.856*** (-3.06)
<i>Accuracy premium = loading on SUE^{Acc} - loading on benchmark SUE:</i>				
Premium (Low Vis)	-0.252 (-0.31)	0.714 (0.41)	0.335 (0.26)	3.054*** (4.00)
Premium (High Vis)	-1.303** (-2.06)	-3.001** (-2.12)	3.382*** (6.17)	7.286*** (10.56)
Premium (High - Low)	-1.052 (-1.07)	-3.715* (-1.72)	3.047** (2.05)	4.232*** (3.93)
Controls / Firm FE / Year-qtr FE	Yes	Yes	Yes	Yes
Observations	9,399	3,705	5,529	5,084
Adjusted R^2	0.088	0.175	0.169	0.142

Notes. This table estimates the accuracy-weighted horse race $CAR_{[-1,+1]}^{Mkt} = \alpha + \beta_1 SUE_{i,q}^{Acc} + \beta_2 SUE_{i,q}^X + \beta_3 SUE_{i,q}^{Acc} \times High_Vis_i + \beta_4 SUE_{i,q}^X \times High_Vis_i + \text{controls} + FE + \varepsilon$ separately within each GPT-capability regime, where $SUE_{i,q}^{Acc}$ is the accuracy-weighted multi-vendor surprise with softmax weights based on each provider's within-firm-quarter z-scored four-quarter rolling MAFE. Regimes track ChatGPT's information-retrieval capability by announcement date: **R0** (before 2022-11-30, Pre-ChatGPT), **R1** (2022-11-30 to 2023-09-26, text-only), **R2** (2023-09-27 to 2024-10-30, browsing / Bing), and **R3** (from 2024-10-31, native ChatGPT Search). **Panel A** uses the equal-weighted five-FDP benchmark SUE^{EW5} ; **Panel B** the I/B/E/S benchmark. The *Accuracy Premium* is the contrast SUE^{Acc} -Benchmark SUE, reported for low- and high-visibility firms (Premium (Low Vis) = $\beta_1 - \beta_2$; Premium (High Vis) = $(\beta_1 + \beta_3) - (\beta_2 + \beta_4)$); the bolded Premium (High - Low) row is the triple difference $\beta_3 - \beta_4$. All specifications include the full control vector, firm fixed effects, and year-quarter fixed effects; the dependent variable is the market-adjusted $CAR_{[-1,+1]}$. Continuous variables are winsorized at the (1, 99) level; standard errors are two-way clustered by firm and calendar year. *t*-statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Variable definitions are in [Appendix A](#).

TABLE 9. LLM Visibility Is Distinct from Classical Visibility

Panel A. Overlap of high-visibility classifications (%)

	LLM Visibility	Firm Size	Analyst Coverage	Inst Ownership
LLM Visibility	100	50.2	49.7	47.1
Firm Size	51.0	100	79.6	40.7
Analyst Coverage	50.6	79.8	100	44.5
Inst Ownership	47.9	40.7	44.5	100

Panel B. Accuracy-premium concentration: Equal-weighted benchmark

Visibility measure	Univariate Δ Premium		Horse-race
	High group	Low group	triple diff.
LLM Visibility	3.263*** (2.59) [11,219]	0.300 (0.25) [11,825]	3.172* (1.80)
Firm Size	0.723 (0.55) [11,901]	2.447** (2.18) [11,143]	0.111 (0.05)
Analyst Coverage	0.109 (0.08) [11,861]	3.076** (2.57) [11,183]	-3.161 (-1.38)
Inst Ownership	3.224** (2.51) [11,593]	0.646 (0.58) [11,451]	3.222* (1.86)
Controls / Firm FE / Year FE	Yes	Yes	Yes
Observations	23,044		23,044

Panel C. Accuracy-premium concentration: I/B/E/S benchmark

Visibility measure	Univariate Δ Premium		Horse-race
	High group	Low group	triple diff.
LLM Visibility	5.743*** (5.67) [11,219]	-0.125 (-0.12) [11,825]	5.841*** (4.03)
Firm Size	2.191 (1.28) [11,901]	2.462*** (2.76) [11,143]	2.010 (0.82)
Analyst Coverage	0.441 (0.29) [11,861]	3.198*** (3.59) [11,183]	-4.099* (-1.77)
Inst Ownership	2.564** (2.46) [11,593]	2.418** (2.08) [11,451]	0.625 (0.41)
Controls / Firm FE / Year FE	Yes	Yes	Yes
Observations	23,044		23,044

Notes. Each measure (LLM visibility; firm size, log market equity; analyst coverage, log average analyst

count; institutional ownership) is split at its cross-firm median; the visibility score is discrete, so firms tied at the median join the Low group, which is why the split is not exactly even. **Panel A:** cell (i, j) is the percentage of firms in measure i 's High group that measure j also classifies as High (diagonal = 100). **Panels B and C:** $\Delta\text{Premium}$ is the post-pre change in the accuracy premium $\beta(SUE^{Acc}) - \beta(SUE^X)$ against benchmark $X =$ the equal-weighted five-FDP consensus (Panel B) or I/B/E/S (Panel C), estimated within each measure's High and Low group; the horse-race triple $(SUE^{Acc} - SUE^X) \times \text{Post}^{\text{Scr}} \times \text{High}_m$ is from one pooled regression including the triple for every measure jointly. Throughout, Post^{Scr} is the post-scraping indicator (announcements on or after the onset of browsing, 2023-09-27), as in the cross-sectional tests (Table 7). All specifications include the full control set, firm and year fixed effects, and firm-clustered standard errors; continuous variables winsorized at (1, 99). t -statistics in parentheses and group observation counts in brackets (each measure's High and Low groups partition the same 23,044 firm-quarters); * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Variable definitions are in [Appendix A](#).

TABLE 10. The Predictive Value of the I/B/E/S Consensus: Measurement and Post-ChatGPT Pricing

Dep. Var.:	Single post-ChatGPT indicator		GPT-capability regimes	
	Mkt CAR (1)	FF4 CAR (2)	Mkt CAR (3)	FF4 CAR (4)
$SUE^{IBES} (\beta_1)$	2.237*** (8.71)	2.241*** (9.02)	2.237*** (8.74)	2.241*** (9.05)
$SUE^{IBES} \times Post (\beta_2)$	0.534 (1.54)	0.565* (1.67)		
$SUE^{IBES} \times \text{Regime 1 (text-only)}$			0.868* (1.91)	1.016** (2.29)
$SUE^{IBES} \times \text{Regime 2 (browsing)}$			0.785 (1.62)	0.591 (1.18)
$SUE^{IBES} \times \text{Regime 3 (native search)}$			0.162 (0.33)	0.268 (0.57)
$SUE^{IBES} \times \text{Predictive Value}^{IBES} (\beta_3)$	2.172*** (4.03)	2.119*** (4.11)	2.174*** (4.05)	2.125*** (4.13)
$SUE^{IBES} \times \text{Predictive Value}^{IBES} \times Post (\beta_4)$	2.049*** (2.68)	2.098*** (2.82)		
$SUE^{IBES} \times \text{Predictive Value}^{IBES} \times \text{Regime 1}$			0.919 (0.93)	0.941 (0.97)
$SUE^{IBES} \times \text{Predictive Value}^{IBES} \times \text{Regime 2}$			1.435 (1.45)	1.777* (1.80)
$SUE^{IBES} \times \text{Predictive Value}^{IBES} \times \text{Regime 3}$			3.540*** (3.08)	3.365*** (3.02)
Lower-order terms / Controls	Yes	Yes	Yes	Yes
Firm FE / Year FE	Yes	Yes	Yes	Yes
Observations	22,356	22,356	22,356	22,356
Adjusted R^2	0.082	0.083	0.082	0.083

Notes. This table measures, for each firm, the predictive value of the I/B/E/S consensus — how closely it has anticipated the firm’s future street numbers — and tests whether the post-ChatGPT strengthening of announcement-window reactions concentrates in firms whose consensus predicts well. $Predictive\ Value_{i,q}^{IBES} = 1 - error_{i,q}$, where $error_{i,q}$ is the consensus’s trailing relative anticipation error: the mean own-basis distance $|actual_{i,q+4} - consensus_{i,q}|$ over up to the twenty most recent source quarters whose year-ahead actuals are public at the announcement (at least twelve required; a stock-split guard applies), scaled by the window mean $|actual|$ so the measure is comparable across firms, and winsorized at (1,99). The measure is continuous (mean 0.56, standard deviation 0.30), higher values meaning the consensus has anticipated the firm’s year-ahead street numbers more closely, and is implementable in real time. Coverage is 22,364 of the 23,833 firm-quarters (8,730 pre-ChatGPT), with no industry restriction; the uncovered remainder are recent listings with fewer than twelve scoreable history quarters. The same construction at the current-quarter horizon yields the consensus’s trailing *accuracy*; across firms the two are highly correlated (Pearson 0.72, Spearman 0.78; 0.82 corrected for split-half measurement error), so the accuracy dimension used throughout the paper carries most of this informativeness content and is preferred there for econometric convenience — the current-quarter error requires no four-quarter maturation and is observable in real time. SUE^{IBES} is the I/B/E/S consensus surprise;

all columns include the Table 5 controls (including dispersion, volatility, and the unexpected-item flag), with firm and year fixed effects and firm-clustered standard errors. The pre-period gradient β_3 reflects the familiar association between predictability and the ERC; the identified claim is the bolded post-ChatGPT increment β_4 , comparable in size to the pre-period gradient — the gradient roughly doubles after ChatGPT (about 0.6 per standard deviation of the measure). The level terms are indistinguishable from zero: high-predictive-value firms earn no different announcement return per se, and the effect operates entirely through the *SUE* loading. Columns (3)–(4) replace the single *Post* indicator with the three GPT-capability regimes of Appendix A (text-only from November 30, 2022; browsing from September 27, 2023; native search from October 31, 2024): the increment rises monotonically across regimes and reaches significance under native search — the same capability profile as the accuracy premium of Table 6, and consistent with the retrieval-intensive nature of a forecast track record, which browsing and native search made accessible. The native-search regime spans only the final five sample quarters, so the regime profile is descriptive and the pooled increment in columns (1)–(2) is the formal test. The increment is not announcement composition in disguise: jointly adding interactions of a large-unexpected-item flag (seven Compustat non-recurring items mapped per [Bochkay et al., 2022](#), top-decile magnitude screen) leaves β_4 unchanged (2.04, $t = 2.67$, market-adjusted; 2.09, $t = 2.83$, FF4), and the two measures are essentially uncorrelated (-0.02). Nor is it selection toward I/B/E/S specifically: in an untabulated decomposition of the surprise into the component common to all providers and the I/B/E/S-specific tilt, both the gradient (2.3, $t = 3.7$) and its post-ChatGPT increment (1.9, $t = 2.0$) price the common component, so the evidence concerns the consensus signal of well-predicted firms rather than provider choice. A within-firm-quarter *four-provider rank* version of the same score — each provider’s own-basis error against its own year-ahead actual — gives a stacked post-period rank-ERC gradient of 0.320 ($t = 4.14$) and a post-pre shift of 0.242 ($t = 2.56$), concentrated in Capital IQ provider-events. t -statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

TABLE 11. Additional Tests: Outage Shock and Information Acquisition

Panel A. ChatGPT outage: accuracy premium on the day after a major outage

Dep. Var.: day-0 AR_0 (market-adj., %)	Benchmark: Equal-Weighted (1)	Benchmark: IBES (2)
$SUE^{Acc} (\beta_1)$	1.631*** (6.93)	1.599*** (7.78)
$SUE^{Acc} \times Outage (\beta_3)$	1.178 (0.80)	-1.685* (-1.85)
Benchmark $SUE (\beta_2)$	0.626*** (3.13)	0.654*** (4.10)
Benchmark $SUE \times Outage (\beta_4)$	-2.608* (-1.79)	1.488** (2.27)
Premium (non-outage)	1.005** (2.48)	0.945*** (2.88)
Premium (outage)	4.792* (1.77)	-2.228* (-1.65)
Δ Premium (Outage-Non)	3.787 (1.37)	-3.173** (-2.29)
Controls / Firm FE / Year FE	Yes	Yes
Observations	14,402	14,402
Adjusted R^2	0.051	0.052

Panel B. Information acquisition: platform brand-search share by regime $\times$ vendor group

Dep. Var.: $S_{brand}/S_{all\ platforms}$	(1) Pre/Post-GPT	(2) Pre + 3-Regime
<i>Post-GPT $\times$ Free Distributor</i>	0.021*** (15.44)	
<i>Post-GPT $\times$ Free FDP</i>	-0.075*** (-63.80)	
<i>Post-GPT $\times$ Paywalled FDP</i>	-0.071*** (-64.91)	
<i>Regime 1 $\times$ Free Distributor</i>		0.021*** (9.35)
<i>Regime 1 $\times$ Free FDP</i>		-0.075*** (-35.57)
<i>Regime 1 $\times$ Paywalled FDP</i>		-0.070*** (-32.85)
<i>Regime 2 $\times$ Free Distributor</i>		0.020*** (10.43)
<i>Regime 2 $\times$ Free FDP</i>		-0.073*** (-33.39)
<i>Regime 2 $\times$ Paywalled FDP</i>		-0.069*** (-38.23)

<i>Regime 3 × Free Distributor</i>		0.022*** (10.96)
<i>Regime 3 × Free FDP</i>		-0.075*** (-42.82)
<i>Regime 3 × Paywalled FDP</i>		-0.073*** (-40.93)
Day-level controls (searches, # earnings, trading day)	Yes	Yes
Calendar-month FE	Yes	Yes
Observations (platform-day)	37,247	37,247
Adjusted R^2	0.046	0.046

Notes. **Panel A** tests whether the accuracy premium collapses on earnings announcements that follow a major ChatGPT outage. The sample is post-ChatGPT firm-quarters; $Outage_{i,q} = 1$ if a ChatGPT incident of at least 240 minutes (from the OpenAI incident log) occurred on the calendar day immediately preceding the announcement (84 outage firm-quarters). The dependent variable is the day-0 market-adjusted abnormal return; both columns include firm and year fixed effects, the full control set, and firm-clustered standard errors. The Accuracy Premium is the loading on SUE^{Acc} minus the loading on the benchmark SUE; the bolded row reports its outage-non-outage difference. **Panel B** reports platform-day evidence on investors' information acquisition: the dependent variable is a platform's share of total platform-specific Google search, regressed on regime $\times$ vendor-group interactions. The 17 platforms are grouped into Free Distributors (Barchart, CNBC, Finviz, Forbes, Investing.com, MarketWatch, Motley Fool, Reuters, Seeking Alpha, TradingView, WSJ, Yahoo Finance), the Free FDP (Zacks), and Paywalled FDPs (Bloomberg, Capital IQ, FactSet, IBES); Free Distributor in the pre-ChatGPT regime is the omitted baseline. Column (1) uses a single Post-GPT indicator; column (2) splits it into three GenAI regimes. Standard errors are HC1. t -statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Variable definitions are in [Appendix A](#).

Appendix A: Variable Definitions

Variable	Definition	Source
Panel A. Across-FDP firm-quarter outcome and surprise variables (indexed by firm i, quarter q).		
$CAR_{i,q,[-1,+1]}^{Mkt}$	Three-day market-adjusted cumulative abnormal return around the earnings-announcement date, $\sum_{t=-1}^{+1} (r_{i,t} - vwret_d_t)$, in percent.	CRSP
$CAR_{i,q,[-1,+1]}^{FFC}$	Three-day FFC four-factor (FF3+UMD) abnormal CAR over $[-1, +1]$ with factor loadings estimated on the $[-260, -11]$ window.	CRSP
$BHAR_{i,q,[-2,+1]}^{Mkt}$	Four-day market-adjusted buy-and-hold abnormal return over $[-2, +1]$.	CRSP
$BHAR_{i,q,[-2,+1]}^{FFC}$	Four-day FFC four-factor buy-and-hold abnormal return over $[-2, +1]$.	CRSP
$SUE_{i,q}^g$	Per-FDP standardized unexpected earnings for provider $g \in \{\text{FactSet, Bloomberg, Zacks, CapIQ, IBES}\}$: $(\text{actual EPS}_{i,q}^g - \text{consensus EPS}_{i,q}^g) / \text{price}_{i,q}$.	FactSet; Bloomberg; Zacks; S&P CapIQ; I/B/E/S
$SUE_{i,q}^{Acc}$	Accuracy-weighted multi-vendor SUE, $\sum_g w_{i,q}^g SUE_{i,q}^g$, where $w_{i,q}^g$ are softmax weights based on each FDP's z -scored four-quarter rolling MAFE.	Calculated
$SUE_{i,q}^{EW5}$	Equal-weighted multi-vendor SUE, $(1/5) \sum_g SUE_{i,q}^g$.	Calculated
$Standardized Pred. Surp_{i,q}^X$	Min-max standardized predicted surprise $(Consensus_{i,q}^{Acc} - Consensus_{i,q}^X) / \text{price}_{i,q}$ for benchmark $X \in \{EW5, IBES\}$.	Calculated
Panel B. Firm and event controls (indexed by firm i, quarter q).		
$Log Market Equity_{i,q-1}$	Natural log of firm i 's market capitalization at the end of quarter $q-1$.	CRSP
$Book-to-Market_{i,q-1}$	Book value of equity divided by market value of equity at quarter-end $q-1$.	Compustat; CRSP
$Leverage_{i,q-1}$	Long-term debt plus debt in current liabilities, scaled by total assets, at $q-1$.	Compustat
$Prior 12-Month Return_{i,q}$	Buy-and-hold raw return over the twelve months ending one month before firm i 's quarter- q announcement.	CRSP
$Log Price_{i,q-1}$	Natural log of share price at quarter-end $q-1$.	CRSP
$Return Volatility_{i,q}$	Standard deviation of daily returns over the $[-260, -11]$ window before the announcement.	CRSP
$Analyst Dispersion_{i,q}$	Standard deviation of analyst EPS forecasts divided by absolute consensus.	I/B/E/S
$Busyness_{i,q}$	Natural log of the number of S&P 1500 firms announcing earnings on firm i 's announcement day.	Compustat

(continued on next page)

Variable	Definition	Source
$ \Delta IB _{i,q}$	Absolute year-over-year change in income before extraordinary items, scaled by lagged total assets.	Compustat
$ \Delta Shares _{i,q}$	Absolute year-over-year change in shares outstanding, scaled by lagged shares.	Compustat
$ SpecItems _{i,q}$	Absolute value of special items, scaled by lagged total assets.	Compustat
$SBC\ High_{i,q}$	Indicator equal to one if stock-based compensation expense exceeds the sample median.	Compustat
$Unexpected\ Item_{i,q}$	Indicator equal to one if the quarter contains a non-recurring item flagged by Compustat.	Compustat
$Stock\ Split_{i,q}$	Indicator equal to one if a stock split was effective within $[-30, +30]$ trading days of the announcement.	CRSP
$Fiscal\ Q4_{i,q}$	Indicator equal to one for fiscal fourth-quarter announcements.	Compustat
$Institutional\ Ownership_{i,q-1}$	Fraction of shares held by institutional investors at $q-1$.	Thomson Reuters 13F
$\#Analysts_{i,q}$	Number of unique analysts in the IBES detail file forecasting EPS for firm-quarter (i, q) .	I/B/E/S

Panel C. Cross-sectional partitioning and identification variables.

$LLM\ Visibility_i$	EPS-extraction indicator $\times \log(1 + \text{unique cited domains})$, averaged across ChatGPT and Perplexity for direct queries about firm i 's quarterly EPS.	DataForSEO
$Cross-FDP\ SUE\ Dispersion_{i,q}$	Inter-quartile range of $SUE_{i,q}^g$ across the five FDPs in firm-quarter (i, q) .	Calculated
$ SpecItems _{i,q} (\%)$	Absolute special items as a percentage of absolute income before extraordinary items, $ SpecItems _{i,q}/ IB _{i,q} \times 100$.	Compustat
$IBES\ Up_i$	Indicator equal to one if IBES's rounded average per-firm-quarter accuracy rank among the five FDPs rose by at least two positions from pre- to post-ChatGPT.	Calculated
$IBES\ Down_i$	Indicator equal to one if IBES's rank fell by at least two positions.	Calculated
$Outage_{i,q}$	Indicator equal to one if a ChatGPT incident of at least 240 minutes occurred on the calendar day immediately preceding firm i 's quarter- q announcement date.	OpenAI
$Large\ Unexpected\ Item\ (LUI_{i,q})$	Indicator equal to one if firm-quarter (i, q) contains a large unexpected item: any of seven Compustat non-recurring items mapped per Bochkay et al. (2022) , with restructuring, acquisition, and special items requiring top-decile magnitude. Refines <i>Unexpected Item</i> (Panel B) with a magnitude screen.	Compustat
$INFO_{i,q}$	Indicator equal to one if firm-quarter (i, q) contains no large unexpected item, $1-LUI_{i,q}$.	Calculated

Panel D. Information-acquisition platform-day variables (indexed by platform p , day t).

(continued on next page)

Variable	Definition	Source
$Brand\ Share_{p,t}$	Platform p 's share of total platform-specific Google search volume on day t , $S_{p,t} / \sum_{p'} S_{p',t}$, summed across the 17 sample platforms.	Google Trends; DataForSEO
$Log\ Total\ Platform\ Searches_t$	Natural log of total Google search volume across the 17 sample platforms on day t .	Google Trends; DataForSEO
$Log\ \# Earnings\ Today_t$	Natural log of one plus the number of S&P 1500 firms announcing earnings on day t .	Compustat
$Trading\ Day_t$	Indicator equal to one on NYSE trading days; zero otherwise.	CRSP
$Free\ Distributor_p$	Indicator equal to one for platforms in the Free Distributors group (Barchart, CNBC, Finviz, Forbes, Investing.com, MarketWatch, Motley Fool, Reuters, Seeking Alpha, TradingView, WSJ, Yahoo Finance).	Classification
$Free\ FDP_p$	Indicator equal to one for Zacks.	Classification
$Paywalled\ FDP_p$	Indicator equal to one for Bloomberg, Capital IQ, FactSet, and IBES.	Classification
$Post-ChatGPT_t$	Indicator equal to one for days on or after 2022-11-30 (ChatGPT public release).	OpenAI
$Regime\ k_t\ (k = 1, 2, 3)$	Mutually exclusive indicators for the three post-GenAI regimes (Regime 1: 2022-11-30 to 2023-09-26; Regime 2: 2023-09-27 to 2024-08-31; Regime 3: 2024-09-01 onward).	OpenAI

Appendix B: Worked Examples of GenAI Aggregation of Analyst Consensus

This appendix shows the prompt and ChatGPT (gpt-4o) responses for two illustrative large-cap firms, IBM and Microsoft. Responses were obtained through the DataForSEO LLM Responses API and represent the unedited output of the model. Each firm subsection ends with a summary table mapping the cited aggregators to their upstream forecast data provider (FDP).

B.0 Prompt

What is the analyst consensus EPS forecast for [TICKER] for the next quarter? Please report the figures compiled by multiple third-party aggregators and cite each source.

B.1 International Business Machines Corporation (IBM)

ChatGPT (gpt-4o) response:

Here's the latest analyst consensus EPS forecast for IBM's next quarter, as compiled from several third-party aggregators.

- **MarketBeat (via Fiscal.ai):** Q2 2026 average EPS estimate is **\$2.96**, based on 2 estimates (low \$2.92, high \$2.99) ([marketbeat.com](https://www.marketbeat.com)).
- **TipRanks:** next quarter's earnings estimate is **\$3.02**, with a range from \$2.92 to \$3.09 ([tipranks.com](https://www.tipranks.com)).
- **Barchart:** next quarter average earnings estimate is **\$2.96**, based on 6 estimates (low \$2.73, high \$3.50) ([barchart.com](https://www.barchart.com)).
- **Yahoo Finance:** Jun 2026 average EPS estimate is **\$3.04**, based on 14 analysts (low \$2.80, high \$3.21) (finance.yahoo.com).

Summary table.

Source	Forecast	Upstream FDP
MarketBeat	\$2.96	Independent aggregator
TipRanks	\$3.02	Independent aggregator
Barchart	\$2.96	Independent aggregator
Yahoo Finance	\$3.04	S&P Capital IQ

B.2 Microsoft Corporation (MSFT)

ChatGPT (gpt-4o) response:

Here are the analyst consensus EPS forecasts for Microsoft (MSFT) for the **next quarter**, as compiled by multiple third-party aggregators. Each figure is accompanied by its source.

- **MarketBeat:** for Q4 FY2026, the consensus EPS estimate is **\$4.04**, based on 5 analyst forecasts ([marketbeat.com](https://www.marketbeat.com)).
- **Zacks Investment Research:** consensus estimate for the next quarter (June 2026) is **\$4.13**, derived from 13 analyst estimates ([zacks.com](https://www.zacks.com)).
- **TipRanks:** next EPS estimate is **\$4.24** ([tipranks.com](https://www.tipranks.com)).

- **ChartMill:** consensus EPS estimate of **\$4.33** for the next earnings (Q4 FY2026), based on 39 analysts (chartmill.com).
- **Seeking Alpha:** an earnings preview lists the consensus EPS at \$4.06 for Q3 FY2026 (already reported), so not applicable for the upcoming quarter (seekingalpha.com).
- **Yahoo Finance:** for the period “Next qtr. (Jun 2026)”, average EPS estimate is **\$4.26**, based on 29 analysts (low \$3.74, high \$5.54) (finance.yahoo.com).
- **Investing.com:** I was unable to locate a clear next-quarter EPS consensus figure (investing.com).

Summary table.

Source	Forecast	Upstream FDP
MarketBeat	\$4.04	Independent aggregator
Zacks	\$4.13	Zacks
TipRanks	\$4.24	Independent aggregator
Yahoo Finance	\$4.26	S&P Capital IQ
ChartMill	\$4.33	Independent aggregator

Appendix C: Numerical Illustration of the Accuracy-Weighted SUE

This appendix illustrates how the accuracy-weighted SUE is constructed for a single firm-quarter. Suppose firm i in quarter q has the following backward-looking accuracy histories ($acc_{i,q}^g = -MAFE_{i,q}^g$ over the prior four quarters, in price-scaled units) across the five FDPs, alongside their current-quarter standardized surprises $SUE_{i,q}^g$ (in percent of price).

Panel A. Weight construction from accuracy history.

Provider g	$acc_{i,q}^g$	$acc_z_{i,q}^g$	$\exp(acc_z_{i,q}^g)$	$w_{i,q}^g$
IBES	-0.0010	+1.5	4.48	0.57
FactSet	-0.0020	+0.5	1.65	0.21
Capital IQ	-0.0030	0.0	1.00	0.13
Zacks	-0.0040	-0.8	0.45	0.06
Bloomberg	-0.0050	-1.2	0.30	0.04
Sum			7.88	1.00

I/B/E/S has the smallest absolute forecast error (largest acc^g) and therefore the highest within-row z -score, which the softmax transform maps to the largest weight (0.57). Bloomberg has the largest absolute forecast error and receives the smallest weight (0.04). The weights sum to one and reduce to the uniform 0.20 vector when the five accuracy scores are equal.

Panel B. Composite construction.

Provider g	$SUE_{i,q}^g$	$w_{i,q}^g \times SUE_{i,q}^g$	$0.20 \times SUE_{i,q}^g$
IBES	+1.20	+0.68	+0.24
FactSet	+0.80	+0.17	+0.16
Capital IQ	+1.00	+0.13	+0.20
Zacks	+1.50	+0.09	+0.30
Bloomberg	-0.50	-0.02	-0.10
Total		$SUE_{i,q}^{Acc} = +1.05$	$SUE_{i,q}^{EW5} = +0.80$

The accuracy-weighted composite $SUE_{i,q}^{Acc} = +1.05$ exceeds the equal-weighted composite $SUE_{i,q}^{EW5} = +0.80$ because the two most-accurate providers in this row (I/B/E/S and FactSet) reported positive surprises, while the least-accurate provider (Bloomberg) was the lone outlier on the downside. The accuracy tilt $SUE_{i,q}^{Acc} - SUE_{i,q}^{EW5} = +0.25$ captures the accuracy-driven correction. In the horse-race regression of Table 6, this tilt is the source of identification between the two composites: when the tilt is informative about announcement-window returns beyond the equal-weighted aggregate, the market is loading on provider-specific accuracy rather than treating all five FDPs symmetrically.

ONLINE APPENDIX

“Synthesizing the Consensus: Generative AI, LLM Visibility, and the Integration of Multi-Provider Consensus Forecasts”

This Online Appendix validates accuracy as the weighting dimension of the multi-vendor synthesis. Table [OA1](#) shows that the accuracy tilt is a weakly dominant design choice: the accuracy-weighted composite is strictly closer to realized earnings than either the equal-weighted or the I/B/E/S-only heuristic in both regimes (Panel A), while the five providers’ consensus slopes for future fundamentals are statistically indistinguishable (Panel B) — so accuracy weighting strictly gains precision at no cost in fundamental informativeness.

TABLE OA1. Validating Accuracy as the Weighting Dimension: Strict Precision Gains, No Informativeness Cost

Panel A. Expectational precision: paired differences in $|SUE|$ against the accuracy-weighted composite.

	Pooled	Pre-ChatGPT	Post-ChatGPT
$ SUE^{EW5} - SUE^{Acc} $	0.033*** (13.30)	0.033*** (10.03)	0.033*** (10.94)
Share with $ SUE^{Acc} $ strictly smaller	52.3%	53.6%	51.5%
$ SUE^{IBES} - SUE^{Acc} $	0.045*** (14.44)	0.049*** (12.10)	0.042*** (11.99)
Share with $ SUE^{Acc} $ strictly smaller	62.0%	62.7%	61.6%
Observations	23,833	9,426	14,407

Panel B. Fundamental informativeness: consensus slopes for four-quarter-ahead outcomes.

	FactSet	Zacks	CapIQ	IBES
Future cash flows (θ_g)	1.016** (2.11)	0.807** (2.08)	0.822** (1.98)	0.845** (2.00)
Wald test of slope equality		$F = 0.80, p = 0.493$		
Future street earnings (θ_g)	1.260** (2.46)	1.098*** (2.66)	1.198** (2.50)	1.282** (2.54)
Wald test of slope equality		$F = 1.75, p = 0.155$		

Notes. This table validates the choice of forecast accuracy as the weighting dimension of the manuscript's composite. **Panel A** reports Diebold–Mariano-style paired tests: the mean difference in absolute standardized surprise between each heuristic composite and the accuracy-weighted composite, estimated on the full firm-quarter panel with firm-clustered standard errors (the formal counterpart of Figure 2). The accuracy-weighted composite is strictly closer to the realized actual in both regimes against both benchmarks. **Panel B** asks what accuracy-based weighting could cost along the provider-level informativeness dimension of [Bochkay et al. \(2022\)](#): each provider's consensus slope for four-quarter-ahead fundamentals, estimated on a stacked provider panel (2015–2025, the four providers with native point-in-time history; provider-specific intercepts, lagged-cash-flow loadings, and industry effects; firm-clustered standard errors). The slopes are statistically indistinguishable across providers for both outcomes, so tilting weights toward accurate providers sacrifices no fundamental informativeness. Bloomberg, excluded above because its native point-in-time history begins in 2020Q2, is indistinguishable from FactSet, CapIQ, and I/B/E/S on the common 2020Q2-onward window (cash-flow slopes 2.16 versus 2.17–2.18); the five-provider equality test on that shorter window rejects only marginally ($p = 0.09$ cash flows, $p = 0.03$ earnings), driven by a lower Zacks slope consistent with the attenuation from its thin, penny-rounded consensus. Together with the persistence of accuracy rankings (Table 4, Spearman $\rho = 0.71$) and their firm-level variation, the panels imply that accuracy weighting weakly dominates: it strictly improves expectational precision, costs nothing in informativeness, and is the only quality dimension observable and exploitable in real time. t -statistics in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.